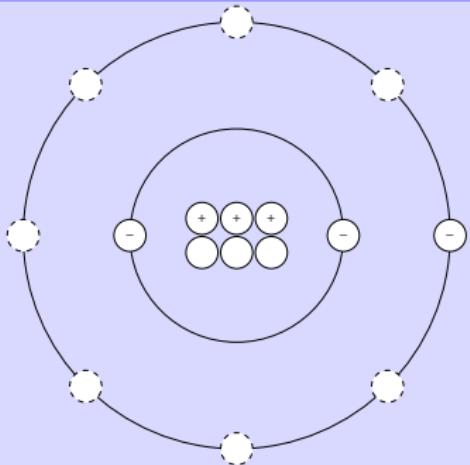


- ▶ The concrete implementation of most digital logic boils down use of **switches**: if the switch is “on” it conducts electricity, if it’s “off” it doesn’t.
- ▶ Higher-level use-cases demand we make the switches as
 1. efficient (in space and time),
 2. robust, and
 3. manufacturableas possible, and thus turn electronics into *micro-electronics*.

Example (lithium atom)

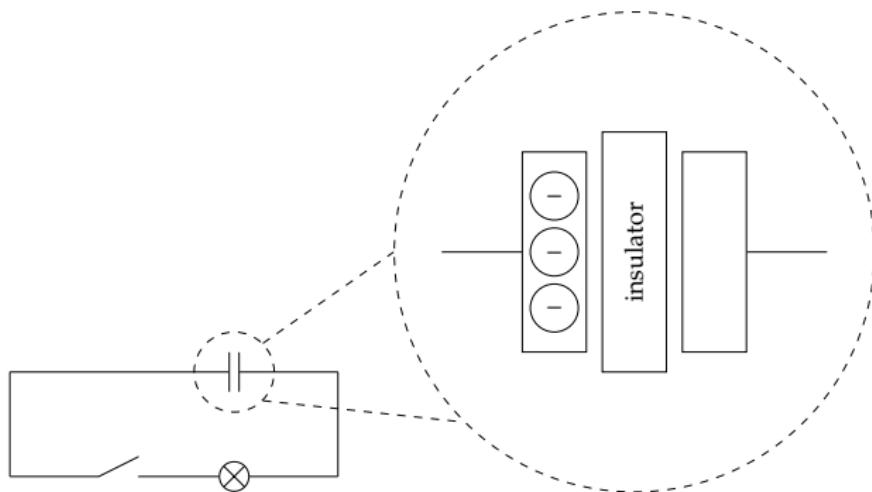


- ▶ An **atom** is built from
 1. a group of nucleons, either **protons** or **neutrons**, called the **nucleus**, and
 2. a cloud of **electrons** arranged in shells.
- ▶ The protons dictate the atomic number; the number of neutrons dictates the isotope.
- ▶ The electron configuration is not arbitrary: the n -th shell can accommodate upto $2n^2$ electrons.
- ▶ An unfilled slot within a shell is called a **hole**.

- ▶ Each sub-atomic particle has a (small) electrical **charge** associated with it
 1. electrons have a negative charge, i.e., $-$,
 2. protons have a positive charge, i.e., $+$, and
 3. neutrons have a neutral charge
- and an **ion** is an atom with a non-zero overall charge.
- ▶ The particles are bound together in a variety of ways ...
- ▶ ... **but** the binding can be disrupted:
 - ▶ As an electron absorbs more energy, it becomes excited; at some threshold, the electron is displaced and becomes free.
 - ▶ A free electron can move between shells **or** between atoms (from an outer shell); roughly speaking they are “attracted” by holes.

“A-level Physics in 10 minutes” (3)

- ▶ An electrical **current** is a flow of electrons, i.e., a flow of charge; free electrons (bound to atoms or not) “move” from low to high potential ...
- ▶ ... this is effectively how a **capacitor** (i.e., a battery) works:



- ▶ A conductivity “rating” says how easily electrons can move:
 - ▶ a **conductor**, e.g., a metal, has high-conductivity (resp. low-resistivity) and allows electrons to move easily, and
 - ▶ an **insulator**, e.g., a vacuum, has low-conductivity (resp. high-resistivity) and does not allow electrons to move easily.

- ▶ Silicon (**Si**) is neat because
 1. it's **abundant**, i.e., there's lots of it and so it doesn't cost much,
 2. it's fairly **inert**, i.e., it's stable enough not to react in weird ways with other things, **and**
 3. it can be “**doped**” with a donor material, e.g.,
 - 3.1 a boron (**B**) or aluminium (**Al**) donor creates extra holes, or
 - 3.2 a phosphorus (**P**) or arsenic (**As**) donor creates extra electronswhich, roughly speaking, allows us to construct materials with the exact sub-atomic properties we want.
- ▶ The result is a *semiconductor*:
 - ▶ A P-type semiconductor has extra holes, N-type has extra electrons.
 - ▶ If we sandwich P-type and N-type layers together, electrons can only move in *one way*, i.e., from an N-type layer to a P-type layer, but *not* in the other direction.

An Aside: History



- ▶ A historically dominant switch technology was the **vacuum tube**:
 - ▶ Looks like a light-bulb with a glass envelope holding a vacuum.
 - ▶ Inside vacuum is an electron producing filament (or cathode) and a metal plate (or anode).
 - ▶ When filament is heated, electrons are produced into vacuum which are attracted by plate.
 - ▶ Reliability of vacuum tubes was reasonably good, but most failed during power-on or power-off ...
 - ▶ ... the terms **bug** and **debug** both allegedly stem (in part) from failure of this sort.

<http://en.wikipedia.org/wiki/File:5651RegulatorTubeInOperation.jpg>

An Aside: History



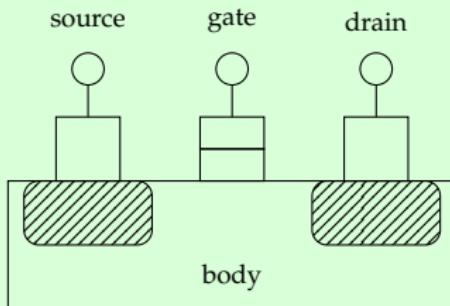
- ▶ The replacement for this generation of technology is the **transistor**.
- ▶ There are *many* types, but a potted overview of the one we'll focus on is:
 - ▶ **1925:** Field Effect Transistor (FET), initially designed and patented by Julius Lilienfeld ...
 - ▶ ... design not really used, due to lack of understanding and means of manufacture.
 - ▶ **1952:** team led by William Shockley at Bell Labs invent junction gate FET (or JFET).
 - ▶ **1960:** Dawon Kahng and Martin Atalla, also at Bell Labs invent the **Metal Oxide Semiconductor FET (MOSFET)**.

<http://en.wikipedia.org/wiki/File:Replica-of-first-transistor.jpg>

Switches and Transistors (1)

- ▶ At a high level, the design of a MOSFET is as follows:

Circuit (MOSFET)

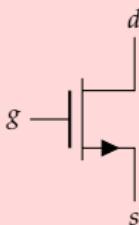


- ▶ The (very) rough idea is ...
 - ▶ Apply a potential difference between gate and source; this controls the size of a **depletion region** between the source and drain.
 - ▶ A large potential difference means a small depletion region, allowing electrons, i.e., current, to flow in a **channel** that links the source and drain.
 - ▶ A small potential difference (smaller than some **threshold**) means a large depletion region which closes the channel.
- ▶ ... a more detailed description needs more Physics!

Switches and Transistors (2)

Definition (N-MOSFET)

- ▶ Written symbolically as

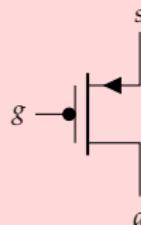


where d , s and g denote the drain, source and gate terminals.

- ▶ Built using **N-type** semiconductor terminals and **P-type** body.
- ▶ A positive potential difference between source and gate widens the depletion channel, meaning source and drain are connected.
- ▶ A negative potential difference between source and gate narrows the depletion channel, meaning source and drain are disconnected.

Definition (P-MOSFET)

- ▶ Written symbolically as



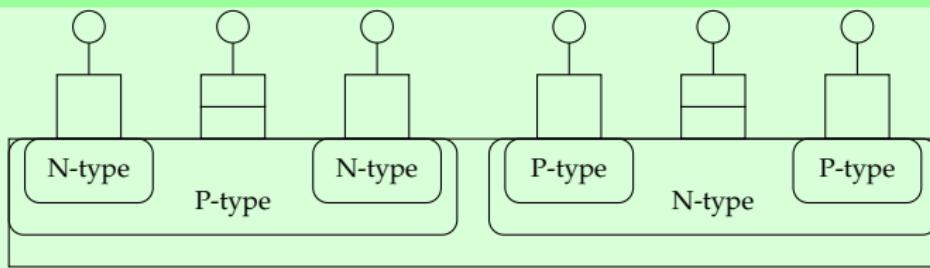
where s , d and g denote the source, drain and gate terminals.

- ▶ Built using **P-type** semiconductor terminals and **N-type** body.
- ▶ A positive potential difference between source and gate narrows the depletion channel, meaning source and drain are disconnected.
- ▶ A negative potential difference between source and gate widens the depletion channel, meaning source and drain are connected.

Switches and Transistors (3)

- ▶ ... **but** they aren't typically used in isolation.
- ▶ A **CMOS** cell combines one N-MOSFET and one P-MOSFET; each transistor in a pair works in a **complementary** way:

Circuit (CMOS cell)



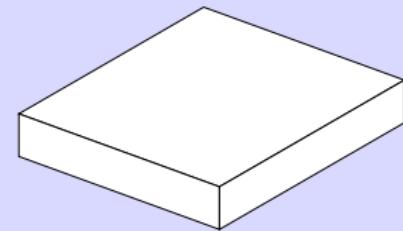
- ▶ We'll see later why this is functionally useful, but behaviourally
 1. only switching from one state to another consumes much power (or **dynamic consumption**) since one transistor does the opposite of the other, thus
 2. there isn't much "leakage" (or **static consumption**) at other times.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

Example (photolithography)



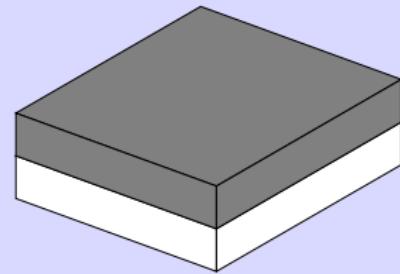
- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

Example (photolithography)



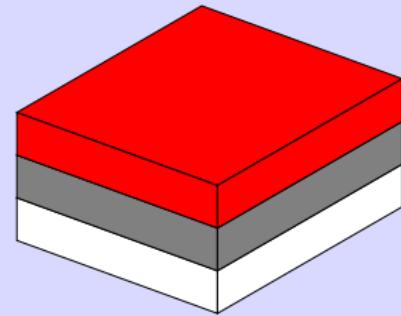
- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

Example (photolithography)



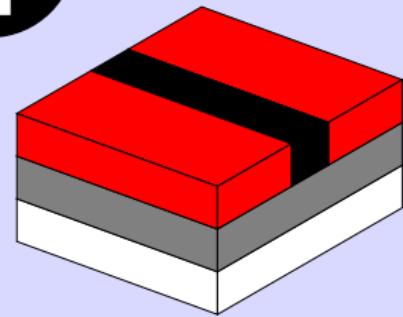
- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

Example (photolithography)



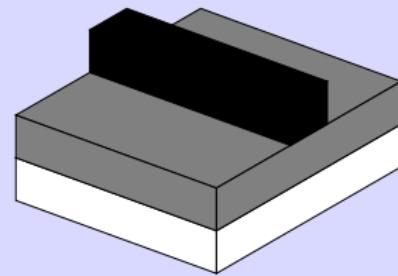
- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

Example (photolithography)



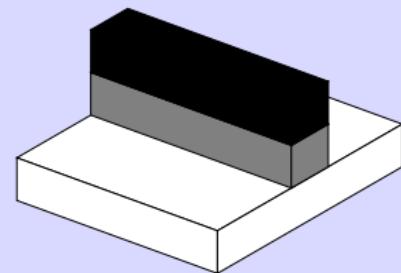
- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

Example (photolithography)



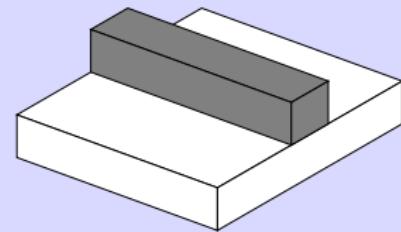
- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (1)

Algorithm (photolithography)

1. Start with a clean, prepared **wafer**.
2. Apply a layer of **substrate** material, e.g., metal or semiconductor.
3. Apply a layer of **photoresist** material.
4. Expose the photoresist to a precise negative or **mask** of design; this hardens the exposed photoresist.
5. Wash away unhardened photoresist.
6. Etch away uncovered substrate.
7. Strip hardened photoresist.

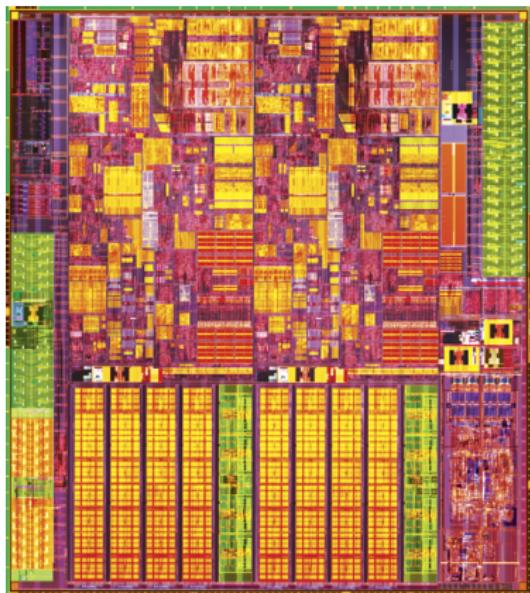
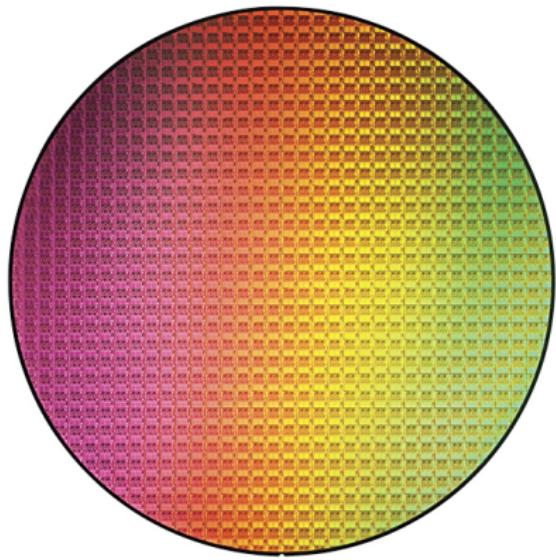
Example (photolithography)



- ▶ Note that:
 - ▶ The algorithm iterates to produce many layers, i.e., the result is 3D not 2D.
 - ▶ Regularity is a major advantage: we can manufacture *many* similar components in a layer using *one* photolithographic process.
 - ▶ The feature size (e.g., 90 nm CMOS) relates to the resolution of this process, e.g., width of wires or density of transistors.

Manufacture (2)

- The end result (left-hand side) is a wafer where ...

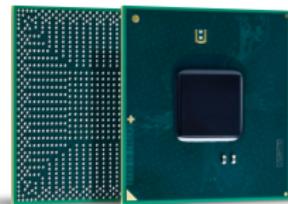


- ... each "block" is an individual component, e.g., a Westmere model Intel processor (right-hand side).

http://www.intel.com/pressroom/archive/releases/2010/20100107comp_sm.htm

Manufacture (3)

- ▶ The components aren't much use in this form; they are **packaged** before use:



- ▶ Note that the packaging acts as
 - ▶ protection against damage, often including a heat sink as well, **and**
 - ▶ an interface between the component and the outside world, using external pins bonded to internal inputs and outputs.

http://www.intel.com/pressroom/archive/releases/2010/20100107comp_sm.htm

Moore's Law: originally an observation

The complexity for minimum component costs has increased at a rate of roughly a factor of two per year. Certainly over the short term this rate can be expected to continue, if not to increase. Over the longer term, the rate of increase is a bit more uncertain, although there is no reason to believe it will not remain nearly constant for at least 10 years. That means by 1975, the number of components per integrated circuit for minimum cost will be 65,000.

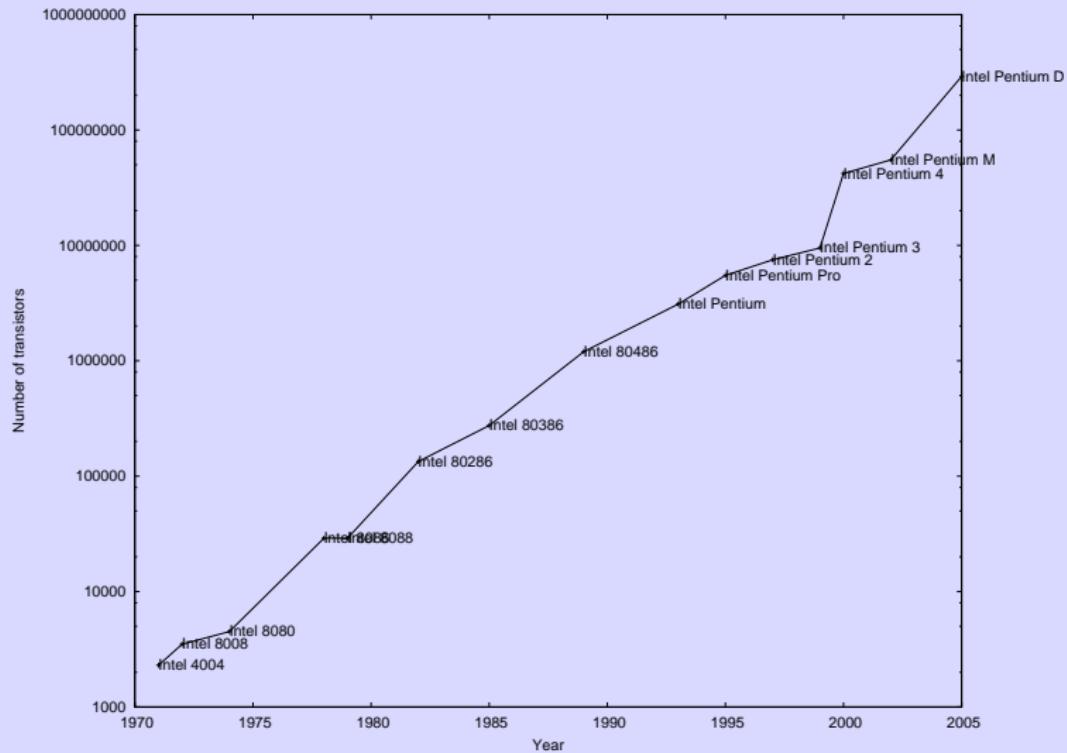
– Moore [5], 1965

and later updated: in short “number of transistors in unit area doubles roughly every two years”.



Manufacture (5)

Example (Moore's law, from 1970 to 2005)



Manufacture (6)

- ▶ Great: with smaller transistors we can
 1. produce more complicated designs that use more transistors, **or**
 2. shrink existing designs, making them smaller and faster

without really thinking about anything!
- ▶ **But**, there's no "free lunch" because (as Moore notes)
 1. **power consumption** and **heat dissipation** become more and more important, and
 2. process variation, which may imply **defects** and reduce yield, starts to increase

so *smarter* design techniques are crucially important.

Conclusions

► Take away points:

1. Semiconductor and transistor design and manufacturing technology lurks in the background of almost *everything* we do ...
2. ... transistors *are* fairly complicated wrt. Physics, but from a conceptual perspective they're *just* switches.
3. Various implications of modern transistor design place tangible constraints on us: at least a high-level understanding of these (e.g., Moore's Law etc.) is valuable.

► Don't panic!

- ▶ A reasonable comment is "WTF?! I thought this was CS, not EE!".
- ▶ The point is to now abstract away the detail:
 - ▶ A micro-processor, at a low-level, is just a complicated arrangement of transistors (i.e., switches).
 - ▶ Now we know how transistors work, we can start to focus more on their function and less on their form.
 - ▶ That is, we'll view them as "black boxes" and build high-level components on top of the functionality they afford.

References and Further Reading

- [1] Wikipedia: Hydraulic analogy.
http://en.wikipedia.org/wiki/Hydraulic_analogy.
- [2] Wikipedia: Moore's law.
http://en.wikipedia.org/wiki/Moore's_law.
- [3] Wikipedia: Semiconductor device fabrication.
http://en.wikipedia.org/wiki/Semiconductor_fabrication.
- [4] Wikipedia: Transistor.
<http://en.wikipedia.org/wiki/Transistor>.
- [5] G.E. Moore.
[Cramming more components onto integrated circuits.](#)
Electronics Magazine, 38(8):114–117, 1965.
- [6] D. Page.
[Chapter 2: Basics of digital logic.](#)
In *A Practical Introduction to Computer Architecture*. Springer-Verlag, 1st edition, 2009.
- [7] W. Stallings.
[Chapter 11: Digital logic.](#)
In *Computer Organisation and Architecture*. Prentice-Hall, 9th edition, 2013.

References and Further Reading

- [8] A.S. Tanenbaum.
Section 3.1: Gates and Boolean algebra.
In *Structured Computer Organisation* [10].
- [9] A.S. Tanenbaum.
Section 3.2: Basic digital logic circuits.
In *Structured Computer Organisation* [10].
- [10] A.S. Tanenbaum.
Structured Computer Organisation.
Prentice-Hall, 6th edition, 2012.