

Instruction Set Architectures (ISAs)

Simon Hollis, COMS12200

What's in an ISA?

- An Instruction Set Architecture (ISA) is a definition of the instructions supported by a particular processor.
- The ISA specifies the exact behaviour of instructions fed to it.
- The ISA should be designed sympathetically for the programming language that will write code for it.

The 'I'

- What exactly is an 'instruction'?
- It's a definition of a unique operation code PLUS some contextual information.
- We say there is an operation plus zero or more arguments.

The 'I'

- We've seen this before with **1 01** being *decodable*.
- **1** is the operation code (*op-code*)
 - Here 1 means '*move*'
- **01** is the argument. Here, it is the literal/constant/immediate value **1**

Decoding

- All instructions must be *uniquely decodable*.
- One-to-one mapping of bit-string to meaning.
- Some ISAs do not define all possible bit-strings.
- This is allowed, but can cause problems if an undefined instruction is encountered.

Creating instructions

- So, to create an ISA we need to produce a list of instructions and their possible arguments.
- What instructions make it to the set are based on experience and need.

Instruction classes

- Most instructions fall naturally into one of a few classes:
 - Arithmetic
 - *e.g. Add, subtract, multiply.*
 - Comparison
 - *e.g. Compare, If, set flags*
 - Memory
 - *e.g. Loads, stores, swaps.*
 - Control flow
 - *e.g. Branches, state changes*

Instruction classes

- Simple processors need only the previous classes, but complex, high-performance ones may employ:
 - Vector processing
 - *e.g. Array manipulation, block operations.*
 - Digital Signal Processing (DSP)
 - *e.g. Fast Fourier Transform (FFT)*

Instruction selection

- How do you go about selecting what instructions should be part of an ISA?
- #1: Think about the tasks in hand
 - Common operations
 - Representative applications
 - Future-proofing?

ISAs

ISA CATEGORISATION: RISC AND CISC

RISC and CISC

- We often categorise ISAs by how close they are to one of two opposing paradigms.
- RISC (Reduced Instruction Set Computer)
- CISC (Complex Instruction Set Computer)

RISC

- RISC (Reduced Instruction Set Computer)
 - The aim is to produce a minimal set of instructions that are *lightweight* and *fast* to execute.
 - Each instruction should have a *similar run time*.
 - Complex operations must be *synthesised* from simple ones.
 - RISC machines run fast but may need lots of instructions to do a complex job.

CISC

- CISC machines are the opposite:
- If there is a specialised operation to be done, add a ***dedicated instruction*** for it (up to a limit!)
- Rich instruction set means ***smaller programs***.
- Mixed-length instructions means ***unpredictability*** and compiler headache.
- Processor ends up ***complex*** and ***slow***.

ISAs

ISA OPTIONS

Instruction choices

- For many operations, you may have a set of choices
 - e.g. `r0 <- r0 + r1`
 - How many ways can we do this?

$$r0 \leftarrow r0 + r1$$

- Can have multiple accumulators:
ACC1, ACC2

1. **ACC1** \leftarrow **r0**

2. **ACC2** \leftarrow **r1**

3. **ACC1** \leftarrow **ACC1** + **ACC2**

4. **r0** \leftarrow **ACC1**

- This takes 4, very simple instructions.

$r0 \leftarrow r0 + r1$

- or, using a stack
 - **PUSH [r0]**
 - **PUSH [r1]**
 - **ADD**
 - **POP [r0]**
- (r0 here assumed to be a memory address here)

$r0 \leftarrow r0 + r1$

- We can get it down to 3 instructions
 - **$ACC \leftarrow r0$: `MOV ACC, r0`**
 - **$ACC \leftarrow ACC + r1$: `ADD r1`**
 - **$r0 \leftarrow ACC$: `MOV r0, ACC`**

$r0 \leftarrow r0 + r1$

- ...or in 2 instructions
 - **MOV ACC, r2**
 - **ADD r0, r0, ACC**

$r0 \leftarrow r0 + r1$

- We can do better
- Set r0 to always be the destination.
 - $r0 \leftarrow r0 + r1 : (\text{ADD } r0, r1)$
- Or, even better -- r0 also always the argument:
 - $r0 \leftarrow r0 + r1 : (\text{ADD } r1)$

Instruction set choice

- So, how do I choose between all the possibilities?
- It's a design choice, taking into account all the tradeoffs

ISA tradeoffs

- Tradeoff 1: more instructions -> more bits
 - More instructions -> greater information
 - But: more space, more energy
- Tradeoff 2: more operands -> more bits
 - More bits -> more work
 - More space, more energy

ISA implications

- Choices of instruction specificity impact how a machine may be programmed.
- Can imply a choice of paradigm:
 - Implicit source & destination: stack machine
 - Implicit destination: accumulator machine
 - Explicit source and destinations: register machine

Picking a set

- We want to cover the useful space of execution needs.
- *No repetition.*
- *No coverage where it is not needed.*
- *Minimum size* to do the job.
- Check if instruction addition is harmful.
 - Recall Amdahl's Law

Amdahl's Law

- Amdahl's Law is a measure of utility of adding an instruction to an ISA.
- What it says is don't add an instruction if it slows down an application overall.

Amdahl's Law

If a new instruction is added to speed up an operation by **s** times, the law gives the maximum overall speedup:

$$\text{speedup} = s / 1 + (f * (s - 1))$$

where **0 ≤ f ≤ 1** is the fraction of time spent *NOT* running the operation in question in the original design

ISAs

ENCODING INSTRUCTIONS

Instruction encodings

- Once we have selected our set of instructions, we need to encode them *unambiguously*.
- A common way is to separate them into code *prefixes* and *operand* areas.

Instruction encodings

- *Space* is a major factor in deciding instruction encodings.
 - Space -> memory size
 - Space -> no of instructions loadable/second
 - Space -> decode stage complexity

Instruction lengths

- Instruction lengths can, therefore vary
 - Some instructions need to provide **more information** than others -> **longer instruction**
 - If instructions can be shorter, that's good.
- All this implies that ***fixed length*** instructions may be sub-optimal.

Instruction lengths

- Consider these two instructions:
- **ADD r0, r0, r1**
 - Op-code + 3 operands
 - 16 registers -> 12 bits + op-code
- **HALT**
 - Only an op-code.
- ***Variable length*** instructions can save space.

Variable or fixed length?

- Variable length instructions:
 - Size efficient
 - Can be extended to support more operands
 - Optimal in some sense.
- BUT: almost infinite amount of possibilities causes decoding to be very complex
 - -> pipeline slowdown

Fixed length instructions

- May be **sub-optimal** in length
- Can get around this by allocating op-codes as part of a **prefix** code.
 - e.g. **100** - *Branch*
 - **110** - *Branch equal to*
 - **111** - *Branch not equal to*
- Major advantage in simplified decode stage.
- Tends to be a winner in most modern designs.

Conclusions

- Creating an ISA is a complicated task
- There are almost infinite ways to do it
- It's a very creative thing
- A choice of ISA has big implications for users of a processor. BUT:
 - History shows that bad ISAs can flourish (e.g. x86)
 - Good ISAs can become bad over time (or vice-versa (e.g. VAX, MIPS))