

Instruction Set Architectures (ISAs)

Simon Hollis, COMS12200

Example ISA

The ARM Thumb V1 ISA

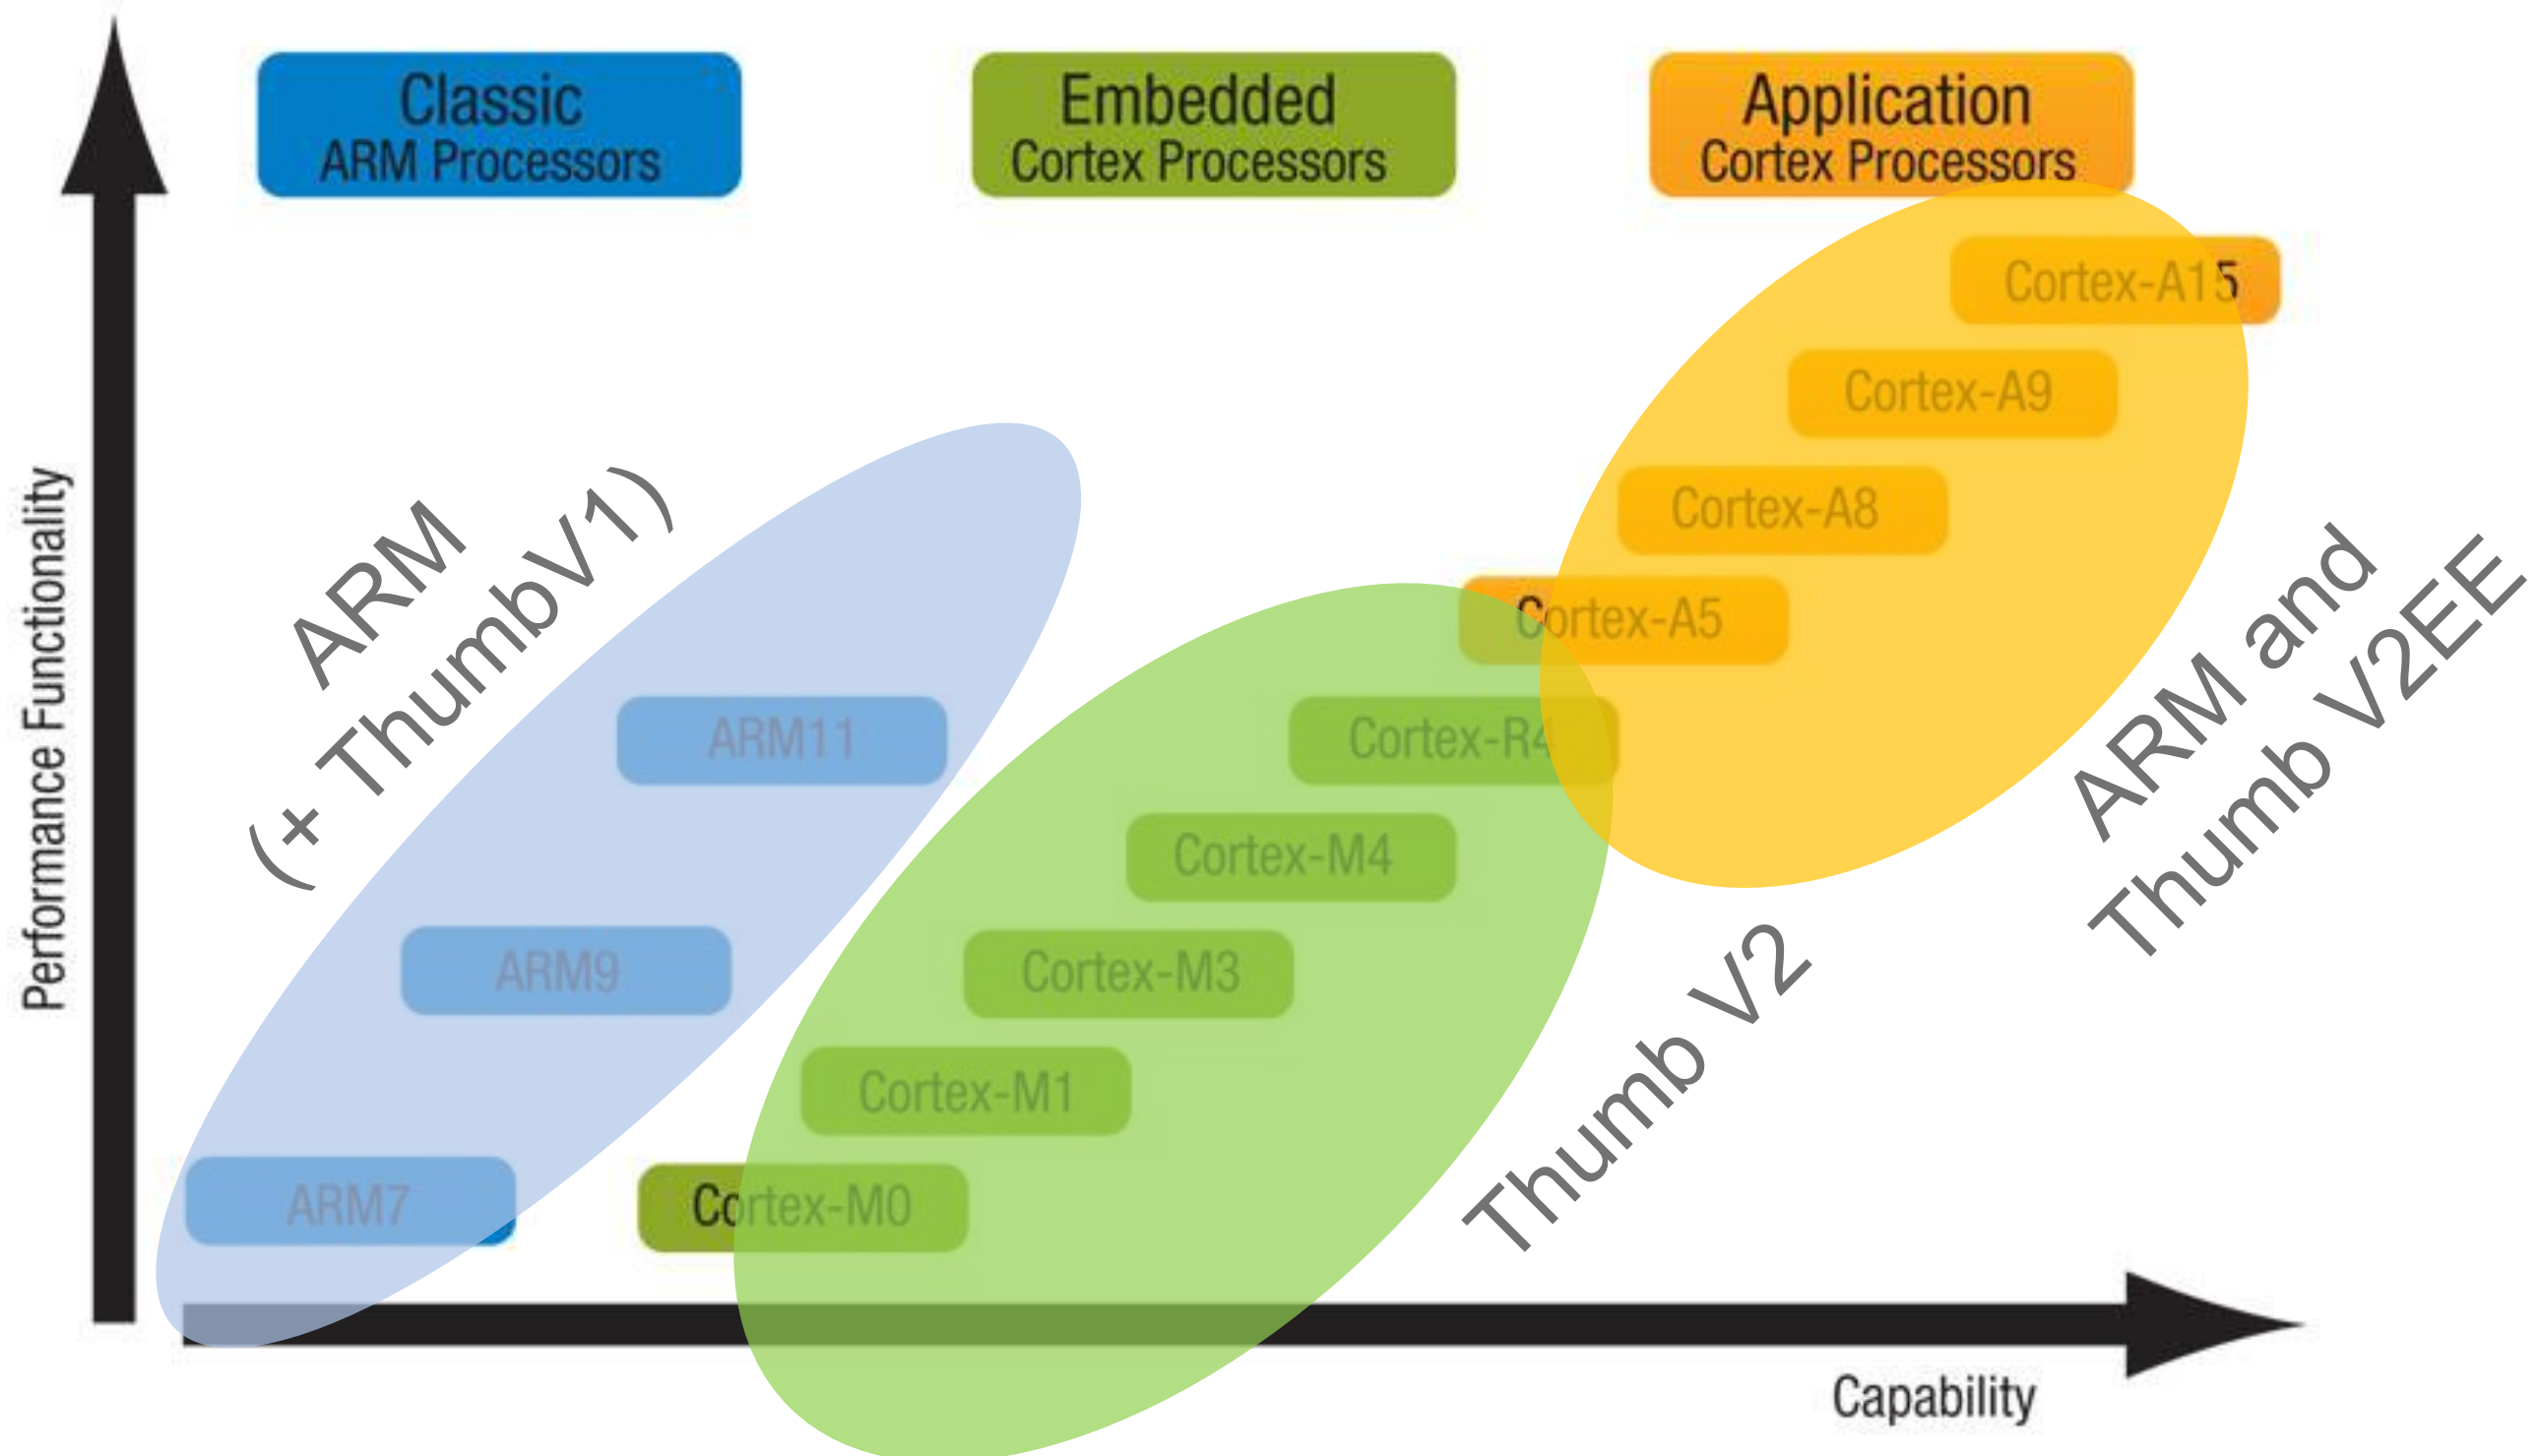
A concrete ISA

- So far, we have talked about ISAs in a generic manner.
- In this lecture and the following one, we will discuss a concrete ISA: the ARM Thumb V1 ISA.

Why Thumb V1?

- ARM is the world's most populous processor
 - Billions shipped per year
- They have multiple device families supporting various ISAs.
 - **ARM** (32-bit)
 - **Thumb V1** (16-bit)
 - **Thumb V2** (16-bit and 32-bit)
 - **Thumb V2EE** (16-bit and 32-bit)

ARM device families



Why Thumb V1?

- Thumb V1 is powerful and compact.
- It is contemporary and delivers good performance for its code size.
- Most importantly for this course, the instructions contained within are relatively simple and few in number.
- All Thumb V1 is also valid Thumb V2.

The ARM architecture

- Before we look at the Thumb ISA, let's look at the ARM architecture at a high-level.
- ARM:
 - Is a register machine, with 16 registers.
 - The PC and LR are part of these registers
 - Is RISC.
 - Has good support for multiple memory addressing modes and is effective with stacks.

ARM M-series registers

Note the
stack pointer
register is
duplicated
(more in
COMS35102!)

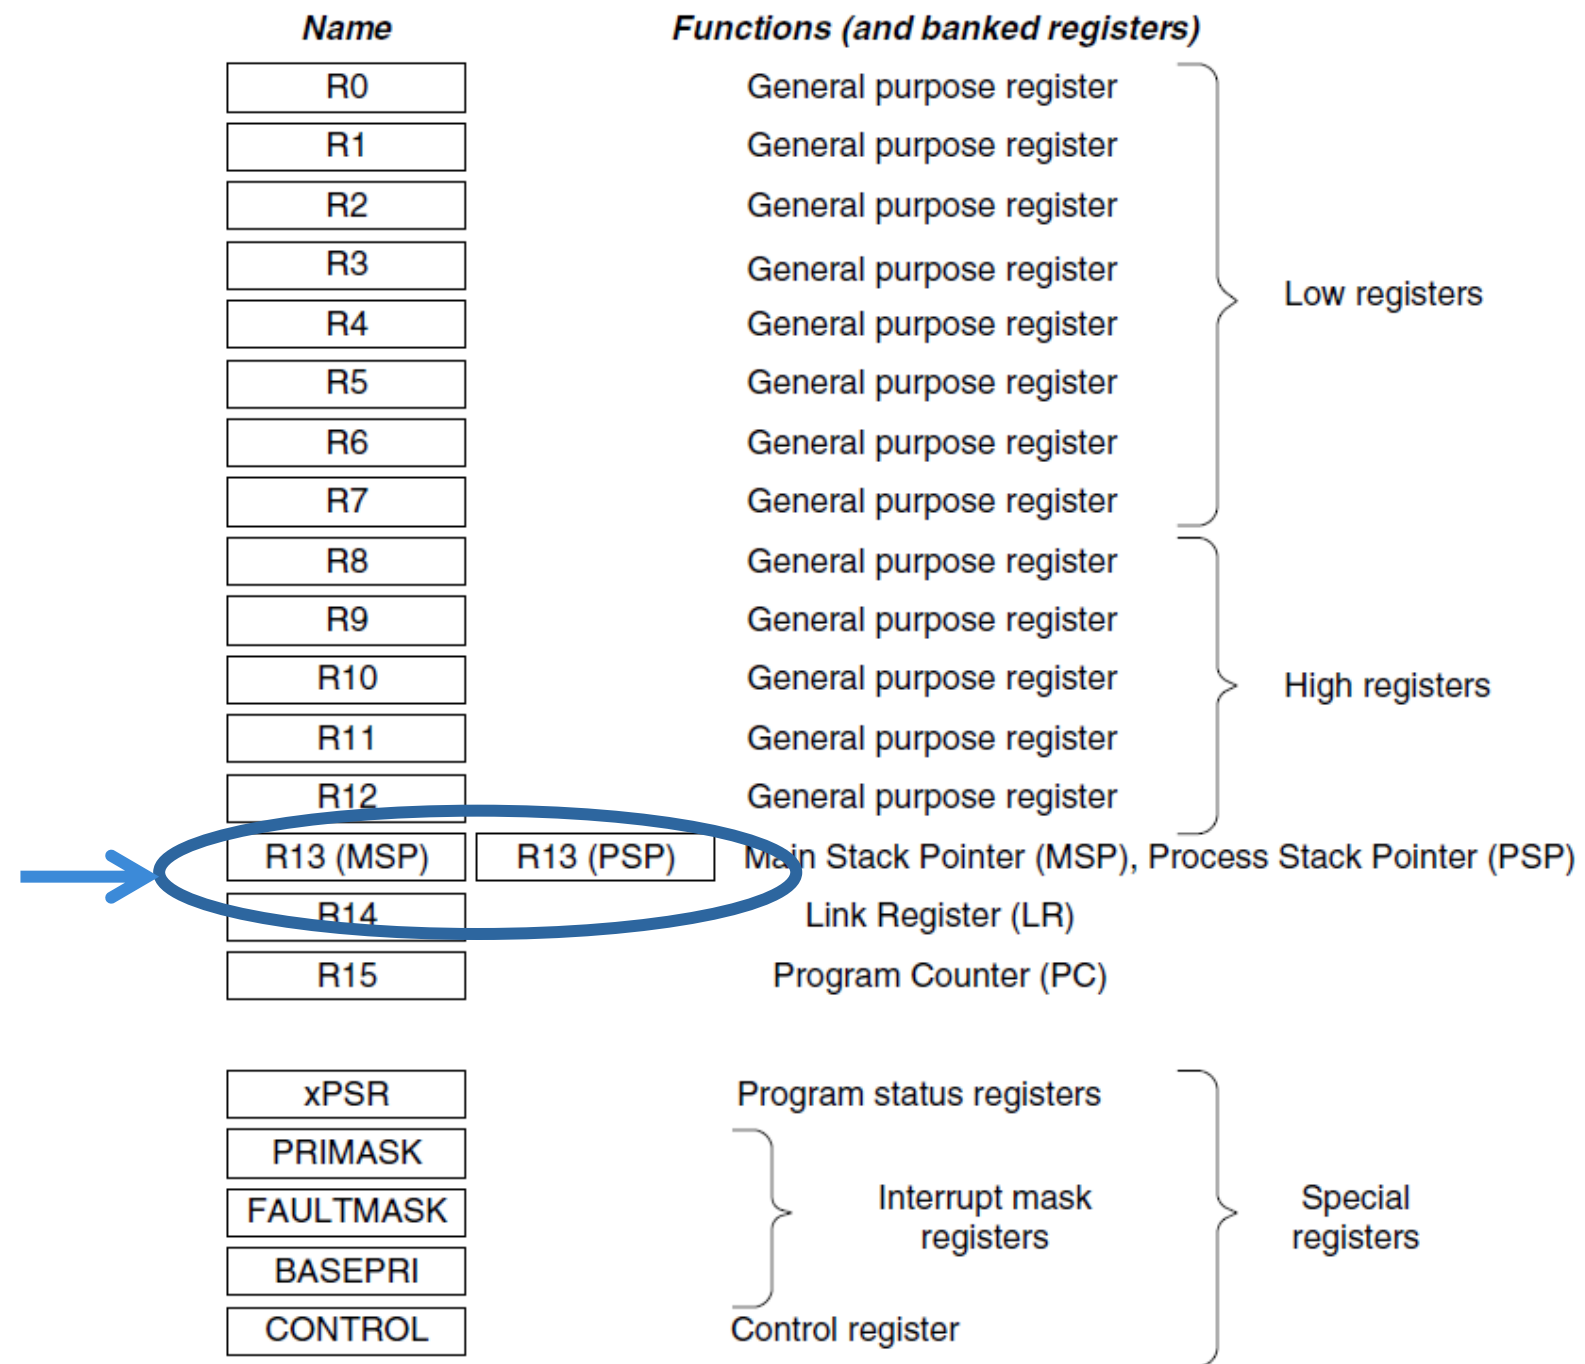


FIGURE 3.1

Registers in the Cortex-M3.

Flexible registers

- The orthogonality of the ARM registers means that *any* combination of them can be used for the source and destination of *any* instruction.
- This gives great flexibility in
 - Register assignment
 - Variable manipulation
- It *includes the Program Counter*
 - Branches can be made as the result of an arithmetic calculation!

Memory manipulation

Most of the addressing modes we saw in previous lectures are supported by the ARM architecture, including

- Register-register
- Register-memory
- Immediate
- Indexed

Memory manipulation

In general, memory addresses are accessed in a two phase approach.

1. Store the memory address to be accessed in a register.
2. Execute a load or store instruction, targeting that register's address.

Memory access example

e.g. to write '42' to memory location 8.

1. **MOV r0, #8** ; *place 8 in r0*

2. **MOV r1, #42** ; *place 42 in r1*

3. **STR r1, [r0]** ; *write value in r1 to
address in r0*

ARM PC

- The ARM's program counter is pre-incremented in an execution cycle.
- It runs ahead of the current instruction by 2 instructions.
- This is because it follows the three stage pipeline model we saw before.

Thumb V1 overview

- The Thumb V1 is a 16-bit instruction set.
- This places some limitations on what it can and cannot do.
 - E.g. encoding three registers from the full ARM set takes 12 bits.
 - There cannot be very many op-codes
 - Op-code length is variable at 5–10 bits

Thumb V1 overview

- Instructions are broadly split into two forms:
 1. Three register operations
 - E.g. `ADD r0, r1, r2`
 2. Two register + one immediate operations
 - E.g. `ADD r0, r1, #10`

Thumb V1 overview

- Most Thumb V1 instructions address one of the 8 ‘low’ registers (i.e. $r0-r7$).
- There are two main forms of register encoding:
 - Three explicit registers
 - e.g. `ADD r0, r1, r2`
 - Two explicit + one implicit
 - e.g. `AND r0, r1`

‘High’ instructions

- Some instructions that need to access registers $> r7$ do allow access to the full register set.
- These instruction encodings are known as ‘high’ instructions.
 - e.g. ADDH, MOVH
 - Often some other functionality of the instruction is sacrificed, such as 3 register format being reduced to 2 register format

Conditionality

- Conditionality in the ARM family is powerful and complex in equal measure.
- Conditional execution occurs in the following way:
 1. Certain instructions' execution has a side-effect of setting conditional bits in the status register.
 2. Subsequent instructions can become conditional on the basis of these set bits.

Thumb conditional example

Here's an example:

1. **ADD r0, #2** ; $r0 = r0 + 2$

2. **BEQ loop** ; *branch to loop if previous result was zero. (i.e. r0 started as -2)*

Thumb conditional

The conditional information was stored in a hardware element called the *status register*.

Here's some important parts:

Bit 31	Bit 30	Bit 29	Bit 28
N	Z	C	V
Negative	Zero (Equal)	Carry	Overflow

Thumb control flow

- Control flow for the Thumb is fairly standard
 - Unconditional branches (BU)
 - Conditional branches (BC)
 - Calls (BL)
- As expected, control flow changes cause a pipeline flush and a execution time penalty.

Thumb control flow

- However, the limited instruction width places restrictions on the reachable distances.
 - BU: 11 bits of immediate
 - BC: 8 bits of immediate
 - BL: 21 + 3 bits of immediate

Double-length BL

- The “Branch with Link” (BL) instruction is unusual.
- It needs greater than 16 bits to encode the jump address.
- So, it is formed of two separate 16 bit instructions, that are concatenated in hardware.
- The first instruction is called “BL1”, the second “BL2”.
- Both are necessary for a successful operation.

ARM stack model

ARM stacks are full, descending

Address	Stack pointer order
0x2010000	0 (initial for stack starting at 0x200ffff)
0x200fffc	1
0x200fff8	2
0x200fff4	3
0x200ffe0	4
0x200ffec	5
0x200ffe8	6
0x200ffe4	7
....

PUSH and POP

The ARM architecture has plenty of support for stack based operation.

PUSH and POP are dedicated instructions to operate on a stack.

- Stack location is defined using the stack pointer (=r13).

PUSH and POP example

This sequence of operations:

1. MOV r0, #2
2. MOV r1, #15
3. MOV r2, #31
4. PUSH r0
5. PUSH r1
6. POP r0
7. PUSH r2
8. PUSH r0
9. POP r0

registers

r0	r1	r2

memory

Stack pointer	Address	Data value
	0x200fffc	
	0x200fff8	
	0x200fff4	
	0x200fff0	
	0x200ffec	
	0x200ffe0	

Thumb V1 quirks

There are certainly some quirks in the Thumb V1 ISA.

- Multiple op-codes for some operations
- Overlapping prefixes
- Hard to reach registers

You'll get to grips with these in the coursework 😊

Summary of Thumb V1

- Thumb V1 is a 16-bit instruction set.
- It predominantly targets only 8 of the 16 registers.
- Conditionality is a combination:
 - Instructions that implicitly set status bits
 - Instructions that are explicitly conditional
- Control flow suffers from restrictions.
- It is, however, compact and efficient.

A word about Thumb V2

- Thumb V2 is newer than V1.
- Many newer ARM processors support it.
- It includes additional instruction formats and features.
- It is much more complex with more variable length instructions.
- The additional complexity is a big disadvantage.
- It does offer some performance benefits over V1.

Relative efficiency

- Thumb v2 benchmarks
 - 74% memory footprint c.f. ARM
 - 98% performance c.f. ARM

