


## Concept: *virtualise* the memory

- ▶ We already virtualised the processor, *but*
    1. *how* do we segregate the processes in memory, and
    2. *why* put up with this restriction?!
  - ▶ In general, several layers of memory management
    1. hardware (RAM, MMU, MPU),
    2. kernel (address space protection and virtualisation), and
    3. user (allocation, deallocation, garbage collection),
- supporting various use-cases, e.g.,
1. process manages memory process uses,
  2. kernel manages memory kernel uses, and
  3. kernel manages memory process uses,

warrant attention ...




## Concept: *virtualise* the memory

- ▶ We already virtualised the processor, *but*
  1. *how* do we segregate the processes in memory, and
  2. *why* put up with this restriction?!
- ▶ In general, several layers of memory management
  1. hardware (RAM, MMU, MPU), and
  2. kernel (address space protection and virtualisation),

supporting various use-cases, e.g.,

  3. kernel manages memory process uses,


warrant attention ...
- ▶ ... *but* we'll consider a narrower remit.



## Concept: *virtualise* the memory

- Specifically, we want processes to
1. *appear* to have dedicated access to the whole physical memory,
  2. have a larger footprint than physical memory *if* required,
  3. be protected wrt. access to their regions of the physical memory,
  4. to share regions of physical memory *if* required,
  5. ...


so, the question is, *how*?



## Concept (1)

### Definition (resolution, relocation etc.)

A (sub-)sequence of standard steps, implemented by a user mode tool-chain plus the kernel, translates a high-level program into a form ready for execution



noting that

- ▶ at various steps we might use an
  - ▶ **abstract address**, e.g., a symbol per
  - ▶ **concrete address**, e.g., a literal per

```
goto foo;
```

```
uint32_t* bar = ( uint32_t* )( 0xDEADBEEF );
```

- ▶ abstract addresses are **resolved** into concrete addresses before execution, and
- ▶ addresses may be **relocated** (or moved, i.e., rewritten) before *or* during execution.

## Concept (2)

### Definition (**address stream** etc.)

Although executed instructions provoke memory accesses, e.g.,




we can often ignore processes themselves, and instead focus on the resulting **address stream** (i.e., a sequence of addresses). An address stream will, *on average*, exhibit various properties:

- ▶ **access locality**, i.e., reuse of the same or “close” addresses, which implies a
- ▶ **working set**, denoted  $\mathcal{W}(P_i)$  for some process  $P_i$ , which captures the set of addresses, or portion of the address space, currently in use.

## Concept (3)

### Definition (address space etc.)

Including a **Memory Management Unit (MMU)** per



allows transparent manipulation of the semantics of addresses and address spaces. Specifically,

- ▶ a **virtual address** relates to the processor view of *a virtual address space* (e.g., associated with a process), whereas
- ▶ a **physical address** relates to the memory view of *the physical address space* (i.e., the actual RAM) noting there is one virtual address space per process, and one physical address space period.

## Concept (4)

► **Goal:** *use* the MMU to realise

1. **translation** of virtual to physical addresses,
  2. **protection** e.g., of the virtual address space associated with one process against access from another, and
  3. **sharing** i.e., controlled non-protection of (or overlap between) address spaces
- and hence *virtualise* the physical memory.


## Mechanism: software (1)

- ▶ **Idea:** *no* MMU!

- ▶ no translation,
- ▶ no protection.

- ▶ **Features:**

address space(s) translated	×
address space(s) protected	×
address space(s) virtualised	×
address space(s) non-contiguous	×
req. hardware support	×
req. software (kernel) support	×
req. software (user) support	×





## Mechanism: software (2)

### ► Idea: *no* MMU!

#### ► still no translation: either

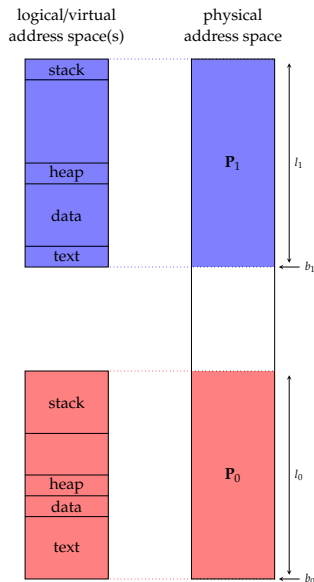
1. linker or
2. loader

relocates each address  $x$  wrt. some base  $b$ ,

#### ► still no protection: assumes process will be “honest” st. $b < x < b + l$ .

### ► Features:

address space(s) translated	×	×
address space(s) protected	×	×
address space(s) virtualised	×	×
address space(s) non-contiguous	×	×
req. hardware support	×	×
req. software (kernel) support	×	✓
req. software (user) support	✓	×



## Mechanism: hardware-based per process segmentation (1)

### ► Idea: per process **segmented memory**.

1. maintain a base and limit register per process,
2. enforce

$$b \leq x < b + l$$

and relocate as before, or

3. enforce

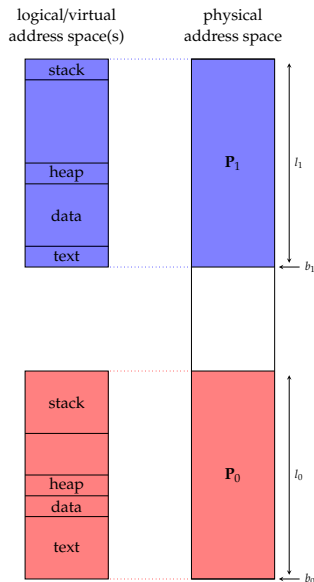
$$0 \leq x < l$$

and translate st.

$$x \mapsto b + x.$$

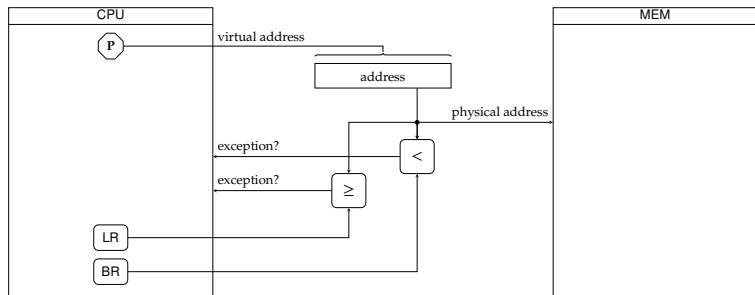
### ► Features:

address space(s) translated	×	✓
address space(s) protected	✓	✓
address space(s) virtualised	×	✓
address space(s) non-contiguous	×	×
req. hardware support	✓	✓
req. software (kernel) support	✓	✓
req. software (user) support	✓	×



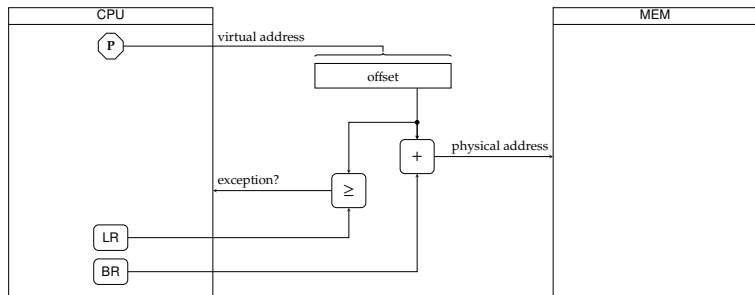
## Mechanism: hardware-based per process segmentation (2)

- An implementation requires MMU-like hardware, e.g.,



## Mechanism: hardware-based per process segmentation (2)

- ▶ An implementation requires MMU-like hardware, e.g.,



## An Aside: allocation, swapping and fragmentation

► **Problem:** where should we load  $\mathcal{P}_i$ .

► **Solution:** we need


1. an allocation algorithm, e.g.,

- first-fit,
- best-fit,
- worst-fit,
- ...


and

2. a data structure to capture the current allocation state.


physical  
address space



file system



swap space



## An Aside: allocation, swapping and fragmentation

► **Problem:** where should we load  $\mathcal{P}_i$ .

► **Solution:** we need


1. an allocation algorithm, e.g.,

- first-fit,
- best-fit,
- worst-fit,
- ...


and

2. a data structure to capture the current allocation state.


physical  
address space



file system



swap space



## An Aside: allocation, swapping and fragmentation

► **Problem:** where should we load  $\mathcal{P}_i$ .

► **Solution:** we need


1. an allocation algorithm, e.g.,

- first-fit,
- best-fit,
- worst-fit,
- ...


and

2. a data structure to capture the current allocation state.


physical  
address space



file system



swap space



## An Aside: allocation, swapping and fragmentation

### ► Problem:

1. cases st.

$$\sum_{i=0}^{i < n} |P_i| > |\text{MEM}|$$

or


2. cases st.

$$\exists i \text{ st. } |P_i| > |\text{MEM}|.$$


### ► Solution(s):

1. **swapping**,
2. some improvement to per process segmentation.


physical  
address space



file system



swap space





## An Aside: allocation, swapping and fragmentation

### ► Problem:

1. cases st.

$$\sum_{i=0}^{i < n} |\mathbf{P}_i| > |\text{MEM}|$$

or


2. cases st.

$$\exists i \text{ st. } |\mathbf{P}_i| > |\text{MEM}|.$$


### ► Solution(s):

1. **swapping**,
2. some improvement to per process segmentation.


physical  
address space



file system



swap space



## An Aside: allocation, swapping and fragmentation

### ► Problem:

1. cases st.

$$\sum_{i=0}^{i < n} |P_i| > |\text{MEM}|$$

or


2. cases st.

$$\exists i \text{ st. } |P_i| > |\text{MEM}|.$$


### ► Solution(s):

1. **swapping**,
2. some improvement to per process segmentation.


physical  
address space



file system



swap space



## An Aside: allocation, swapping and fragmentation

### ► Problem:

1. cases st.

$$\sum_{i=0}^{i < n} |P_i| > |\text{MEM}|$$

or


2. cases st.

$$\exists i \text{ st. } |P_i| > |\text{MEM}|.$$


### ► Solution(s):

1. **swapping**,
2. some improvement to per process segmentation.


physical  
address space



file system



swap space



## An Aside: allocation, swapping and fragmentation

### ► Problem:

1. cases st.

$$\sum_{i=0}^{i < n} |P_i| > |\text{MEM}|$$

or


2. cases st.

$$\exists i \text{ st. } |P_i| > |\text{MEM}|.$$


### ► Solution(s):

1. **swapping**,
2. some improvement to per process segmentation.


physical  
address space



file system




swap space




## An Aside: allocation, swapping and fragmentation

- **Problem: fragmentation**, namely
  1. **internal** (i.e., *within* allocations), or
  2. **external** (i.e., *between* allocations).
- **Solution(s):**
  - **compaction**,
  - some improvement to per process segmentation.


physical  
address space



file system



swap space



## Mechanism: hardware-based per segment segmentation (1)

### ► Idea: per segment **segmented memory**.

- maintain a **segment table**  $T$  per process,
- let  $t = \log_2(|T|)$ , check

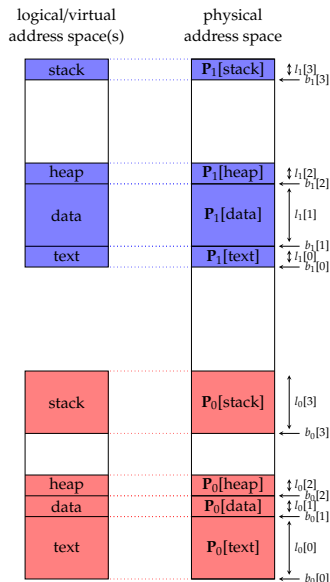
$$0 \leq \text{LSB}_{w-t}(x) < l[\text{MSB}_t(x)],$$

and translate st.

$$x \mapsto b[\text{MSB}_t(x)] + \text{LSB}_{w-t}(x).$$

### ► Features:

address space(s) translated	✓
address space(s) protected	✓
address space(s) virtualised	✓
address space(s) non-contiguous	✓
req. hardware support	✓
req. software (kernel) support	✓
req. software (user) support	✓



## Mechanism: hardware-based per segment segmentation (1)

### ► Idea: per segment **segmented memory**.

- maintain a **segment table**  $T$  per process,
- let  $t = \log_2(|T|)$ , check

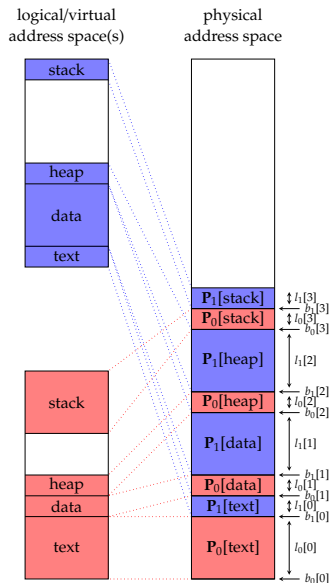
$$0 \leq \text{LSB}_{w-t}(x) < l[\text{MSB}_t(x)],$$

and translate st.

$$x \mapsto b[\text{MSB}_t(x)] + \text{LSB}_{w-t}(x).$$

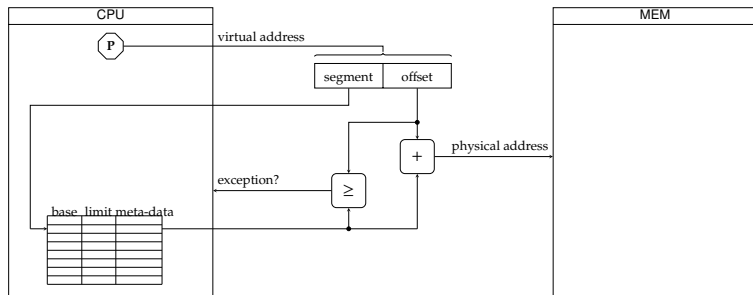
### ► Features:

address space(s) translated	✓
address space(s) protected	✓
address space(s) virtualised	✓
address space(s) non-contiguous	✓
req. hardware support	✓
req. software (kernel) support	✓
req. software (user) support	✓



## Mechanism: hardware-based per segment segmentation (2)

- ▶ An implementation requires MMU-like hardware, e.g.,



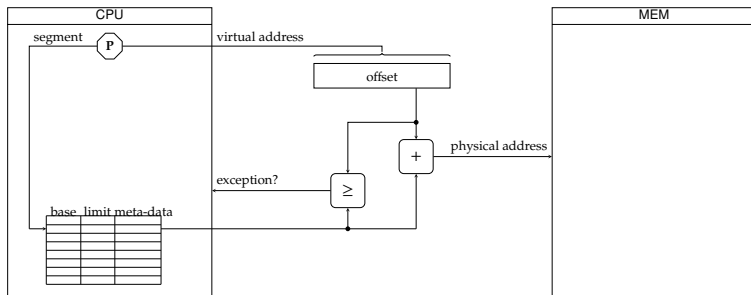
noting we *could* opt to index into the table via

1. one address, i.e., split address into a segment identifier and offset.



## Mechanism: hardware-based per segment segmentation (2)

- ▶ An implementation requires MMU-like hardware, e.g.,



noting we *could* opt to index into the table via

1. one address, i.e., split address into a segment identifier and offset, or
2. two address, i.e., a dedicated segment identifier and offset.

## Mechanism: hardware-based paging (1)

### ► Idea: **paged memory**.

- fix  $l = \rho$ , and divide
  - virtual address space(s) into **pages**,
  - physical address space into **page frames**

of  $l$  bytes in each case,

- maintain a **page table**  $T$  per process,
- let  $t = \log_2(|T|)$ , and translate st.

$$x \mapsto b[\text{MSB}_t(x)] \cdot l + \text{LSB}_{w-t}(x).$$

noting no check is required since


$$0 \leq \text{LSB}_{w-t}(x) < l$$

by definition.



### ► Features:

address space(s) translated	✓
address space(s) protected	✓
address space(s) virtualised	✓
address space(s) non-contiguous	✓
req. hardware support	✓
req. software (kernel) support	✓
req. software (user) support	✓

logical/virtual  
address space(s)



physical  
address space



## Mechanism: hardware-based paging (1)

### ► Idea: paged memory.

- fix  $l = \rho$ , and divide
  - virtual address space(s) into **pages**,
  - physical address space into **page frames**

of  $l$  bytes in each case,

- maintain a **page table**  $T$  per process,
- let  $t = \log_2(|T|)$ , and translate st.

$$x \mapsto b[\text{MSB}_t(x)] \cdot l + \text{LSB}_{w-t}(x).$$

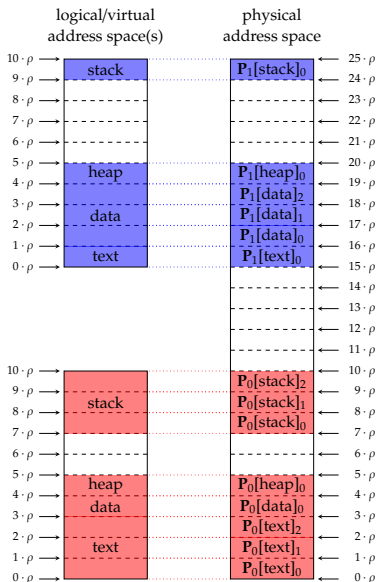
noting no check is required since

$$0 \leq \text{LSB}_{w-t}(x) < l$$

by definition.

### ► Features:

address space(s) translated	✓
address space(s) protected	✓
address space(s) virtualised	✓
address space(s) non-contiguous	✓
req. hardware support	✓
req. software (kernel) support	✓
req. software (user) support	✓



## Mechanism: hardware-based paging (1)

### ► Idea: paged memory.

- fix  $l = \rho$ , and divide
  - virtual address space(s) into **pages**,
  - physical address space into **page frames**

of  $l$  bytes in each case,

- maintain a **page table**  $T$  per process,
- let  $t = \log_2(|T|)$ , and translate st.

$$x \mapsto b[\text{MSB}_t(x)] \cdot l + \text{LSB}_{w-t}(x).$$

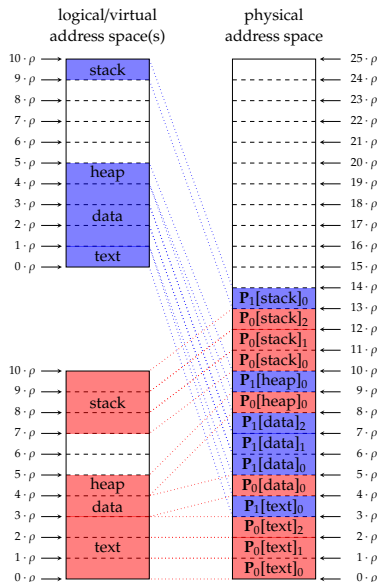
noting no check is required since

$$0 \leq \text{LSB}_{w-t}(x) < l$$

by definition.

### ► Features:

address space(s) translated	✓
address space(s) protected	✓
address space(s) virtualised	✓
address space(s) non-contiguous	✓
req. hardware support	✓
req. software (kernel) support	✓
req. software (user) support	✓



# Mechanism: hardware-based paging (1)

## ► Idea: paged memory.

- fix  $l = \rho$ , and divide
  - virtual address space(s) into **pages**,
  - physical address space into **page frames**

of  $l$  bytes in each case,

- maintain a **page table**  $T$  per process,
- let  $t = \log_2(|T|)$ , and translate st.

$$x \mapsto b[\text{MSB}_t(x)] \cdot l + \text{LSB}_{w-t}(x).$$


noting no check is required since

$$0 \leq \text{LSB}_{w-t}(x) < l$$

by definition.


## ► Features:

address space(s) translated	✓
address space(s) protected	✓
address space(s) virtualised	✓
address space(s) non-contiguous	✓
req. hardware support	✓
req. software (kernel) support	✓
req. software (user) support	✓



## Mechanism: hardware-based paging (2)


- An implementation requires MMU-like hardware, e.g.,



noting the page table consists of **Page Table Entries (PTEs)**.

## Mechanism: hardware-based paging (3)


- **Improvement #1:** since the page table is *large*, we could



1. store the page table in memory,
  2. point at the page table with a **Page Table Register (PTR)**, and
  3. use a  $\tau$ -entry **Translation Look-aside Buffer (TLB)** to cache the page table, noting
    - flush the TLB during a context switch, or
    - include a process identifier as a disambiguation tag,
- st. cached PTEs for one process cannot be used by another.

## Mechanism: hardware-based paging (3)

- **Improvement #1:** since the page table is *large*, we could

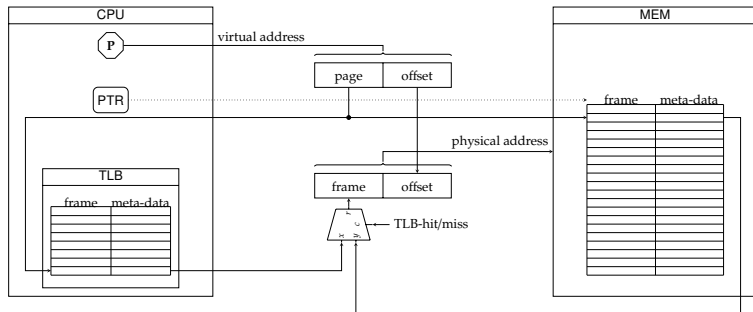


1. store the page table in memory,
  2. point at the page table with a **Page Table Register (PTR)**, and
  3. use a  $\tau$ -entry **Translation Look-aside Buffer (TLB)** to cache the page table, noting
    - flush the TLB during a context switch, or
    - include a process identifier as a disambiguation tag,
- st. cached PTEs for one process cannot be used by another.



## Mechanism: hardware-based paging (3)


- **Improvement #1:** since the page table is *large*, we could




1. store the page table in memory,
  2. point at the page table with a **Page Table Register (PTR)**, and
  3. use a  $\tau$ -entry **Translation Look-aside Buffer (TLB)** to cache the page table, noting
    - flush the TLB during a context switch, or
    - include a process identifier as a disambiguation tag,
- st. cached PTEs for one process cannot be used by another.

## Mechanism: hardware-based paging (4)

- **Improvement #2:** since the page table is *sparse*, we could



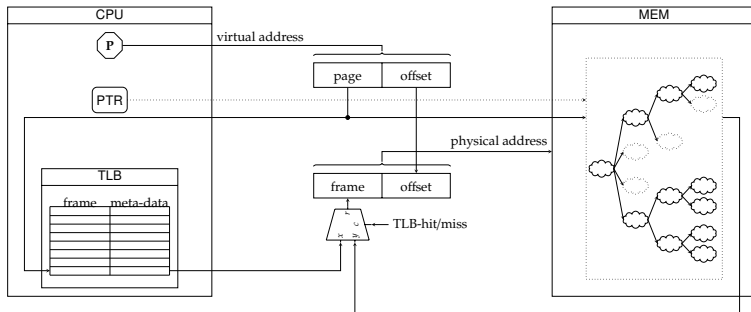
1. store the page table as a  $\lambda$ -level tree (vs. a list),
2. decompose original page number to index into each level, i.e.,




3. use a valid flag to indicate whether or not a sub-tree exists.

## Mechanism: hardware-based paging (4)

- **Improvement #2:** since the page table is *sparse*, we could



1. store the page table as a  $\lambda$ -level tree (vs. a list),
2. decompose original page number to index into each level, i.e.,

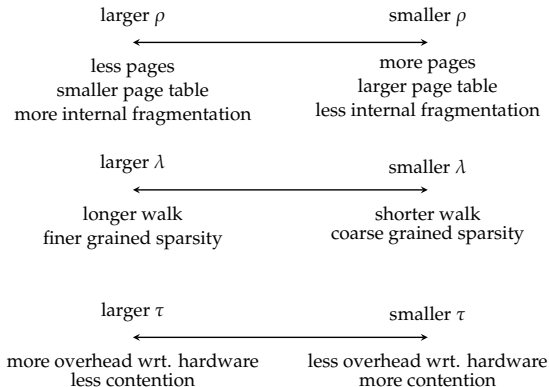


3. use a valid flag to indicate whether or not a sub-tree exists.

## Mechanism: hardware-based paging (5)

► ... *but*, we need to

1. select various (non-independent) parameters,
2. consider how to interface with the wider memory hierarchy, and
3. handle various exceptions appropriately.



## Mechanism: hardware-based paging (5)

- ▶ ... *but*, we need to
  - 1. select various (non-independent) parameters,
  - 2. consider how to interface with the wider memory hierarchy, and
  - 3. handle various exceptions appropriately.
- ▶ any given cache could potentially be placed before (i.e., deal with virtual addresses) *or* after (i.e., deal with physical addresses) the MMU,
- ▶ it can make sense to align the page size with the swap space (i.e., disk) transfer size.

## Mechanism: hardware-based paging (5)

► ... *but*, we need to

1. select various (non-independent) parameters,
2. consider how to interface with the wider memory hierarchy, and
3. handle various exceptions appropriately.

soft TLB miss	{ page is in memory page table entry isn't in TLB
hard TLB miss	{ page isn't in memory page table entry isn't in TLB
invalid page fault	{ page isn't in memory page isn't valid in page table
soft page fault	{ page is in memory page isn't valid in page table
hard page fault	{ page isn't in memory page is valid in page table
access fault	{ fails some check wrt. meta-data

## Implementation: ARMv7-A (1)

- ▶ ARMv7-A supports *two* (very flexible) mechanisms via
  1. the **Protected Memory System Architecture (PMSA)** [5, Chapter B5] and
  2. the **Virtual Memory System Architecture (VMSA)** [5, Chapter B3]both of which are controlled via the co-processor interface [5, Chapters B4 and B6].

## Implementation: ARMv7-A (2) VMSA

### ► Some details:

#### 1. It supports

- a 32-bit virtual address space, and
- *upto* a 40-bit physical address space

with the latter realised via the **Large Physical Address Extension (LPAE)** ...

#### 2. ... and so two PTE formats [5, Section B3.3], namely

long	⇒	64-bit PTE	$\left\{ \begin{array}{l} \text{upto } \lambda = 3 \text{ levels} \\ \text{translates 32-bit to 40-bit address spaces at 4KiB granularity} \end{array} \right.$
short	⇒	32-bit PTE	$\left\{ \begin{array}{l} \text{upto } \lambda = 2 \text{ levels} \\ \text{translates 32-bit to 32-bit address spaces at 4KiB granularity} \end{array} \right.$
short	⇒	32-bit PTE	$\left\{ \begin{array}{l} \text{upto } \lambda = 2 \text{ levels} \\ \text{translates 32-bit to 40-bit address spaces at 16MiB granularity} \end{array} \right.$

plus per-level variants of each.



## Implementation: ARMv7-A (2) VMSA

- Some details:

- 3. It supports four page sizes [5, Section B3.3]

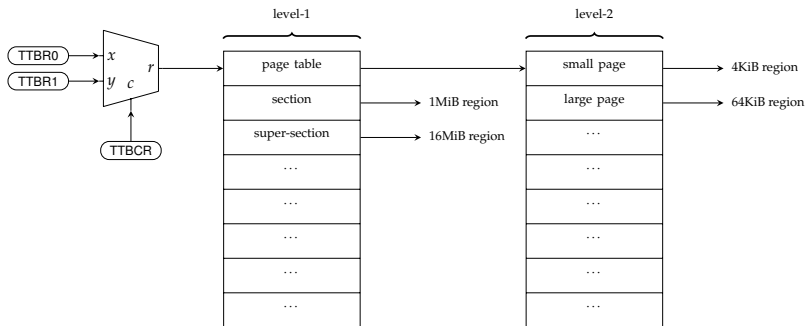
small page	$\Rightarrow$	$\rho = 4\text{KiB}$	$\leadsto$	12-bit offsets
large page	$\Rightarrow$	$\rho = 64\text{KiB}$	$\leadsto$	16-bit offsets
section	$\Rightarrow$	$\rho = 1\text{MiB}$	$\leadsto$	20-bit offsets
super-section	$\Rightarrow$	$\rho = 16\text{MiB}$	$\leadsto$	24-bit offsets

- 4. It uses two PTRs named TTBR0 and TTBR1, selecting one via TTBCR.

## Example

Consider a (simple) example where we set  $\lambda = 2$ , utilise short PTEs only, utilise small pages only, and ignore functionality such as ASID.

The (general) 2-level page table organisation can be described as follows




although in this (specific) example, all level-1 PTEs will point to a level-2 page table, and all level-2 PTEs will point to a small page by definition.

## Example

Consider a (simple) example where we set  $\lambda = 2$ , utilise short PTEs only, utilise small pages only, and ignore functionality such as ASID.


The format of (short) PTEs

- ▶ at level-1 [5, Figure B3-4] is



and

- ▶ at level-2 [5, Figure B3-5] is



### Example

Consider a (simple) example where we set  $\lambda = 2$ , utilise short PTEs only, utilise small pages only, and ignore functionality such as ASID.

To load from some virtual address  $x$ , we proceed as follows:

```
1  if PTE  $E$  for  $x$  is resident in the appropriate TLB(s) then
2    if access control check for  $x$  and  $E$  passes then
3      | load from  $\text{MEM}[E[PA] + x_{11,\dots,0}]$ 
4    else
5      | raise exception
6    end
7  else
8    if  $\text{MSB}_n(x) = 0$  then
9      | load level-1 entry  $E_1$  from  $\text{MEM}[\text{TTBR0} + x_{31,\dots,20}]$ 
10   else
11     | load level-1 entry  $E_1$  from  $\text{MEM}[\text{TTBR1} + x_{31,\dots,20}]$ 
12   end
13   if  $E_1$  is invalid or access control check fails then raise exception
14   load level-2 entry  $E_2$  from  $\text{MEM}[E_1[PA] + x_{19,\dots,12}]$ 
15   if  $E_2$  is invalid or access control check fails then raise exception
16   load from  $\text{MEM}[E_2[PA] + x_{11,\dots,0}]$ 
17   update TLB(s)
18 end
```

Continued in next lecture ...

## References

- [1] Wikipedia: Memory management.  
[https://en.wikipedia.org/wiki/Memory\\_management](https://en.wikipedia.org/wiki/Memory_management).
- [2] Wikipedia: Memory segmentation.  
[https://en.wikipedia.org/wiki/Memory\\_segmentation](https://en.wikipedia.org/wiki/Memory_segmentation).
- [3] Wikipedia: Translation Look-aside Buffer (TLB).  
[https://en.wikipedia.org/wiki/Translation\\_lookaside\\_buffer](https://en.wikipedia.org/wiki/Translation_lookaside_buffer).
- [4] Wikipedia: Virtual memory.  
[https://en.wikipedia.org/wiki/Virtual\\_memory](https://en.wikipedia.org/wiki/Virtual_memory).
- [5] ARM Limited.  
[ARM Architecture Reference Manual: ARMv7-A and ARMv7-R edition](#).  
Technical Report DDI-0406C, 2014.  
<http://infocenter.arm.com/help/topic/com.arm.doc.ddi0406c/index.html>.
- [6] M. Gorman.  
[Understanding the Linux Virtual Memory Manager](#).  
Prentice Hall, 2004.  
<http://www.kernel.org/doc/gorman/>.
- [7] A. Silberschatz, P.B. Galvin, and G. Gagne.  
[Chapter 8: Memory management strategies](#).  
In *Operating System Concepts* [9].

## References

- [8] A. Silberschatz, P.B. Galvin, and G. Gagne.  
[Chapter 9: Virtual-memory management.](#)  
In *Operating System Concepts* [9].
- [9] A. Silberschatz, P.B. Galvin, and G. Gagne.  
[Operating System Concepts.](#)  
Wiley, 9th edition, 2014.
- [10] A. N. Sloss, D. Symes, and C. Wright.  
[ARM System Developer's Guide: Designing and Optimizing System Software.](#)  
Elsevier, 2004.
- [11] A. N. Sloss, D. Symes, and C. Wright.  
[Chapter 13: Memory protection units.](#)  
In *ARM System Developer's Guide: Designing and Optimizing System Software* [10].
- [12] A. N. Sloss, D. Symes, and C. Wright.  
[Chapter 14: Memory management units.](#)  
In *ARM System Developer's Guide: Designing and Optimizing System Software* [10].
- [13] A.S. Tanenbaum and H. Bos.  
[Chapter 3: Memory managment.](#)  
In *Modern Operating Systems*. Pearson, 4th edition, 2015.