


Continued from last lecture ...

- ▶ ... so far so good, *but* we need to explore
 1. how the file system supports our assumed access model, *and*
 2. how the underling **storage device** supports the file system.

Mechanism: (mass storage) devices \leadsto blocks (1)

► **Example: optical disks** (inc. CDs and DVDs).




noting that

- one might attempt
 1. **random access**, and/or
 2. **sequential access**
- and
- efficiency of said access is limited by

$$\begin{array}{ll} \text{transfer rate} & \propto \text{read/write head performance} \\ \text{positioning latency} & \approx \text{seek latency} + \text{rotational latency} \end{array}$$

Mechanism: (mass storage) devices \leadsto blocks (1)

► Example: magnetic disks.



noting that

- one might attempt
 1. **random access**, and/or
 2. **sequential access**
- and
- efficiency of said access is limited by

$$\begin{array}{ll} \text{transfer rate} & \propto \text{read/write head performance} \\ \text{positioning latency} & \approx \text{seek latency} + \text{rotational latency} \end{array}$$


Mechanism: (mass storage) devices \leadsto blocks (2)

- ▶ We add structure to the medium via several steps

medium =


Mechanism: (mass storage) devices \leadsto blocks (2)

- We add structure to the medium via several steps



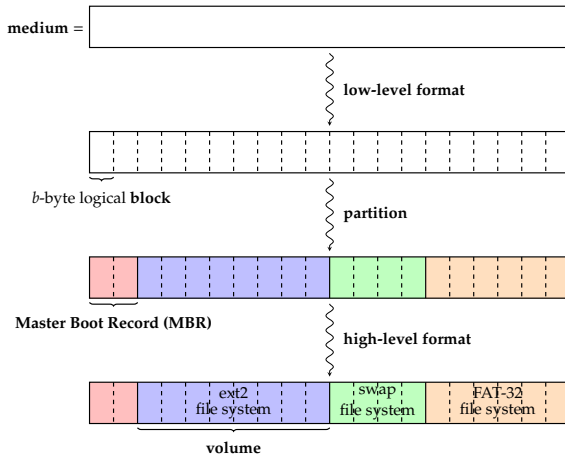
Mechanism: (mass storage) devices \leadsto blocks (2)

- We add structure to the medium via several steps



Mechanism: (mass storage) devices \leadsto blocks (2)


- We add structure to the medium via several steps



then ...


Mechanism: (mass storage) devices \leadsto blocks (3)

- ... assume an interface as previously described, e.g.




Mechanism: (mass storage) devices \leadsto blocks (3)

- ... assume an interface as previously described, e.g.



Mechanism: (mass storage) devices \leadsto blocks (3)

- ... assume an interface as previously described, e.g.

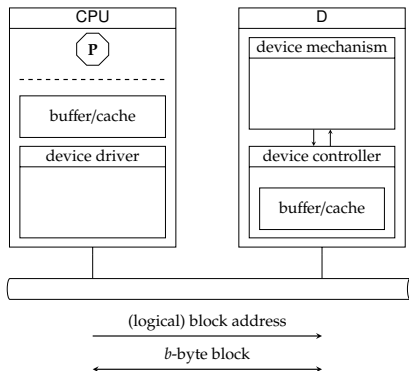


but, since efficiency is crucial we (typically) also

1. amortise overhead by fixing transferring b -byte blocks,
2. use **Logical Block Addressing (LBA)**, forcing translation by the device controller.

Mechanism: (mass storage) devices \leadsto blocks (3)

- ... assume an interface as previously described, e.g.



but, since efficiency is crucial we (typically) also

1. amortise overhead by fixing transferring b -byte blocks,
2. use **Logical Block Addressing (LBA)**, forcing translation by the device controller, and
3. buffer and/or cache accesses (in various layers).

Mechanism: blocks \leadsto files (1)

► **Challenge:** given a device with

- fixed number of logical blocks and
- fixed sized logical blocks,

realise (hierarchical) file system supporting


- representation and
- manipulation

of

- variable number of files, and
- variable sized files.


► **Solution:** we need

1. an allocation algorithm, and
2. a data structure to capture the current allocation state.




Mechanism: blocks \leadsto files (1)

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Idea: contiguous allocation.**
 - allocation is more challenging,
 - + sequential *and* random access is efficient,
 - internal *and* external fragmentation,
 - + no storage overhead.




Mechanism: blocks \leadsto files (1)

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Idea: linked allocation.**
 - + allocation is less challenging,
 - + sequential access is efficient,
 - random access is inefficient,
 - + internal fragmentation only,
 - some storage overhead due to pointers.



Mechanism: blocks \leadsto files (1)

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Idea: linked allocation.**
 - + allocation is less challenging,
 - + sequential access is efficient,
 - random access is inefficient,
 - + internal fragmentation only,
 - some storage overhead due to pointers.



Mechanism: blocks \leadsto files (1)

► **Challenge:** given a device with

- fixed number of logical blocks and
- fixed sized logical blocks,

realise (hierarchical) file system supporting


- representation and
- manipulation

of

- variable number of files, and
- variable sized files.


► **Idea: indexed allocation.**

- + allocation is less challenging,
- + sequential *and* random access is efficient,
- + internal fragmentation only,
- some storage overhead due to pointers.



Mechanism: blocks \leadsto files (1)

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Problem:** what if the index block is full?
- ▶ **Solution(s):** use *multiple* index blocks, e.g., via
 1. linked list,
 2. linked tree, or
 3. various hybrid(s) ...



Mechanism: blocks \leadsto files (2)

► **Problem:** larger storage capacity means more logical blocks, so

1. larger logical block addresses,
2. decreased access locality, and
3. greater overhead (in time and space) wrt. allocation.

► **Solution:** use a hybrid part contiguous, part non-contiguous approach, e.g.,

1. **cluster** \simeq fixed size, contiguous group of logical blocks:

- e.g., divide w -bit logical block address into two

$$\text{logical block group address} = \begin{array}{|c|c|} \hline \overset{w-1}{\text{offset}} & \overset{t}{\text{00...0}} \overset{t-1}{\text{0}} \overset{0}{\text{0}} \\ \hline \end{array}$$

- each offset now addresses a contiguous group of 2^t logical blocks.

2. **extent** \simeq variable size, contiguous group of logical blocks:

- e.g., divide w -bit logical block address into two

$$\text{logical block group address} = \begin{array}{|c|c|} \hline \overset{w-1}{\text{offset}} & \overset{t}{\text{length}} \overset{t-1}{\text{length}} \overset{0}{\text{length}} \\ \hline \end{array}$$

- each offset now addresses a contiguous group of logical blocks whose length is given by the t LSBs,
- this is more flexible, *but* yields complications wrt. seeking and allocation.

although from here on we ignore this option.

Mechanism: blocks \leadsto files (3)

- **Problem:** each write requires one or more of

1. update the allocation state,
2. update the file meta-data, *and*
3. update the file data

which *must* be **atomic**: if not, the file system can become inconsistent.

- **Solution:**

- describe update in write-ahead log (or journal),
- commit update to file system iff. write to log is complete

Mechanism: blocks \leadsto files (4)

- ▶ **Question:** what *is* a directory?
- ▶ **Answer:** a mapping, e.g.,

identifier \mapsto (meta-data, data)

or

(identifier, meta-data) \mapsto data

which also hint at

1. options for where meta-data should reside, and
2. the fact a file might not *itself* have an identifier!

Mechanism: blocks \leadsto files (4)

- ▶ **Question:** what *is* a directory?
- ▶ **Answer:** a mapping, e.g.,

identifier \mapsto (meta-data, data)

or

(identifier, meta-data) \mapsto data

which also hint at

1. options for where meta-data should reside, and
2. the fact a file might not *itself* have an identifier!

- ▶ **Problem:** *how* should a directory be represented?
- ▶ **Solution:**
 1. list,
 2. tree,
 3. hash table,
 4. ...

Mechanism: blocks \rightsquigarrow files (4)

- ▶ **Question:** what *is* a directory?
- ▶ **Answer:** a mapping, e.g.,

identifier \mapsto (meta-data, data)

or

(identifier, meta-data) \mapsto data

which also hint at

1. options for where meta-data should reside, and
2. the fact a file might not *itself* have an identifier!

- ▶ **Problem:** *where* should a directory representation be stored?
- ▶ **Solution:**


1. as a file, plus special-purpose rules for access,
2. as a special-purpose structure,
3. ...

i.e., unified or segregated wrt. the rest of the file system.

Implementation: devices \leadsto file systems (1)

Windows-centric: FAT

- **Idea: File Allocation Table (FAT) \approx fancy linked allocation.**



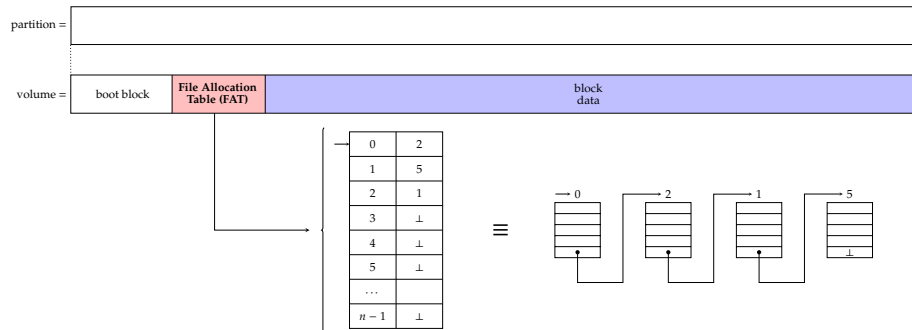
noting that

- + resolves issue of random access wrt. linked allocation,
- need to retain FAT in memory ... which, for n logical blocks, can be large!

Implementation: devices \leadsto file systems (1)

Windows-centric: FAT

- **Idea: File Allocation Table (FAT) \approx fancy linked allocation.**




noting that

- + resolves issue of random access wrt. linked allocation,
- need to retain FAT in memory ... which, for n logical blocks, can be large!

Implementation: devices \leadsto file systems (2)

UNIX-centric: UFS

- **Idea: Unix File System (UFS)** [15, Section 4] \approx fancy indexed allocation.



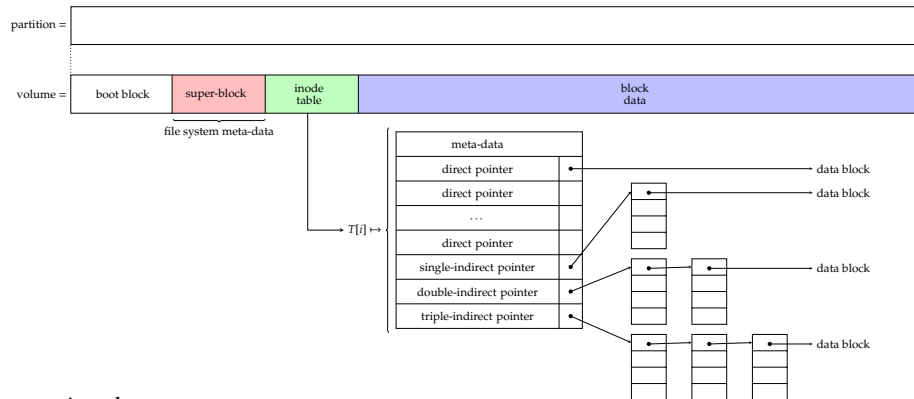
noting that

- + inodes are of (small) fixed size, so indexing into the inode table is efficient,
- linked representation of free space (for inodes *and* blocks).

Implementation: devices \leadsto file systems (2)

UNIX-centric: UFS

- **Idea: Unix File System (UFS)** [15, Section 4] \approx fancy indexed allocation.




noting that

- + inodes are of (small) fixed size, so indexing into the inode table is efficient,
- linked representation of free space (for inodes *and* blocks).

Implementation: devices \leadsto file systems (3)

UNIX-centric: FFS

- **Idea: Fast File System (FFS) [8]** \approx UFS + larger block size ($\geq 4\text{KiB}$).




noting that

- + bit-map representation of free space (for inodes *and* blocks),
- + redundant copy of super-block improves fault tolerance (for overhead of space),
- + cylinder groups increase access locality,
- + includes additional features such as soft links.

Implementation: devices \leadsto file systems (4)

UNIX-centric: ext2

- **Idea: Second Extended File System (ext2) [9] \approx FFS + caching.**



noting that

- + improved directory representation via hash tables [9, Section 4.2],
- + caching and asynchronous writes improve performance ...
- ... but with disadvantages wrt. coherence (viz. robustness).

► Take away points:

- This is a broad and complex topic: it involves (at least)
 1. a hardware aspect:
 - an interrupt controller,
 - a block device
 2. a low(er)-level software aspect:
 - an interrupt handler,
 - a device driver,
 - a file system driver
 3. a high(er)-level software aspect:
 - some data structures (e.g., mount and file descriptor tables),
 - any relevant POSIX system calls (e.g., `write`)
- Keep in mind that, even then,
 - we've excluded and/or simplified various (sub-)topics,
 - there are numerous trade-offs involved, meaning it is often hard to identify one ideal solution.

References

- [1] Wikipedia: Disk formatting.
https://en.wikipedia.org/wiki/Disk_formatting.
- [2] Wikipedia: Disk partitioning.
https://en.wikipedia.org/wiki/Disk_partitioning.
- [3] Wikipedia: File system.
https://en.wikipedia.org/wiki/File_system.
- [4] Wikipedia: HTree.
<https://en.wikipedia.org/wiki/HTree>.
- [5] Wikipedia: Magnetic storage.
http://en.wikipedia.org/wiki/Magnetic_storage.
- [6] Wikipedia: Master Boot Record (MRB).
https://en.wikipedia.org/wiki/Master_boot_record.
- [7] Wikipedia: Optical disk.
http://en.wikipedia.org/wiki/Optical_disc.
- [8] M.K. McKusick, W.N. Joy, S.J. Leffler, and R.S. Fabry.
[A fast file system for UNIX](#).
ACM Transactions on Computer Systems, 2(3):181–197, 1984.

References

- [9] D. Poirier.
[Second extended file system.](#)
<http://www.nongnu.org/ext2-doc/>.
- [10] V. Prabhakaran, A.C. Arpaci-Dusseau, and R.H. Arpaci-Dusseau.
[Analysis and evolution of journaling file systems.](#)
In *USENIX Annual Technical Conference (ATEC)*, pages 105–120, 2005.
- [11] A. Silberschatz, P.B. Galvin, and G. Gagne.
[Chapter 11: Implementing file systems.](#)
In *Operating System Concepts* [13].
- [12] A. Silberschatz, P.B. Galvin, and G. Gagne.
[Chapter 12: Mass storage structure.](#)
In *Operating System Concepts* [13].
- [13] A. Silberschatz, P.B. Galvin, and G. Gagne.
[Operating System Concepts.](#)
Wiley, 9th edition, 2014.
- [14] A.S. Tanenbaum and H. Bos.
[Chapter 4: File systems.](#)
In *Modern Operating Systems*. Pearson, 4th edition, 2015.

References

- [15] K. Thompson.
[UNIX implementation.](#)
Bell System Technical Journal, 57(6):1931–1946, 1978.
- [16] S.C. Tweedie.
[Journaling the Linux ext2fs filesystem.](#)
In *LinuxExpo*, 1998.