

# Concurrent Computing

---

Lecturers: Prof. Majid Mirmehdi majid@cs.bris.ac.uk)  
Dr. Tilo Burghardt tilo@cs.bris.ac.uk  
Dr. Daniel Page page@cs.bris.ac.uk

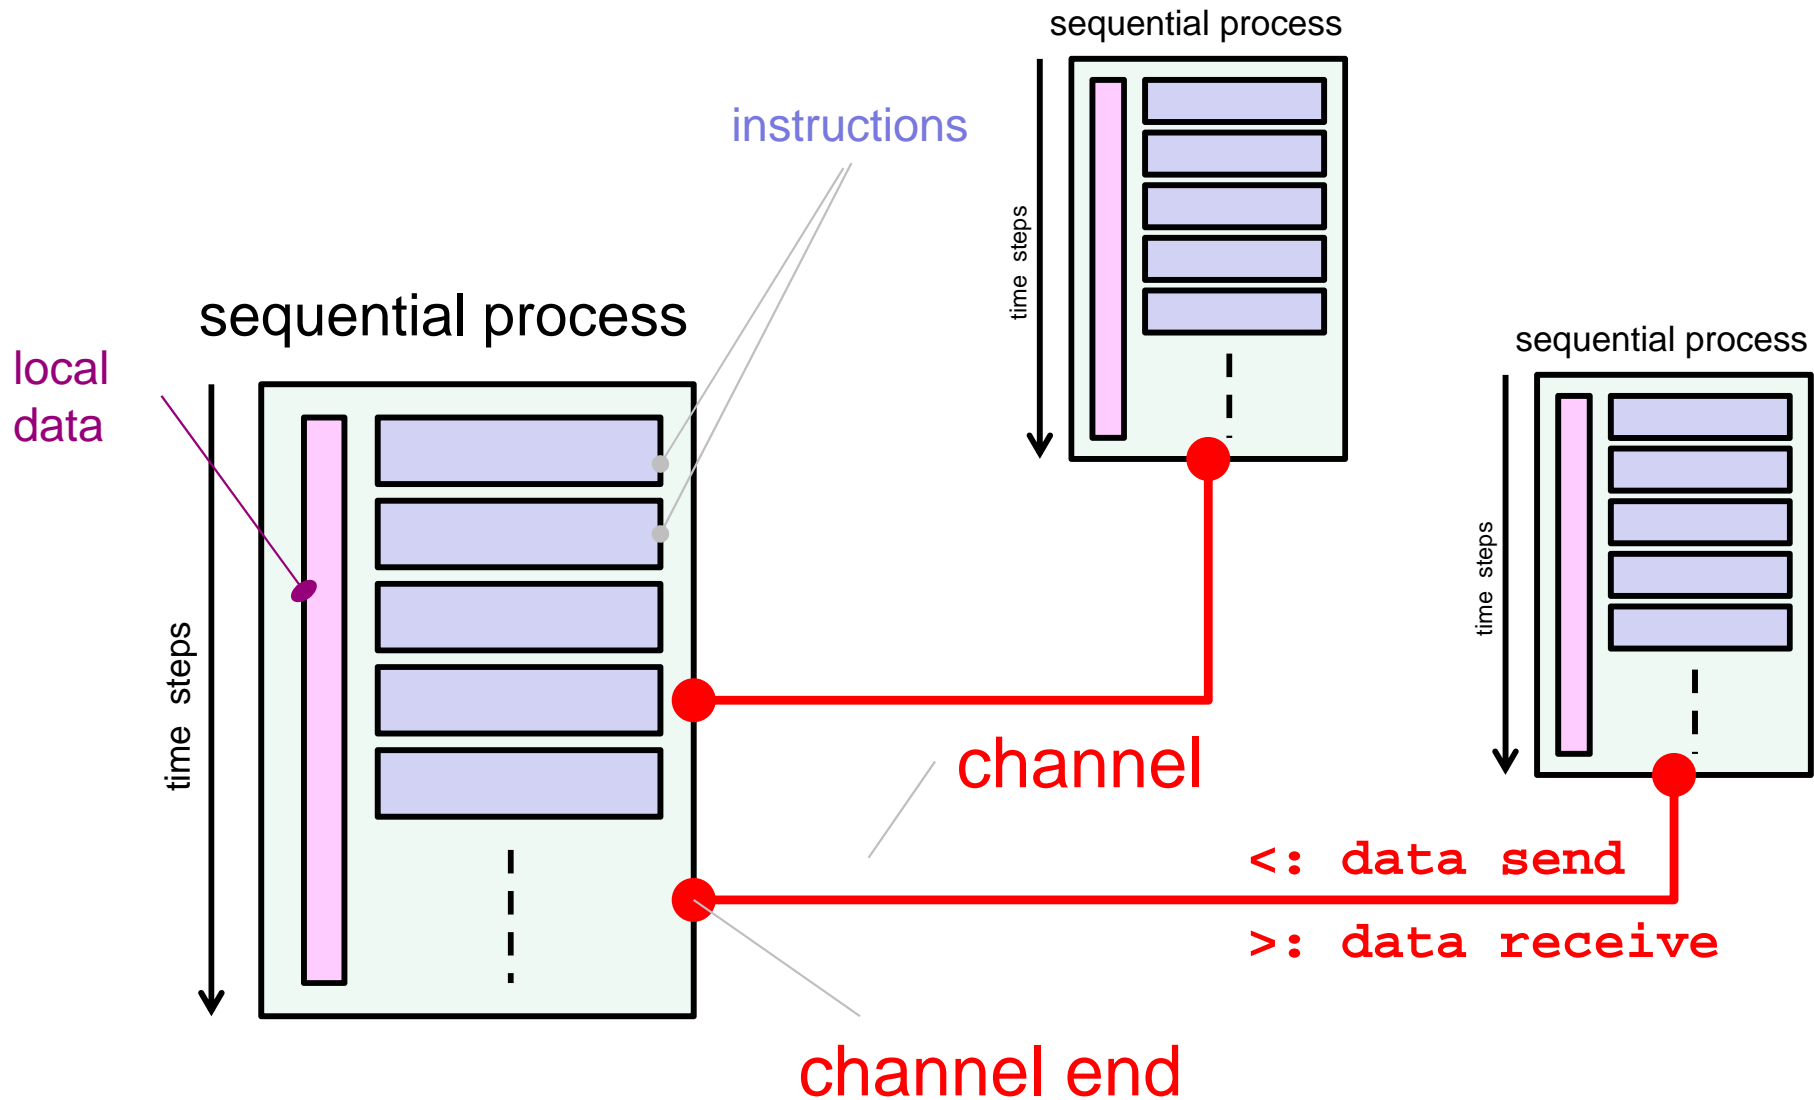
Web: <http://www.cs.bris.ac.uk/Teaching/Resources/COMS20001>



## LECTURE MM4

### *BASICS OF CONCURRENT SYSTEM DESIGN*

# The Building Blocks: Processes and Channels




# Stages in Designing a Concurrent System

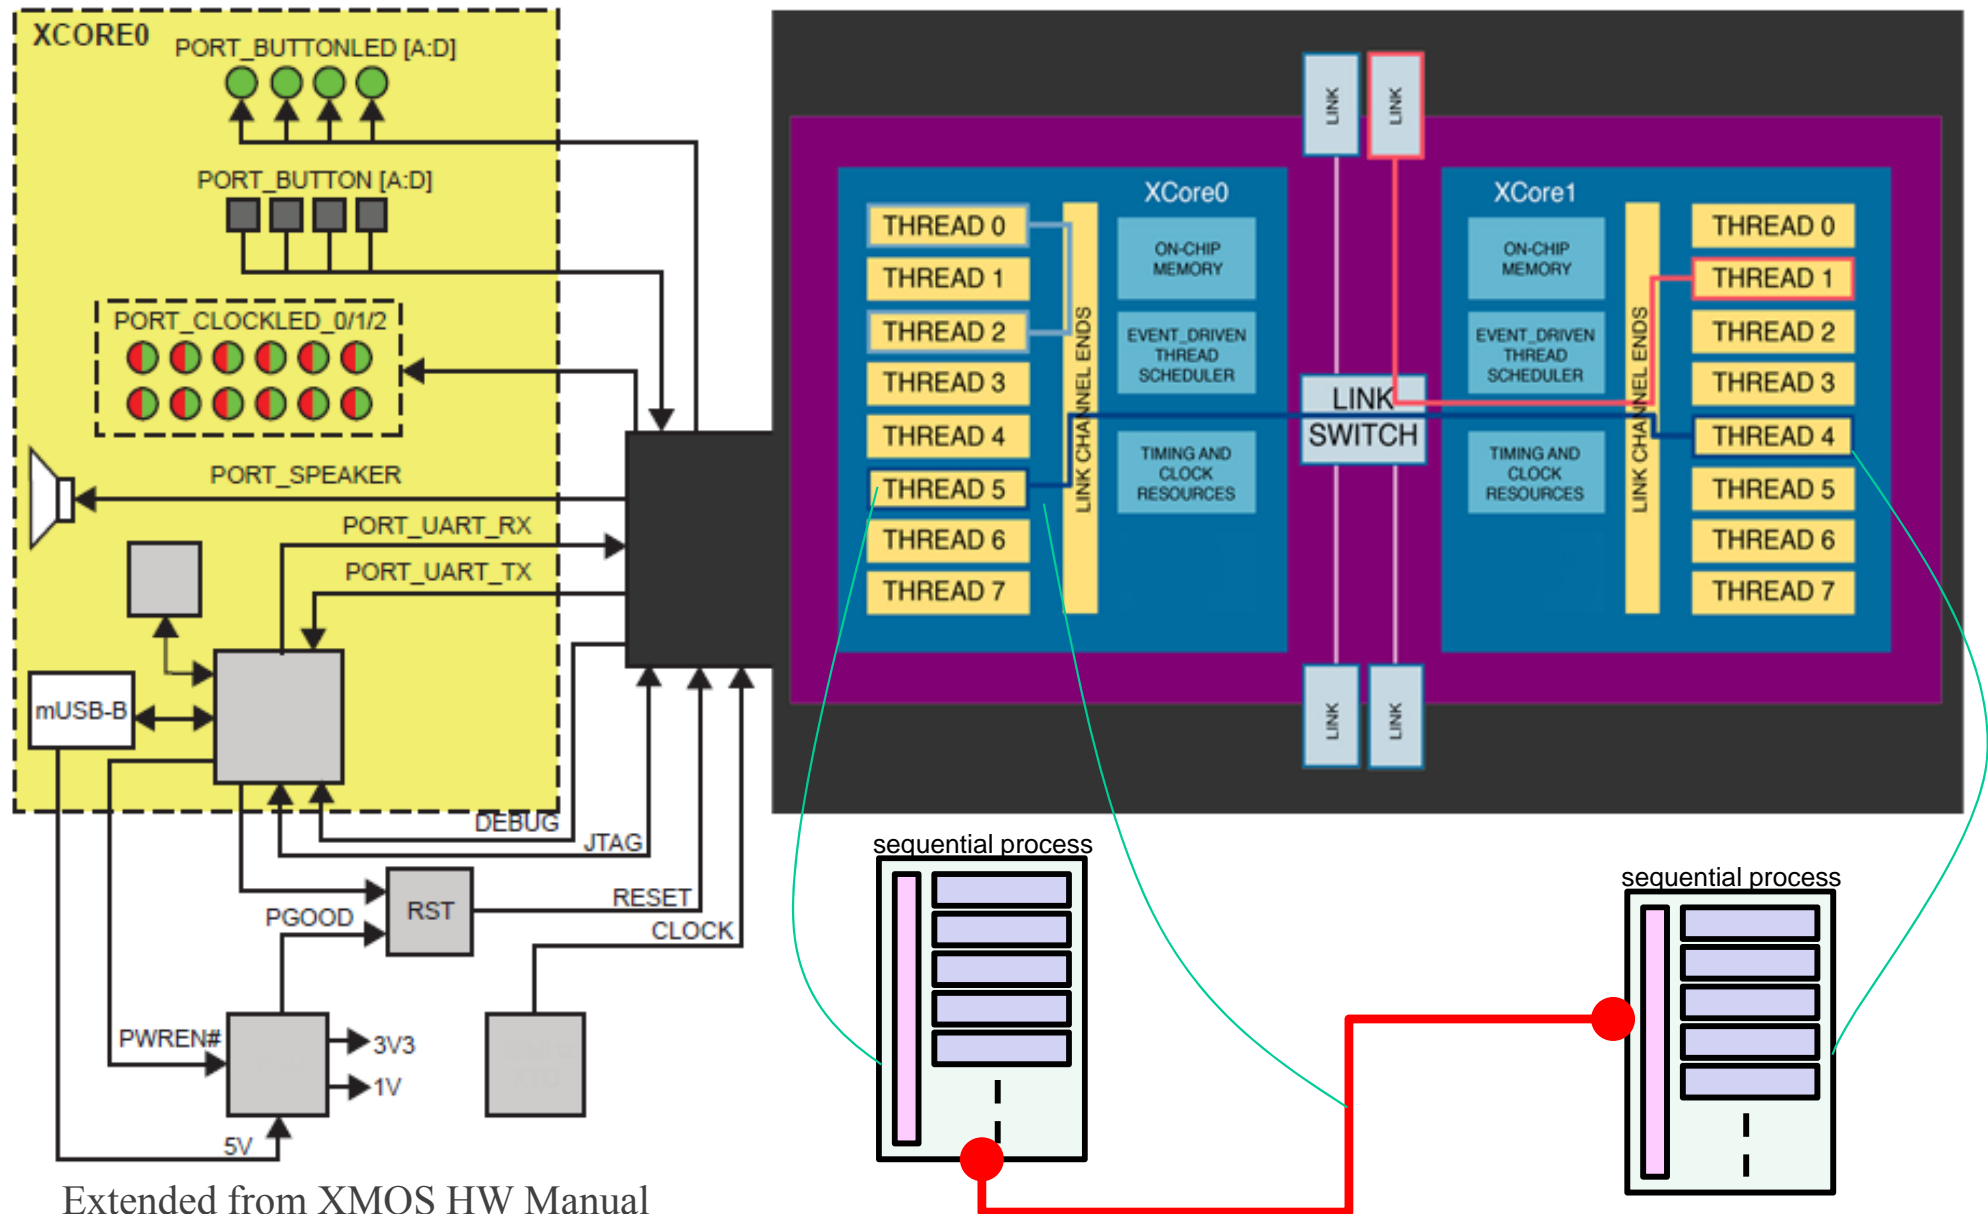
## Three fundamental stages:

- **Assignment** of application tasks to processes
- **Orchestration** of processes into a concurrent program
- **Mapping** of program onto available resources

# Stages in Designing a Concurrent System



# Mapping to Hardware on XMOS XS-1



# Basic Example: Distance Evaluation

- Assume that the task is to design a concurrent system that uses two processes efficiently in order to calculate distances between:
  - patrol locations stored in an array **A**
  - emergency locations stored in an array **B**and finally outputs the 2D result array **C**

```
// SEQUENTIAL SOLUTION (FRAGMENT)
...
typedef point...;           //define point structure
point    A[4]={...};       //initial locations A
point    B[4]={...};       //initial locations B
int       C[4][4]={...};   //initial result matrix C

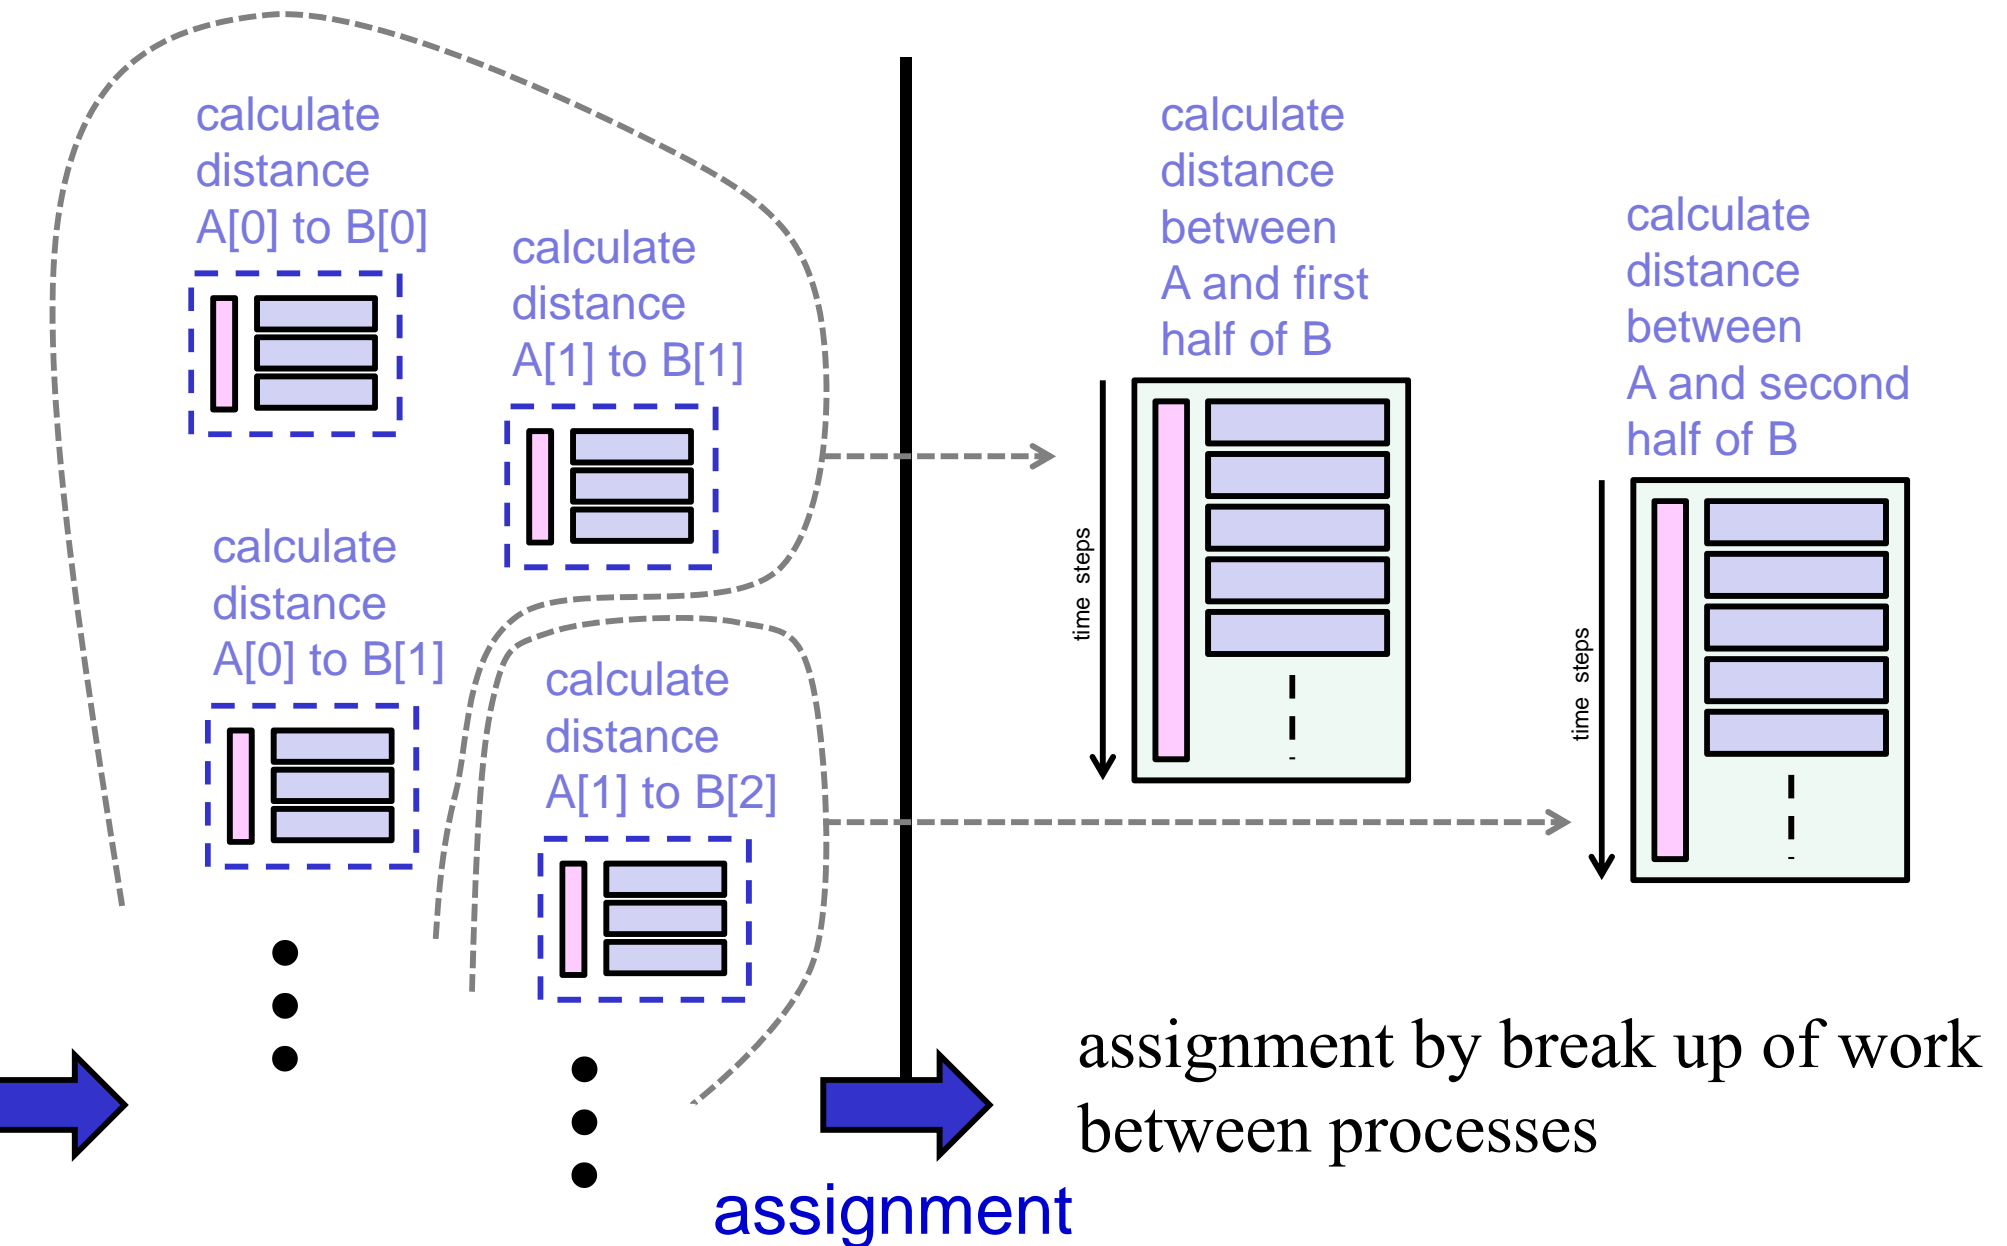
// calculate all distances
for (int i = 0; i<4; i++)
    for (int j = 0; j<4; j++)
        C[i][j] = dist(A[i], B[j]);

// output result C
...
```

# Problem Decomposition & Task Assignment

- **Creation of Process Candidates**  
(i.e. identifying tasks and partitioning the task space into processes by bundling tasks)
  - **Criteria:** minimum coupling, maximum compactness
    - Goal: clustering tasks of similar functionality demanding similar data portions
    - Analogy: design of objects in OO programming
  - **Target:** balance work and reduce communication during runtime
    - application often hints at good process candidates

# Example Task Assignment

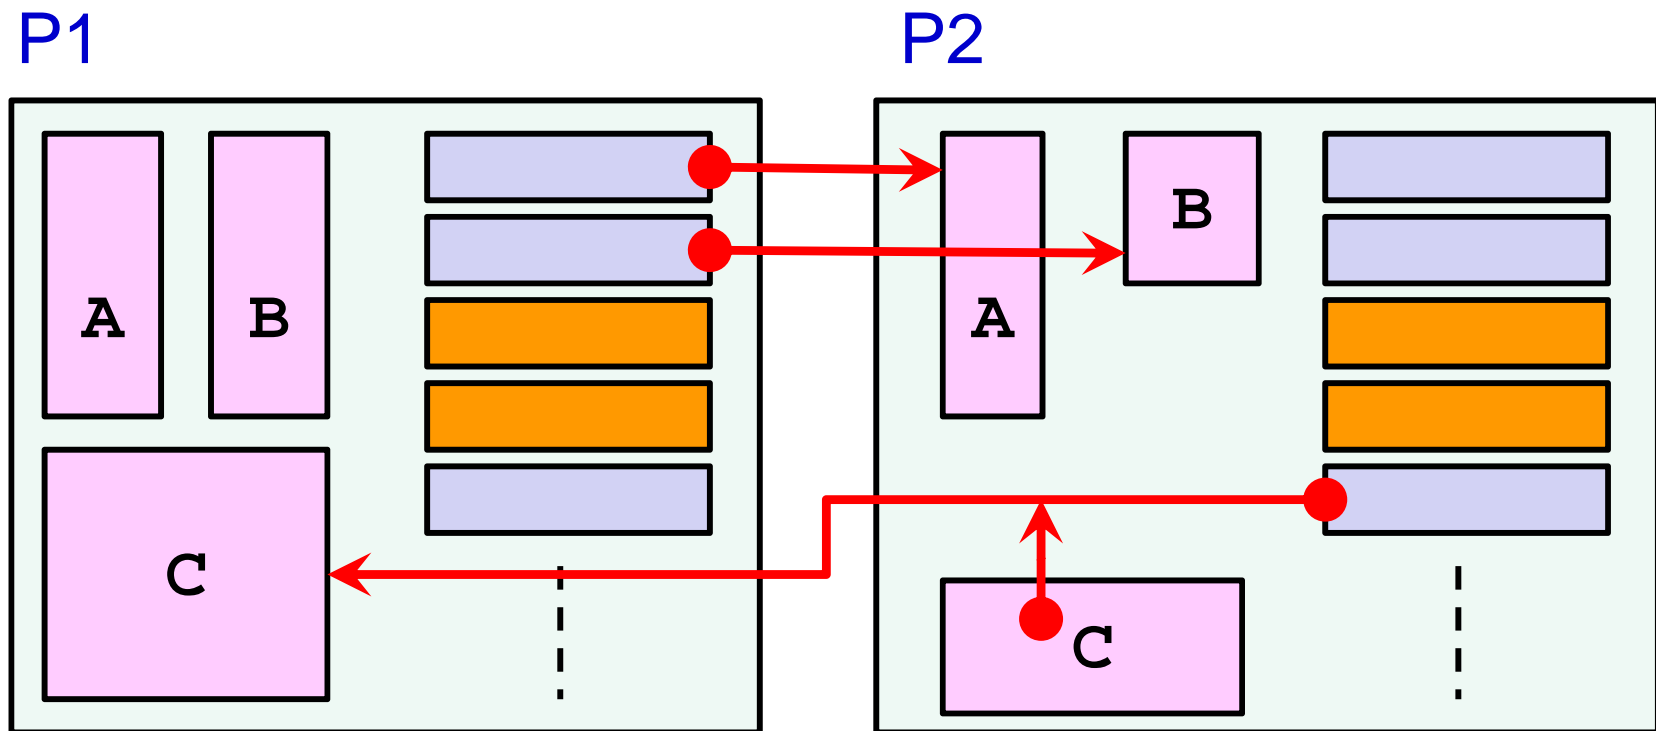




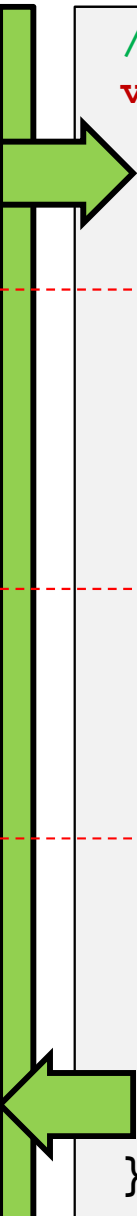
# Basic Example: Process Orchestration

- **Orchestration:** determine communication strategy and channels

- e.g. P1 sends data to P2
- P1 and P2 compute
- P2 sends output to P1



# Concurrent XC Fragment showing Data Flow



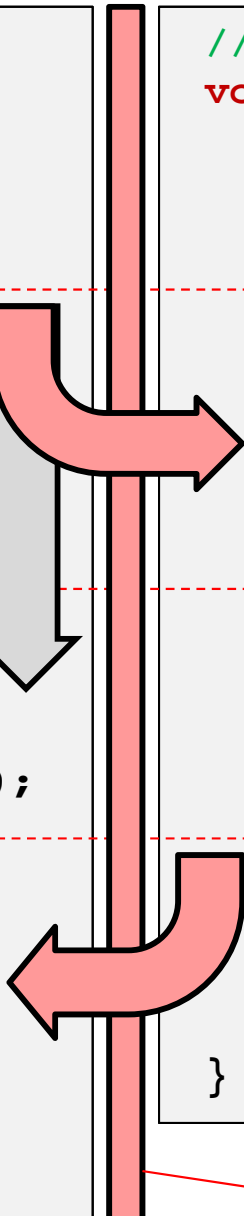
```
// process P1 (fragment)
void processOne(chanend ch) {
    point A[4]={...}; //init A
    point B[4]={...}; //init B
    int    C[4][4]={...}; //init C

    // send data to process two
    for (int i = 0; i<4; i++)
        ch <: A[i];
    for (int i = 0; i<2; i++)
        ch <: B[i];

    // calculate some distances
    for (int i = 0; i<4; i++)
        for (int j = 2; j<4; j++)
            C[i][j] = dist(A[i], B[j]);

    // receive results
    for (int i = 0; i<4; i++)
        for (int j = 0; j<2; j++)
            ch :> C[i][j];

    ... //output result C
}
```



```
// process P2 (fragment)
void processTwo(chanend ch) {
    point A[4]; //allocate local memory
    point B[2];
    int    C[4][2]={...};

    // receive data from process one
    for (int i = 0; i<4; i++)
        ch :> A[i];
    for (int i = 0; i<2; i++)
        ch :> B[i];

    // calculate some distances
    for (int i = 0; i<4; i++)
        for (int j = 0; j<2; j++)
            C[i][j] = dist(A[i], B[j]);

    // send results
    for (int i = 0; i<4; i++)
        for (int j = 0; j<2; j++)
            ch <: C[i][j];
}
```


channel ch

# Trade-off: Parallelisation vs. Communication


- in the example, distance calculations between points are independent of each other (...and thus, can be parallelised easily... ‘embarrassingly’ parallel)
  - Dividing work over  $N$  processes does not lead to  $N$ -times speedup as there are communication demands:
    - e.g. copy of **A** needs to be sent to all processors + copy of subset of **B** to each processor + results need collecting
- **COMMUNICATION is NOT free...**  
**...it creates overheads and takes time!**

# Impact of Granularity: Fine vs. Coarse

- computation is typically separated by periods of communication
- **Fine-grain System**
  - Small computation to communication ratio
  - Little computational work between communication stages
  - Less opportunity for computation enhancement
  - High communication overhead
  - Good for reactive & dynamic problems
- **Coarse-grain System**
  - Large computation to communication ratio
  - Big blocks of computational work between communications
  - More opportunity for performance increase due to large blocks
  - Hard to load balance efficiently



compute  
communicate



# Dependencies: Limits to Performance Scalability

- So far, we have focussed on problems that are easily divided into parallel threads ... an important design issue is **Data Dependency!**
- Example:

**A = B + C**


**D = A + A**

**E = D / C**


One computation can not proceed until another is complete.

- Dependencies between (interim) results enforce waiting at synchronisation points...
- Ordering and balancing of calculations can help minimise idle waiting for results...


# Analysing Data Dependency: Reduction Graphs



# Serial Reduction vs. Tree Reduction

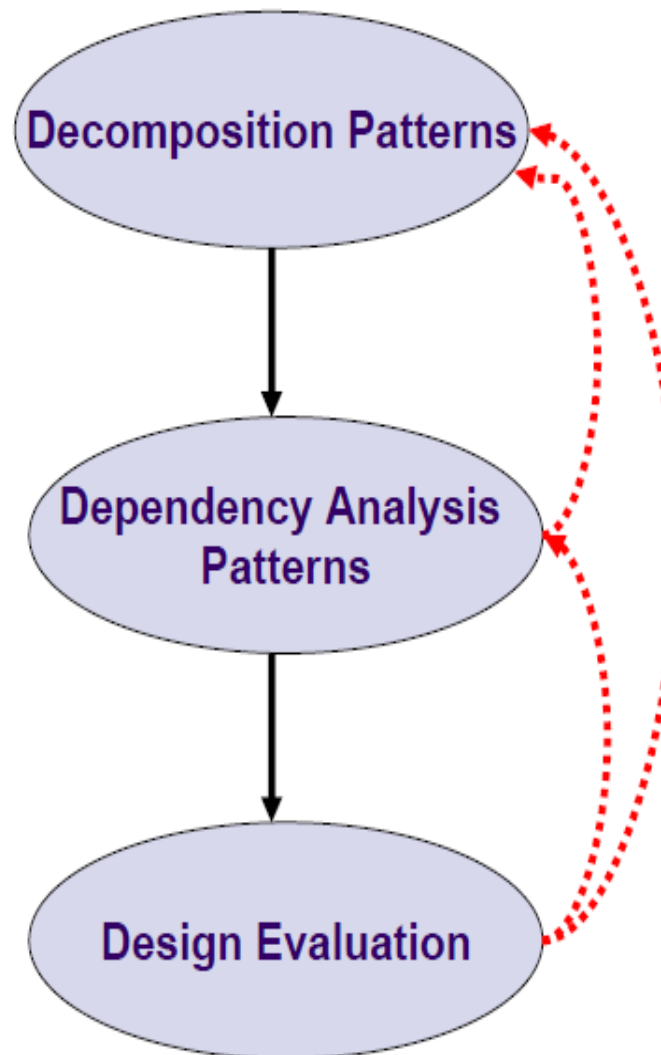



- Good for non-associative combination of interim values
- Usually followed by broadcast of result



- Good for associative combination of interim values
- $N$  steps for  $2^N$  elements
- Especially efficient when only one task needs result

# Design Space for Concurrent Systems




- Processes formed are at right *granularity*?
- Processes are *compact* operating on local data?
- Communication *dependencies* are well *reduced*?
- Expected message *traffic is low*?
- Hardware *supports* approach well?



# Bernstein's Criteria for Task Independency

(What can be parallelised?)



A J Bernstein



$R_i \rightarrow$  set of memory locations read (**input**) by task  $T_i$

$W_j \rightarrow$  set of memory locations written to (**output**) by task  $T_j$

**Two tasks  $T_1$  and  $T_2$  are independent if:**

- **input** to  $T_1$  is not part of **output** from  $T_2$
- **input** to  $T_2$  is not part of **output** from  $T_1$
- **outputs** from  $T_1$  and  $T_2$  do not overlap

**Example:**



$R_1 = \{ x, y \}$   
 $W_1 = \{ a \}$

$R_2 = \{ x, z \}$   
 $W_2 = \{ b \}$

$$R_1 \cap W_2 = \phi$$

$$R_2 \cap W_1 = \phi$$

$$W_1 \cap W_2 = \phi$$