

# Concurrent Computing

---

Lecturers:                      Prof. Majid Mirmehdi   [majid@cs.bris.ac.uk](mailto:majid@cs.bris.ac.uk)  
                                        Dr. Tilo Burghardt        [tilo@cs.bris.ac.uk](mailto:tilo@cs.bris.ac.uk)  
                                        Dr. Daniel Page           [page@cs.bris.ac.uk](mailto:page@cs.bris.ac.uk)

Web:                                <http://www.cs.bris.ac.uk/Teaching/Resources/COMS20001>



## LECTURE MM5

### *PARADIGMS OF PARALLELISM*


# Lecture Overview

- **Paradigms of Hardware Parallelism**
  - Flynn's Taxonomy
  - SISD, SIMD, MISD, MIMD
- **Paradigms of Logical Parallelism**
  - Independence: Bernstein's Condition
  - Geometric Parallelisation
  - Parallelisation via Farming
  - Algorithmic Parallelisation
- **Process Networks**
- **Performance issues**

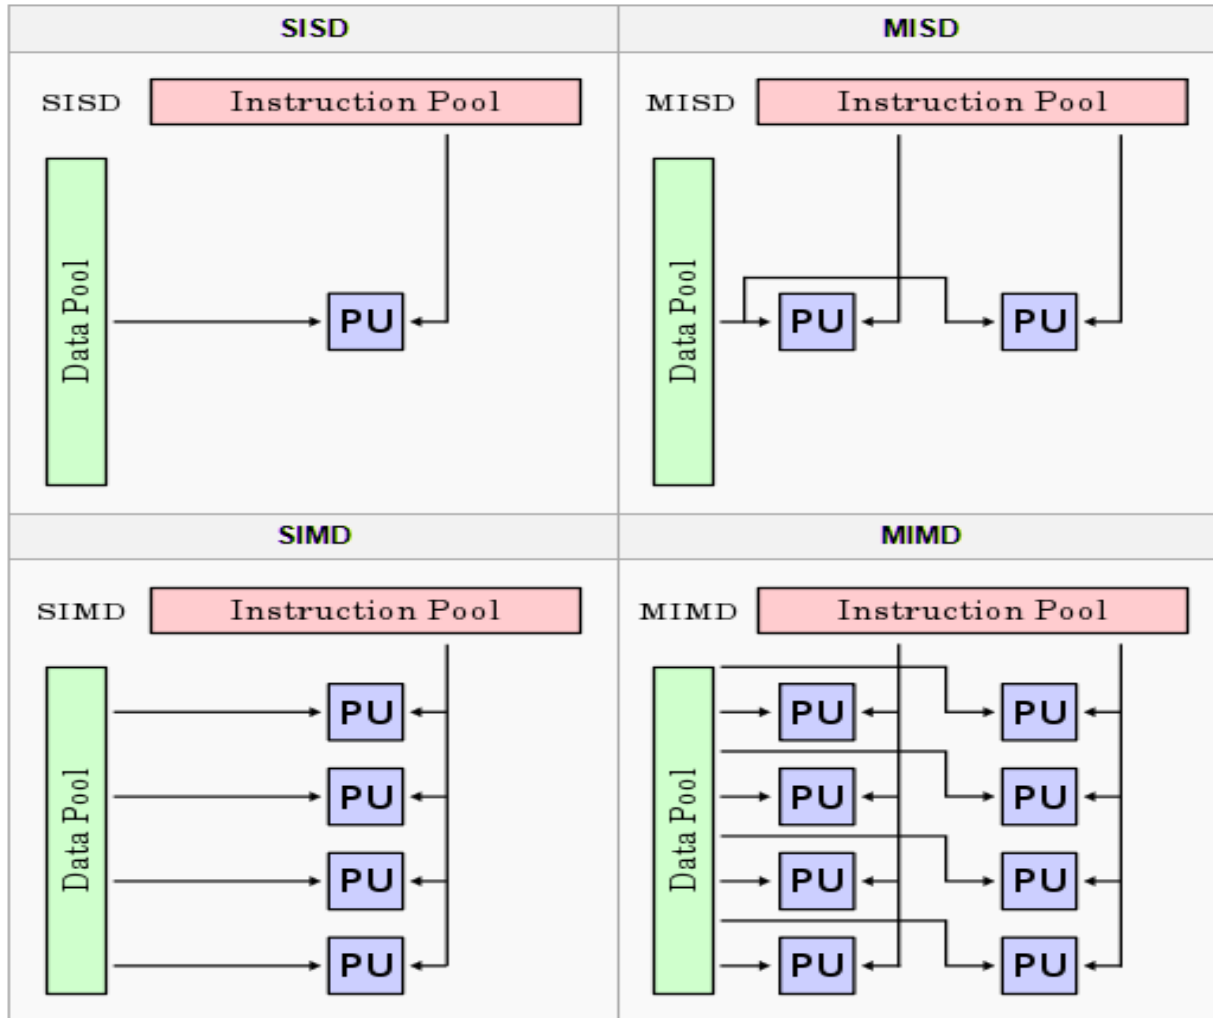


Blue Crystal, Bristol

# Flynn's Taxonomy



Michael J. Flynn




Flynn identifies **four classes** based on the number of concurrent instruction (or control) and data streams available in an architecture.

# Paradigms of Hardware Parallelism

(What fundamental types of parallelisation exist?)

## SISD

...sequential computer, no parallelism. Single control unit fetches **single instruction** stream from memory. Instruction manipulates **single data** item at a time., e.g. *1990s Pentium*




## SIMD

...exploits **multiple data** streams against **single instruction** stream, e.g. *MMX technology for multimedia data types on Pentiums*



## MISD

...**multiple instructions** operate on a **single data** stream; often used for fault tolerance., e.g. Space Shuttle flight controllers



## MIMD

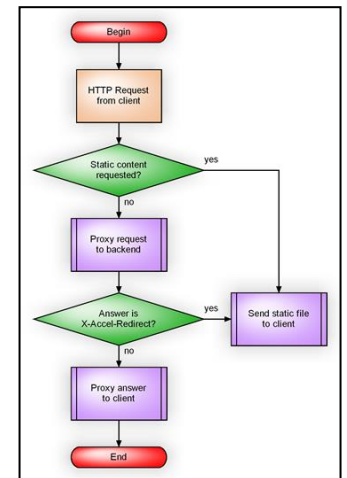

...**multiple** autonomous processors simultaneously execute different **instructions** on different **data**, e.g. *Intel i7, XMOS XC-1A*




# Paradigms of Logical Parallelism

(How can it be parallelised?)

- **Geometric:** i.e. data is divided up
  - processes are quasi-independent but identical
  - each process operates on a portion of data
  - interaction with neighbours (according to geometry)
- **Farming:** i.e. data is divided up
  - single source (farmer) distributes work (data) to workers
  - work is done independently (at own pace) by worker(s)
  - workers send results to harvester
- **Algorithmic:** i.e. algorithm is parallelised
  - quasi-independent processes execute sections of algorithm
  - processes may be non-identical
  - data is passed between processes according to algorithm
- (and mixtures of the above)



# Geometric Parallelism (static data division)




- **Geometric: data is proportioned equally**

Useful when there is a natural way to decompose the data set into smaller "chunks" and allocate them to individual processors. Replication can often serve as the enabling programming concept.

**Each processor executes the same code but clearly operates on a subset of the total data.**

## Wall building analogy




 Distribute bricks amongst a number of bricklayers

 Each bricklayer responsible for a vertical section of wall.

 Communication required with those building neighbouring sections.



# Geometric Parallelism (static data division)



- **Geometric Parallelism**

Useful for  
the data is  
them often

Each processor  
on a different


Wall time

Discrete

Each

Computation

sections.



# Case Study: Geometric Parallelism

## Matrix Multiplication

matrix  $A$  of order  $m \times p$  with elements  $a_{ik}$ , and  
matrix  $B$  of order  $p \times n$  with elements  $b_{ik}$ .

Product is matrix  $C$  of order  $m \times n$  with elements  $c_{ik}$  where...

$$c_{ik} = \sum_{j=1}^p a_{ij} b_{jk}$$


Each element  $c_{ik}$  is product of  $i^{th}$  row of  $A$  with  $k^{th}$  column of  $B$ .



# Matrix Multiplication (sequential)

```
// sequential solution, fragment
...
for (int i = 1; i < m; i++)
  for (int k = 1; k < m; k++) {
    int tot = 0;
    for (int j = 1; j < m; j++)
      tot := tot + a[i,j]*b[j,k];
    c[i,k] = tot;
  }
...
```

$$c_{ik} = \sum_{j=1}^m a_{ij} b_{jk}$$



# Matrix Multiplication: Cell Evaluation

Each cell communicates with neighbours via vertical and horizontal input and output channels.

```
// CODE FRAGMENT: cell evaluation

void multCell(chanend left, chanend top, chanend right, chanend bottom) {

    int result = 0;
    int a, b;
    for (int p=0; p<3; p++) {
        par {
            top :> a;
            left :> b;
        }
        par {
            result = result + (a * b);
            bottom <: a;
            right <: b;
        }
    }
}
```

# Matrix Multiplication: Main Harness

```
// CODE FRAGMENT: Harness for setting up the multiplier cells in grid

// dimensions of M x N matrix
#define M 3 // no of rows
#define N 3 // no of columns

// main harness
int main(void) {
    chan hor[M][N+1]; // horizontal channels
    chan ver[N][M+1]; // vertical channels
    // add code here to feed grid with values
    par (int i=0;i<M;i++) { // go through rows
        par (int j=0;j<N;j++) { // go through columns
            multCell(hor[i][j],ver[j][i],hor[i][j+1],ver[j][i+1]);
        }
    }
    // add code here to bleed grid with output results
    return 0;
}
```

# Farming (dynamic data division)



- **Farming: processors work on data subsets when ready**

Involves a farmer distributing work.

- Workers receive work from farmer,
- process the work or pass it on to next worker,
- result is returned to harvester (e.g. farmer).

**Used for computation intensive calculations, e.g. ray tracing.**

## Wall building analogy




Each bricklayer lays a brick at a time.



Bricks and mortar are taken from a common pool.




Communication of interim results not particularly necessary.




# Farming (dynamic data division)




# Common Topologies of Processor Farms



Simple Star



Hierarchical Tree (with co-worker pairs)

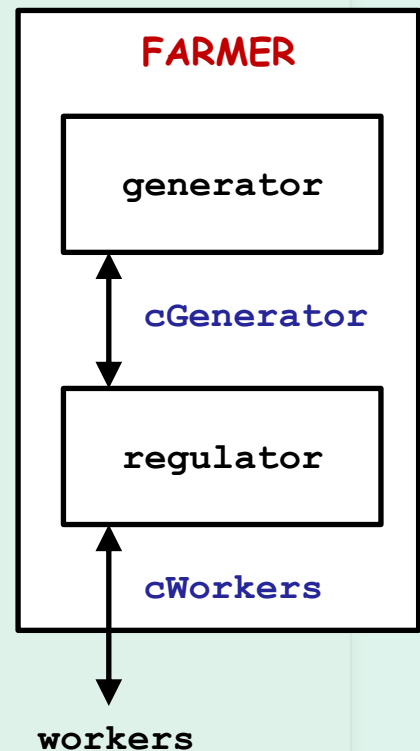


Deep Stumps (star of worker chains)



# XC Regulator Fragment for Processor Farm

```
// fragment code for farm regulator
void regulator(chanend cWorkers, chanend cGenerator, int workerNum) {
    ...
    int idle = workerNum;
    while (farmerRunning) {
        if (idle > 0)
            select { //some workers free, results or data allowed
                case cWorkers :> result:
                    cGenerator <: result;
                    idle++;
                    break;
                case cGenerator :> dataPacket:
                    cWorkers <: dataPacket;
                    idle--;
                    break;
                ...
            }
        else
            select { //all workers busy, only results allowed
                case cWorkers :> result:
                    cGenerator <: result;
                    idle++;
                    break;
                ...
            }
    }
    ...
}
```



# The Price Tag: Communication Overheads

For example in farming...

- Adding more workers to a farm can enhance performance.

**However,**

**... the farm can become swamped with overheads of routing data to/from workers.**

- This takes processing time away from the real work! Serious performance degradation if too many workers are added.  
→ Communication bottleneck!
- Finding the right balance between communication overheads and time spent processing can be very hard!

...can also occur in geometric parallelism...

# Algorithm Parallelisation



- **Structure of algorithm is decomposed/parallelised**

Stream of data pass through a pipeline.

Each stage of the pipeline performs some operation.

Video processing pipeline with 4 stages

1. Smoothing a raw image
2. Feature extraction from smoothed image
3. Object recognition from feature descriptions
4. Graphics display of recognized objects

**Example**

**Wall building analogy**



Each bricklayer lays a row of bricks.



Work is done simultaneously with others, but offset in time.



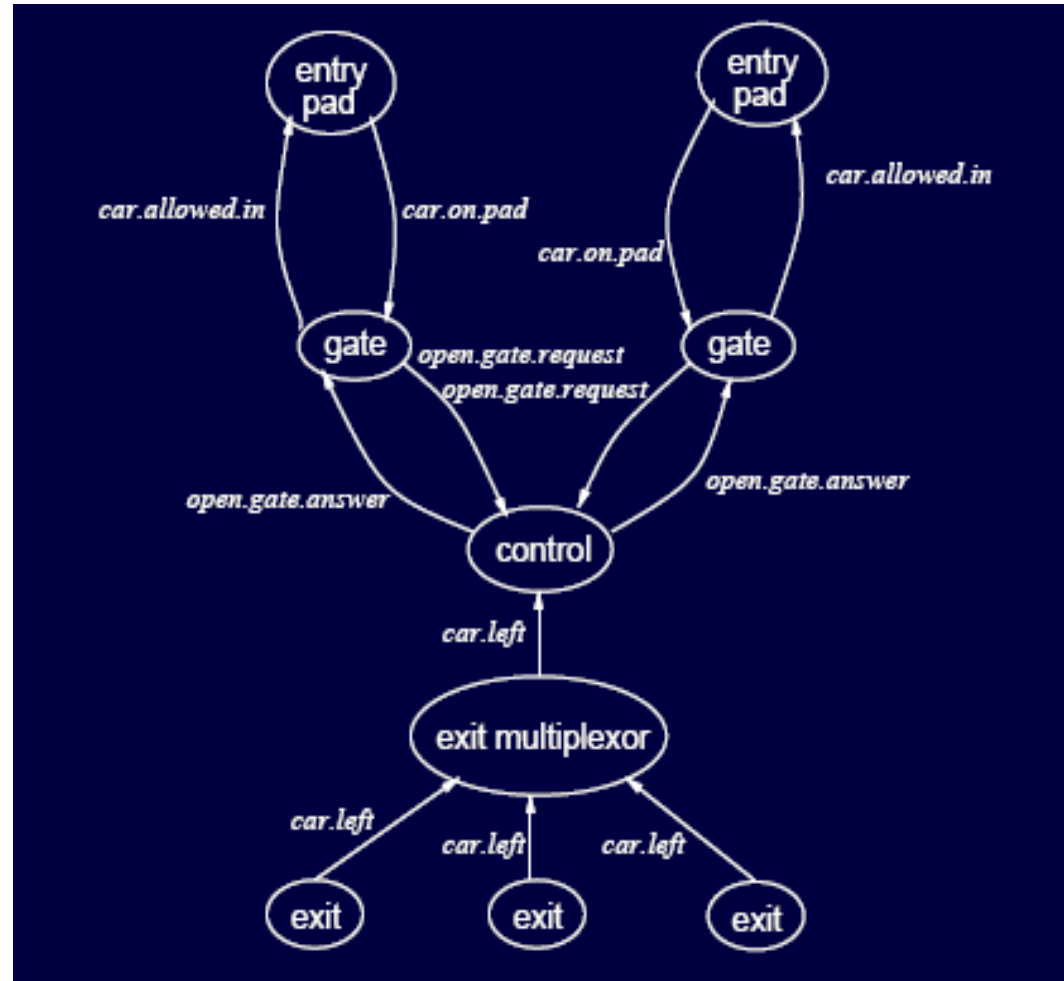
Work at pace of slowest worker, some workers idle at beginning/end.

# Generic Topologies: Process Networks

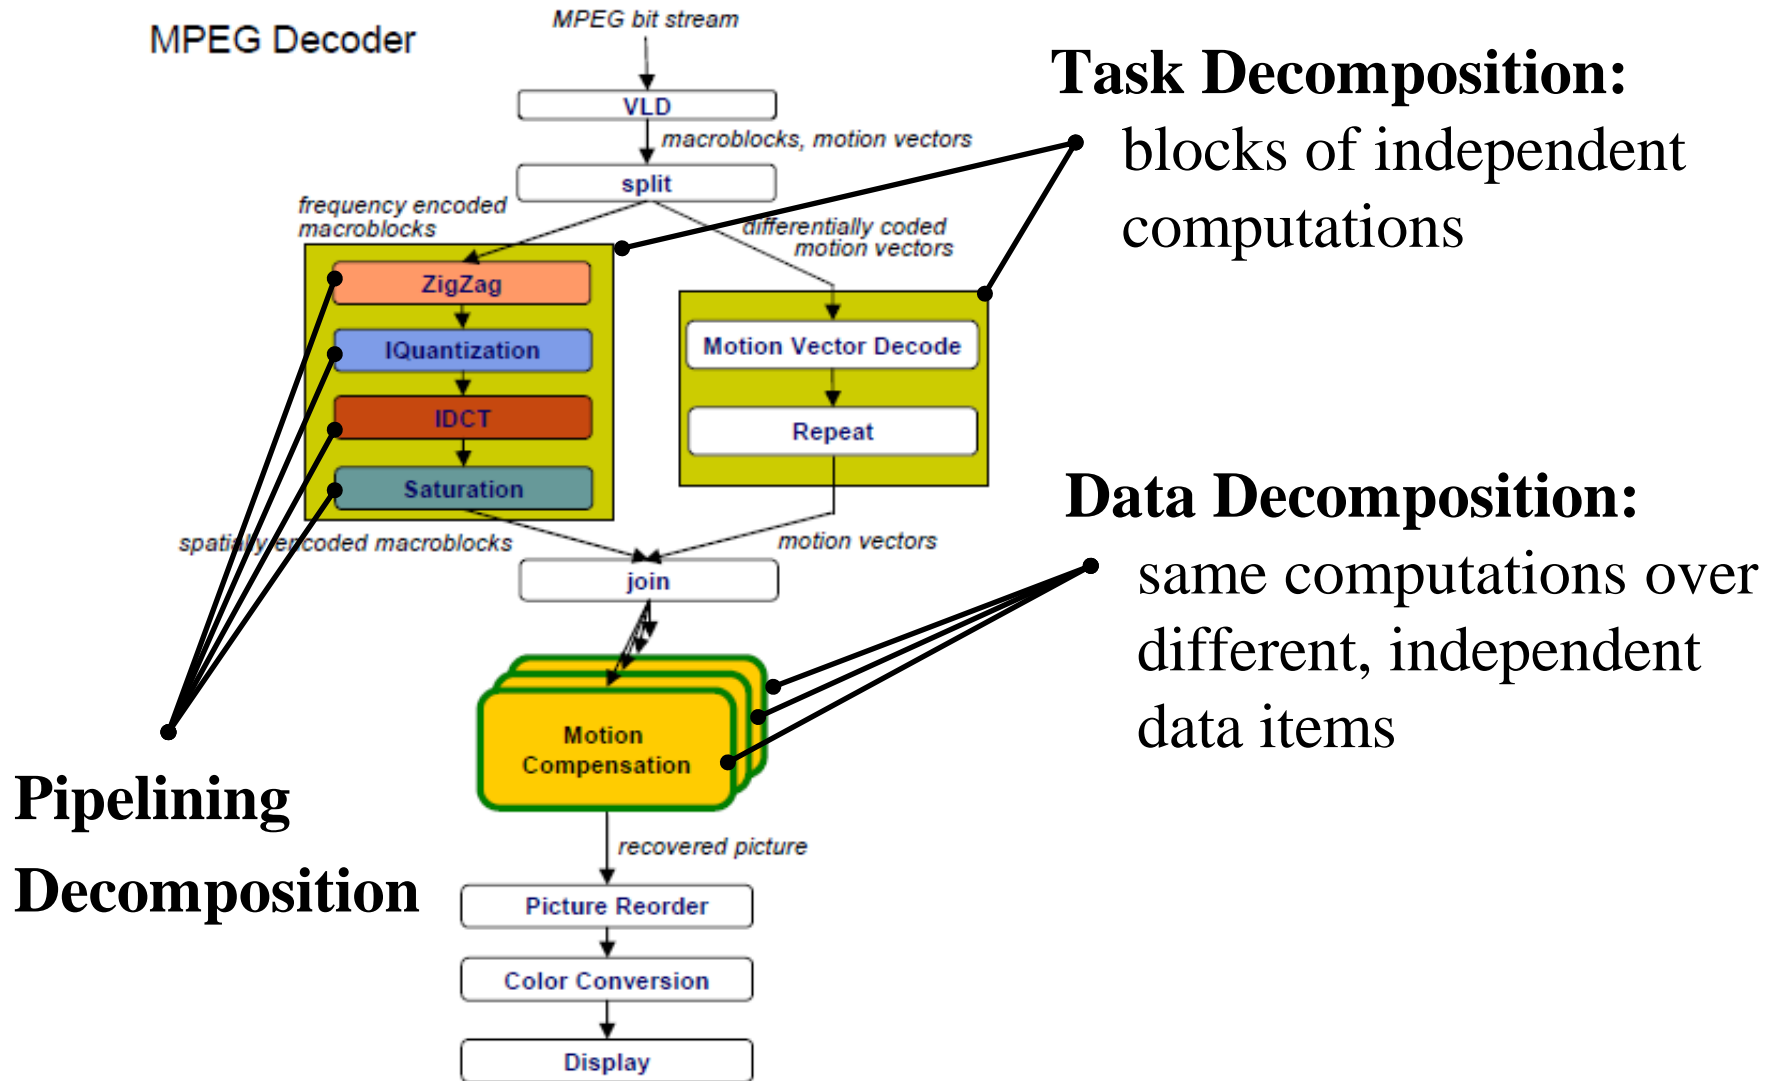
Networks provide the most flexible process topology.

## Example


Car park with two entry pads (sensors) controlling a gate each, and three exit gates.



# Example: Decomposition of MPEG Decoder

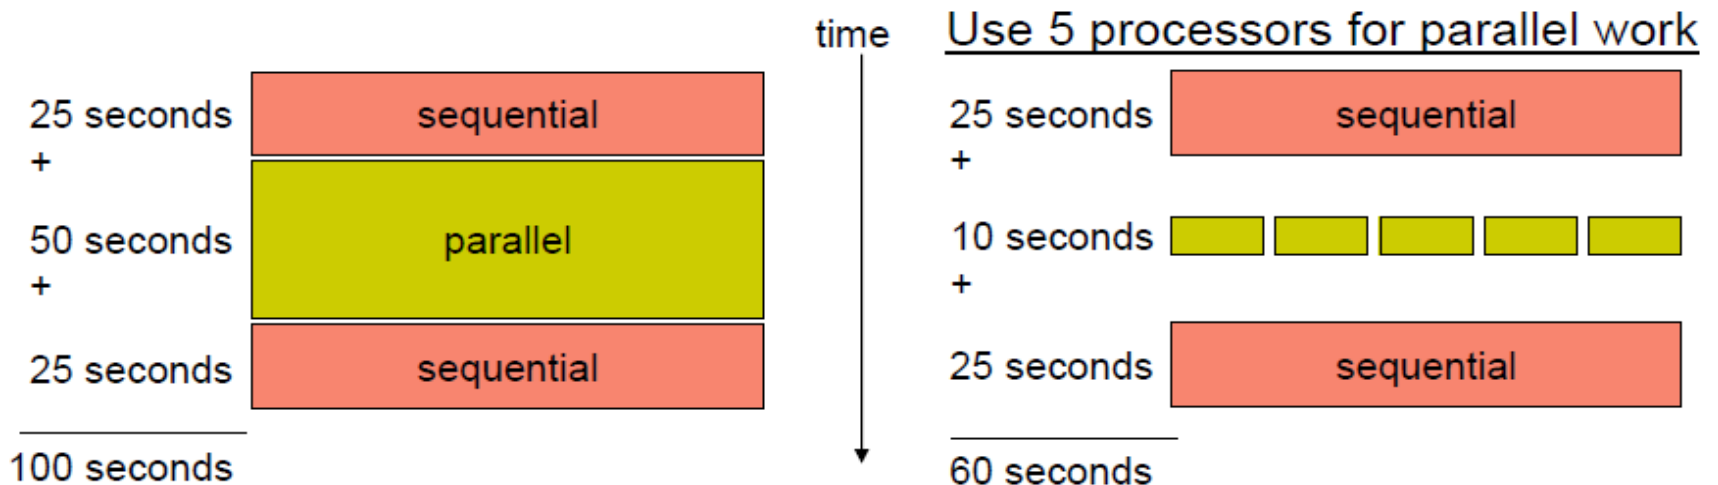


# Speedup Limit: Amdahl's Law




Gene Amdahl

The speedup of a program using multiple processors in parallel computing is limited by the time needed for the sequential fraction of the program.





# Speedup Limit: Amdahl's Law



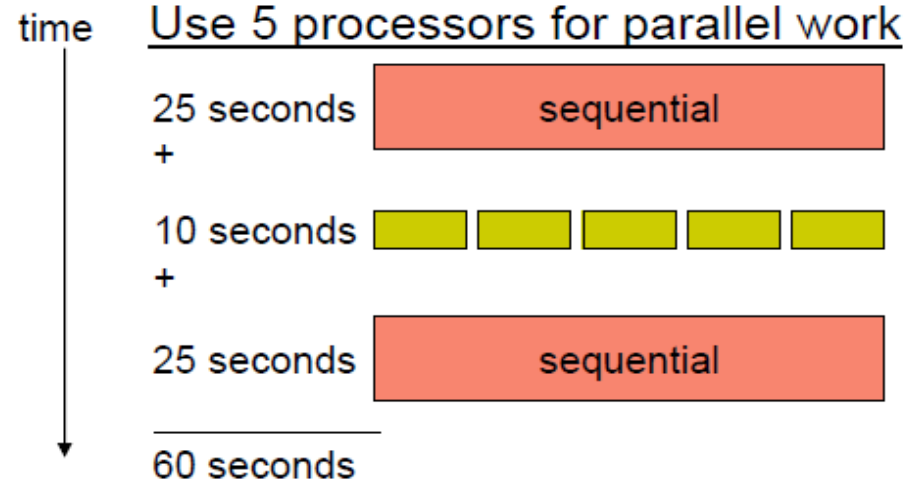
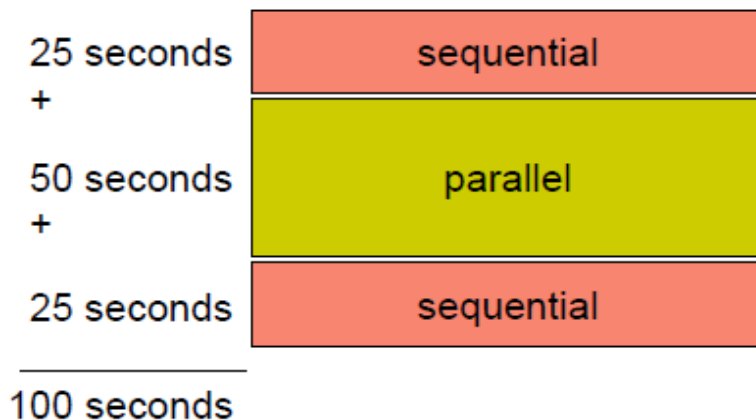
Gene Amdahl

$p$  = fraction of work that can be parallelized  
 $n$  = the number of processor

$$\text{speedup} = \frac{\text{old running time}}{\text{new running time}}$$
$$= \frac{1}{(1-p) + \frac{p}{n}}$$

fraction of time to complete sequential work

fraction of time to complete parallel work



# Factors that Impact on Performance...

- The *number* and *throughput* of the processing nodes.
- The *bandwidth* between the processing nodes. Measured in Mbytes/sec (or Gbytes/sec etc.) this determines the rate at which data can be sent from one node to another.
- The *latency* in the node-to-node connections measures (in seconds) their delays in transmission - influenced by the bit rate, the time taken to route data, and the transmission protocol used.
- The *Flow Control* strategy used influences performance - nodes must be prevented from flooding each other, or the routing fabric, with traffic.

# Factors that Impact on Performance...

- The amount of **Local Memory**, e.g. in MB, often limits local computation without communication
- The **Granularity** of the system describes how coarsely/finely a computation and its data is partitioned into concurrently computing processes – this directly impacts on communication and load balancing
- The **Speedup factor** for a parallel computer can be defined as:

$$S(N) = \frac{\text{Execution time using one processor}}{\text{Execution time using } N \text{ processors}}$$

- The **Efficiency** is then defined as:

$$E(N) = \frac{S(N)}{N}$$