

# Concurrent Computing

Lecturers:

Prof. Majid Mirmehdi ( majid@cs.bris.ac.uk)

Dr. Tilo Burghardt ( tilo@cs.bris.ac.uk)

Dr. Simon Hollis( simon@cs.bris.ac.uk)

Dr. Daniel Page ( page@cs.bris.ac.uk)

Web:

<http://www.cs.bris.ac.uk/Teaching/Resources/COMS20001>



## LECTURE 10 / 11

*Liveness  
&  
Deadlock*

[Many thanks to Kerstin Eder, major parts of these lecture slides are taken from or based on materials originally prepared by her.]

# Recap: Labelled Transitions and Traces

Traces can be extracted from a transition diagram!

 Labelled transitions capture same information as transition diagrams!

Relationship between traces of a process and its labelled transitions:

$$traces(P) = \{\langle \rangle\} \cup \{\langle a \rangle \frown tr \mid P \xrightarrow{a} Q \text{ and } tr \in traces(Q)\}$$

where  $\frown$  denotes concatenation, i.e.

$$\langle a_1, \dots, a_m \rangle \frown \langle b_1, \dots, b_n \rangle = \langle a_1, \dots, a_m, b_1, \dots, b_n \rangle$$

Behaviour of new CSP operators can be defined through labelled transition rules.  Use above relationship to calculate the traces of processes defined in terms of new operators.

# Recap: Refusals and Failures

We write  $P/tr$  for the process whose behaviour is whatever  $P$  could do after the trace  $tr$  has been observed.

## Failures of a process:

$$\text{failures}(P) = \{(tr, X) \mid tr \in \text{traces}(P) \text{ and } X \in \text{refusals}(P/tr)\}$$

- $P = a \rightarrow b \rightarrow STOP$  with  $\alpha(P) = \{a, b\}$

Transition Diagram of  $P$ :



$$\text{traces}(P) = \{\langle \rangle, \langle a \rangle, \langle a, b \rangle\}$$

$$\text{refusals}(P/\langle \rangle) = \{\{\}, \{b\}\}$$

$$\text{refusals}(P/\langle a \rangle) = \{\{\}, \{a\}\}$$

$$\text{refusals}(P/\langle a, b \rangle) = \{\{\}, \{a\}, \{b\}, \{a, b\}\}$$

$$\begin{aligned} \text{failures}(P) &= \{(\langle \rangle, \{\}), (\langle \rangle, \{b\}), \\ &\quad (\langle a \rangle, \{\}), (\langle a \rangle, \{a\}), \\ &\quad (\langle a, b \rangle, \{\}), (\langle a, b \rangle, \{a\}), (\langle a, b \rangle, \{b\}), (\langle a, b \rangle, \{a, b\})\} \end{aligned}$$

# Failure Refinement I

Failure refinement is defined in a similar way to trace refinement:

$$P \sqsubseteq_F Q \text{ if and only if } \text{failures}(Q) \subseteq \text{failures}(P)$$

(Pronounce: " $P$  is failure refined by  $Q$ ")

**Failure refinement in specifications:**

- $SPEC = a \rightarrow b \rightarrow SPEC$ 
  - 💡 Use  $SPEC$  with trace refinement, get a *safety specification*!
- Find some processes  $P$  which satisfy  $SPEC \sqsubseteq_T P$ .  
 $P = STOP, P = a \rightarrow STOP, P = a \rightarrow b \rightarrow STOP, \dots$
- What effect has  $SPEC \sqsubseteq_F P$ ? 💡 First, calculate  $\text{failures}(SPEC)$ !

# Failure Refinement II

$$\begin{aligned} \text{failures}(SPEC) = & \{(\langle a, b \rangle^n \setminus \langle a \rangle, \emptyset) \mid n \geq 0\} \\ & \cup \{(\langle a, b \rangle^n \setminus \langle a \rangle, \{a\}) \mid n \geq 0\} \\ & \cup \{(\langle a, b \rangle^n, \emptyset) \mid n \geq 0\} \\ & \cup \{(\langle a, b \rangle^n, \{b\}) \mid n \geq 0\} \end{aligned}$$

To find out whether  $SPEC \sqsubseteq_F STOP$ , calculate:

$$\text{failures}(STOP) = \{(\langle \rangle, \emptyset), (\langle \rangle, \{a\}), (\langle \rangle, \{b\}), (\langle \rangle, \{a, b\})\}$$

Pairs  $(\langle \rangle, \{a\})$  and  $(\langle \rangle, \{a, b\})$  are failures of  $STOP$ , but not of  $SPEC$ .  
Hence,  $SPEC \not\sqsubseteq_F STOP$ .

Now, consider  $P = a \rightarrow STOP$ .

$$\text{failures}(P) = \{(\langle \rangle, \emptyset), (\langle \rangle, \{b\}), (\langle a \rangle, \emptyset), (\langle a \rangle, \{a\}), (\langle a \rangle, \{b\}), (\langle a \rangle, \{a, b\})\}$$

Failure pairs  $(\langle a \rangle, \{b\})$  and  $(\langle a \rangle, \{a, b\})$  are failures of  $P$  but not of  $SPEC$ ; so again  $SPEC \not\sqsubseteq_F P$ .

# Liveness (guaranteed execution of some behaviour)

$SPEC \sqsubseteq_F P$  is a *liveness* specification which requires  $P$  to do certain events.

■ Which definitions of  $P$  satisfy  $SPEC = a \rightarrow b \rightarrow SPEC$ ?

Obviously  $P = a \rightarrow b \rightarrow P$  does.

 It is (in this case) the only process satisfying this specification!

(Specification is too tight; pins down implementation precisely.)

# Liveness Specification Example & Hiding

Process  $P$  with alphabet  $\{a, b, c\}$ .

- Want to specify that  $P$  must be able to do an infinite sequence of alternating  $a$  and  $b$  events, starting with  $a$ .
  - We do not care about  $c$  events.
- Use process  $ALT = a \rightarrow b \rightarrow ALT$  as before.  
Allow  $c$  events to occur freely through hiding:  $ALT \sqsubseteq_F (P \setminus \{c\})$
- Definitions of  $P$  satisfying this specification include:  $P = a \rightarrow b \rightarrow P$ ,  
 $P = c \rightarrow a \rightarrow c \rightarrow c \rightarrow b \rightarrow P$ ,  $P = a \rightarrow b \rightarrow c \rightarrow P$ .
- 💡 All are the same as  $ALT$  when  $c$  is hidden!
- Definitions of  $P$  not satisfying this specification include:  $P = STOP$ ,  
 $P = a \rightarrow b \rightarrow (P \square a \rightarrow c \rightarrow STOP)$

# Safety Spec vs. Liveness Spec

Saying that  $tr \in traces(P)$  is a *positive* statement.

💡 Describes something that  $P$  can do!

$SPEC \sqsubseteq_T P$  puts limit on traces that  $P$  can do; restricts behaviour.

💡  $P$  may fail a *safety* (trace) specification by doing too much.

Saying that  $(tr, X) \in failures(P)$  is a *negative* statement.

💡 Describes something that  $P$  cannot do!

$SPEC \sqsubseteq_F P$  puts limit on what  $P$  can fail to do.

⇒ Requires  $P$  to accept at least a certain range of behaviours.

💡  $P$  may fail a *liveness* (failure) specification by refusing too much, i.e. by not doing enough.

# Example: Moving Furniture

- Two furniture movers need to move a table and a piano. Each requires two people to lift it.

*PETE* = *lift.piano* → *PETE*  
    □ *lift.table* → *PETE*

*DAVE* = *lift.piano* → *DAVE*  
    □ *lift.table* → *DAVE*

*TEAM* = *PETE* || *DAVE*

- 💡 Both Pete and Dave make their decisions independently! (□)
  - If both make same choice, they can cooperate in moving an object.

# Deadlock Example: Moving Furniture

- If their choices are different, ...

$PETE \xrightarrow{\tau} lift.piano \rightarrow PETE$

$DAVE \xrightarrow{\tau} lift.table \rightarrow DAVE$

💡  $lift.piano \rightarrow PETE \parallel lift.table \rightarrow DAVE$  cannot do anything.  
(It is equivalent to the process *STOP!*)

A state of a process is deadlocked if it can **refuse** to do every event.  
*STOP* is the simplest deadlocked process.

# Deadlock Example: Children Painting

■ Ella and Kate share a paint box and an easel.

*ELLA* = *ella.get.box* → *ella.get.easel* → *ella.paint* →  
*ella.put.box* → *ella.put.easel* → *ELLA*

*KATE* = *kate.get.easel* → *kate.get.box* → *kate.paint* →  
*kate.put.easel* → *kate.put.box* → *KATE*

*EASEL* = *ella.get.easel* → *ella.put.easel* → *EASEL*  
□ *kate.get.easel* → *kate.put.easel* → *EASEL*

*BOX* = *ella.get.box* → *ella.put.box* → *BOX*  
□ *kate.get.box* → *kate.put.box* → *BOX*

Combination of two children, box and easel:

*PAINTING* = *ELLA||KATE||EASEL||BOX*

(Assume synchronisation on (intersection of) individual alphabets.)

# Conditions for Deadlock

Coffman, Elphick and Shoshani identified 4 *necessary and sufficient* conditions for deadlock [System Deadlocks. ACM Computing Surveys 3, 2 (June), p. 67-78, 1971.]

1. Agents claim exclusive control of the resources they require.  
⇒ “**Mutual exclusion**” condition
2. Agents hold resources already allocated to them while waiting for additional resources.  
⇒ “**Wait for**” condition
3. Resources cannot be forcibly removed from the agent holding them until the resources are used to completion.  
⇒ “**No preemption**” condition
4. A circular chain of agents exists, s.t. each agent holds one or more resources that are being requested by the next task in the chain.  
⇒ “**Circular wait**” condition

# Breaking Deadlock

Aim: System in which possibility of deadlock is excluded a priori.

💡 Ensure that at least one of the conditions is not satisfied!

⇒ Constrain the way in which requests for resources are made.

- Usually “**Mutual exclusion**” condition **cannot be denied**.
- Each agent must request all its required resources at once and cannot proceed until all have been granted. 💡 Make it **atomic**!  
⇒ “**Wait for**” condition **denied**.
- If an agent holding certain resources is denied a further request, that agent must release its original resources and, if necessary, request them again together with the original resources.  
⇒ “**No preemption**” condition **denied**.
- Imposition of a *linear* ordering of resource types.  
⇒ “**Circular wait**” condition **denied**.

# Example: Breaking Deadlock

- ✿ If an agent holding certain resources is denied a further request, that agent must release its original resources and, if necessary, request them again together with the original resources.  
⇒ “No preemption” condition denied.

- Let Ella return tools before they have been used, rather than wait indefinitely for them to be available.

$ELLA' = ella.get.box \rightarrow (ella.put.box \rightarrow ELLA'$   
 $\quad \square ella.get.easel \rightarrow ella.paint \rightarrow$   
 $\quad \quad ella.put.box \rightarrow ella.put.easel \rightarrow ELLA')$   
 $\quad \square ella.get.easel \rightarrow (ella.put.easel \rightarrow ELLA'$   
 $\quad \quad \square ella.get.box \rightarrow ella.paint \rightarrow$   
 $\quad \quad \quad ella.put.easel \rightarrow ella.put.box \rightarrow ELLA')$

Will this get rid of the deadlock? Yes, but it introduces a "livelock"!

# Summary

- **Deadlock**
  - ...failure model using trace-refusal pairs
  - ...conditions for deadlock
  - ...breaking deadlock
- **Trace Refinement:** safety specification – (do nothing wrong)
- **Failure Refinement:** liveness specification – (do something right)

**Next Time:** Livelock Freedom – (actually progress a task)