



- Recall: we know how to realise



st. hosts can transmit TCP segments to each other.

- **Recall:** we know how to realise



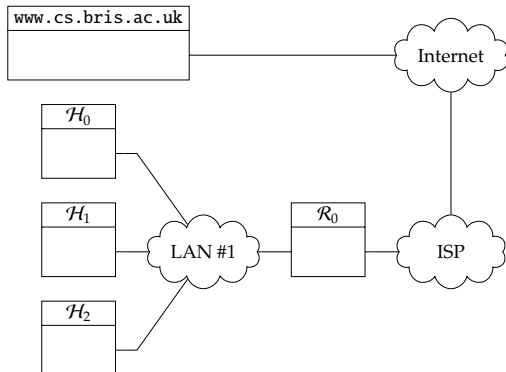
st. hosts can transmit TCP segments to each other ...

- ... *but*

1. how *could* both LANs use the same private IP address block, and
2. how does the client know the IP address of **www.cs.bris.ac.uk**?


Problem #1 \leadsto NAT (1)

- **Problem:** consider the following inter-network



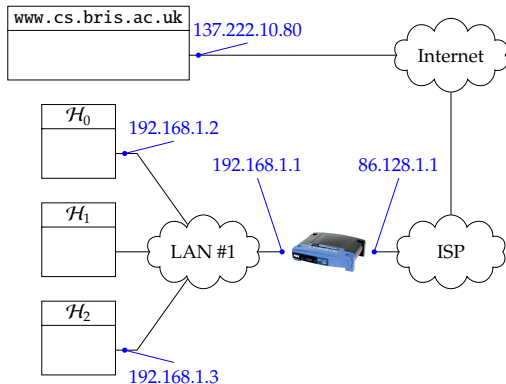
Problem #1 \leadsto NAT (1)

- **Problem:** consider the following inter-network



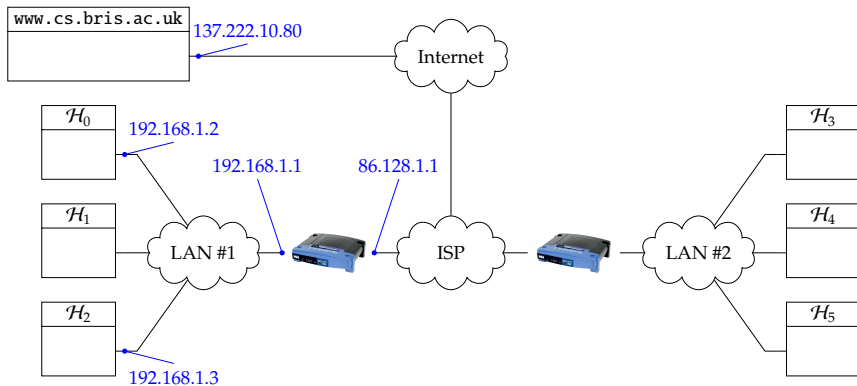
Problem #1 \leadsto NAT (1)

- **Problem:** consider the following inter-network



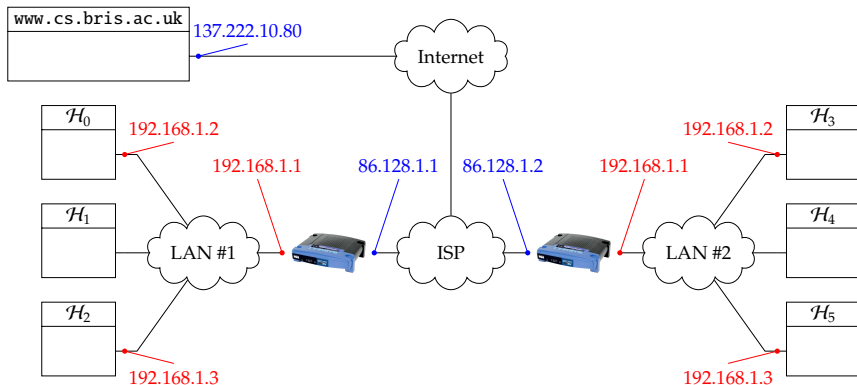
Problem #1 \leadsto NAT (1)

- **Problem:** consider the following inter-network



Problem #1 \leadsto NAT (1)

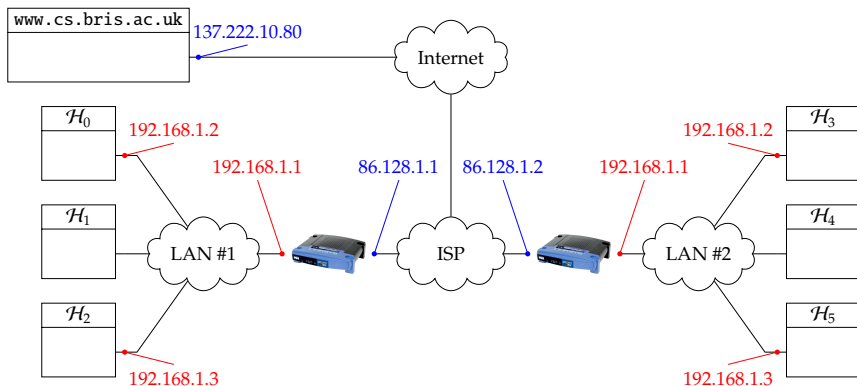
- **Problem:** consider the following inter-network



where LANs #1 and #2 (legitimately) use the private address block `192.168.0.0/16`, so the host IP addresses conflict.

Problem #1 \leadsto NAT (1)

- **Problem:** consider the following inter-network




where LANs #1 and #2 (legitimately) use the private address block `192.168.0.0/16`, so the host IP addresses conflict.

- (A) **solution:** Network Address Translation (NAT) [7].

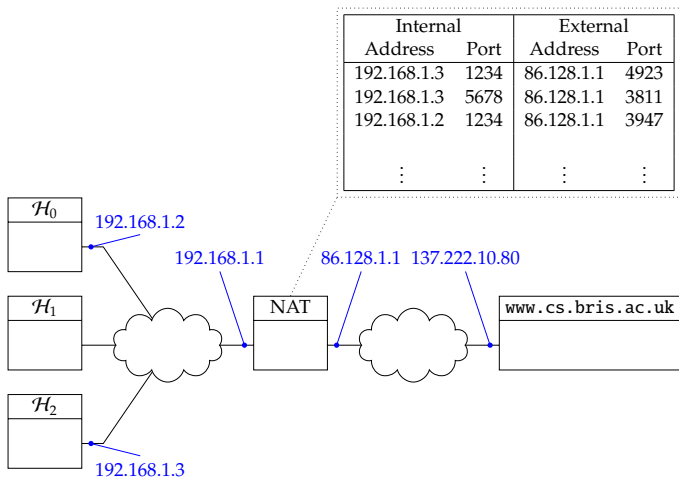
Problem #1 \leadsto NAT (2)

► Example:



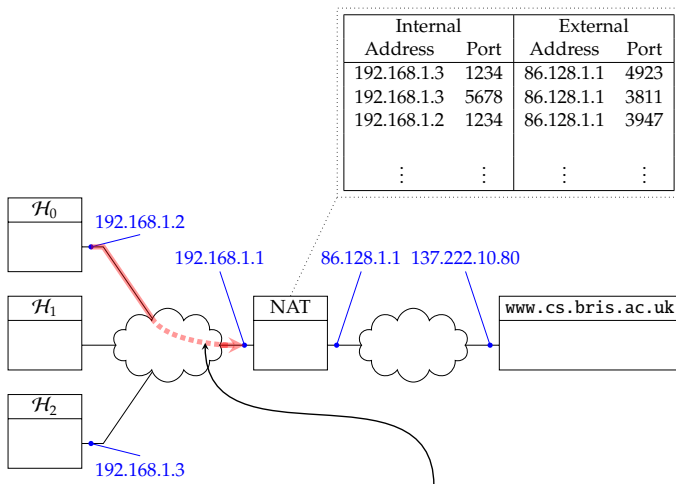
Problem #1 \leadsto NAT (2)

► Example:



Problem #1 \leadsto NAT (2)

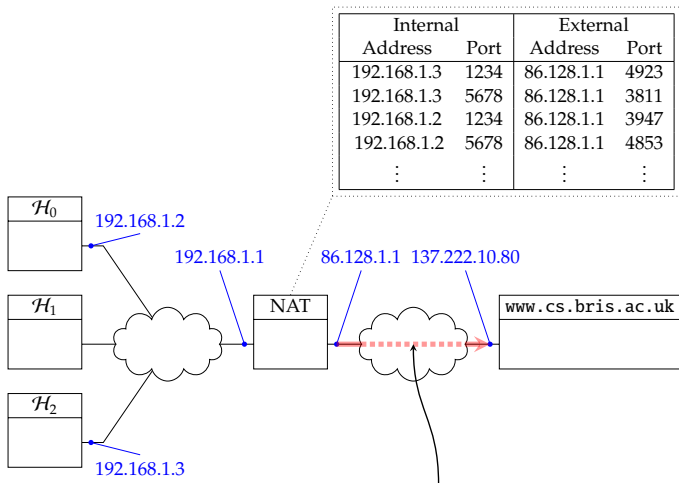
► Example:



$H_{IPv4}[\text{src} = 192.168.1.2, \text{dst} = 137.222.10.80] \parallel H_{TCP}[\text{syn} = \text{true}, \text{src port} = 5678, \text{dst port} = 80]$

Problem #1 \leadsto NAT (2)

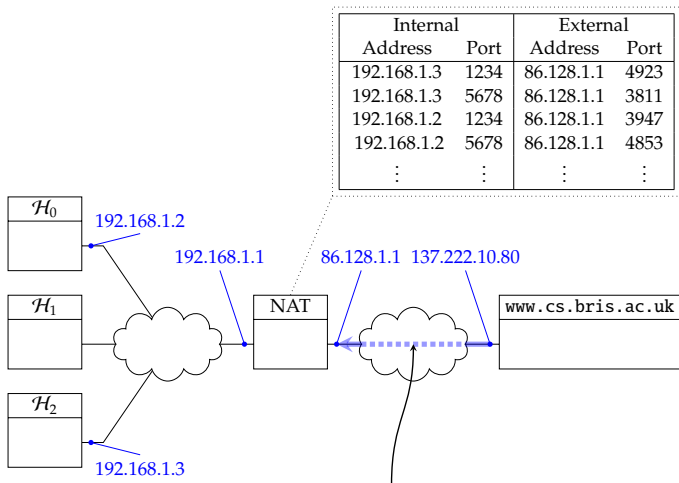
► Example:



$H_{IPv4}[\text{src} = 86.120.1.1, \text{dst} = 137.222.10.80] \parallel H_{TCP}[\text{syn} = \text{true}, \text{src port} = 4853, \text{dst port} = 80]$

Problem #1 \leadsto NAT (2)

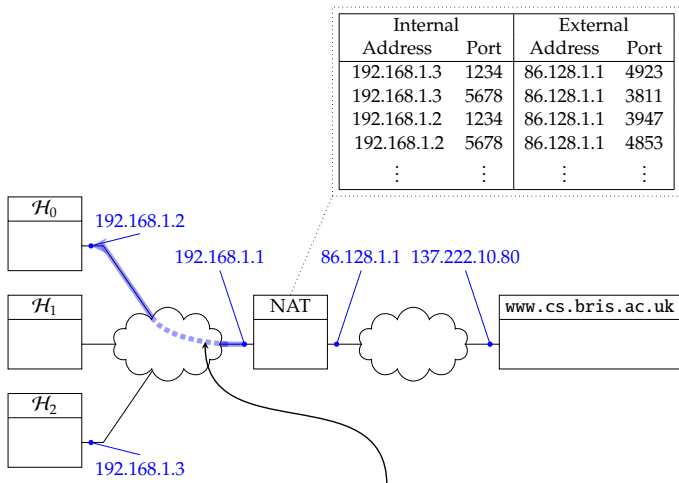
► Example:



$H_{IPv4}[\text{src} = 137.222.10.80, \text{dst} = 86.128.1.1] \parallel H_{TCP}[\text{syn} = \text{true}, \text{ack} = \text{true}, \text{src port} = 80, \text{dst port} = 4853]$

Problem #1 \leadsto NAT (2)

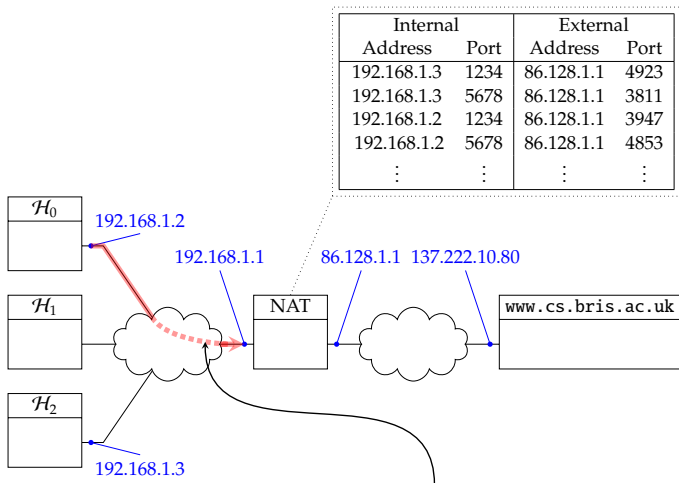
► Example:



$H_{IPv4}[\text{src} = 137.222.10.80, \text{dst} = 192.168.1.2] \parallel H_{TCP}[\text{syn} = \text{true}, \text{ack} = \text{true}, \text{src port} = 80, \text{dst port} = 5678]$

Problem #1 \leadsto NAT (2)

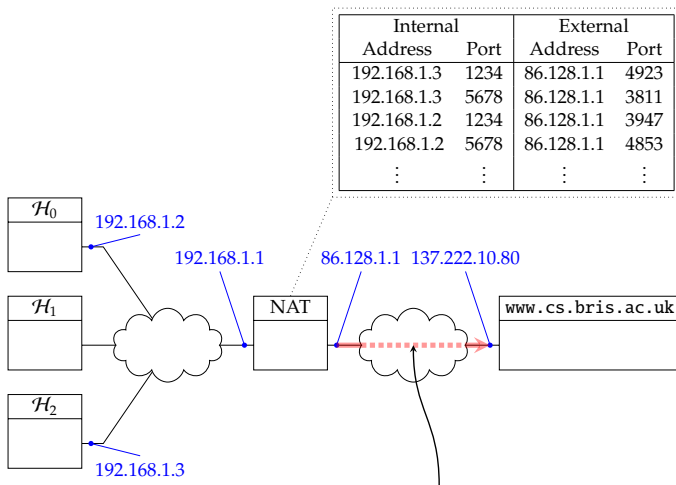
► Example:



$H_{IPv4}[\text{src} = 192.168.1.2, \text{dst} = 137.222.10.80] \parallel H_{TCP}[\text{ack} = \text{true}, \text{src port} = 5678, \text{dst port} = 80]$

Problem #1 \leadsto NAT (2)

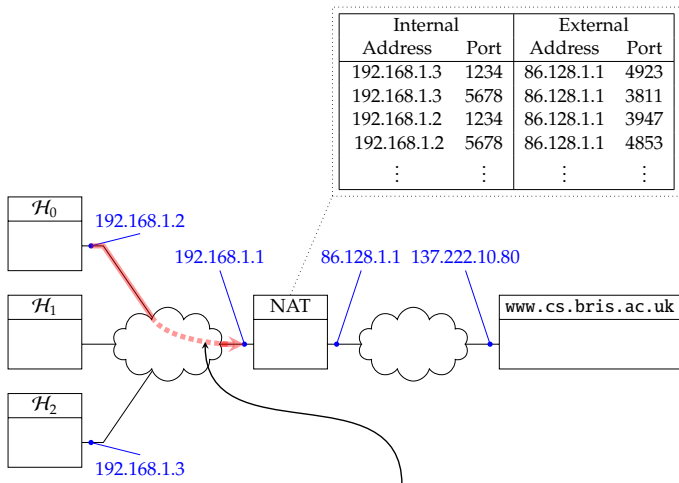
► Example:



$H_{IPv4}[\text{src} = 86.120.1.1, \text{dst} = 137.222.10.80] \parallel H_{TCP}[\text{ack} = \text{true}, \text{src port} = 4853, \text{dst port} = 80]$

Problem #1 \leadsto NAT (2)

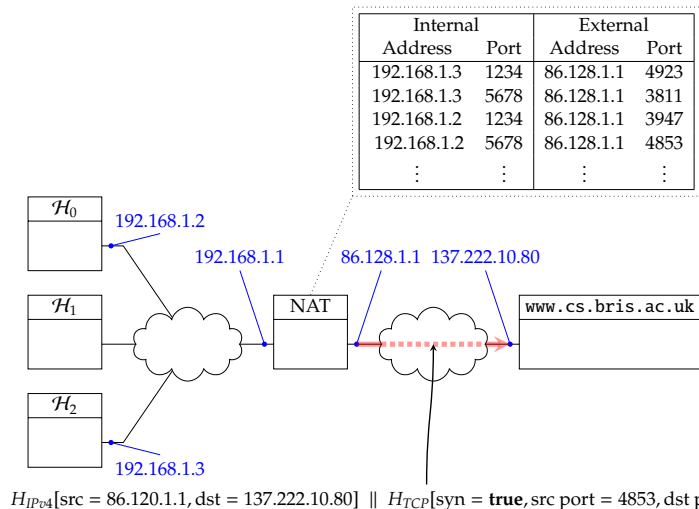
► Example:



$H_{IPv4}[\text{src} = 192.168.1.2, \text{dst} = 137.222.10.80] \parallel H_{TCP}[\text{syn} = \text{true}, \text{src port} = 5678, \text{dst port} = 80] \parallel \langle . \rangle$

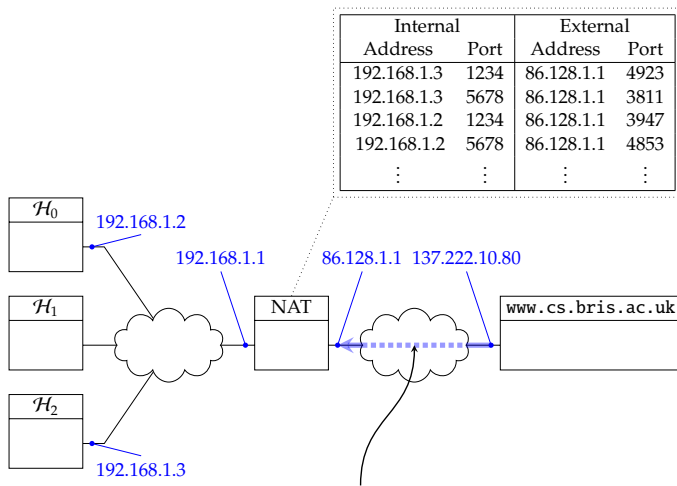
Problem #1 \leadsto NAT (2)

► Example:



Problem #1 \leadsto NAT (2)

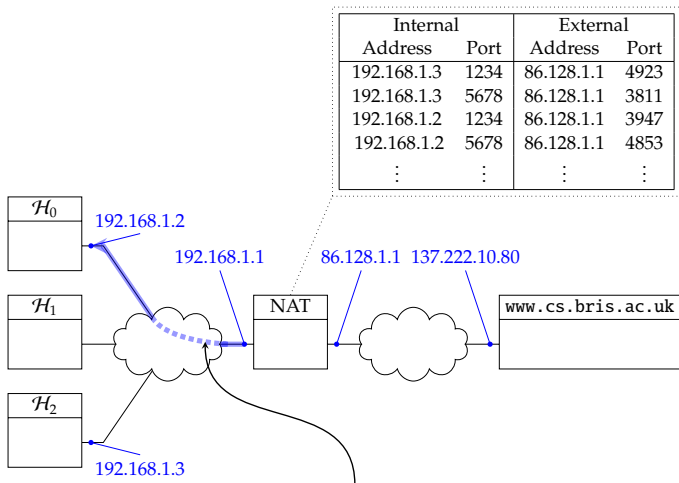
► Example:



$H_{IPv4}[\text{src} = 137.222.10.80, \text{dst} = 86.120.1.1] \parallel H_{TCP}[\text{syn} = \text{true}, \text{ack} = \text{true}, \text{src port} = 80, \text{dst port} = 4853] \parallel \langle . \rangle$

Problem #1 \leadsto NAT (2)

► Example:



$H_{IPv4}[\text{src} = 137.222.10.80, \text{dst} = 192.168.1.2] \parallel H_{TCP}[\text{syn} = \text{true}, \text{ack} = \text{true}, \text{src port} = 80, \text{dst port} = 5678] \parallel \langle . \rangle$

Problem #1 \leadsto NAT (3)

- ▶ NAT appliances are examples of **middlebox** [6]:
 - ▶ exist “in” network (like routers), but
 - ▶ operate in more than network layer (i.e., don’t *just* forward packets, like hosts), so
 - ▶ break layered model and connectivity, and can result in strange side-effects

hence, NAT specifically

- ▶ **Good:**
 - ▶ helps mitigate IPv4 address scarcity problem, and
 - ▶ can be easily, centrally deployed
- ▶ **Bad:**
 - ▶ is difficult to extend beyond TCP,
 - ▶ requires initial outgoing connection, and
 - ▶ potentially fails if application layer depends on and/or exposes (internal) IP address

plus, internal hosts are anonymised to some extent since their traffic flows are merged externally.


Problem #2 \leadsto DNS (1)

► Problem:

Definition (**name**, **address** and **resolution**)

- a (human-readable) **name** is used to *identify* the resource,
- an (machine-readable) **address** is used to *locate* the resource, and
- **resolution** (resp. **reverse resolution**) maps a name to an address (resp. address to a name).

so, we need some mechanism to perform



where

- the LHS is a **domain name**, and
- the RHS is an IP address.

Problem #2 \leadsto DNS (2)

► Solutions:

1. fully centralised: pre \sim 1984 via ARPANET `hosts.txt` [8], e.g.,

```
NET : 128.54.0.0 : UCSD :
```

and

```
HOST : 128.54.0.1 : SDCSVAX,UCSD : VAX-11/780 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,UDP :
```

2. partially decentralised: via

- in `/etc/networks`, e.g.,

```
loopback 127.0.0.0
```

and

- in `/etc/hosts`, e.g.,

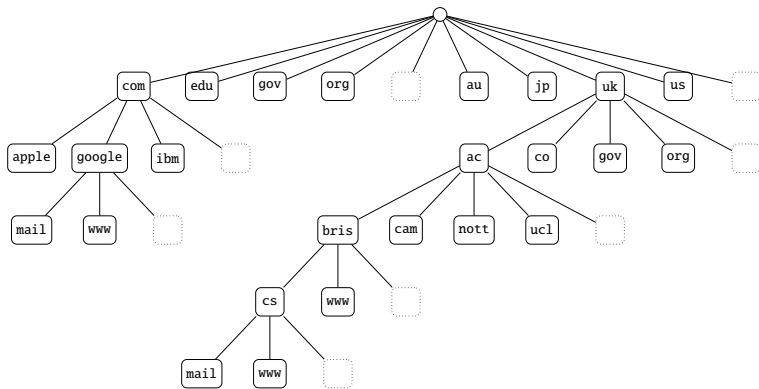
```
127.0.0.1 localhost
```

on a POSIX-style OS (e.g., Linux),

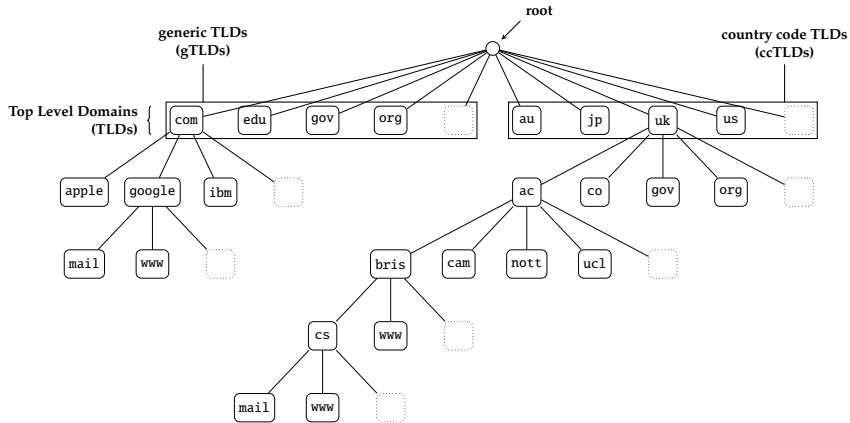
3. fully decentralised: post \sim 1984 via **Domain Name System (DNS)** [9, 10] which includes

- 3.1 a managed, hierarchical **name space**,
- 3.2 a protocol, used to communicate queries and responses, plus
- 3.3 an infrastructure, comprised of **name servers**, used to host the (distributed) database and respond to queries.

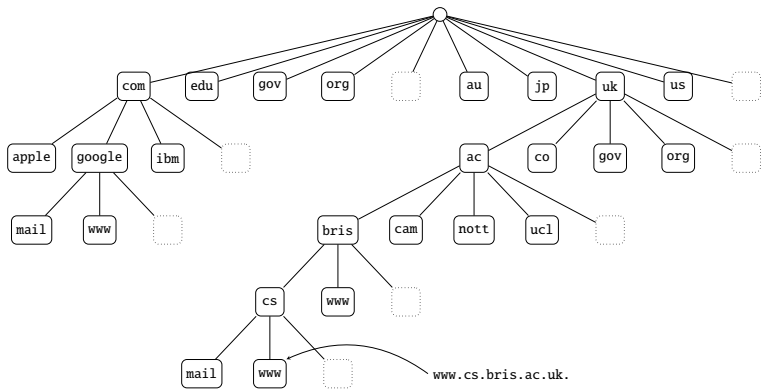
Problem #2 \leadsto DNS (3)



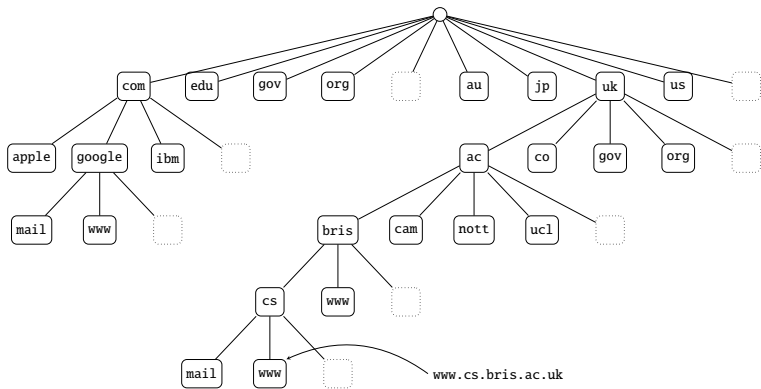
Problem #2 \leadsto DNS (3)



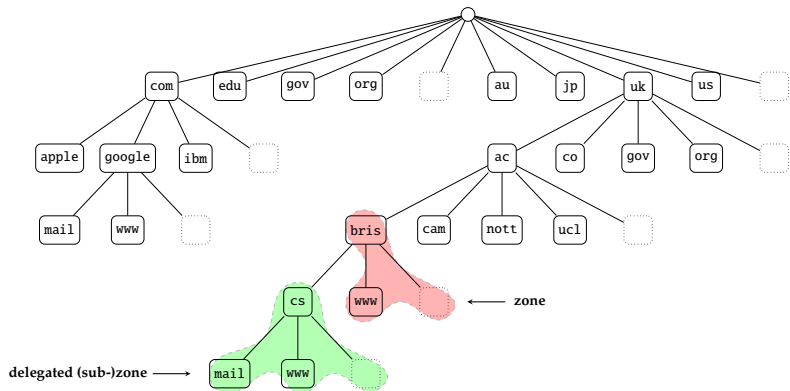
Problem #2 \leadsto DNS (3)



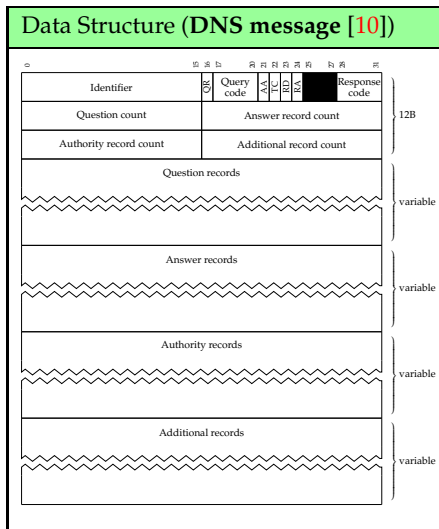
Problem #2 \leadsto DNS (3)



Problem #2 \leadsto DNS (3)



Problem #2 \leadsto DNS (4)




The data structure includes:

- ▶ A 16-bit identifier used to link query and response messages.
- ▶ 4-bit query and response codes.
- ▶ A set of flags, including
 - ▶ 1-bit **Query/Response (QR)** flag, to disambiguate queries from responses,
 - ▶ 1-bit **Authoritative Answer (AA)** flag, which marks responses that are authoritative,
 - ▶ 1-bit **Truncation Flag (TC)** which marks responses longer than the UDP 512B limit,
 - ▶ 1-bit **Recursion Desired (RD)** flag, used to request recursive resolution,
 - ▶ 1-bit **Recursion Available (RA)** flag, used to signal availability of recursive resolution.
- ▶ Some number of records (of each type).

Problem #2 \leadsto DNS (4)

Data Structure (DNS question record [10])

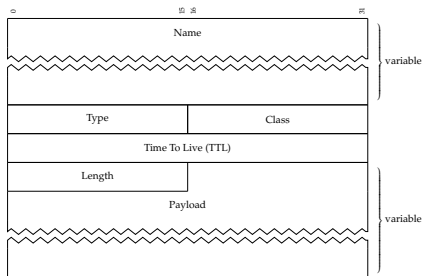


The data structure includes:

- ▶ A variable length, length-prefixed name string.
- ▶ A 16-bit record type.
- ▶ A 16-bit record class.

Problem #2 \leadsto DNS (4)

Data Structure (DNS resource record [10])




The data structure includes:

- ▶ A variable length, length-prefixed name string.
- ▶ A 16-bit record type.
- ▶ A 16-bit record class.
- ▶ A 32-bit **Time To Live (TTL)** field which controls cache retention.
- ▶ A 16-bit record length, and the record-specific payload, e.g.,
 1. a 32-bit IP address for A records,
 2. a variable length name for NS records, and
 3. a variable length name for CNAME records.

Problem #2 \leadsto DNS (5)

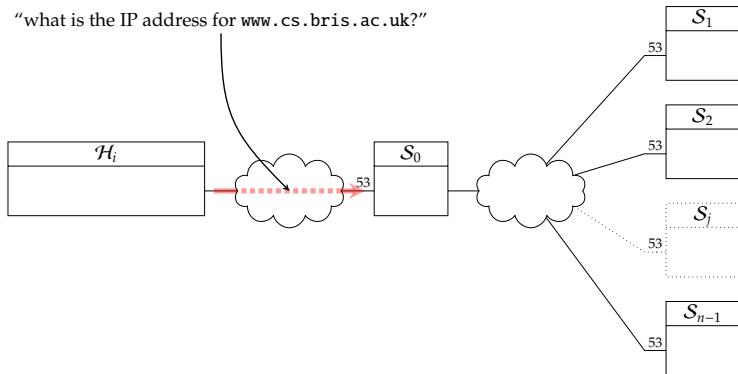
► Example:



Problem #2 \leadsto DNS (5)

► Example:


“what is the IP address for `www.cs.bris.ac.uk`?”



Problem #2 \leadsto DNS (5)

► Example:


"what is the IP address for `www.cs.bris.ac.uk`?"



Problem #2 \leadsto DNS (5)

► Example:


“ask ns1.nic.uk, whose IP address is 195.66.240.130”



Problem #2 \leadsto DNS (5)

► Example:


"what is the IP address for `www.cs.bris.ac.uk`?"



Problem #2 \leadsto DNS (5)

► Example:


“ask ns0.ja.net whose IP address is 193.63.94.20”



Problem #2 \leadsto DNS (5)


► Example:

"what is the IP address for `www.cs.bris.ac.uk`?"




Problem #2 \leadsto DNS (5)

► Example:




Problem #2 \leadsto DNS (5)

► Example:



Problem #2 \leadsto DNS (5)

► Example:



- ▶ Take away points:
 - ▶ The “glue” protocols outlined here solve issues that stem from practical deployment and use.
 - ▶ DNS in particular is interesting, since
 - ▶ it highlights the tricky compromise enforced by lack of middle-ground between UDP and TCP,
 - ▶ like DHCP *users* consider it part of the network but in reality it is an application level mechanism,
 - ▶ although not *too* old, it already highlights the need for flexibility to meet changing requirements and use-cases,
 - ▶ versus other glue protocols, there is a much stronger efficiency requirement, plus
 - ▶ tangential requirements (e.g., security [5], cf. DNSSEC [4] etc.) make it a *much* larger and more challenging problem.

References

- [1] Wikipedia: Domain name.
http://en.wikipedia.org/wiki/Domain_name.
- [2] Wikipedia: Domain Name System (DNS).
http://en.wikipedia.org/wiki/Domain_Name_System.
- [3] Wikipedia: Network address translation.
http://en.wikipedia.org/wiki/Network_address_translation.
- [4] R. Arends, R. Austein, M. Larson, D. Massey, and S. Rose.
[DNS security introduction and requirements](#).
Internet Engineering Task Force (IETF) Request for Comments (RFC) 4033, 2005.
<http://tools.ietf.org/html/rfc4033>.
- [5] D. Atkins and R. Austein.
[Threat analysis of the Domain Name System \(DNS\)](#).
Internet Engineering Task Force (IETF) Request for Comments (RFC) 3833, 2004.
<http://tools.ietf.org/html/rfc3833>.
- [6] B. Carpenter and S. Brim.
[Middleboxes: Taxonomy and issues](#).
Internet Engineering Task Force (IETF) Request for Comments (RFC) 3234, 2002.
<http://tools.ietf.org/html/rfc3234>.

References

- [7] K. Egevang and P. Francis.
[The IP Network Address Translator \(NAT\).](#)
Internet Engineering Task Force (IETF) Request for Comments (RFC) 1631, 1994.
<http://tools.ietf.org/html/rfc1631>.
- [8] E. Feinler, K. Harrenstien, Z.-S. Su, and V. White.
[DoD Internet host table specification.](#)
Internet Engineering Task Force (IETF) Request for Comments (RFC) 810, 1982.
<http://tools.ietf.org/html/rfc810>.
- [9] P. Mockapetris.
[Domain names - concepts and facilities.](#)
Internet Engineering Task Force (IETF) Request for Comments (RFC) 882, 1983.
<http://tools.ietf.org/html/rfc882>.
- [10] P. Mockapetris.
[Domain names - implementation and specification.](#)
Internet Engineering Task Force (IETF) Request for Comments (RFC) 883, 1983.
<http://tools.ietf.org/html/rfc883>.