

Concurrent Computing (Computer Networks)

Daniel Page

Department of Computer Science,
University Of Bristol,
Merchant Venturers Building,
Woodland Road,
Bristol, BS8 1UB. UK.
<Daniel.Page@bristol.ac.uk>


March 14, 2016

Keep in mind there are *two* PDFs available (of which this is the latter):

1. a PDF of examinable material used as lecture slides, and
2. a PDF of non-examinable, extra material:
 - ▶ the associated notes page may be pre-populated with extra, written explanation of material covered in lecture(s), plus
 - ▶ anything with a “grey’ed out” header/footer represents extra material which is useful and/or interesting but out of scope (and hence not covered).

Notes:


Notes:



- **Goal:** *finally* investigate the **application layer**, e.g.,
 - the (mainly OS-based) network stack implementation,
 - the interface between application and network stack, i.e.,
 1. a **raw socket**, or
 2. the **POSIX sockets API**,
- and
- examples of how *you* can use all this!

Notes:

- Using a raw socket is basically a way to bypass layers of the network stack, and use lower layers directly instead. `traceroute` may need to directly manipulate IP header fields, for example, rather than have the IPv4 layer act for it.
- The POSIX sockets API is effectively a standard version of the original TCP/IP implementation, namely **Berkeley sockets** (or **BSD sockets**) which first appeared in 1983. As a result, there's a tendency to think it as being "UNIX only". In reality, however, almost every OS uses roughly the same interface; **Winsock**, the Windows socket API, roughly does so for example.



- $F = H_{802.11} \parallel H_{IPv4} \parallel H_{TCP} \parallel \langle 0, 1, 1, 0, 1, 0, 0, 1 \rangle \parallel T_{802.11}$
Goal: *finally* investigate the **application layer**, e.g.,
 - the (mainly OS-based) network stack implementation,
 - the interface between application and network stack, i.e.,
 1. a **raw socket**, or
 2. the **POSIX sockets API**,
- and
- examples of how *you* can use all this!

Notes:

- Using a raw socket is basically a way to bypass layers of the network stack, and use lower layers directly instead. `traceroute` may need to directly manipulate IP header fields, for example, rather than have the IPv4 layer act for it.
- The POSIX sockets API is effectively a standard version of the original TCP/IP implementation, namely **Berkeley sockets** (or **BSD sockets**) which first appeared in 1983. As a result, there's a tendency to think it as being "UNIX only". In reality, however, almost every OS uses roughly the same interface; **Winsock**, the Windows socket API, roughly does so for example.

POSIX sockets API (1) – The Interface

	Function	Description	Blocking?
UDP {	socket	Form the data structure used to describe communication end-point	×
	bind	Associate socket data structure with (local) address	×
	close	Close socket and stop using it	×
	shutdown	Close socket and stop using it, with control over how	×
	getsockopt	Get or set options for a socket, i.e., control how it functions	×
	setsockopt		
	sendto	Transmit a datagram to (remote) address	✓
	recvfrom	Receive a datagram from (remote) address	✓
	listen	Mark socket as passive, i.e., for incoming connections	×
	accept	Wait for a connection to be established	✓
TCP {	connect	Actively establish a connection with (remote) address	✓
	send	Transmit a segment via connection	✓
	recv	Receive a segment via connection	✓
	select	Wait for activity that would allow non-blocking access	✓
	poll		

Notes:

- The Linux implementation of the socket API is well documented: see, for example,

`man -s 2 bind`

and so on. Note that some of the functions listed aren't strictly part of the (standard) API: rather, they represent "helper" functions which make using it easier.

- You may see `gethostbyname` or `gethostbyaddr` used in example code; these are deprecated analogues of `getnameinfo` and `getaddrinfo`, which are now preferred.
- The fact *some* functions are blocking versus non-blocking means it is important not to misunderstand their name: `recvfrom`, for example, should be read as "register to receive data some time in the future" rather than "receive data *now*".
- Strictly speaking use of `connect` within the context of UDP can make sense: rather than establish a connection, it simply specifies the default address for subsequent `sendto` and `recvfrom` invocations.
- You can think about the socket API via analogy with making telephone calls:
 - `socket` is like obtaining a telephone to use,
 - `bind` is like making your telephone number available st. incoming callers know it,
 - `listen` is like plugging in the telephone, and turning on ringer so you know when an incoming call occurs,
 - `connect` is like placing a call using the destinations number,
 - `accept` is like picking up the receiver to answer incoming call, with caller ID st. the destination knows the source number,
 - `close` is like putting down receiver to terminate an active call.

In the same context, you could think about DNS as being analogous to the telephone book: it allows one to look a number based on a description.

POSIX sockets API (1) – The Interface

Function	Description
getnameinfo	Convert an internal, machine-readable data structure into a host name
getaddrinfo	Convert a host name into an internal, machine-readable data structure
inet_aton	Convert a dotted-decimal address into a binary, machine-readable address
inet_ntoa	Convert a binary, machine-readable address into a dotted-decimal address
htons/htons	Convert a 16/32-bit host order integer into network order
ntohs/ntohs	Convert a 16/32-bit network order integer into host order

Notes:

- The Linux implementation of the socket API is well documented: see, for example,

`man -s 2 bind`

and so on. Note that some of the functions listed aren't strictly part of the (standard) API: rather, they represent "helper" functions which make using it easier.

- You may see `gethostbyname` or `gethostbyaddr` used in example code; these are deprecated analogues of `getnameinfo` and `getaddrinfo`, which are now preferred.
- The fact *some* functions are blocking versus non-blocking means it is important not to misunderstand their name: `recvfrom`, for example, should be read as "register to receive data some time in the future" rather than "receive data *now*".
- Strictly speaking use of `connect` within the context of UDP can make sense: rather than establish a connection, it simply specifies the default address for subsequent `sendto` and `recvfrom` invocations.
- You can think about the socket API via analogy with making telephone calls:
 - `socket` is like obtaining a telephone to use,
 - `bind` is like making your telephone number available st. incoming callers know it,
 - `listen` is like plugging in the telephone, and turning on ringer so you know when an incoming call occurs,
 - `connect` is like placing a call using the destinations number,
 - `accept` is like picking up the receiver to answer incoming call, with caller ID st. the destination knows the source number,
 - `close` is like putting down receiver to terminate an active call.

In the same context, you could think about DNS as being analogous to the telephone book: it allows one to look a number based on a description.

POSIX sockets API (2) – An Implementation (in Linux)

► Some (rough) design goals might include

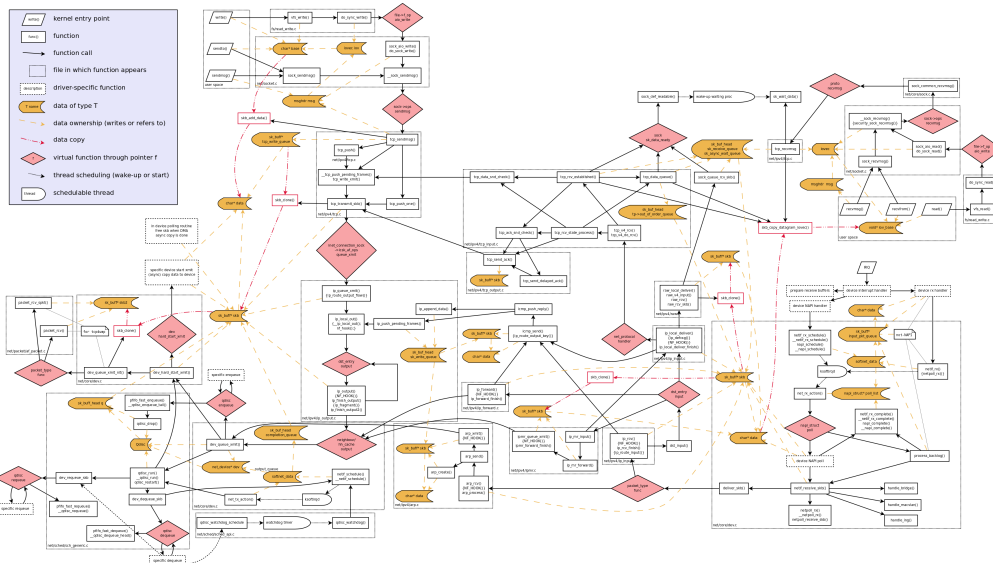
1. offer POSIX-compliant interface,
2. offer RFC-compliant implementation,
3. maximise efficiency (e.g., low-latency, effective use of bandwidth),
4. maximise flexibility (e.g., general- not special-purpose),
5. allow configurability,
6. ...

which lead to some underlying golden rules, e.g.,

- make use of all possible hardware support,
- make use of effective data structures,
- minimise copying,
- optimise for common-case,
- ensure correctness for corner-cases,
- ...

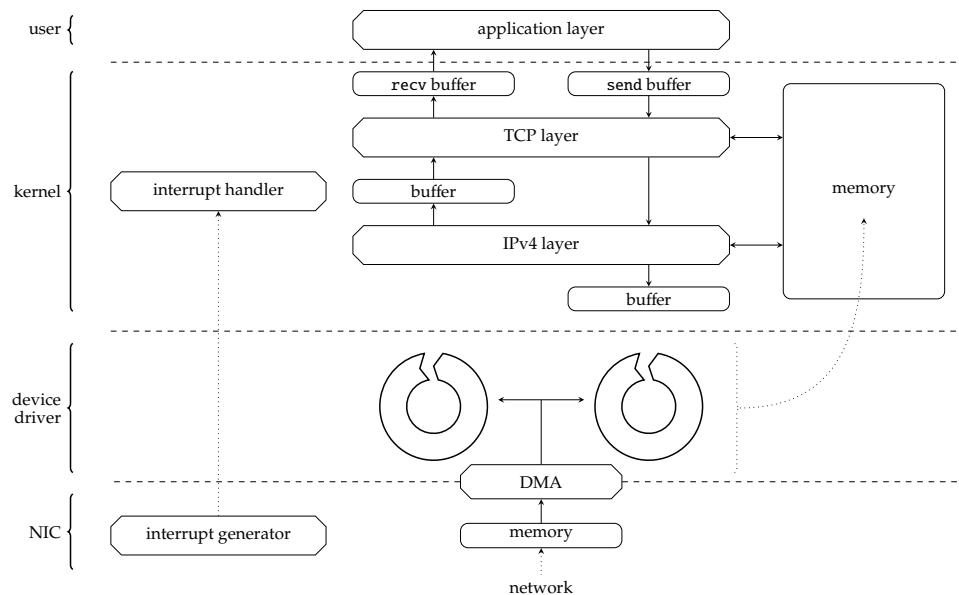
Notes:

POSIX sockets API (3) – An Implementation (in Linux)



Notes:

POSIX sockets API (4) – An Implementation (in Linux)



Notes:

- A challenge wrt. describing “the” Linux implementation is that it is continually evolving! On one hand, this is positive: it continues to improve so as to meet new demands and use-case. On the other hand, the moving target makes it harder to offer an accurate, general overview. Various documents try to do so, however, at various levels of detail. A good example is

<http://datatag.web.cern.ch/datatag/papers/tr-datatag-2004-1.pdf>

which stems from an EU-funded research project.

POSIX sockets API (5) – An Implementation (in Linux)

- **Fact:** ports are basically buffers within network stack.
- **Implication #1:**
 - packets and segments might be received out-of-order, *but*
 - buffering enforces in-order delivery to the application.

Notes:

- The “some sort of time-out timer” might seem a little vague, but we’ve already encountered a concrete example: Nagle’s algorithm can be seen as triggering transmission in this way.

POSIX sockets API (5) – An Implementation (in Linux)

- ▶ **Fact:** ports are basically buffers within network stack.
- ▶ **Implication #2:** send and transmission are decoupled ...
- ▶ ... transmission *could* occur
 1. when a complete segment is accumulated, or
 2. when transmission is forced, e.g., via
 - ▶ use of the PSH flag, or
 - ▶ some sort of time-out timer

so basically needs to realise a trade-off:

- ▶ less efficient wrt. latency (wait more time) but more efficient wrt. bandwidth (transmit complete segments more often), or
- ▶ more efficient wrt. latency (wait less time) but less efficient wrt. bandwidth (transmit complete segments less often).

Notes:

- The “some sort of time-out timer” might seem a little vague, but we’ve already encountered a concrete example: Nagle’s algorithm can be seen as triggering transmission in this way.

POSIX sockets API (5) – An Implementation (in Linux)

- ▶ **Fact:** ports are basically buffers within network stack.
- ▶ **Implication #3:** send and recv are decoupled ...
- ▶ ... any one of

send () ~~~~~> recv (. .) = 


is possible.

Notes:

- The “some sort of time-out timer” might seem a little vague, but we’ve already encountered a concrete example: Nagle’s algorithm can be seen as triggering transmission in this way.

POSIX sockets API (5) – An Implementation (in Linux)

- **Fact:** ports are basically buffers within network stack.
- **Implication #3:** send and recv are decoupled ...
- ... any one of




is possible.

Notes:

- The “some sort of time-out timer” might seem a little vague, but we’ve already encountered a concrete example: Nagle’s algorithm can be seen as triggering transmission in this way.

POSIX sockets API (5) – An Implementation (in Linux)

- **Fact:** ports are basically buffers within network stack.
- **Implication #3:** send and recv are decoupled ...
- ... any one of




is possible.

Notes:


- The “some sort of time-out timer” might seem a little vague, but we’ve already encountered a concrete example: Nagle’s algorithm can be seen as triggering transmission in this way.

Using POSIX sockets (1) – UDP



Notes:

Using POSIX sockets (2) – TCP



Notes:

Using POSIX sockets (3) – An “echo uppercase” TCP client/server

Listing (C)

```
1 #include <sys/socket.h>
2 #include <arpa/inet.h>
3 #include <unistd.h>
4
5 void handle( int cs ) {
6     char t[ 1024 ];
7
8     while( true ) {
9         // terminal -> t
10        fgets( t, 1024, stdin );
11        // server <- t
12        send( cs, t, strlen( t ), 0 );
13        // server -> t'
14        t[ recv( cs, t, 1023, 0 ) ] = '\0';
15        // terminal <- t'
16        fputs( t, stdout );
17    }
18
19    // close connection
20    close( cs );
21 }
```

Listing (C)

```
1 int main( int argc, char* argv[] ) {
2     struct sockaddr_in sa; socklen_t sl = sizeof( sa );
3     struct sockaddr_in ca; socklen_t cl = sizeof( ca );
4
5     memset( &sa, 0, sl );
6
7     sa.sin_family = AF_INET;
8     sa.sin_addr.s_addr = inet_addr( argv[ 1 ] );
9     sa.sin_port = htons( atoi( argv[ 2 ] ) );
10
11    // open socket
12    int cs = socket( AF_INET, SOCK_STREAM, 0 );
13    // open connection
14    connect( cs, ( struct sockaddr* )( &sa ), sl );
15    // handle connection
16    handle( cs );
17
18    return 0;
19 }
```

Notes:

- This code has a number of caveats and limitations; some compromises have been made to fit it on the slide(s), so try to treat it as an example rather than a reference.
 - There is no error checking *at all*, which is clearly not ideal! More or less *every* function, bind for example, returns an error code that *should* be checked carefully; the global variable `errno` captures more detailed information about any error that occurs.
 - Both the client and server loop indefinitely, i.e., neither allow a way for the user to terminate the connection without forcibly terminating the associated process: this might not be a great design strategy in general.
 - It should *really* use a type cast around `&ca`, i.e., use `(struct sockaddr*)(&ca)`, when calling `accept`. This highlights a subtle design choice, stemming in part from how C structures work. If you look at the definitions of `sockaddr` and `sockaddr_in`, it's clear you can cast the latter into the former: they share the same first field, so in a sense `sockaddr` is just a more general version of `sockaddr_in`.
 - An important issue with the code relates to `send`: it doesn't in fact *guarantee* to transmit all the data you give to it as input. As a result, it is common to use an auxiliary function of the form

```
void sendall( int s, uint8_t* x, int n ) {
    int l = n;

    while( n > 0 ) {
        if( ( n -= send( s, x + l - n, n, 0 ) ) < 1 ) {
            break;
        }
    }
}
```

to give a more useful outcome.

- The IP address and port are specified via the command-line, but there are alternatives. For example, the following

IP address	Port	Semantics
hard-coded, <code>INADDR_ANY</code>	hard-coded, zero	kernel chooses IP address, kernel chooses port
hard-coded, <code>INADDR_ANY</code>	user-supplied, non-zero	kernel chooses IP address, user specifies port
user-supplied	hard-coded zero	user specifies IP address, kernel chooses port
user-supplied	user-supplied, non-zero	user specifies IP address, user specifies port

shows one can defer to the kernel to some extent. Even then, however, IP addresses are demanded over domain names: it is of course possible to perform DNS resolution via the (local) resolver with appropriate additions.

Using POSIX sockets (3) – An “echo uppercase” TCP client/server

Listing (C)

```
1 #include <sys/socket.h>
2 #include <arpa/inet.h>
3 #include <unistd.h>
4
5 void handle( int cs ) {
6     char t[ 1024 ];
7
8     while( true ) {
9         // client -> t
10        t[ recv( cs, t, 1023, 0 ) ] = '\0';
11        // t' = toupper( t )
12        for( int i = 0; i < strlen( t ); i++ ) {
13            t[ i ] = toupper( t[ i ] );
14        }
15        // client <- t'
16        send( cs, t, strlen( t ), 0 );
17    }
18
19    // close connection
20    close( cs );
21 }
```

Listing (C)

```
1 int main( int argc, char* argv[] ) {
2     struct sockaddr_in sa; socklen_t sl = sizeof( sa );
3     struct sockaddr_in ca; socklen_t cl = sizeof( ca );
4
5     memset( &sa, 0, sl );
6
7     sa.sin_family = AF_INET;
8     sa.sin_addr.s_addr = inet_addr( argv[ 1 ] );
9     sa.sin_port = htons( atoi( argv[ 2 ] ) );
10
11    // open socket
12    int ss = socket( AF_INET, SOCK_STREAM, IPPROTO_IP );
13    // bind socket
14    bind( ss, ( struct sockaddr* )( &sa ), sl );
15    // listen for connections
16    listen( ss, 10 );
17
18    while( true ) {
19
20
21        // open connection
22        int cs = accept( ss, &ca, &cl );
23        // handle connection
24        handle( cs );
25    }
26
27    // close socket
28    close( ss );
29
30    return 0;
31 }
```

Notes:

- This code has a number of caveats and limitations; some compromises have been made to fit it on the slide(s), so try to treat it as an example rather than a reference.
 - There is no error checking *at all*, which is clearly not ideal! More or less *every* function, bind for example, returns an error code that *should* be checked carefully; the global variable `errno` captures more detailed information about any error that occurs.
 - Both the client and server loop indefinitely, i.e., neither allow a way for the user to terminate the connection without forcibly terminating the associated process: this might not be a great design strategy in general.
 - It should *really* use a type cast around `&ca`, i.e., use `(struct sockaddr*)(&ca)`, when calling `accept`. This highlights a subtle design choice, stemming in part from how C structures work. If you look at the definitions of `sockaddr` and `sockaddr_in`, it's clear you can cast the latter into the former: they share the same first field, so in a sense `sockaddr` is just a more general version of `sockaddr_in`.
 - An important issue with the code relates to `send`: it doesn't in fact *guarantee* to transmit all the data you give to it as input. As a result, it is common to use an auxiliary function of the form

```
void sendall( int s, uint8_t* x, int n ) {
    int l = n;

    while( n > 0 ) {
        if( ( n -= send( s, x + l - n, n, 0 ) ) < 1 ) {
            break;
        }
    }
}
```

to give a more useful outcome.

- The IP address and port are specified via the command-line, but there are alternatives. For example, the following

IP address	Port	Semantics
hard-coded, <code>INADDR_ANY</code>	hard-coded, zero	kernel chooses IP address, kernel chooses port
hard-coded, <code>INADDR_ANY</code>	user-supplied, non-zero	kernel chooses IP address, user specifies port
user-supplied	hard-coded zero	user specifies IP address, kernel chooses port
user-supplied	user-supplied, non-zero	user specifies IP address, user specifies port

shows one can defer to the kernel to some extent. Even then, however, IP addresses are demanded over domain names: it is of course possible to perform DNS resolution via the (local) resolver with appropriate additions.

Using POSIX sockets (3) – An “echo uppercase” TCP client/server

Listing (C)

```
1 #include <sys/socket.h>
2 #include <arpa/inet.h>
3 #include <unistd.h>
4
5 void* handle( void* __cs ) {
6     char t[ 1024 ];
7
8     int cs = *( int* )( __cs );
9
10    while( true ) {
11        // client -> t
12        t[ recv( cs, t, 1023, 0 ) ] = '\0';
13        // t' = toupper( t )
14        for( int i = 0; i < strlen( t ); i++ ) {
15            t[ i ] = toupper( t[ i ] );
16        }
17        // client <- t'
18        send( cs, t, strlen( t ), 0 );
19    }
20
21    // close connection
22    close( cs );
23
24    return NULL;
25 }
```

Listing (C)

```
1 int main( int argc, char* argv[ ] ) {
2     struct sockaddr_in sa; socklen_t sl = sizeof( sa );
3     struct sockaddr_in ca; socklen_t cl = sizeof( ca );
4
5     memset( &sa, 0, sl );
6
7     sa.sin_family = AF_INET;
8     sa.sin_addr.s_addr = inet_addr( argv[ 1 ] );
9     sa.sin_port = htons( atoi( argv[ 2 ] ) );
10
11    // open socket
12    int ss = socket( AF_INET, SOCK_STREAM, IPPROTO_IP );
13    // bind socket
14    bind( ss, ( struct sockaddr* )( &sa ), sl );
15    // listen for connections
16    listen( ss, 10 );
17
18    while( true ) {
19        pthread_t id;
20
21        // open connection
22        int cs = accept( ss, &ca, &cl );
23        // handle connection
24        pthread_create( &id, NULL, &handle, &cs );
25    }
26
27    // close socket
28    close( ss );
29
30    return 0;
31 }
```

Notes:

- This code has a number of caveats and limitations; some compromises have been made to fit it on the slide(s), so try to treat it as an example rather than a reference.
 - There is no error checking *at all*, which is clearly not ideal! More or less *every* function, bind for example, returns an error code that *should* be checked carefully; the global variable `errno` captures more detailed information about any error that occurs.
 - Both the client and server loop indefinitely, i.e., neither allow a way for the user to terminate the connection without forcibly terminating the associated process: this might not be a great design strategy in general.
 - It should *really* use a type cast around `&ca`, i.e., use `(struct sockaddr*)(&ca)`, when calling `accept`. This highlights a subtle design choice, stemming in part from how C structures work. If you look at the definitions of `sockaddr` and `sockaddr_in`, it's clear you can cast the latter into the former: they share the same first field, so in a sense `sockaddr` is just a more general version of `sockaddr_in`.
 - An important issue with the code relates to `send`: it doesn't in fact *guarantee* to transmit all the data you give to it as input. As a result, it is common to use an auxiliary function of the form

```
void sendall( int s, uint8_t* x, int n ) {
    int l = n;

    while( n > 0 ) {
        if( ( n -= send( s, x + l - n, n, 0 ) ) < 1 ) {
            break;
        }
    }
}
```

to give a more useful outcome.

- The IP address and port are specified via the command-line, but there are alternatives. For example, the following

IP address	Port	Semantics
hard-coded, <code>INADDR_ANY</code>	hard-coded, zero	kernel chooses IP address, kernel chooses port
hard-coded, <code>INADDR_ANY</code>	user-supplied, non-zero	kernel chooses IP address, user specifies port
user-supplied	hard-coded zero	user specifies IP address, kernel chooses port
user-supplied	user-supplied, non-zero	user specifies IP address, user specifies port

shows one can defer to the kernel to some extent. Even then, however, IP addresses are demanded over domain names: it is of course possible to perform DNS resolution via the (local) resolver with appropriate additions.

Notes:

Using POSIX sockets (4) – An “echo uppercase” TCP client/server

Listing (Python)

```
1 import socket, sys
2
3 def handle( cs ) :
4     while( True ) :
5         # terminal -> t
6         t = sys.stdin.readline()
7         # server <- t
8         cs.send( t )
9         # server -> t'
10        t = cs.recv( 1024 )
11        # terminal <- t'
12        sys.stdout.write( t )
13
14    if( __name__ == "__main__" ) :
15        host = sys.argv[ 1 ]
16        port = int( sys.argv[ 2 ] )
17
18        # open socket
19        cs = socket.socket( socket.AF_INET, socket.SOCK_STREAM )
20        # open connection
21        cs.connect( ( host, port ) )
22        # handle connection
23        handle( cs )
24        # close connection
25        cs.close()
```

Listing (Python)

```
1 import socket, sys
2
3 def handle( cs, ca ) :
4     while( True ) :
5         # client -> t
6         t = cs.recv( 1024 )
7         # t' = toupper( t )
8         t = t.upper()
9         # client <- t'
10        cs.send( t )
11
12    # close connection
13    cs.close()
14
15 if( __name__ == "__main__" ) :
16     host = sys.argv[ 1 ]
17     port = int( sys.argv[ 2 ] )
18
19    # open socket
20    ss = socket.socket( socket.AF_INET, socket.SOCK_STREAM )
21    # bind socket
22    ss.bind( ( host, port ) )
23    # listen for connections
24    ss.listen( 10 )
25
26    while( True ) :
27        # open connection
28        ( cs, ca ) = ss.accept()
29        # handle connection
30        handle( cs, ca )
31
32    # close socket
33    ss.close()
```

Notes:

Listing (Python)

```
1 import socket, sys, thread
2
3 def handle( cs, ca ) :
4     while( True ) :
5         # client -> t
6         t = cs.recv( 1024 )
7         # t' = toupper( t )
8         t = t.upper()
9         # client <- t'
10        cs.send( t )
11
12    # close connection
13    cs.close()
14
15 if( __name__ == "__main__" ) :
16     host = sys.argv[ 1 ]
17     port = int( sys.argv[ 2 ] )
18
19    # open socket
20    ss = socket.socket( socket.AF_INET, socket.SOCK_STREAM )
21    # bind socket
22    ss.bind( ( host, port ) )
23    # listen for connections
24    ss.listen( 10 )
25
26    while( True ) :
27        # open connection
28        ( cs, ca ) = ss.accept()
29        # handle connection
30        thread.start_new_thread( handle, ( cs, ca ) )
31
32    # close socket
33    ss.close()
```

Notes:

Listing (Python)

```

1 import SocketServer, sys
2
3 class server( SocketServer.BaseRequestHandler ):
4     def handle( self ) :
5         while( True ) :
6             # client -> t
7             t = self.request.recv( 1024 )
8             # t' = toupper( t )
9             t = t.upper()
10            # client <- t'
11            self.request.sendall( t )
12
13 if( __name__ == "__main__" ) :
14     host = sys.argv[ 1 ]
15     port = int( sys.argv[ 2 ] )
16
17     SocketServer.TCPServer( ( host, port ), server ).serve_forever()

```

Notes:

Conclusions

► Take away points:

- Ultimately, the POSIX sockets API is an abstraction of the network ...
- ... even so, it's hard to argue you can totally avoid having to understand the underlying technology.
- As with any design, it has **good** and **bad** features: for example,
 - it offers a uniform interface to analogous concepts (cf. **domain sockets**, for IPC),
 - it allows special-case implementation choices such as use of **TCP offload**,
 - the abstraction offered is still low-level so can be hard to use (directly),
 - numerous requirements have changed over time (e.g., network vs. host order, new protocols, new use-cases), but the API hasn't,
 - ...

Notes:

References

- [1] Wikipedia: Berkeley sockets.
http://en.wikipedia.org/wiki/Berkeley_sockets.
- [2] Wikipedia: Raw socket.
http://en.wikipedia.org/wiki/Raw_socket.
- [3] Wikipedia: TCP offload engine.
http://en.wikipedia.org/wiki/TCP_offload_engine.
- [4] Wikipedia: UNIX domain socket.
http://en.wikipedia.org/wiki/Unix_domain_socket.
- [5] Wikipedia: Winsock.
<http://en.wikipedia.org/wiki/Winsock>.
- [6] Standard for information technology - portable operating system interface (POSIX).
[Institute of Electrical and Electronics Engineers \(IEEE\) 1003.1, 2008.](http://standards.ieee.org/findstds/standard/1003.1-2008.html)
<http://standards.ieee.org/findstds/standard/1003.1-2008.html>.
- [7] W.R. Stevens, B. Fenner, and A.M. Rudoff.
[*UNIX Network Programming Volume 1: The Sockets Networking API*](#).
Addison Wesley, 3rd edition, 2003.

Notes: