


- **Goal:** *finally* investigate the **application layer**, e.g.,
  - the (mainly OS-based) network stack implementation,
  - the interface between application and network stack, i.e.,
    1. a **raw socket**, or
    2. the **POSIX sockets API**,
- and
- examples of how *you* can use all this!



- **Goal:** *finally* investigate the **application layer**, e.g.,
- the (mainly OS-based) network stack implementation,
  - the interface between application and network stack, i.e.,
    1. a **raw socket**, or
    2. the **POSIX sockets API**,
 and
  - examples of how *you* can use all this!

## POSIX sockets API (1) – The Interface

	Function	Description	Blocking?
UDP {	<b>socket</b>	Form the data structure used to describe communication end-point	×
	<b>bind</b>	Associate socket data structure with (local) address	×
	<b>close</b>	Close socket and stop using it	×
	<b>shutdown</b>	Close socket and stop using it, with control over how	×
	<b>getsockopt</b>	Get or set options for a socket, i.e., control how it functions	×
	<b>setsockopt</b>		
	<b>sendto</b>	Transmit a datagram to (remote) address	✓
	<b>recvfrom</b>	Receive a datagram from (remote) address	✓
	<b>listen</b>	Mark socket as passive, i.e., for incoming connections	×
	<b>accept</b>	Wait for a connection to be established	✓
TCP {	<b>connect</b>	Actively establish a connection with (remote) address	✓
	<b>send</b>	Transmit a segment via connection	✓
	<b>recv</b>	Receive a segment via connection	✓
	<b>select</b>	Wait for activity that would allow non-blocking access	✓
	<b>poll</b>		

## POSIX sockets API (1) – The Interface

Function	Description
<b>getnameinfo</b>	Convert an internal, machine-readable data structure into a host name
<b>getaddrinfo</b>	Convert a host name into an internal, machine-readable data structure
<b>inet_aton</b>	Convert a dotted-decimal address into a binary, machine-readable address
<b>inet_ntoa</b>	Convert a binary, machine-readable address into a dotted-decimal address
<b>htons/htons</b>	Convert a 16/32-bit host order integer into network order
<b>ntohs/ntohs</b>	Convert a 16/32-bit network order integer into host order

## POSIX sockets API (2) – An Implementation (in Linux)

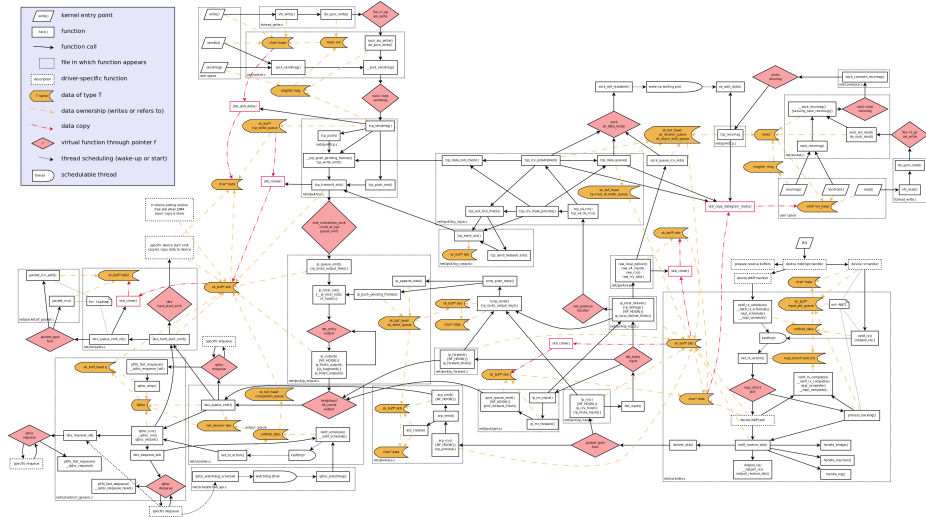
- ▶ Some (rough) **design goals** might include

1. offer POSIX-compliant interface,
2. offer RFC-compliant implementation,
3. maximise efficiency (e.g., low-latency, effective use of bandwidth),
4. maximise flexibility (e.g., general- not special-purpose),
5. allow configurability,
6. ...

which lead to some underlying **golden rules**, e.g.,

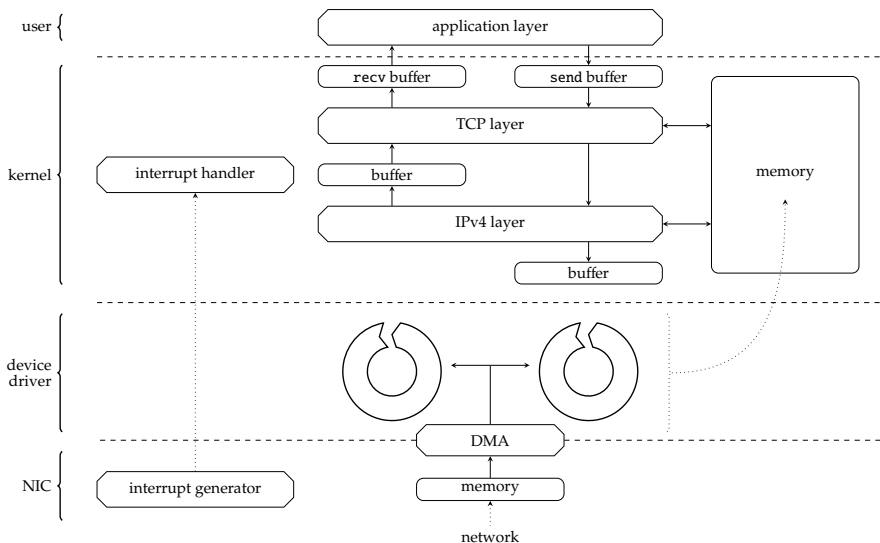
- ▶ make use of all possible hardware support,
- ▶ make use of effective data structures,
- ▶ minimise copying,
- ▶ optimise for common-case,
- ▶ ensure correctness for corner-cases,
- ▶ ...

# POSIX sockets API (3) – An Implementation (in Linux)



[http://www.linuxfoundation.org/collaborate/workgroups/networking/kernel\\_flow](http://www.linuxfoundation.org/collaborate/workgroups/networking/kernel_flow)

## POSIX sockets API (4) – An Implementation (in Linux)



## POSIX sockets API (5) – An Implementation (in Linux)

- ▶ **Fact:** ports are basically buffers within network stack.
- ▶ **Implication #1:**
  - ▶ packets and segments might be received out-of-order, *but*
  - ▶ buffering enforces in-order delivery to the application.



## POSIX sockets API (5) – An Implementation (in Linux)

- ▶ **Fact:** ports are basically buffers within network stack.
- ▶ **Implication #2:** send and transmission are decoupled ...
- ▶ ... transmission *could* occur
  1. when a complete segment is accumulated, or
  2. when transmission is forced, e.g., via
    - ▶ use of the PSH flag, or
    - ▶ some sort of time-out timer

so basically needs to realise a trade-off:

- ▶ less efficient wrt. latency (wait more time) but more efficient wrt. bandwidth (transmit complete segments more often), or
- ▶ more efficient wrt. latency (wait less time) but less efficient wrt. bandwidth (transmit complete segments less often).

## POSIX sockets API (5) – An Implementation (in Linux)


- ▶ **Fact:** ports are basically buffers within network stack.
- ▶ **Implication #3:** send and recv are decoupled ...
- ▶ ... any one of



is possible.

## POSIX sockets API (5) – An Implementation (in Linux)


- ▶ **Fact:** ports are basically buffers within network stack.
- ▶ **Implication #3:** send and recv are decoupled ...
- ▶ ... any one of



is possible.


## POSIX sockets API (5) – An Implementation (in Linux)

- **Fact:** ports are basically buffers within network stack.
- **Implication #3:** send and recv are decoupled ...
- ... any one of




is possible.

## Using POSIX sockets (1) – UDP



## Using POSIX sockets (2) – TCP



## Using POSIX sockets (3) – An “echo uppercase” TCP client/server

### Listing (C)

```
1 #include <sys/socket.h>
2 #include <arpa/inet.h>
3 #include <unistd.h>
4
5 void handle( int cs ) {
6     char t[ 1024 ];
7
8     while( true ) {
9         // terminal -> t
10        fgets( t, 1024, stdin );
11        // server <- t
12        send( cs, t, strlen( t ), 0 );
13        // server -> t'
14        t[ recv( cs, t, 1023, 0 ) ] = '\0';
15        // terminal <- t'
16        fputs( t, stdout );
17    }
18
19    // close connection
20    close( cs );
21 }
```

### Listing (C)

```
1 int main( int argc, char* argv[] ) {
2     struct sockaddr_in sa; socklen_t sl = sizeof( sa );
3     struct sockaddr_in ca; socklen_t cl = sizeof( ca );
4
5     memset( &sa, 0, sl );
6
7     sa.sin_family      = AF_INET;
8     sa.sin_addr.s_addr = inet_addr( argv[ 1 ] );
9     sa.sin_port        = htons( atoi( argv[ 2 ] ) );
10
11    // open socket
12    int cs = socket( AF_INET, SOCK_STREAM, 0 );
13    // open connection
14    connect( cs, ( struct sockaddr* )( &sa ), sl );
15    // handle connection
16    handle( cs );
17
18    return 0;
19 }
```

## Using POSIX sockets (3) – An “echo uppercase” TCP client/server

### Listing (C)

```
1 #include <sys/socket.h>
2 #include <arpa/inet.h>
3 #include <unistd.h>
4
5 void handle( int cs ) {
6     char t[ 1024 ];
7
8     while( true ) {
9         // client -> t
10        t[ recv( cs, t, 1023, 0 ) ] = '\0';
11        // t' = toupper( t )
12        for( int i = 0; i < strlen( t ); i++ ) {
13            t[ i ] = toupper( t[ i ] );
14        }
15        // client <- t'
16        send( cs, t, strlen( t ), 0 );
17    }
18
19    // close connection
20    close( cs );
21
22 }
23
24 }
25 }
```

### Listing (C)

```
1 int main( int argc, char* argv[] ) {
2     struct sockaddr_in sa; socklen_t sl = sizeof( sa );
3     struct sockaddr_in ca; socklen_t cl = sizeof( ca );
4
5     memset( &sa, 0, sl );
6
7     sa.sin_family = AF_INET;
8     sa.sin_addr.s_addr = inet_addr( argv[ 1 ] );
9     sa.sin_port = htons( atoi( argv[ 2 ] ) );
10
11    // open socket
12    int ss = socket( AF_INET, SOCK_STREAM, IPPROTO_IP );
13    // bind socket
14    bind( ss, ( struct sockaddr* )( &sa ), sl );
15    // listen for connections
16    listen( ss, 10 );
17
18    while( true ) {
19
20        // open connection
21        int cs = accept( ss, &ca, &cl );
22        // handle connection
23        handle( cs );
24    }
25
26    // close socket
27    close( ss );
28
29    return 0;
30
31 }
```



## Using POSIX sockets (3) – An “echo uppercase” TCP client/server

### Listing (C)

```
1 #include <sys/socket.h>
2 #include <arpa/inet.h>
3 #include <unistd.h>
4
5 void* handle( void* __cs ) {
6     char t[ 1024 ];
7
8     int cs = *( int* )( __cs );
9
10    while( true ) {
11        // client -> t
12        t[ recv( cs, t, 1023, 0 ) ] = '\0';
13        // t' = toupper( t )
14        for( int i = 0; i < strlen( t ); i++ ) {
15            t[ i ] = toupper( t[ i ] );
16        }
17        // client <- t'
18        send( cs, t, strlen( t ), 0 );
19    }
20
21    // close connection
22    close( cs );
23
24    return NULL;
25 }
```

### Listing (C)

```
1 int main( int argc, char* argv[] ) {
2     struct sockaddr_in sa; socklen_t sl = sizeof( sa );
3     struct sockaddr_in ca; socklen_t cl = sizeof( ca );
4
5     memset( &sa, 0, sl );
6
7     sa.sin_family      = AF_INET;
8     sa.sin_addr.s_addr = inet_addr( argv[ 1 ] );
9     sa.sin_port        = htons( atoi( argv[ 2 ] ) );
10
11    // open socket
12    int ss = socket( AF_INET, SOCK_STREAM, IPPROTO_IP );
13    // bind socket
14    bind( ss, ( struct sockaddr* )( &sa ), sl );
15    // listen for connections
16    listen( ss, 10 );
17
18    while( true ) {
19        pthread_t id;
20
21        // open connection
22        int cs = accept( ss, &ca, &cl );
23        // handle connection
24        pthread_create( &id, NULL, &handle, &cs );
25    }
26
27    // close socket
28    close( ss );
29
30    return 0;
31 }
```

## ► Take away points:

- Ultimately, the POSIX sockets API is an abstraction of the network ...
- ... even so, it's hard to argue you can totally avoid having to understand the underlying technology.
- As with any design, it has **good** and **bad** features: for example,
  - it offers a uniform interface to analogous concepts (cf. **domain sockets**, for IPC),
  - it allows special-case implementation choices such as use of **TCP offload**,
  - the abstraction offered is still low-level so can be hard to use (directly),
  - numerous requirements have changed over time (e.g., network vs. host order, new protocols, new use-cases), but the API hasn't,
  - ...

## References

- [1] Wikipedia: Berkeley sockets.  
[http://en.wikipedia.org/wiki/Berkeley\\_sockets](http://en.wikipedia.org/wiki/Berkeley_sockets).
- [2] Wikipedia: Raw socket.  
[http://en.wikipedia.org/wiki/Raw\\_socket](http://en.wikipedia.org/wiki/Raw_socket).
- [3] Wikipedia: TCP offload engine.  
[http://en.wikipedia.org/wiki/TCP\\_offload\\_engine](http://en.wikipedia.org/wiki/TCP_offload_engine).
- [4] Wikipedia: UNIX domain socket.  
[http://en.wikipedia.org/wiki/Unix\\_domain\\_socket](http://en.wikipedia.org/wiki/Unix_domain_socket).
- [5] Wikipedia: Winsock.  
<http://en.wikipedia.org/wiki/Winsock>.
- [6] Standard for information technology - portable operating system interface (POSIX).  
[Institute of Electrical and Electronics Engineers \(IEEE\) 1003.1, 2008.](http://standards.ieee.org/findstds/standard/1003.1-2008.html)  
<http://standards.ieee.org/findstds/standard/1003.1-2008.html>.
- [7] W.R. Stevens, B. Fenner, and A.M. Rudoff.  
*UNIX Network Programming Volume 1: The Sockets Networking API*.  
Addison Wesley, 3rd edition, 2003.