

COMS20001 lab. worksheet: week #15

- Both the hardware and software in MVB-2.11 is managed by the IT Services Zone E team. If you encounter a problem (e.g., a workstation that fails to boot, an error when you try to use some software, or you just cannot log into your account), they can help: either talk to them in room MVB-3.41, submit a service request online via

<http://servicedesk.bristol.ac.uk>

or talk to the dedicated CS Teaching Technologist, Richard Grafton, in room MVB-2.07.

- We intend this worksheet to be attempted, at least partially, in the associated lab. session. Your attendance is important, since this session represents a central form of feedback and help for COMS20001. Perhaps more so than in units from earlier years, *you* need to actively ask questions of and seek help from either the lectures and/or lab. demonstrators present: passively expecting them to provide solutions is less ideal.
- The questions are roughly classified as either L (for coursework related questions that should be completed in the lab. session), or A (for additional questions that are entirely optional). Keep in mind that we only *expect* you to complete the first class of questions: the additional content has been provided *purely* for your benefit and/or interest, so there is no problem with nor penalty for totally ignoring it (since it is not directly assessed).

Before you start work, download and unarchive^a the file

http://www.ole.bris.ac.uk/bbcswebdav/courses/COMS20001_2015/csdsp/os/sheet/lab/lab-3_q.tar.gz

somewhere secure^b in your file system: it is intended to act as a starting point for your own work, and will be referred to in what follows.

^aExecute the command `tar xvfz lab-3_q.tar.gz` from a BASH shell (e.g., in a terminal window), or use the archive manager GUI (available by using the menu Applications→Accessories→Archive Manager or directly executing `file-roller`) if you prefer.

^bFor example, the Private sub-directory in your home directory.

Q1[L]. This question develops a simple yet functioning operating system kernel, which is based on the concept of cooperative multi-tasking. The goal is to provide a concrete introduction to representation of processes, and, crucially, realisation of a basic context switching mechanism. Doing so requires support for system calls, and, in particular, a `yield` system call that allows a process to cooperatively yield control of the processor.

Although the worksheet concludes with a set of challenges relating to experimental exploration, the fundamental goal throughout is improved understanding of the material. Put another way, the work you do will be biased toward reading and understanding rather than programming at this point. It is crucial *not* to view this as optional effort: carefully working through what is, admittedly, a detailed worksheet will allow you to more easily and rapidly engage with the intellectual vs. engineering challenges involved.

Q1–§1 Explore the archive content

As shown by Figure 1, the content and structure of the archived material provided matches the worksheet from week #13; the only difference is some additional files within the (previously empty) user directory which demand explanation.

Normally, the kernel would be totally separate from any given user program: the idea is the kernel can (dynamically) retrieve the program image from some storage medium (e.g., a disk), then use it to populate the address space (e.g., the text segment) of a process that captures the state of execution. However, doing so is complex; there must actually be a storage medium to start with, for example, and even then the kernel must implement components such as a device driver and file system. Such demands are *too* complex at this stage, so we make two fairly significant simplifications: we

- assume the whole system (i.e., the kernel *plus* any user programs) is (statically) compiled to form a single kernel image, *and*
- assume each such user program is executed once and once only, and never terminates: this means we opt for some special- vs. general-purpose mechanisms, e.g., by pre-allocating some of the resources (e.g., PCBs, stack space) each process requires.

In support of both, the additional files are

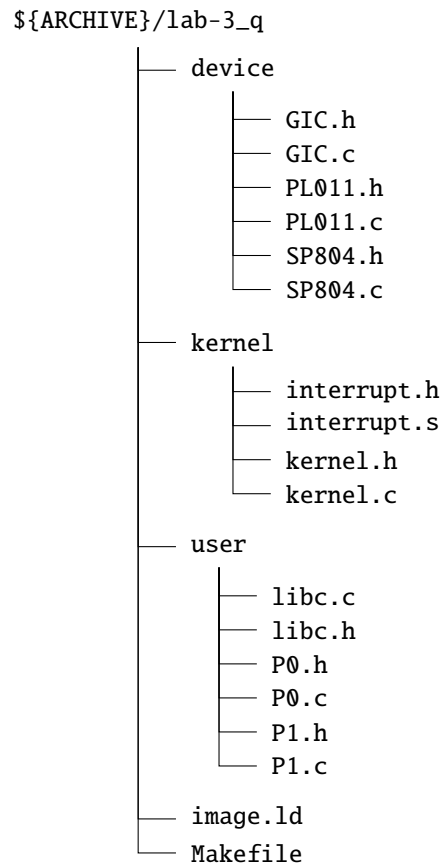


Figure 1: A diagrammatic description of the material in `lab-3_q.tar.gz`.

- `P0.[ch]` and `P1.[ch]`, which are definitions of the two user programs (statically) compiled into the kernel image, and
- `libc.[ch]`, which represents a (very limited) standard library against which the programs are linked: the role of this library is similar to the (full) C standard library, in the sense it abstracts interaction with the kernel via two system call wrappers.

Based on the above, the kernel itself has a fairly limited remit: it needs to a) include an appropriate process table that captures the state of execution (for each executed user program), and b) implement a context switching mechanism and scheduling algorithm so said processes can be suspended and resumed as need be.

Q1–§2 Understand the archive content

image.ld Figure 4 illustrates the linker script `image.ld`. It controls how `ld` produces the kernel image from object files, which, in turn, stem from compilation of the source code files; the resulting layout in memory is illustrated by Figure 2.

interrupt.[sh] Figure 5 and Figure 7 illustrate the header file `interrupt.h` and source code `interrupt.s`: the latter is subtly, but *significantly* more complex than similar files in previous worksheets. This complexity stems from the fact that when an interrupt occurs, we now expect a *user* mode process to be executing. As a result, the execution context must be a) preserved at the interrupt handler entry point and b) restored at the interrupt handler exit point. These differences are limited to the top part, and so the two functions outlined below:

- `handler_svc` can be viewed as three parts:
 - a 5-instruction prologue,
 - invocation of `kernel_handler_svc`, the high-level handler in `kernel.c`, and
 - a 5-instruction epilogue.

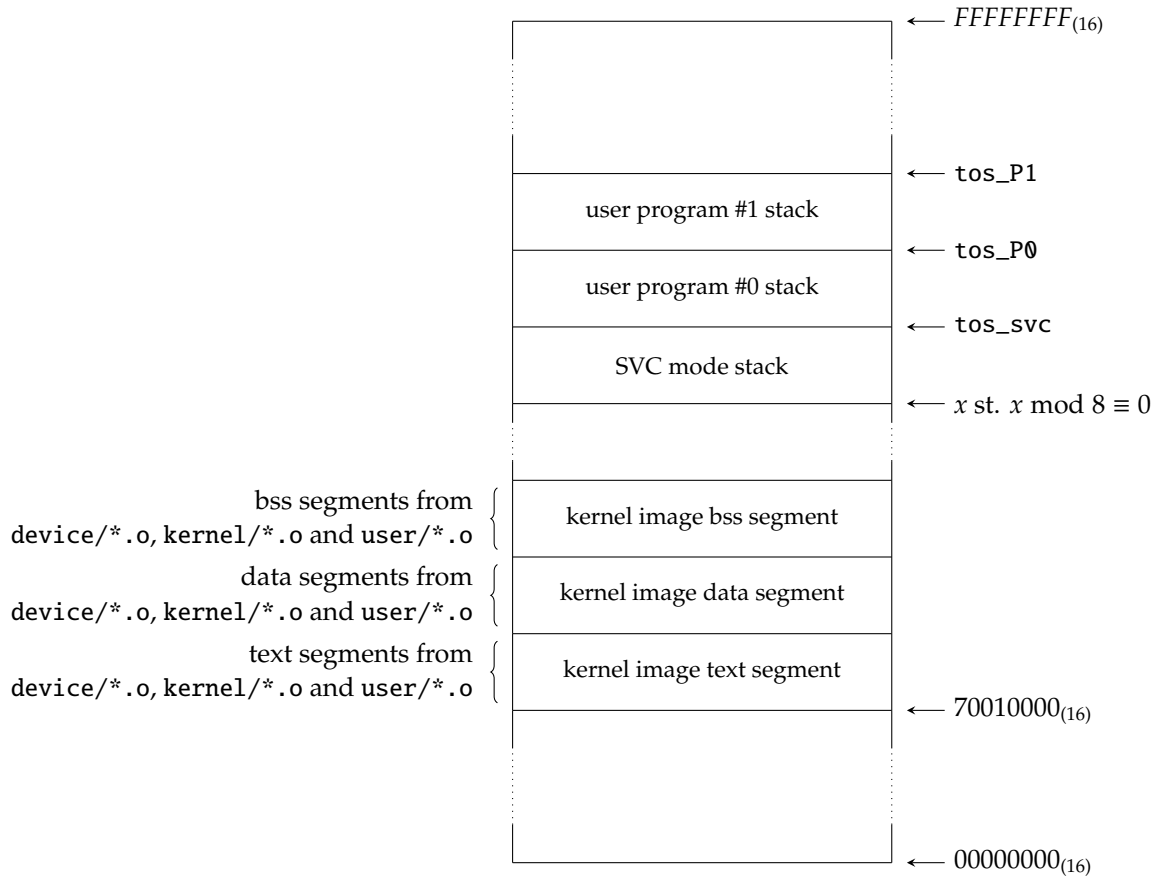


Figure 2: A diagrammatic description of the memory layout realised by `image.ld`.

To understand how this works, imagine the processor is executing a user mode process (in USR mode) and executes an `svc` instruction. This raises a supervisor call interrupt, and causes the interrupt handling mechanism to invoke `handler_svc` (with the processor now in SVC mode). The kernel must avoid corrupting the execution context associated with the user process, because otherwise would be impossible to resume (i.e., recommence execution of) the the user process later. To see how a lack of care *could* cause problems, note that only certain registers are banked: the SVC mode `r0` register holds a value "owned" by USR mode, for example. So, it preserves said context on the SVC mode stack:

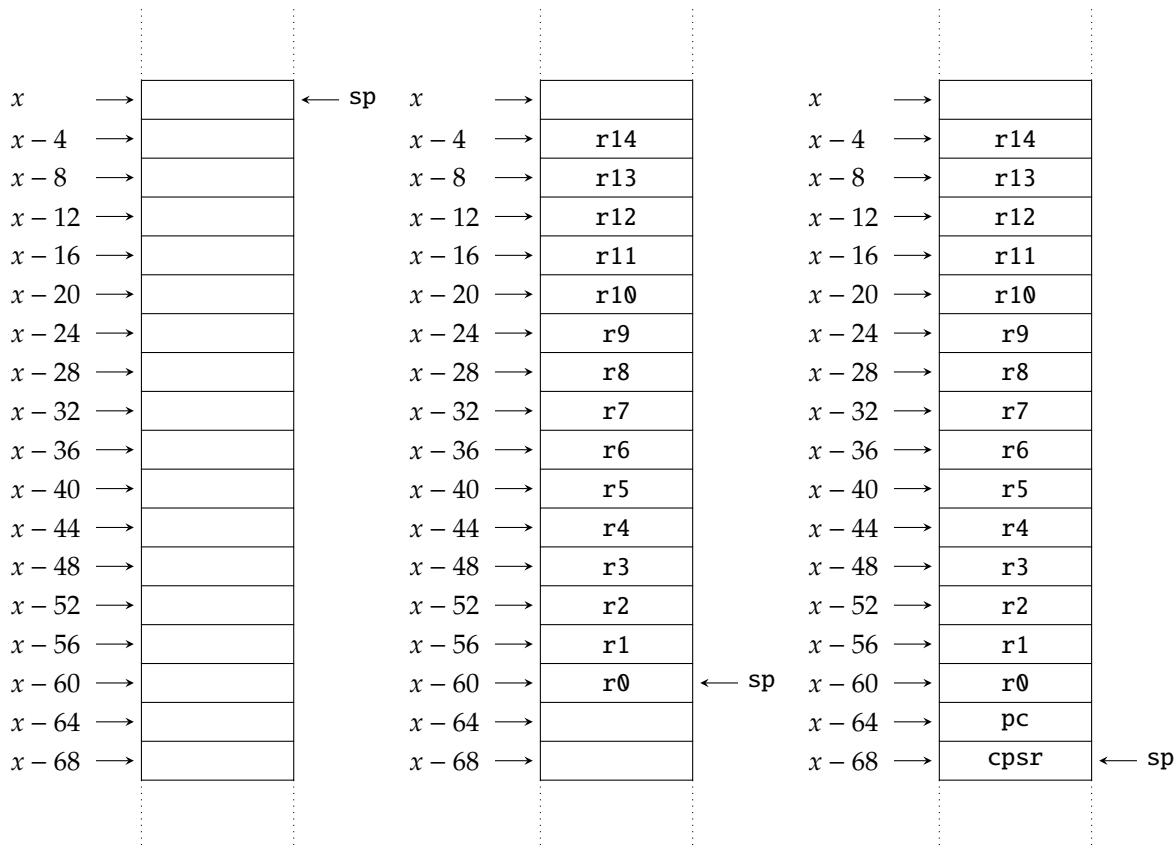
- Line #23 allocates space on the SVC mode stack for the execution context.
- Line #24 preserve USR mode `r0` through to `r12`, plus `sp` (i.e., `r13`) and `lr` (i.e., `r14`).
- Lines #25 and #26 preserve USR mode CPSR (which was copied into SVC mode SPSR by the interrupt), plus USR mode `sp` (i.e., `r13`, which was copied into SVC mode `lr` by the interrupt).

Figure 3 illustrates the stack content at each step in this process. So, with the USR mode execution context preserved on the SVC mode stack, `kernel_handler_svc` is invoked:

- Line #28 set argument #0 equal to the SVC mode `sp`.
- Line #29 and #30 set argument #1 equal to the 24-bit immediate value encoded in the `svc` instruction that raised the interrupt now being handled.
- Line #31 invokes `kernel_handler_svc` with the arguments as specified above.

Since argument #0 is effectively a pointer to the preserved execution context, `kernel_handler_svc` can use/update this as it sees fit. However, it must consider one fact, namely that whatever the preserved execution context contains when the function returns will be restored by the `handler_svc` epilogue:

- Lines #33 and #34 restore USR mode CPSR into SVC mode SPSR, plus USR mode `sp` into SVC mode `lr`.
- Line #35 restore USR mode `r0` through to `r12`, plus `sp` (i.e., `r13`) and `lr` (i.e., `r14`).
- Line #36 deallocates space on the SVC mode stack for the execution context.



(a) Stack layout at function entry point (b) Stack layout after `sub` and `stmia` (c) Stack layout after `stmdb` (line #31), (i.e., line #27). (lines #28 and #29). capturing the preserved USR mode state.

Figure 3: An example showing how `handler_svc` preserves (lines #28 to #31) the USR mode state on the SVC mode stack ready for use by `kernel_handler_svc` and then subsequent restoration (lines #38 to #41).

- Line #37 returns control at the now restored program counter value held in SVC mode `lr`; note that simultaneously, SVC mode SPSR is copied into USR mode CPSR which is what invokes the processor mode switch.
- `handler_rst` is a lot easier to explain:
 - As with the worksheet from week #14, lines # 6 to #9 initialise the interrupt vector table and SVC mode stack.
 - Lines #16 and #20 are identical to the prologue in `handler_svc`: they restore an execution context from the SVC mode stack once `kernel_handler_rst` returns.
 - Unlike `handler_svc`, however, there is nothing sane to preserve: this interrupt relates to initialisation after reset, so there can be no currently executing user process. As a result, line #11 simply allocates the right amount of space on the SVC mode stack for an execution context, but does not fill it with anything.
 - Lines #13 and #14 invoke `kernel_handler_rst`, with `sp` as an argument so the function can update the execution context: it *must* do so, otherwise the restoration steps will result in “junk” being written into all the registers.

It is vital to remember this implementation represents a trade-off in favour of simplicity, but at the expense of efficiency. For example, preserving the *entire* execution context on *every* interrupt allows for a clean and uniform separation between low- and high-level interrupt handlers, *but* could be inefficient if/when the same execution context is restored as preserved: in this case, we may have avoided the overhead of the former (and could do so, more easily at least, by writing more of the interrupt handler at a low-level where control over use of registers is easier and more explicit).

kernel.[ch] Figure 6 and Figure 8 illustrate the header file `kernel.h` and source code `kernel.c`. When compared to previous worksheets, the former has some additional content which needs explanation:

- Lines #31 to #33 define a type `ctx_t` that captures an execution context (i.e., the processor state). In addition to noting that each pertinent component is present, keep in mind that `ctx_t` uses a specific order for the components: reading the content from `sp` upward, it matches Figure 3 st. elements of said content. So if we have pointer `ctx` of type `ctx_t*` equal to `sp`, we can access values of the preserved execution context via fields in `ctx`.
- Line #35 define a type `pid_t` that captures an integer Process IDentifier (PID),
- Lines #37 to #40 define a type `pcb_t` that captures a Process Control Block (PCB), instances of which form entries in the process table: given the limited remit here, such entry simply includes a PID and execution context.

These type definitions are used by the former, which implements the kernel itself:

- Line #12 defines a fixed-size array of `pcb_t` instances for use as the process table: there is one entry for each user process, with a pointer into this table maintained so it is clear which PCB is active (i.e., which process is currently executing).
- Lines #14 to #25 implement the function scheduler, which is the scheduler algorithm. We know there will always be two processes, so the implementation is (very) special-purpose: it simply checks which process is active, and performs a context switch to suspend it and resume the other. You could think of this as round-robin scheduling, but with two processes clearly there is only one choice! `memcpy` is used to copy the associated execution contexts into place, before updating current to reflect the (newly) active PCB.
- Line #27 to #55 implement the high-level interrupt handler `kernel_handler_rst`. First it initialises the process table by copying information into two PCBs, one for each user program; in each case the PCB is first zero'ed by `memset`, and then components such as the program counter value are initialised. In short, these steps reflect (automatic) execution of the two user programs and hence formation of two associated user processes. With the process table initialised, the function then selects a PCB entry and activates it: `memcpy` is used to copy the execution context to addresses pointed at by `ctx` (as allocated by `handler_rst`).

Note `handler_rst` will restore the execution context, and resume execution once `kernel_handler_rst` returns: the initialised value of CPSR has a subtle but important role. More specifically, we the (preserved) CPSR value is initialised so it reflects the processor in USR mode with IRQ interrupts enabled. By doing so, when the preserved execution context is restored, the processor is automatically switched to USR mode and IRQ interrupts enabled: there is no need to perform either explicitly, as there is in `kernel_handler_rst` say.

- Lines #57 to #86 implement the high-level interrupt handler `kernel_handler_svc`: recall this is invoked by `handler_svc` with the two arguments formed as
 - a pointer to the preserved execution context (on the SVC mode stack, i.e., the value of `sp`), and
 - the immediate operand of the `svc` instruction that raised the interrupt being handled: this is used as the system call identifier.

The function uses the system call identifier to decide how the interrupt should be handled, and so what the kernel-facing side of the system call API should do:

- Lines #67 to #70 deal with identifier `00(16)`, i.e., the `yield` system call: when one of the user processes invokes `yield` as defined in `libc.c`, the resulting supervisor call interrupt is eventually handled by this case. The semantics of this system call are that the process scheduler should be invoked, i.e., that the currently executing user process has yielded control of the processor (thus permitting the other to execute).
- Lines #71 to #82 Identifier `01(16)` is the `write` system call: when one of the user processes invokes `write` as defined in `libc.c`, the resulting supervisor call interrupt is eventually handled by this case. The semantics of this system call are that some `n`-element sequence of bytes pointed to by `x` are written to file descriptor `fd`; the kernel ignores `fd`, and simply writes those bytes to the `PL011_t` instance `UART0`.

libc.[ch] Figure 13 and Figure 14 illustrate the header file `libc.h` and source code `libc.c`. The idea here is to capture the user-facing side of the system call API. Put another way, the *kernel-facing* side (i.e., the *implementation*) is described above; `libc.h` and `libc.c` support use of that implementation, offering a more appropriate abstraction by wrapping low-level `svc` instructions in high(er)-level functions. These wrappers abstract both the fact assembly language is required, *and* the actual calling convention used. Various choices exist: here we opt to use the `svc` instruction itself to communicate a system call identifier. However, we could equally have defined the API so this is communicated via a register in the same way as any arguments.

P[01].[ch] Figure 9 and Figure 10 illustrate the header file `P0.h` and source code `P0.c`, which describe user program #0; the files associated with user program #1 are essentially the same, so we focus on the former alone as an example.

- Line #4 defines a 20-character string `x` user to identify the program (this is the only difference in user program #1).
- Lines #6 to #8 implement an infinite `while` loop, which implies that the function never returns. Each iteration makes two system calls, namely
 - a call to `write` that instructs the kernel to write `x` to the file descriptor whose identifier is 0: per the above, this is interpreted as transmitting it via the `PL011_t` instance `UART0`, and
 - a call to `yield` that instructs the kernel to invoke the scheduler: per the above this means a context switch occurs.

Q1–§3 Experiment with the archive content

Following the same approach as in the worksheet from week #13, first launch QEMU then `gdb`. Issue the

`continue`

command to `gdb` in the debugging terminal so the kernel image is executed. You should observe the two strings “hello world, I’m P0” and “hello world, I’m P1” being written to the emulation terminal. At an intuitive level, the behaviour of this kernel is easy to summarise. It should be clear that when the processor is reset, each of the two user programs are (automatically) executed and form associated processes. Over time, the processor is shared between said processes; this is achieved by the kernel switching between them when the process currently executing (cooperatively) relinquishes control via the `yield` system call. Each process does so, using the `write` system call to identify itself beforehand (so as to *demonstrate* it is executing), meaning the result is a sequence of identifier strings written to the emulator terminal as observed.

Q1–§4 Next steps

There are various things you could (optionally) do next: here are some ideas.

- a An effective way to gain insight into an example of this type, in which the timing of events is crucial, is by drawing a time-line to illustrate said events. For example, by including a) when events (e.g., interrupts) occur, b) what the processor and memory (e.g., stack) state is, and c) what instructions (i.e., what function) is being executed at different periods of time, one can align the observed behaviour with the source code. Try to construct such a diagram for this example, limiting the duration to the first three context switches at most.
- b Add an *additional*, third user program: initially this could be more or less identical to the other two, but the goal would be to explore questions like
 - i what happens if it behaves in a non-cooperative (i.e., does not invoke `yield`) or semi-cooperative (i.e., does invoke `yield`, just less often than the other two user programs) manner, or
 - ii whether/how the new user program can exploit the the lack of protection implemented by this kernel, and what behaviour might result when it does so.
- c Add an *additional* system call by updating both to the user- and kernel-facing side of the system call API. A good candidate is `read`, since this acts as the natural counterpart to `write` by allowing processes to read bytes from the emulated UART.
- d As outlined above, there are various possible choices wrt. design of the system call API: investigate how system call invocation is implemented in Linux, then alter this design to match.

```

1  SECTIONS {
2      /* assign address (per QEMU) */
3      . = 0x70010000;
4      /* place text segment(s) */
5      .text : { kernel/interrupt.o(.text) *(.text .rodata) }
6      /* place data segment(s) */
7      .data : { *(.data) }
8      /* place bss segment(s) */
9      .bss : { *(.bss) }
10     /* align address (per AAPCS) */
11     . = ALIGN(8);
12     /* allocate stack for svc mode */
13     . = . + 0x00001000;
14     tos_svc = .;
15     /* allocate stack for P0 */
16     . = . + 0x00001000;
17     tos_P0 = .;
18     /* allocate stack for P1 */
19     . = . + 0x00001000;
20     tos_P1 = .;
21 }

```

Figure 4: image.ld

```

1  #ifndef __INTERRUPT_H
2  #define __INTERRUPT_H
3
4  // enable IRQ interrupts
5  extern void irq_enable();
6  // disable IRQ interrupts
7  extern void irq_unable();
8
9  #endif

```

Figure 5: kernel/interrupt.h

```

1  #ifndef __KERNEL_H
2  #define __KERNEL_H
3
4  #include <stddef.h>
5  #include <stdint.h>
6
7  #include "GIC.h"
8  #include "PL011.h"
9  #include "SP804.h"
10
11 #include "interrupt.h"
12
13 // Include functionality from newlib, the embedded standard C library.
14
15 #include <string.h>
16
17 // Include definitions relating to the 2 user programs.
18
19 #include "P0.h"
20 #include "P1.h"
21
22 /* The kernel source code is made simpler by three type definitions:
23  *
24  * - a type that captures each component of an execution context (i.e.,
25  *   processor state) in a compatible order wrt. the low-level handler
26  *   preservation and restoration prologue and epilogue,
27  * - a type that captures a process identifier, which is basically an
28  *   integer, and
29  * - a type that captures a process PCB.
30  */
31
32 typedef struct {
33     uint32_t cpsr, pc, gpr[ 13 ], sp, lr;
34 } ctx_t;
35
36 typedef int pid_t;
37
38 typedef struct {
39     pid_t pid;
40     ctx_t ctx;
41 } pcb_t;
42
43 #endif

```

Figure 6: kernel/kernel.h


```

1  /* Each of the following is a low-level interrupt handler: each one is
2  * tasked with handling a different interrupt type, and acts as a sort
3  * of wrapper around a high-level, C-based handler.
4  */
5
6  handler_rst: bl    table_copy          @ initialise interrupt vector table
7
8              msr    cpsr, #0xD3        @ enter SVC mode with no interrupts
9              ldr    sp, =tos_svc       @ initialise SVC mode stack
10
11             sub    sp, sp, #68         @ initialise dummy context
12
13             mov    r0, sp              @ set C function arg. = SP
14             bl     kernel_handler_rst  @ invoke C function
15
16             ldmia  sp!, { r0, lr }     @ load  USR mode PC and CPSR
17             msr    spsr, r0            @ set  USR mode      CPSR
18             ldmia  sp, { r0-r12, sp, lr }^ @ load  USR mode registers
19             add    sp, sp, #60         @ update SVC mode SP
20             movs   pc, lr              @ return from interrupt
21
22 handler_svc: sub    lr, lr, #0          @ correct return address
23             sub    sp, sp, #60         @ update SVC mode stack
24             stmia  sp, { r0-r12, sp, lr }^ @ store  USR registers
25             mrs    r0, spsr           @ get   USR      CPSR
26             stmdb  sp!, { r0, lr }     @ store  USR PC and CPSR
27
28             mov    r0, sp              @ set  C function arg. = SP
29             ldr    r1, [ lr, #-4 ]     @ load          svc instruction
30             bic    r1, r1, #0xFF000000 @ set  C function arg. = svc immediate
31             bl     kernel_handler_svc  @ invoke C function
32
33             ldmia  sp!, { r0, lr }     @ load  USR mode PC and CPSR
34             msr    spsr, r0            @ set  USR mode      CPSR
35             ldmia  sp, { r0-r12, sp, lr }^ @ load  USR mode registers
36             add    sp, sp, #60         @ update SVC mode SP
37             movs   pc, lr              @ return from interrupt
38
39 /* The following captures the interrupt vector table, plus a function
40 * to copy it into place (which is called on reset): note that
41 *
42 * - for interrupts we don't handle an infinite loop is realised (to
43 * to approximate halting the processor), and
44 * - we copy the table itself, *plus* the associated addresses stored
45 * as static data: this preserves the relative offset between each
46 * ldr instruction and wherever it loads from.
47 */
48
49 table_data: ldr    pc, address_rst      @ reset          vector -> SVC mode
50             b      .                   @ undefined instruction vector -> UND mode
51             ldr    pc, address_svc      @ supervisor call  vector -> SVC mode
52             b      .                   @ abort (prefetch)   vector -> ABT mode
53             b      .                   @ abort (data)       vector -> ABT mode
54             b      .                   @ reserved
55             b      .                   @ IRQ                vector -> IRQ mode
56             b      .                   @ FIQ                vector -> FIQ mode
57
58 address_rst: .word handler_rst
59 address_svc: .word handler_svc
60
61 table_copy: mov    r0, #0               @ set destination address
62             ldr    r1, =table_data      @ set source      address
63             ldr    r2, =table_copy      @ set source      limit
64
65 table_loop: ldr    r3, [ r1 ], #4        @ load word, inc. source      address
66             str    r3, [ r0 ], #4        @ store word, inc. destination address
67
68             cmp    r1, r2               @ loop if address != limit
69             bne    table_loop
70
71             mov    pc, lr               @ return
72
73 /* These function enable and disable IRQ interrupts respectively, by
74 * toggling the 7-th bit of CPSR to either 0 or 1.
75 */
76
77 .global irq_enable
78 .global irq_unable
79
80 irq_enable: mrs    r0, cpsr             @ enable IRQ interrupts
81             bic    r0, r0, #0x80
82             msr    cpsr_c, r0
83
84             mov    pc, lr
85
86 irq_unable: mrs    r0, cpsr             @ disable IRQ interrupts
87             orr    r0, r0, #0x80
88             msr    cpsr_c, r0
89
90             mov    pc, lr

```

Figure 7: kernel/interrupt.s


```

1  #include "kernel.h"
2
3  /* Since we *know* there will be 2 processes, stemming from the 2 user
4   * programs, we can
5   *
6   * - allocate a fixed-size process table (of PCBs), and use a pointer
7   *   to keep track of which entry is currently executing, and
8   * - employ a fixed-case of round-robin scheduling: no more processes
9   *   can be created, and neither is able to complete.
10  */
11
12  pcb_t pcb[ 2 ], *current = NULL;
13
14  void scheduler( ctx_t* ctx ) {
15      if ( current == &pcb[ 0 ] ) {
16          memcpy( &pcb[ 0 ].ctx, ctx, sizeof( ctx_t ) );
17          memcpy( ctx, &pcb[ 1 ].ctx, sizeof( ctx_t ) );
18          current = &pcb[ 1 ];
19      }
20      else if ( current == &pcb[ 1 ] ) {
21          memcpy( &pcb[ 1 ].ctx, ctx, sizeof( ctx_t ) );
22          memcpy( ctx, &pcb[ 0 ].ctx, sizeof( ctx_t ) );
23          current = &pcb[ 0 ];
24      }
25  }
26
27  void kernel_handler_rst( ctx_t* ctx ) {
28      /* Initialise PCBs representing processes stemming from execution of
29       * the two user programs. Note in each case that
30       *
31       * - the CPSR value of 0x50 means the processor is switched into USR
32       *   mode, with IRQ interrupts enabled, and
33       * - the PC and SP values match the entry point and top of stack.
34       */
35
36      memset( &pcb[ 0 ], 0, sizeof( pcb_t ) );
37      pcb[ 0 ].pid = 0;
38      pcb[ 0 ].ctx.cpsr = 0x50;
39      pcb[ 0 ].ctx.pc = ( uint32_t )( entry_P0 );
40      pcb[ 0 ].ctx.sp = ( uint32_t )( &tos_P0 );
41
42      memset( &pcb[ 1 ], 0, sizeof( pcb_t ) );
43      pcb[ 1 ].pid = 1;
44      pcb[ 1 ].ctx.cpsr = 0x50;
45      pcb[ 1 ].ctx.pc = ( uint32_t )( entry_P1 );
46      pcb[ 1 ].ctx.sp = ( uint32_t )( &tos_P1 );
47
48      /* Once the PCBs are initialised, we (arbitrarily) select one to be
49       * restored (i.e., executed) when the function then returns.
50       */
51
52      current = &pcb[ 0 ]; memcpy( ctx, &current->ctx, sizeof( ctx_t ) );
53
54      return;
55  }
56
57  void kernel_handler_svc( ctx_t* ctx, uint32_t id ) {
58      /* Based on the identified encoded as an immediate operand in the
59       * instruction,
60       *
61       * - read the arguments from preserved usr mode registers,
62       * - perform whatever is appropriate for this system call,
63       * - write any return value back to preserved usr mode registers.
64       */
65
66      switch( id ) {
67          case 0x00 : { // yield()
68              scheduler( ctx );
69              break;
70          }
71          case 0x01 : { // write( fd, x, n )
72              int fd = ( int )( ctx->gpr[ 0 ] );
73              char* x = ( char* )( ctx->gpr[ 1 ] );
74              int n = ( int )( ctx->gpr[ 2 ] );
75
76              for( int i = 0; i < n; i++ ) {
77                  PL011_putc( UART0, *x++ );
78              }
79
80              ctx->gpr[ 0 ] = n;
81              break;
82          }
83          default : { // unknown
84              break;
85          }
86      }
87
88      return;
89  }

```

Figure 8: kernel/kernel.c

```

1  #ifndef __P0_H
2  #define __P0_H
3
4  #include <stddef.h>
5  #include <stdint.h>
6
7  #include "libc.h"
8
9  // define symbols for P0 entry point and top of stack
10 extern void (*entry_P0)();
11 extern uint32_t tos_P0;
12
13 #endif

```

Figure 9: user/P0.h

```

1  #include "P0.h"
2
3  void P0() {
4      char* x = "hello world, I'm P0\n";
5
6      while( 1 ) {
7          write( 0, x, 20 ); yield();
8      }
9  }
10
11 void (*entry_P0)() = &P0;

```

Figure 10: user/P0.c

```

1  #ifndef __P1_H
2  #define __P1_H
3
4  #include <stddef.h>
5  #include <stdint.h>
6
7  #include "libc.h"
8
9  // define symbols for P1 entry point and top of stack
10 extern void (*entry_P1)();
11 extern uint32_t tos_P1;
12
13 #endif

```

Figure 11: user/P1.h

```

1  #include "P1.h"
2
3  void P1() {
4      char* x = "hello world, I'm P1\n";
5
6      while( 1 ) {
7          write( 0, x, 20 ); yield();
8      }
9  }
10
11 void (*entry_P1)() = &P1;

```

Figure 12: user/P1.c

```

1  #ifndef __LIBC_H
2  #define __LIBC_H
3
4  #include <stddef.h>
5  #include <stdint.h>
6
7  // cooperatively yield control of processor, i.e., invoke the scheduler
8  void yield();
9
10 // write n bytes from x to the file descriptor fd
11 int write( int fd, void* x, size_t n );
12
13 #endif

```

Figure 13: user/libc.h

Q2[A]. There is a fantastic (short) film at

<http://www.youtube.com/watch?v=Q07PhW5sCEk>

focusing on the Compatible Time-Sharing System (CTTS), developed at MIT in the early 1960s. CTTS was *highly* influential in terms of later technology: a range of subsequent systems such as Multics and Unix refining concepts and approaches *it* pioneered. However, the film is arguably more remarkable because a) it offers such a brilliant explanation of what, at the time, was a revolutionary concept, and b) gives an insight into the challenges faced by computing in that era: it neatly illustrates the cost associated with maintaining one computer, for example, and so the value in maximising utilisation of it!

```
1 #include "libc.h"
2
3 void yield() {
4     asm volatile( "svc #0      \n" );
5 }
6
7 int write( int fd, void* x, size_t n ) {
8     int r;
9
10    asm volatile( "mov r0, %1 \n"
11                  "mov r1, %2 \n"
12                  "mov r2, %3 \n"
13                  "svc #1      \n"
14                  "mov %0, r0 \n"
15                  : "=r" (r)
16                  : "r" (fd), "r" (x), "r" (n)
17                  : "r0", "r1", "r2" );
18
19    return r;
20 }
```

Figure 14: user/libc.c