

Data Structures and Algorithms – COMS21103

2015/2016

Dynamic Search Structures

Self-balancing Trees and Skip Lists

Benjamin Sach

Dynamic Search Structures

A **dynamic search structure**,
stores a set of elements

*Each element x must have a **unique** key - $x.key$*

The following operations are supported:

INSERT(x, k) - inserts x with key $k = x.key$

FIND(k) - returns the (unique) element x with $x.key = k$
(or reports that it doesn't exist)

DELETE(k) - deletes the (unique) element x with $x.key = k$
(or reports that it doesn't exist)

Dynamic Search Structures

A **dynamic search structure**,
stores a set of elements

*Each element x must have a **unique** key - $x.key$*

The following operations are supported:

INSERT(x, k) - inserts x with key $k = x.key$

FIND(k) - returns the (unique) element x with $x.key = k$
(or reports that it doesn't exist)

DELETE(k) - deletes the (unique) element x with $x.key = k$
(or reports that it doesn't exist)

We would also like it to support (among others):

PREDECESSOR(k) - returns the (unique) element x
with the largest key such that $x.key < k$

RANGEFIND(k_1, k_2) - returns every element x with $k_1 \leq x.key \leq k_2$

Using a Linked List as a Dynamic Search Structure

There are many ways in which we could implement a search structure. . .

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

Using a Linked List as a Dynamic Search Structure

There are many ways in which we could implement a search structure. . .

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,

- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*

Using a Linked List as a Dynamic Search Structure


There are many ways in which we could implement a search structure...

but they aren't all efficient

Let n denote the number of elements stored in the structure

- our goal is to implement a structure with operations which scale well as n grows

We could implement a Dynamic Search Structure using an unsorted linked list:



INSERT is very efficient,


- add the new item to the head of the list in $O(1)$ time

FIND and **DELETE** are very *inefficient*, they take $O(n)$ time

- we have to look through the entire linked list to find an item *(in the worst case)*


Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



FIND(21)


Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left



FIND(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left
21 = 21 found it!



FIND(21)


Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node



- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node


- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right



DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right



DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left



DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left



h

DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left
21 = 21 found it!



DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys

We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left
21 = 21 found it!
now delete it!



h

DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height


The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



21 < 27 so go left
21 > 20 so go right
21 < 23 so go left
21 = 21 found it!
now delete it!



DELETE(21)

Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys



We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Recall that in a binary search tree, for any node

- all the nodes in the left subtree have smaller keys
- all the nodes in the right subtree have larger keys


We perform a **FIND** operation by following a path from the root...

this takes $O(h)$ time - where h is the height

The **INSERT** and **DELETE** operations are similar and also take $O(h)$ time


Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...




How big is h ?

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements




How big is h ?

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?


It might be as small as $\log_2 n$ (if the tree is perfectly balanced)

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?

It might be as small as $\log_2 n$ (if the tree is perfectly balanced)


but it might be as big as n (if the tree is completely unbalanced)

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?

It might be as small as $\log_2 n$ (if the tree is perfectly balanced)

but it might be as big as n (if the tree is completely unbalanced)


In particular, each **INSERT** could increase h by one

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?

It might be as small as $\log_2 n$ (if the tree is perfectly balanced)

but it might be as big as n (if the tree is completely unbalanced)


In particular, each **INSERT** could increase h by one

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?

It might be as small as $\log_2 n$ (if the tree is perfectly balanced)
but it might be as big as n (if the tree is completely unbalanced)


In particular, each **INSERT** could increase h by one

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?

It might be as small as $\log_2 n$ (if the tree is perfectly balanced)

but it might be as big as n (if the tree is completely unbalanced)


In particular, each **INSERT** could increase h by one

Binary Search Trees

A classic choice for a dynamic search structure is
a binary search tree...



n is the number
of elements



How big is h ?

It might be as small as $\log_2 n$ (if the tree is perfectly balanced)

but it might be as big as n (if the tree is completely unbalanced)

In particular, each **INSERT** could increase h by one

how can we overcome this?

Part one


Self-balancing trees

*inspired by slides by Inge Li Gørtz
in turn inspired by slides by Kevin Wayne*

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*


Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length *(always, all the time)*




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length *(always, all the time)*



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length *(always, all the time)*



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length *(always, all the time)*



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length *(always, all the time)*



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)



2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys


*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

2-3-4 Trees

Key idea: Nodes can have between 2 and 4 children *(hence the name)*

Perfect balance - every path from the root to a leaf has the same length
(always, all the time)




2-node: 2 children and 1 key
3-node: 3 children and 2 keys
4-node: 4 children and 3 keys

*Like in a binary search tree,
the keys held at a node determine
the contents of its subtrees*

The ● are “dummy leaves” (they don’t do or contain anything)

The FIND operation


Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...



*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation


Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...



*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13
found it!

FIND(12)




*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13
found it!

FIND(12)



*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*


*What is the time complexity of the **FIND** operation?*

The FIND operation

Just like in a binary search tree,
we perform a **FIND** operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13
found it!

FIND(12)



*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

*What is the time complexity of the **FIND** operation?*


It's $O(h)$ again

The FIND operation

Just like in a binary search tree,
we perform a FIND operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13
found it!

FIND(12)



*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

What is the time complexity of the FIND operation?

It's $O(h)$ again


(each step down the path takes $O(1)$ time)

The FIND operation

Just like in a binary search tree,
we perform a FIND operation by following a path from the root...

12 is between 11 and 18
12 is smaller than 13
found it!

FIND(12)



*decisions are made by inspecting the key(s) at the current node
and following the appropriate edge*

What is the time complexity of the FIND operation?

It's $O(h)$ again

(each step down the path takes $O(1)$ time)


What is the height, h of a 2-3-4 tree?

The height of a 2-3-4 tree


Perfect balance - every path from the root to a leaf has the same length
(we'll justify this as we go along)

This implies that the height, h of a 2-3-4 tree with n nodes is

Best case: $\log_4 n = \frac{\log_2 n}{2}$ (all 4-nodes)



Worst case: $\log_2 n$ (all 2-nodes)




h is between 10 and 20 for a million nodes

The height of a 2-3-4 tree


Perfect balance - every path from the root to a leaf has the same length
(we'll justify this as we go along)

This implies that the height, h of a 2-3-4 tree with n nodes is

Best case: $\log_4 n = \frac{\log_2 n}{2}$ (all 4-nodes)



Worst case: $\log_2 n$ (all 2-nodes)



h is between 10 and 20 for a million nodes


The time complexity of the **FIND** operation is $O(h)$

The height of a 2-3-4 tree


Perfect balance - every path from the root to a leaf has the same length
(we'll justify this as we go along)

This implies that the height, h of a 2-3-4 tree with n nodes is

Best case: $\log_4 n = \frac{\log_2 n}{2}$ (all 4-nodes)




Worst case: $\log_2 n$ (all 2-nodes)



h is between 10 and 20 for a million nodes


The time complexity of the **FIND** operation is $O(h) = O(\log n)$

The INSERT operation



To perform $\text{INSERT}(x, k)$,

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

The INSERT operation


INSERT($x, 16$)



To perform INSERT(x, k),

Step 1: Search for the key k as if performing FIND(k).


The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform INSERT(x, k),

Step 1: Search for the key k as if performing FIND(k).

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform INSERT(x, k),

Step 1: Search for the key k as if performing FIND(k).

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation




To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

Step 4: If the leaf is a 4-node,

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

Step 4: If the leaf is a 4-node, ???

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

Step 4: If the leaf is a 4-node, ???


We will make sure this *never* happens

SPLITTING 4-nodes




*We can SPLIT any 4-node into two 2-nodes
if it's parent isn't a 4-node*

SPLITTING 4-nodes




We can **SPLIT** any 4-node into two 2-nodes
if it's parent isn't a 4-node

BEFORE




AFTER

SPLITTING 4-nodes



We can **SPLIT** any 4-node into two 2-nodes
if its parent isn't a 4-node


BEFORE




AFTER

The extra key is pushed up to the parent
(so it won't work if the parent is a 4-node)

SPLITTING 4-nodes




BEFORE




AFTER

SPLITTING 4-nodes




BEFORE




AFTER

The extra key is pushed up to the parent
(so it won't work if the parent is a 4-node)

SPLITTING 4-nodes



BEFORE




AFTER


The extra key is pushed up to the parent
(so it won't work if the parent is a 4-node)

no path lengths have changed

SPLITTING 4-nodes



BEFORE




AFTER


The extra key is pushed up to the parent
(so it won't work if the parent is a 4-node)

no path lengths have changed
(if it was **perfectly balanced**, it still is)

SPLITTING 4-nodes



BEFORE




AFTER

The extra key is pushed up to the parent
(so it won't work if the parent is a 4-node)

no path lengths have changed
(if it was **perfectly balanced**, it still is)

SPLIT takes $O(1)$ time

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation

SPLIT this!



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation

SPLIT this!



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation

SPLIT this!



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation

SPLIT this!



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation

SPLIT this!



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation

SPLIT this!



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,


Step 1: Search for the key k as if performing $\text{FIND}(k)$.

SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

OK, one more thing...

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

OK, one more thing...

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

OK, one more thing...

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

OK, one more thing... what happens when we SPLIT the root?

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node

Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

OK, one more thing... what happens when we SPLIT the root?

The INSERT operation



To perform $\text{INSERT}(x, k)$,

Step 1: Search for the key k as if performing $\text{FIND}(k)$.


SPLIT 4-nodes as we go down

Step 2: If the leaf is a 2-node,
insert (x, k) , converting it into a 3-node


Step 3: If the leaf is a 3-node,
insert (x, k) , converting it into a 4-node

OK, one more thing... what happens when we SPLIT the root?

The INSERT operation



The INSERT operation




SPLITTING the root increases the height of the tree
and increases the length of all root-leaf paths by one

So it maintains the **perfect balance** property

- i.e every path from the root to a leaf has the same length

The INSERT operation




SPLITTING the root increases the height of the tree
and increases the length of all root-leaf paths by one

So it maintains the **perfect balance** property

- i.e every path from the root to a leaf has the same length

This is the only way **INSERT** can affect the length of paths
so it also maintains the **perfect balance** property

The INSERT operation




SPLITTING the root increases the height of the tree
and increases the length of all root-leaf paths by one

So it maintains the **perfect balance** property
- i.e every path from the root to a leaf has the same length

This is the only way **INSERT** can affect the length of paths
so it also maintains the **perfect balance** property

As each **SPLIT** takes $O(1)$ time, overall **INSERT** takes $O(\log n)$ time

The INSERT operation



To perform INSERT(x , k),

Step 1: Search for the key k as if performing FIND(k).


SPLIT 4-nodes as we go down

Step 2: If the bottom node is a 2-node,
insert (x , k), converting it into a 3-node

Step 3: If the bottom node is a 3-node,
insert (x , k), converting it into a 4-node


As each SPLIT takes $O(1)$ time, overall INSERT takes $O(\log n)$ time

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)


The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).


The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.


The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.


The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).


The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

The DELETE operation




To perform $\text{DELETE}(k)$ **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using $\text{FIND}(k)$.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation




To perform $\text{DELETE}(k)$ **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using $\text{FIND}(k)$.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation




To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation




To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation




To perform **DELETE(k)** on a leaf (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation




To perform **DELETE(k)** on a leaf (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation




To perform **DELETE(k)** on a leaf (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

Step 4: If the leaf is a 2-node,

The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using **FIND(k)**.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

Step 4: If the leaf is a 2-node, ???

The DELETE operation



To perform $\text{DELETE}(k)$ **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using $\text{FIND}(k)$.

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node


Step 4: If the leaf is a 2-node, ??? *We will make sure this never happens*

FUSING 2-nodes




We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

FUSING 2-nodes




We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

BEFORE




AFTER

FUSING 2-nodes



We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node


BEFORE



AFTER

The extra key is pulled down from the parent
(so it won't work if the parent is a 2-node)


FUSING 2-nodes



We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

This is the opposite of a
SPLIT operation


BEFORE



AFTER

The extra key is pulled down from the parent
(so it won't work if the parent is a 2-node)


FUSING 2-nodes



We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

This is the opposite of a
SPLIT operation

BEFORE




AFTER

The extra key is pulled down from the parent
(so it won't work if the parent is a 2-node)

these subtrees haven't changed


FUSING 2-nodes



We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

This is the opposite of a
SPLIT operation

BEFORE




AFTER

The extra key is pulled down from the parent
(so it won't work if the parent is a 2-node)

no path lengths have changed

these subtrees haven't changed


FUSING 2-nodes



We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

This is the opposite of a
SPLIT operation

BEFORE



AFTER


The extra key is pulled down from the parent
(so it won't work if the parent is a 2-node)

no path lengths have changed

(if it was **perfectly balanced**, it still is)

these subtrees haven't changed


FUSING 2-nodes



We can **FUSE** two 2-nodes (with the same parent)
into a 4-node
if that parent isn't a 2-node

This is the opposite of a
SPLIT operation

BEFORE




AFTER

The extra key is pulled down from the parent
(so it won't work if the parent is a 2-node)

no path lengths have changed
(if it was **perfectly balanced**, it still is)

FUSE takes $O(1)$ time

TRANSFERRING keys




If there is a 2-node and a 3-node

(with the same parent)

we can perform a TRANSFER

(even if the parent is the root)

TRANSFERRING keys




If there is a 2-node and a 3-node

(with the same parent)

*we can perform a **TRANSFER***


(even if the parent is the root)

BEFORE



AFTER

TRANSFERRING keys




If there is a 2-node and a 3-node

(with the same parent)

*we can perform a **TRANSFER***

(even if the parent is the root)


BEFORE



AFTER

The keys have been rearranged

TRANSFERING keys




If there is a 2-node and a 3-node

(with the same parent)

*we can perform a **TRANSFER***

(even if the parent is the root)

BEFORE




AFTER

The keys have been rearranged

these subtrees haven't changed


TRANSFERING keys



*If there is a 2-node and a 3-node
(with the same parent)*

*we can perform a **TRANSFER**
(even if the parent is the root)*

BEFORE




AFTER

The keys have been rearranged

no path lengths have changed

these subtrees haven't changed


TRANSFERING keys



*If there is a 2-node and a 3-node
(with the same parent)*

*we can perform a **TRANSFER**
(even if the parent is the root)*

BEFORE



AFTER


The keys have been rearranged

no path lengths have changed

*(if it was **perfectly balanced**, it still is)*

these subtrees haven't changed


TRANSFERRING keys



*If there is a 2-node and a 3-node
(with the same parent)*

*we can perform a **TRANSFER**
(even if the parent is the root)*

BEFORE



AFTER


The keys have been rearranged

no path lengths have changed

*(if it was **perfectly balanced**, it still is)*

TRANSFER takes $O(1)$ time

TRANSFERRING keys



If there is a 2-node and a 3-node


(with the same parent)

we can perform a TRANSFER

(even if the parent is the root)

TRANSFER also works with a 2-node and a 4-node

BEFORE



AFTER

The keys have been rearranged


no path lengths have changed

*(if it was **perfectly balanced**, it still is)*

these subtrees haven't changed

TRANSFER takes $O(1)$ time

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using $\text{FIND}(k)$.

use FUSE and TRANSFER to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using **FIND(k)**.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using $\text{FIND}(k)$.

use FUSE and TRANSFER to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using **FIND(k)**.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using **FIND(k)**.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ on a leaf (we'll deal with other nodes later)

Step 1: Search for the key k using $\text{FIND}(k)$.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE(k)** **on a leaf** (*we'll deal with other nodes later*)


Step 1: Search for the key k using **FIND(k)**.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,
delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,
delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ on a leaf (we'll deal with other nodes later)

Step 1: Search for the key k using $\text{FIND}(k)$.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ on a leaf (we'll deal with other nodes later)

Step 1: Search for the key k using $\text{FIND}(k)$.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ on a leaf (we'll deal with other nodes later)

Step 1: Search for the key k using $\text{FIND}(k)$.

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform $\text{DELETE}(k)$ **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using $\text{FIND}(k)$.

use FUSE and TRANSFER to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,


delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation

Now we can DELETE the 5



To perform $\text{DELETE}(k)$ on a leaf (we'll deal with other nodes later)

Step 1: Search for the key k using $\text{FIND}(k)$.

use FUSE and TRANSFER to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,


delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation

Now we can DELETE the 5



To perform $\text{DELETE}(k)$ on a leaf (we'll deal with other nodes later)

Step 1: Search for the key k using $\text{FIND}(k)$.

use FUSE and TRANSFER to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down


Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,


delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

OK, one more thing...

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,

delete (x, k) , converting it into a 2-node


Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node

OK, one more thing... what happens when we **FUSE** the root?

FUSING the root


We said that we could only **FUSE** two 2-nodes if the parent was not a 2-node. . .
we make an exception for the root



FUSING the root can decrease the height of the tree
which in turn decreases the length of all root-leaf paths by one

FUSING the root

We said that we could only **FUSE** two 2-nodes if the parent was not a 2-node. . .
we make an exception for the root




FUSING the root can decrease the height of the tree
which in turn decreases the length of all root-leaf paths by one

So it maintains the **perfect balance** property
- *i.e every path from the root to a leaf has the same length*

FUSING the root

We said that we could only **FUSE** two 2-nodes if the parent was not a 2-node. . .
we make an exception for the root




FUSING the root can decrease the height of the tree
which in turn decreases the length of all root-leaf paths by one

So it maintains the **perfect balance** property
- *i.e every path from the root to a leaf has the same length*

This is the only way **DELETE** can affect the length of paths
so it also maintains the **perfect balance** property

FUSING the root

We said that we could only **FUSE** two 2-nodes if the parent was not a 2-node. . .
we make an exception for the root




FUSING the root can decrease the height of the tree
which in turn decreases the length of all root-leaf paths by one

So it maintains the **perfect balance** property
- i.e every path from the root to a leaf has the same length

This is the only way **DELETE** can affect the length of paths
so it also maintains the **perfect balance** property

As each **FUSE** or **TRANSFER** takes $O(1)$ time, overall **DELETE** takes $O(\log n)$ time

The DELETE operation



To perform **DELETE**(k) **on a leaf** (*we'll deal with other nodes later*)

Step 1: Search for the key k using **FIND**(k).

use **FUSE** and **TRANSFER** to convert 2-nodes as we go down

Step 2: If the leaf is a 3-node,


delete (x, k) , converting it into a 2-node

Step 3: If the leaf is a 4-node,

delete (x, k) , converting it into a 3-node


As each **FUSE** or **TRANSFER** takes $O(1)$ time, overall **DELETE** takes $O(\log n)$ time

The DELETE operation



What if we want to DELETE something other than a leaf?

The DELETE operation




What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

The DELETE operation




What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

The DELETE operation




What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

The DELETE operation




What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

The DELETE operation




What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

The DELETE operation




What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

The DELETE operation



What if we want to DELETE something other than a leaf?


Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k' such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf

The DELETE operation



What if we want to DELETE something other than a leaf?


Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k' such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf

The DELETE operation



What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)


- that's the element with the largest key k'
such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf

Step 3: Overwrite k with another copy of k'

The DELETE operation



What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)


- that's the element with the largest key k'
such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf

Step 3: Overwrite k with another copy of k'

The DELETE operation



What if we want to **DELETE** something other than a leaf?

Step 1: Find the **PREDECESSOR** of k (this is essentially the same as **FIND**)


- that's the element with the largest key k'
such that $k' < k$

Step 2: Call **DELETE**(k')

- fortunately k' is *always* a leaf

Step 3: Overwrite k with another copy of k'

The DELETE operation



What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)


- that's the element with the largest key k'
such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf

Step 3: Overwrite k with another copy of k'

The DELETE operation



What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)


- that's the element with the largest key k'
such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf

Step 3: Overwrite k with another copy of k'

The DELETE operation



What if we want to DELETE something other than a leaf?

Step 1: Find the PREDECESSOR of k (this is essentially the same as FIND)

- that's the element with the largest key k'
such that $k' < k$

Step 2: Call DELETE(k')

- fortunately k' is *always* a leaf


Step 3: Overwrite k with another copy of k'

This also takes $O(\log n)$ time

2-3-4 tree summary

A 2-3-4 is a data structure based on a tree structure

which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$




each of these operations takes *worst case* $O(\log n)$ time

2-3-4 tree summary

A 2-3-4 is a data structure based on a tree structure

which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$




each of these operations takes *worst case* $O(\log n)$ time

*Unfortunately, 2-3-4 trees are awkward to implement
because the nodes don't all have the same number of children*

2-3-4 tree summary

A 2-3-4 is a data structure based on a tree structure

which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$




each of these operations takes *worst case* $O(\log n)$ time

*Unfortunately, 2-3-4 trees are awkward to implement
because the nodes don't all have the same number of children*

So, what is used in practice?

Red-Black tree summary


A Red-Black tree is a data structure based on a **binary** tree structure which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



each of these operations takes *worst case* $O(\log n)$ time

Red-Black tree summary

A Red-Black tree is a data structure based on a **binary** tree structure
which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$




each of these operations takes *worst case* $O(\log n)$ time

The root is **black**

Red-Black tree summary

A Red-Black tree is a data structure based on a **binary** tree structure which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$




each of these operations takes *worst case* $O(\log n)$ time

The root is **black**

All root-to-leaf paths have the same number of **black** nodes

Red-Black tree summary

A Red-Black tree is a data structure based on a **binary** tree structure which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



each of these operations takes *worst case* $O(\log n)$ time


The root is **black**

All root-to-leaf paths have the same number of **black** nodes

Red nodes cannot have **red** children

Red-Black tree summary

A Red-Black tree is a data structure based on a **binary** tree structure which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



each of these operations takes *worst case* $O(\log n)$ time

The root is **black**


All root-to-leaf paths have the same number of **black** nodes

Red nodes cannot have **red** children

If these are used in practice, why did you waste our time with 2-3-4 trees?

Red-Black tree summary

A Red-Black tree is a data structure based on a **binary** tree structure which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



each of these operations takes *worst case* $O(\log n)$ time

The root is **black**

All root-to-leaf paths have the same number of **black** nodes

Red nodes cannot have **red** children


If these are used in practice, why did you waste our time with 2-3-4 trees?

1. 2-3-4 trees are conceptually much nicer

Red-Black tree summary

A Red-Black tree is a data structure based on a **binary** tree structure

which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



each of these operations takes *worst case* $O(\log n)$ time

The root is **black**

All root-to-leaf paths have the same number of **black** nodes



Red nodes cannot have **red** children

If these are used in practice, why did you waste our time with 2-3-4 trees?

1. 2-3-4 trees are conceptually much nicer
2. they are secretly the same :)

2-3-4 trees vs. Red-Black trees



Any 2-3-4 tree can be converted into a Red-Black tree (and visa-versa)



The operations on 2-3-4 trees also have equivalent operations on a Red-Black tree
(the details of the Red-Black tree operations are in CLRS Chapter 13)

2-3-4 trees vs. Red-Black trees



Any 2-3-4 tree can be converted into a Red-Black tree (and visa-versa)



The operations on 2-3-4 trees also have equivalent operations on a Red-Black tree
(the details of the Red-Black tree operations are in CLRS Chapter 13)

2-3-4 trees vs. Red-Black trees



Any 2-3-4 tree can be converted into a Red-Black tree (and visa-versa)



The operations on 2-3-4 trees also have equivalent operations on a Red-Black tree
(the details of the Red-Black tree operations are in CLRS Chapter 13)

2-3-4 trees vs. Red-Black trees


Any 2-3-4 tree can be converted into a Red-Black tree (and visa-versa)



The operations on 2-3-4 trees also have equivalent operations on a Red-Black tree
(the details of the Red-Black tree operations are in CLRS Chapter 13)

2-3-4 trees vs. Red-Black trees


Any 2-3-4 tree can be converted into a Red-Black tree (and visa-versa)




The operations on 2-3-4 trees also have equivalent operations on a Red-Black tree
(the details of the Red-Black tree operations are in CLRS Chapter 13)

2-3-4 trees vs. Red-Black trees

Any 2-3-4 tree can be converted into a Red-Black tree (and visa-versa)



You can think of a Red-Black tree as
a way to implement a 2-3-4 tree



The operations on 2-3-4 trees also have equivalent operations on a Red-Black tree
(the details of the Red-Black tree operations are in CLRS Chapter 13)

Dynamic Search Structure Summary

A **dynamic search structure** supports (at least) the following three operations

$\text{DELETE}(k)$ - deletes the (unique) element x with $x.\text{key} = k$

$\text{INSERT}(x, k)$ - inserts x with key $k = x.\text{key}$

$\text{FIND}(k)$ - returns the (unique) element x with $x.\text{key} = k$

Here are the *worst case* time complexities of the structures we have seen. . .

	INSERT	DELETE	FIND
Unsorted Linked List	$O(1)$	$O(n)$	$O(n)$
Binary Search Tree	$O(n)$	$O(n)$	$O(n)$
2-3-4 Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$
Red-Black Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$

End of part one

Part two

Skip lists

inspired by slides by Ashley Montanaro

Dynamic Search Structures

A **dynamic search structure**,
stores a set of elements

*Each element x must have a **unique** key - $x.key$*

The following operations are supported:

INSERT(x, k) - inserts x with key $k = x.key$

FIND(k) - returns the (unique) element x with $x.key = k$
(or reports that it doesn't exist)

DELETE(k) - deletes the (unique) element x with $x.key = k$
(or reports that it doesn't exist)

We would also like it to support:

PREDECESSOR(k) - returns the (unique) element x
with the largest key such that $x.key < k$

RANGEFIND(k_1, k_2) - returns every element x with $k_1 \leq x.key \leq k_2$

Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?



The bottleneck is **FIND**, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and **DELETE** also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the **FIND** operation?

Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the FIND operation?


Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the FIND operation?


Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the FIND operation?


Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the FIND operation?

Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the FIND operation?

Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the FIND operation?

Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?

FIND(18)



The bottleneck is FIND, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and DELETE also take $O(n)$ time *but only because they rely on FIND*


How can we speed up the FIND operation?

Using a Linked List as a Dynamic Search Structure (again)

Earlier we briefly considered using an unsorted Linked List as a dynamic search structure



What about using a *sorted* Linked List?



The bottleneck is **FIND**, which is *very inefficient*,

- we have to look through the entire linked list to find an item
(in the worst case)

INSERT and **DELETE** also take $O(n)$ time *but only because they rely on FIND*

How can we speed up the **FIND** operation?


Making Shortcuts

How about adding some shortcuts?




Making Shortcuts

How about adding some shortcuts?




Making Shortcuts

How about adding some shortcuts?



Making Shortcuts


How about adding some shortcuts?



Making Shortcuts

How about adding some shortcuts?


We've attached a second linked list containing only some of the keys. . .



Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
and go right until we come to a key $k' > k$
then we move down to the bottom list
and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform **FIND(k)** we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform `FIND(k)` we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform `FIND(k)` we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform **FIND(k)** we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

How long does this take?

Making Shortcuts

How about adding some shortcuts?

We've attached a second linked list containing only some of the keys. . .




To perform $\text{FIND}(k)$ we start in the top list
 and go right until we come to a key $k' > k$
 then we move down to the bottom list
 and go right until we find k

How long does this take?

That depends on where we place the shortcuts

Linked Lists with two levels


Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)

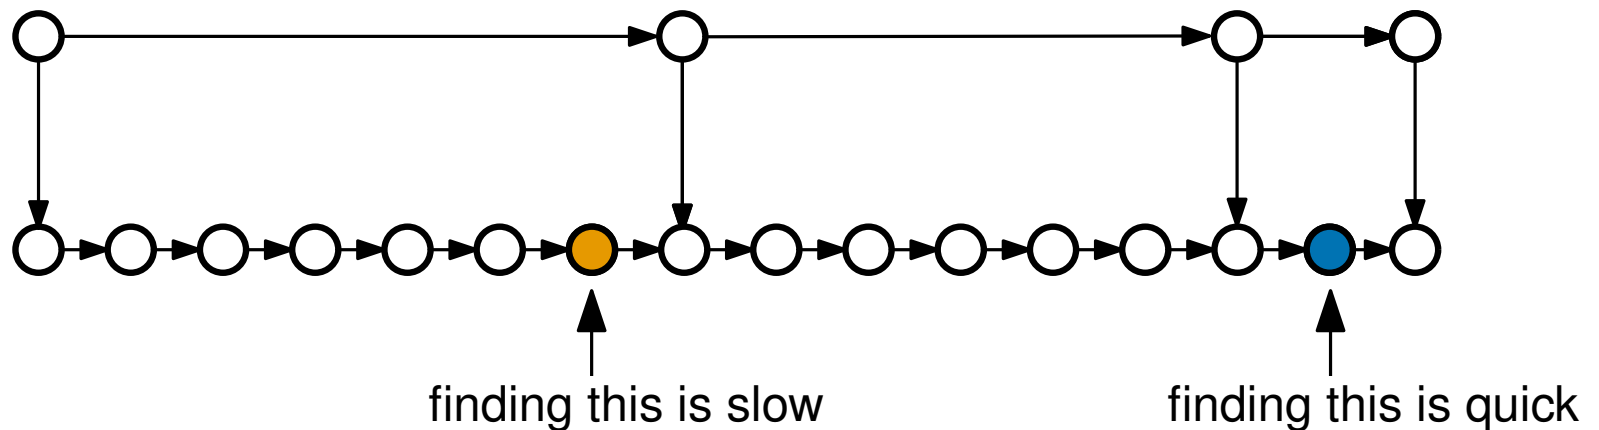
where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)


where should we put them to minimise
the *worst case* time for a **FIND** operation?



Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)

where should we put them to minimise
the *worst case* time for a **FIND** operation?




If we spread out the m keys in the top list evenly . . .

Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)

where should we put them to minimise
the *worst case* time for a **FIND** operation?




If we spread out the m keys in the top list evenly . . .

the *worst case* time for a **FIND** operation becomes $O(m + n/m)$

Linked Lists with two levels

Imagine that we decide to place m keys in the top list. . .
(the bottom list always contains all n keys)

where should we put them to minimise
the *worst case* time for a **FIND** operation?




If we spread out the m keys in the top list evenly . . .

the *worst case* time for a **FIND** operation becomes $O(m + n/m)$

Linked Lists with two levels

Imagine that we decide to place m keys in the top list...
(the bottom list always contains all n keys)

where should we put them to minimise
the *worst case* time for a **FIND** operation?




If we spread out the m keys in the top list evenly ...

the *worst case* time for a **FIND** operation becomes $O(m + n/m)$

Linked Lists with two levels

Imagine that we decide to place m keys in the top list...
(the bottom list always contains all n keys)

where should we put them to minimise
the *worst case* time for a **FIND** operation?



If we spread out the m keys in the top list evenly ...

the *worst case* time for a **FIND** operation becomes $O(m + n/m)$

By setting $m = \sqrt{n}$, we get the *worst case* time for a **FIND** operation is $O(\sqrt{n})$

Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below


They are chosen to be as evenly spread as possible

Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible




Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible




Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible




Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible




Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible




Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible




Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible

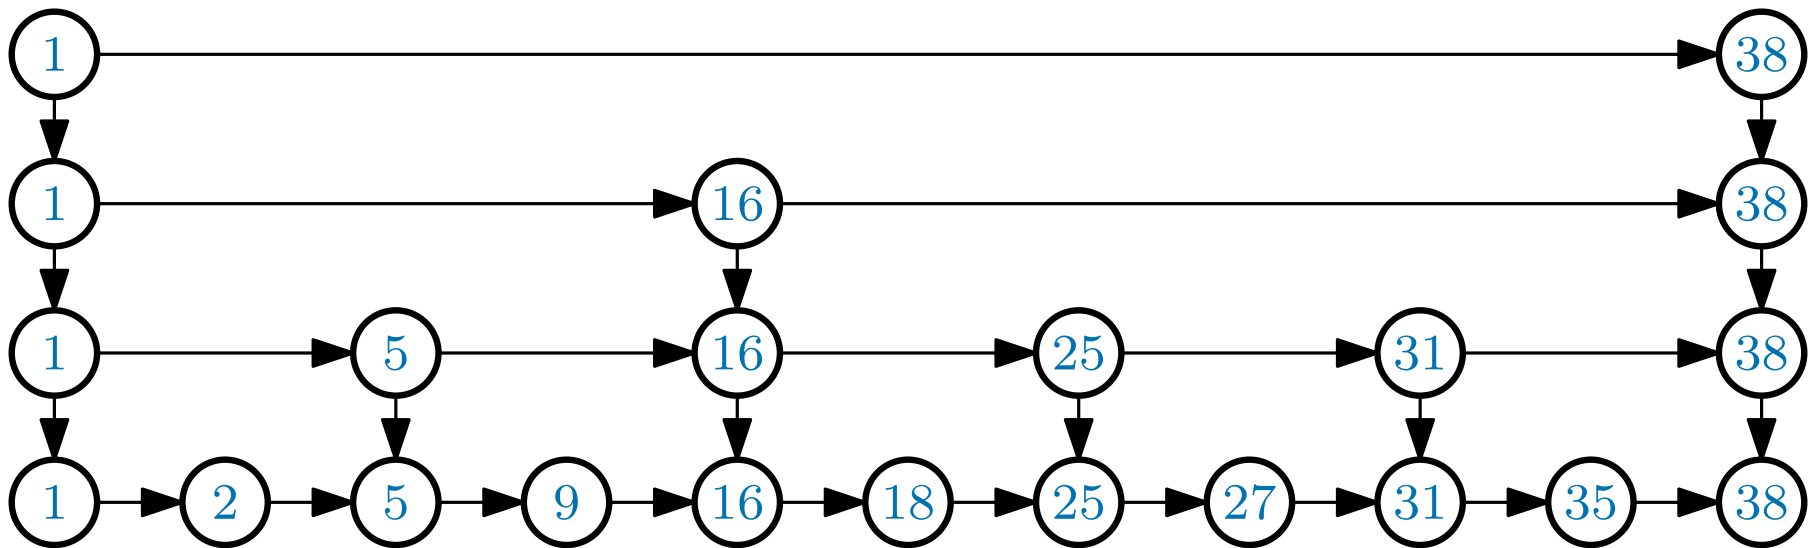


Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible



The bottom level contains every key


and every level contains the leftmost and rightmost keys

Linked Lists with many levels

How about adding even more lists? *(each list is called a **level**)*

Each **level** will now contain **half** of the keys *(rounding up)* from the level below

They are chosen to be as evenly spread as possible



The bottom level contains every key

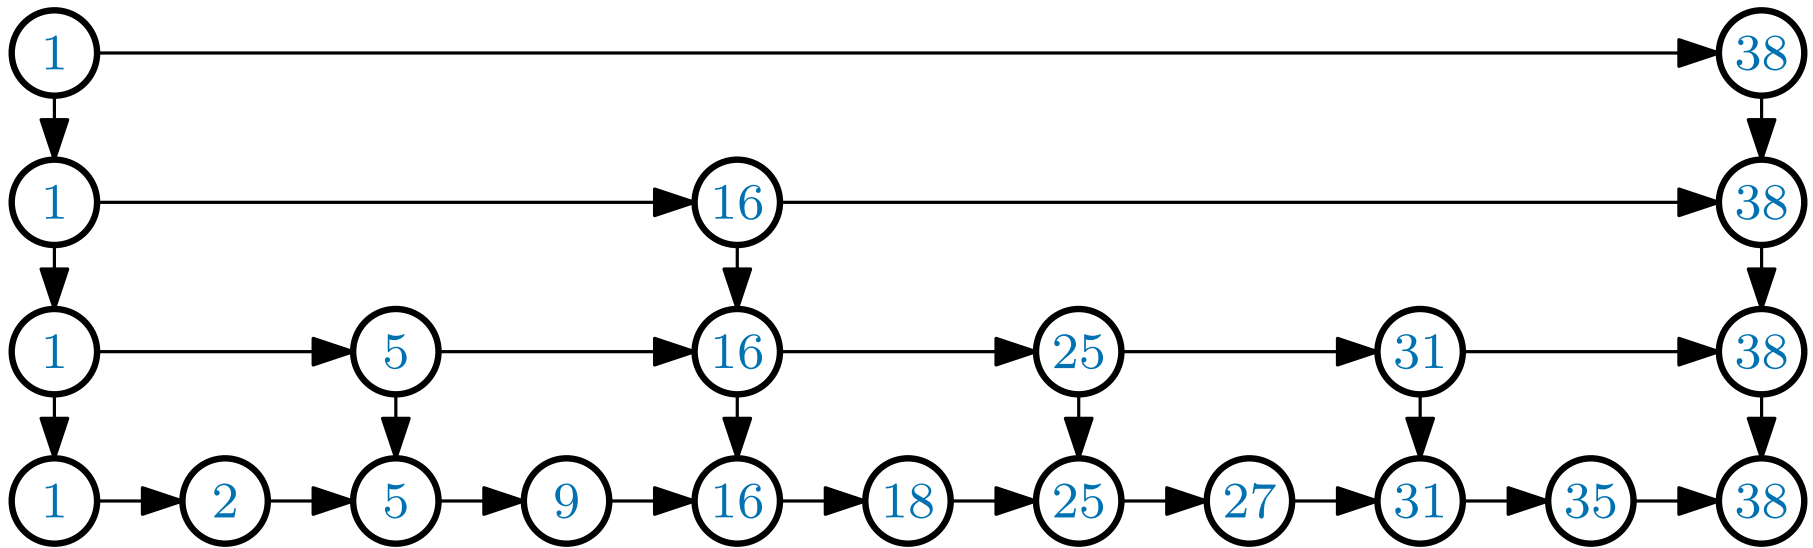
and every level contains the leftmost and rightmost keys

As each level contains half of the keys from the level below, there are $O(\log n)$ levels

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?


(essentially just like before)



FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left *(the head of the top level)*

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

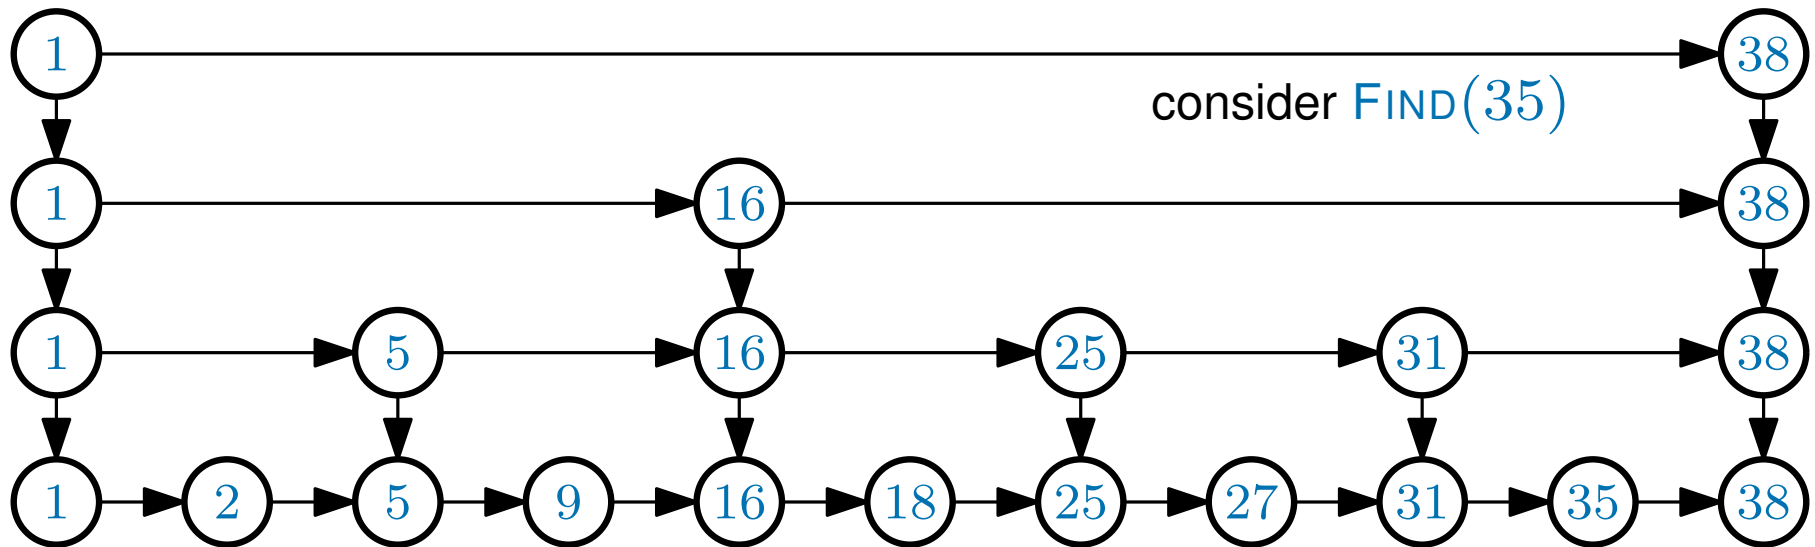
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left *(the head of the top level)*

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

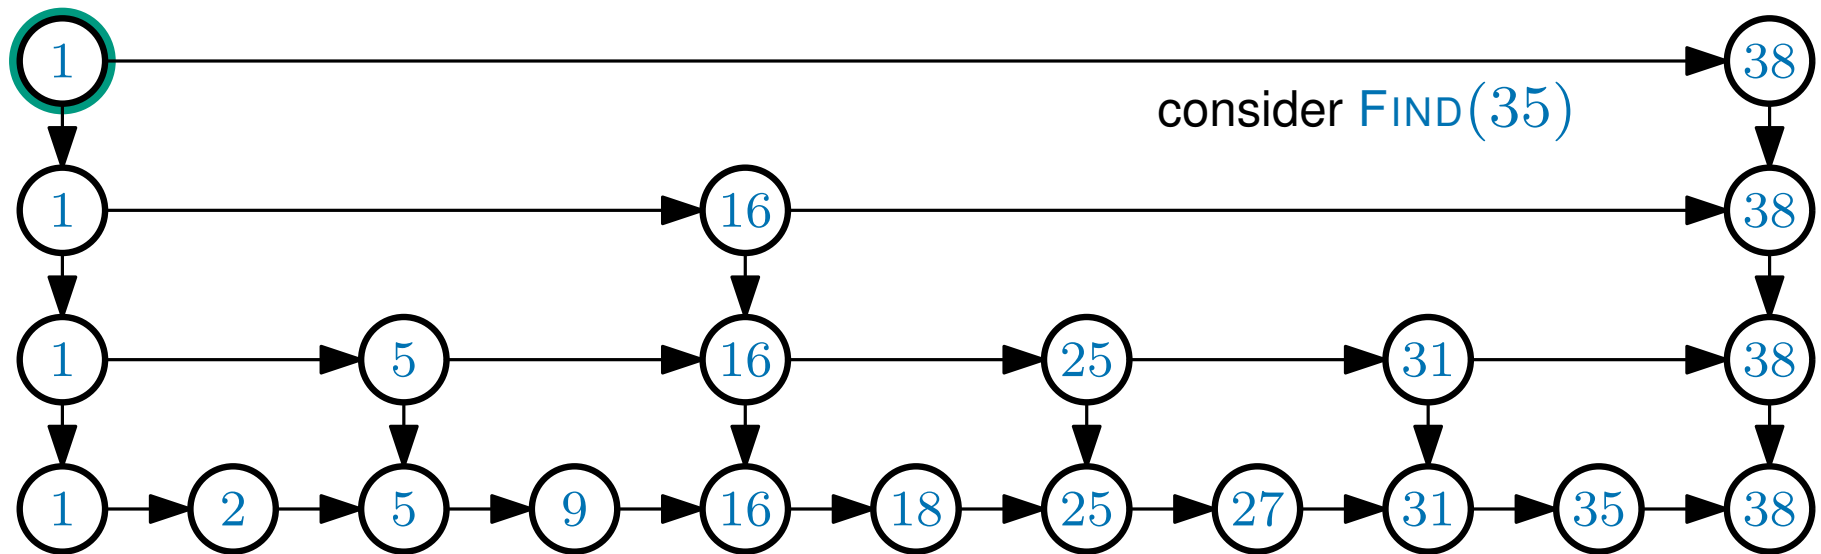
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (*the head of the top level*)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

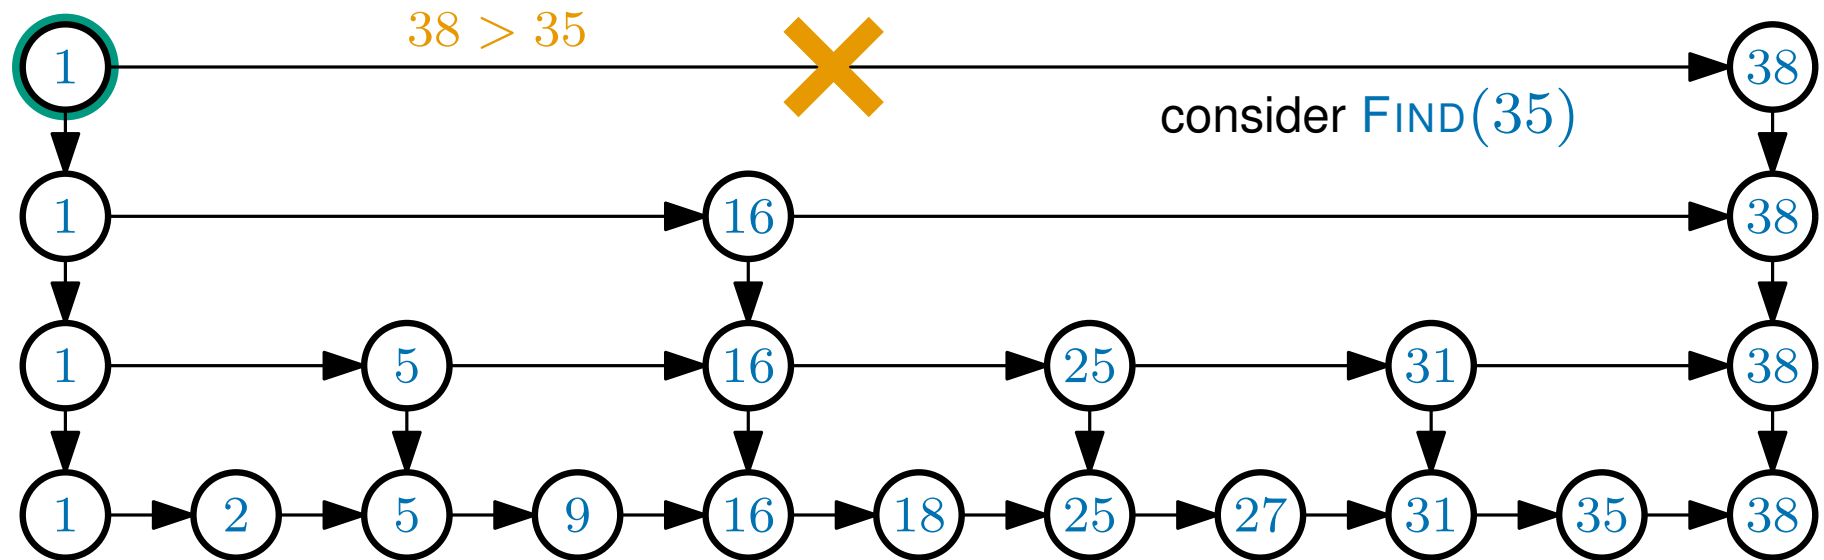
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

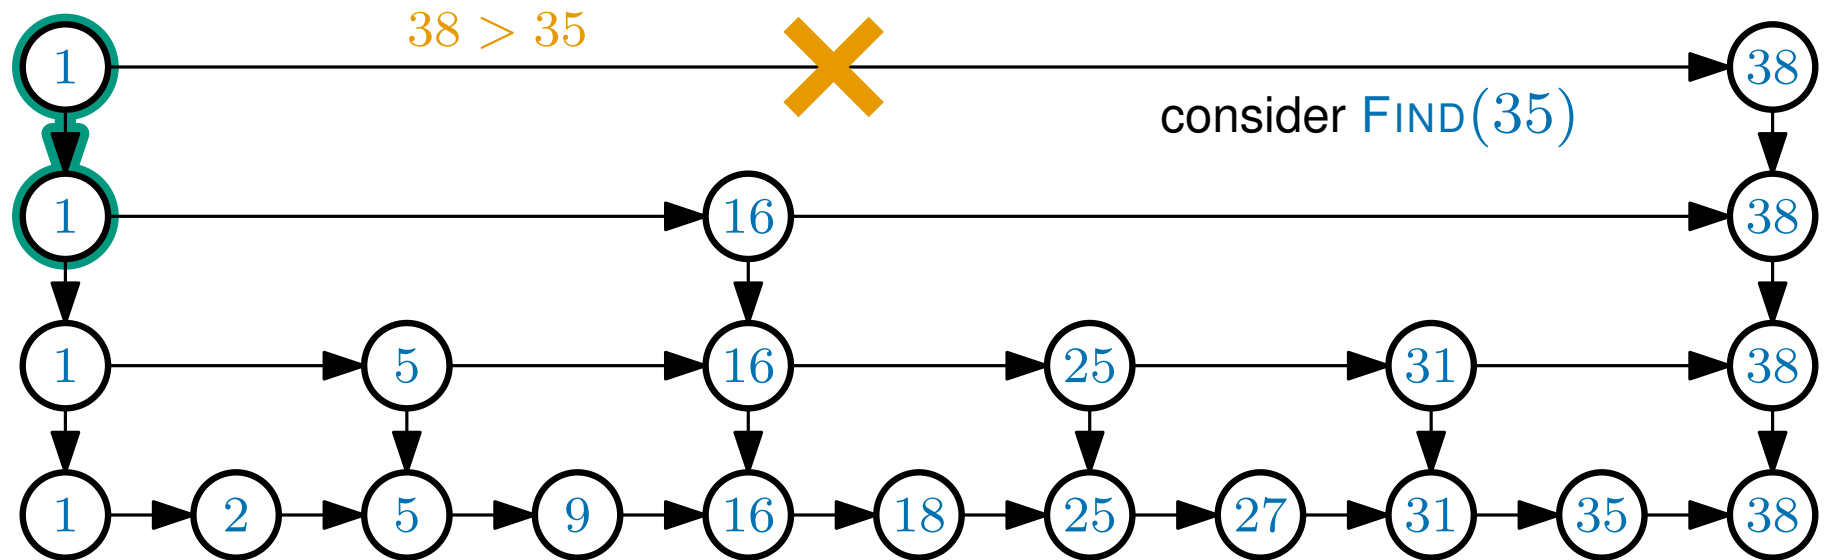
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right


Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

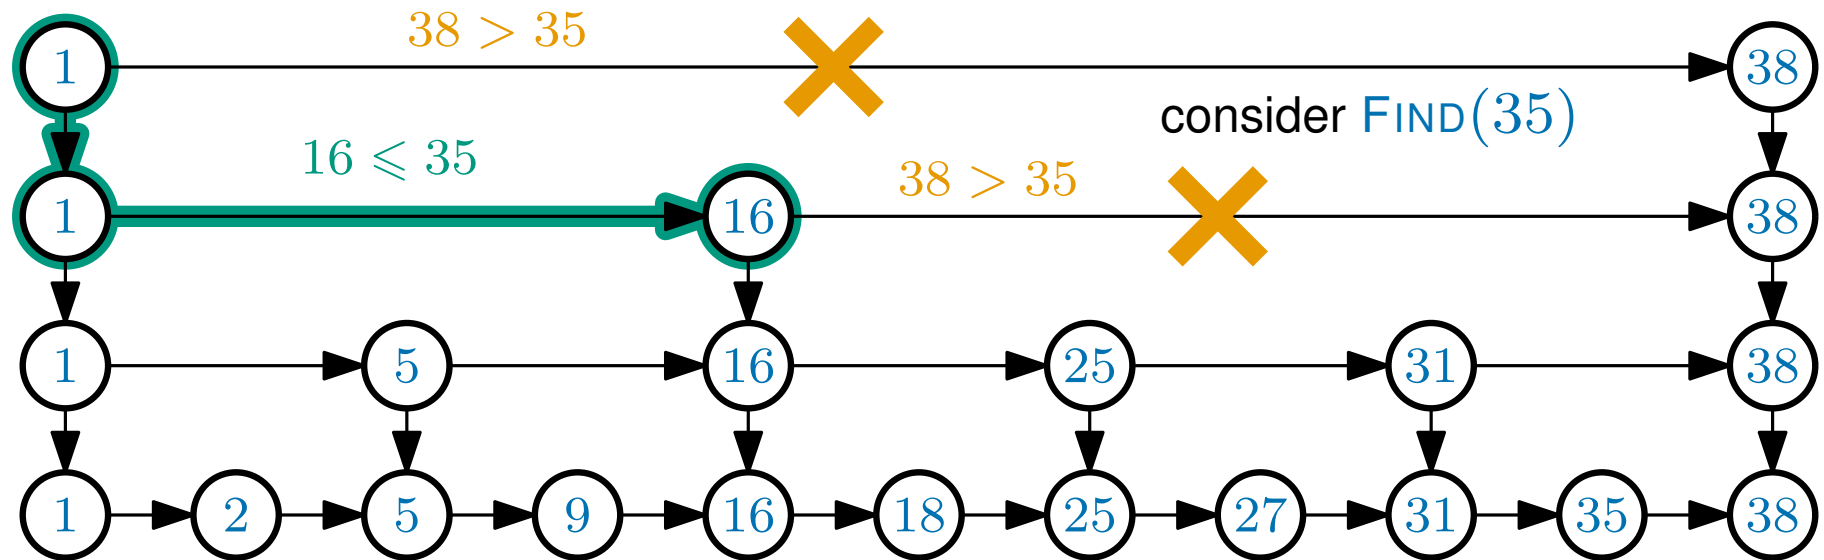
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right


Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (*the head of the top level*)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right


Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right


Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right


Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

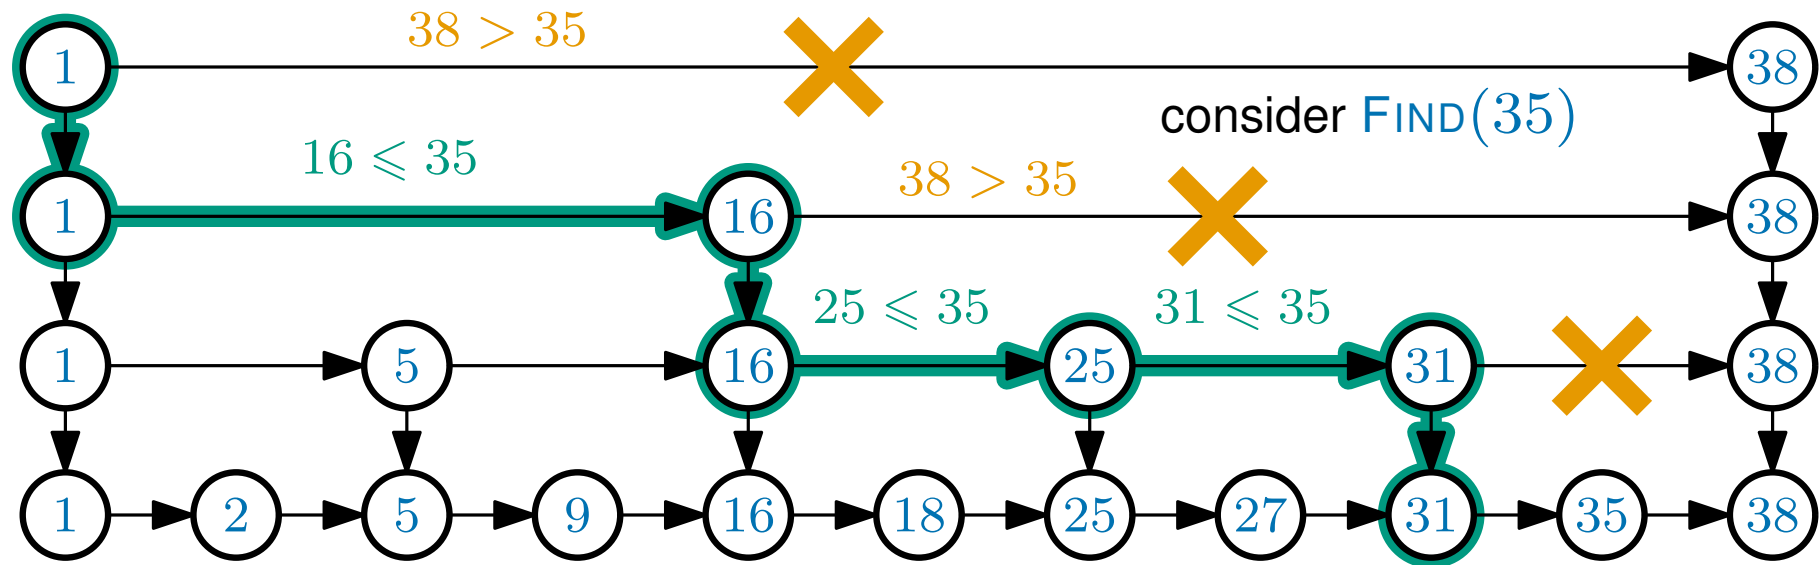
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

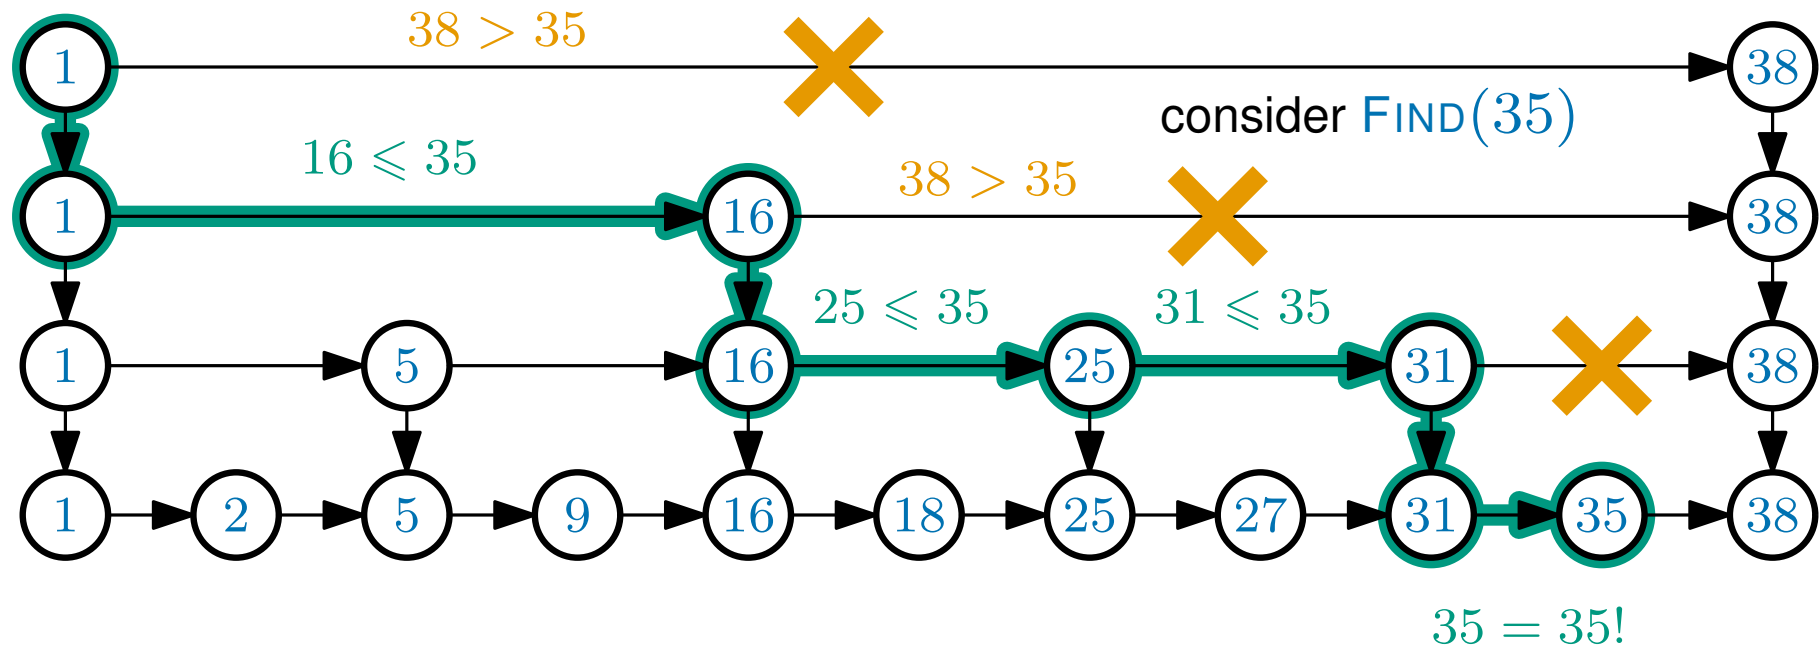
Else

Move down

FIND in multi-level linked lists

How do we perform $\text{FIND}(k)$ in multi-level linked list?

(essentially just like before)



To perform $\text{FIND}(k)$,

Start at the top-left (the head of the top level)

While you haven't found k :

If the node to the right's key, $k' \leq k$

Move right

Else

Move down

The complexity of FIND

How long does $\text{FIND}(k)$ take in a multi-level linked list?

The complexity of FIND

How long does $\text{FIND}(k)$ take in a multi-level linked list?


Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right** **at most 2 times** on any level $i + 1$
because we stopped moving right on level i

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right** **at most 2 times** on any level $i + 1$
*because we stopped **moving right** on level i*

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right at most 2 times** on any level $i + 1$
because we stopped moving right on level i

The complexity of FIND

How long does $\text{FIND}(k)$ take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right** **at most 2 times** on any level $i + 1$
because we stopped moving right on level i

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right** **at most 2 times** on any level $i + 1$
*because we stopped **moving right** on level i*

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right** **at most 2 times** on any level $i + 1$
because we stopped moving right on level i

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly




Observation 3 We only **move right** at most 2 times on any level $i + 1$
because we stopped moving right on level i

The complexity of **FIND**

How long does **FIND**(k) take in a multi-level linked list?

Observation 1 We only **move down** at most $O(\log n)$ times
because there are only $O(\log n)$ levels

Observation 2 Between any two nodes on level i ,
 there are **at most 2 nodes** on level $i + 1$
because we took half the nodes and spread them evenly



Observation 3 We only **move right** at most 2 times on any level $i + 1$
because we stopped moving right on level i


Fact We only **move** at most $O(\log n)$ times while performing a **FIND**

Multi-level Linked Lists

If we had a **multi-level linked list** with $O(\log n)$ levels

where each level contained **half** of the keys **from the level below**
and the keys were **evenly spread** as possible

then we could perform **FIND** in $O(\log n)$ time




Multi-level Linked Lists

If we had a **multi-level linked list** with $O(\log n)$ levels

where each level contained **half** of the keys **from the level below**
and the keys were **evenly spread** as possible


then we could perform **FIND** in $O(\log n)$ time



How are we going to do **INSERTS** and **DELETES**?

Multi-level Linked Lists

If we had a **multi-level linked list** with $O(\log n)$ levels
 where each level contained **half** of the keys **from the level below**
and the keys were **evenly spread** as possible
then we could perform **FIND** in $O(\log n)$ time




How are we going to do **INSERTS** and **DELETES**?

Which levels should we put an **INSERTED** key into?

Multi-level Linked Lists

If we had a **multi-level linked list** with $O(\log n)$ levels
 where each level contained **half** of the keys **from the level below**
and the keys were **evenly spread** as possible
then we could perform **FIND** in $O(\log n)$ time




How are we going to do **INSERTS** and **DELETES**?

Which levels should we put an **INSERTED** key into?

How can we keep a good spread of keys at each levels?

Multi-level Linked Lists

If we had a **multi-level linked list** with $O(\log n)$ levels
 where each level contained **half** of the keys **from the level below**
 and the keys were **evenly spread** as possible
then we could perform **FIND** in $O(\log n)$ time



How are we going to do **INSERTS** and **DELETES**?

Which levels should we put an **INSERTED** key into?

How can we keep a good spread of keys at each levels?

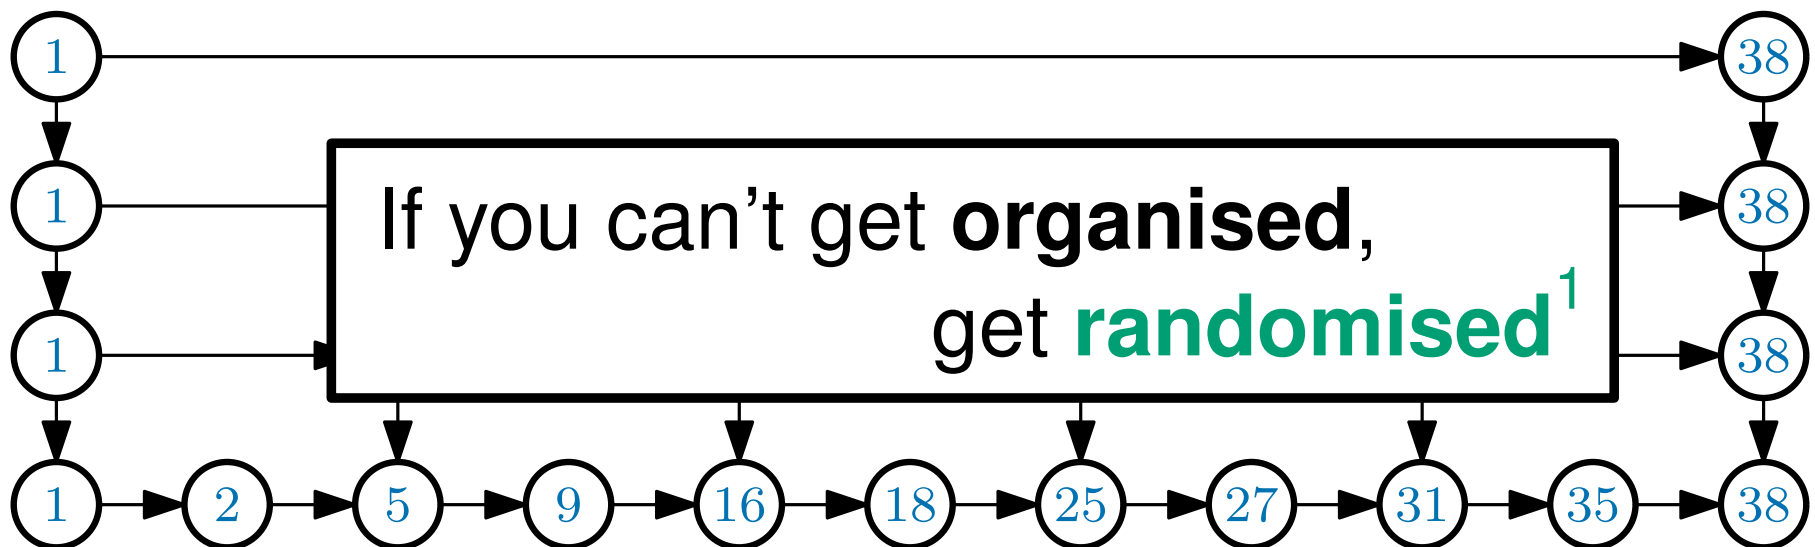
*especially when we don't know what will be **INSERTED** and **DELETED** in the future*

Multi-level Linked Lists

If we had a **multi-level linked list** with $O(\log n)$ levels

where each level contained **half** of the keys **from the level below**
and the keys were **evenly spread** as possible

then we could perform **FIND** in $O(\log n)$ time



How are we going to do **INSERTS** and **DELETES**?

Which levels should we put an **INSERTED** key into?

How can we keep a good spread of keys at each levels?

*especially when we don't know what will be **INSERTED** and **DELETED** in the future*

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List. . .
by flipping coins



(we still always include the smallest and largest keys in every level)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...

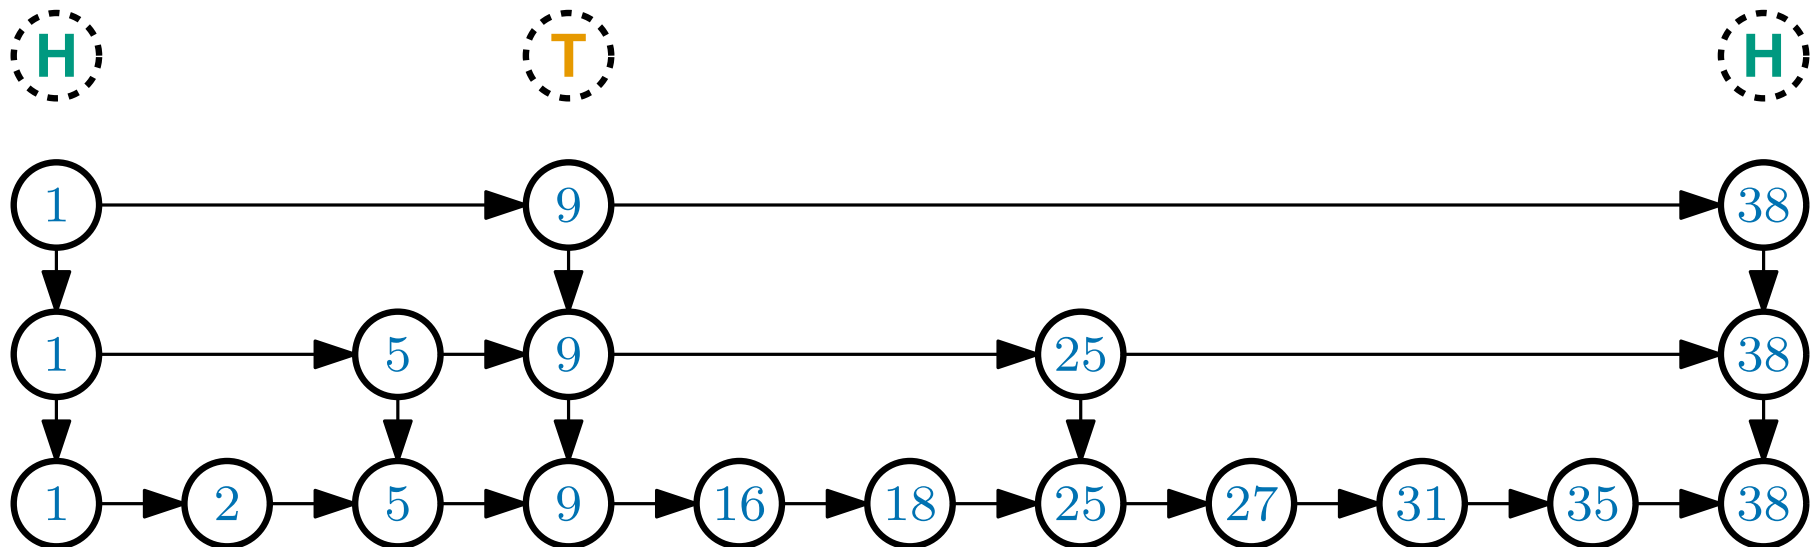
For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Flip one coin for each key...


For each key that got a head, put it in the new top level

Repeat with the keys from the new top level

(stop when the top level contains only the smallest and largest keys)

Building Multi-level Linked Lists by flipping coins


Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins




(we still always include the smallest and largest keys in every level)

This doesn't look quite perfect but actually, it's very good with high probability
(more on this later)

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins




(we still always include the smallest and largest keys in every level)

This doesn't look quite perfect but actually, it's very good with high probability
(more on this later)

The intuition is that n coin flips contain about $\frac{n}{2}$ heads and about $\frac{n}{2}$ tails

Building Multi-level Linked Lists by flipping coins

Before we formally introduce Skip Lists, we let's rewind
and try building another Multi-level Linked List...
by flipping coins



(we still always include the smallest and largest keys in every level)

This doesn't look quite perfect but actually, it's very good with high probability
(more on this later)


The intuition is that n coin flips contain about $\frac{n}{2}$ heads and about $\frac{n}{2}$ tails
and the heads are roughly evenly spread out

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...




Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform **INSERT**(x, k),

Step 1: Use **FIND**(k) to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

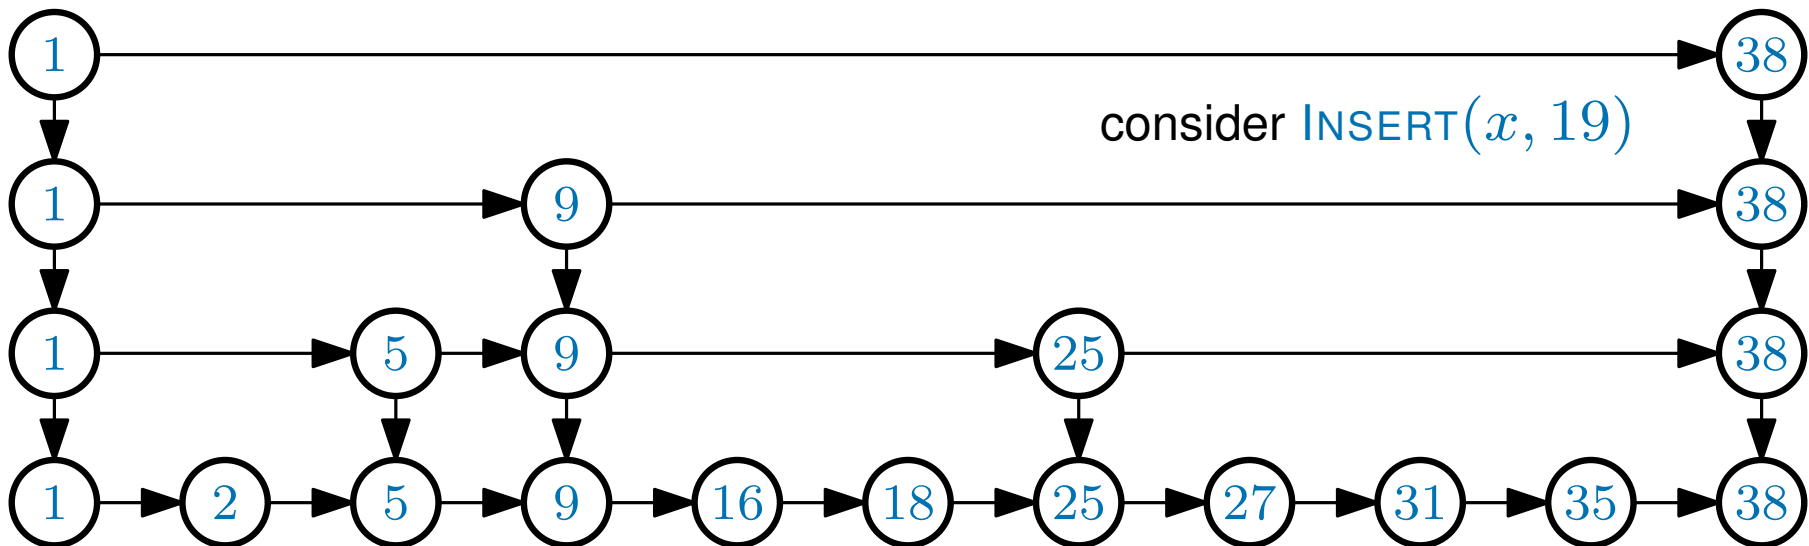
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

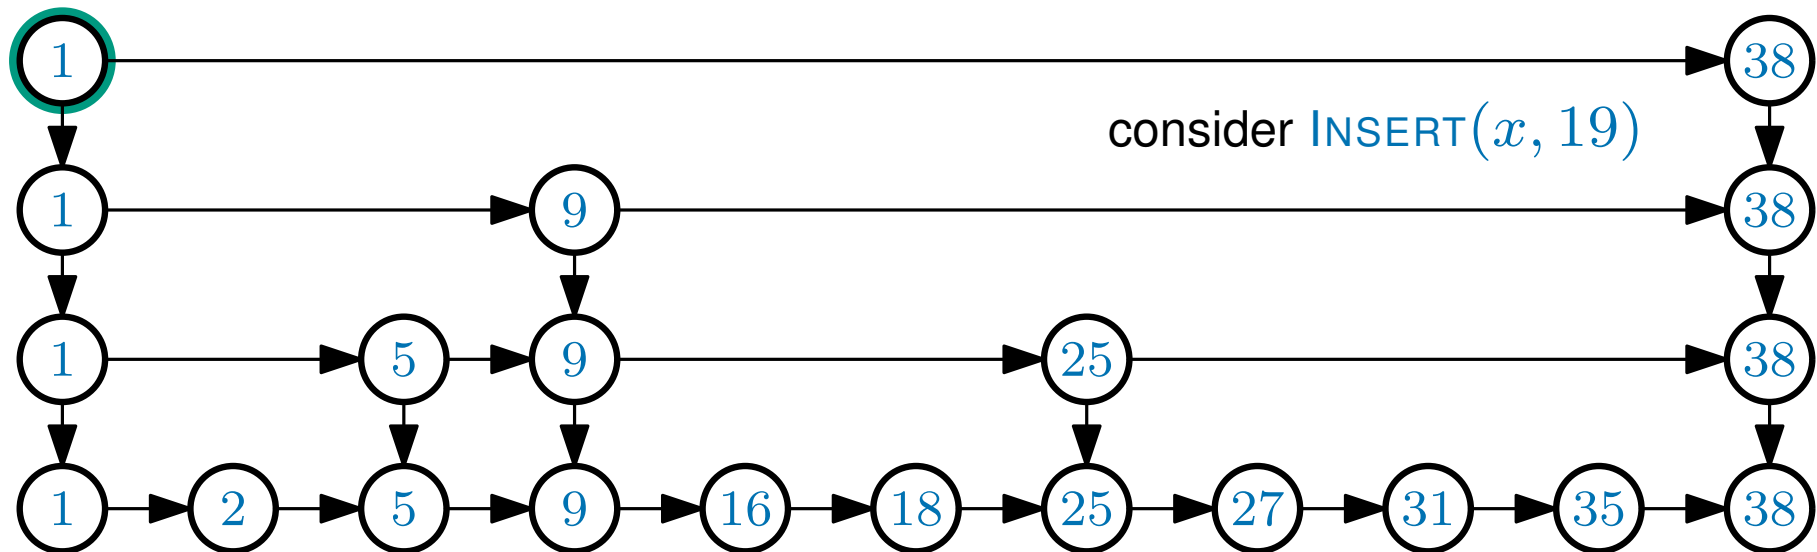
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

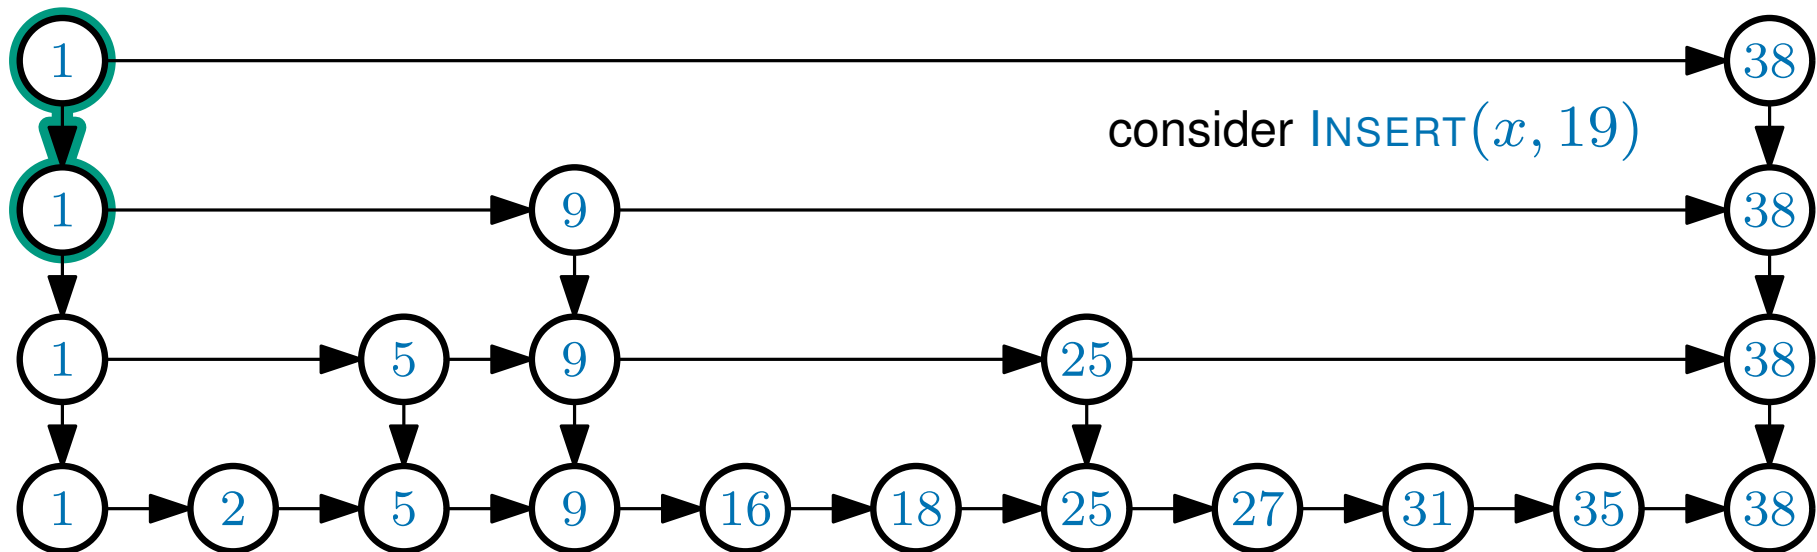
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

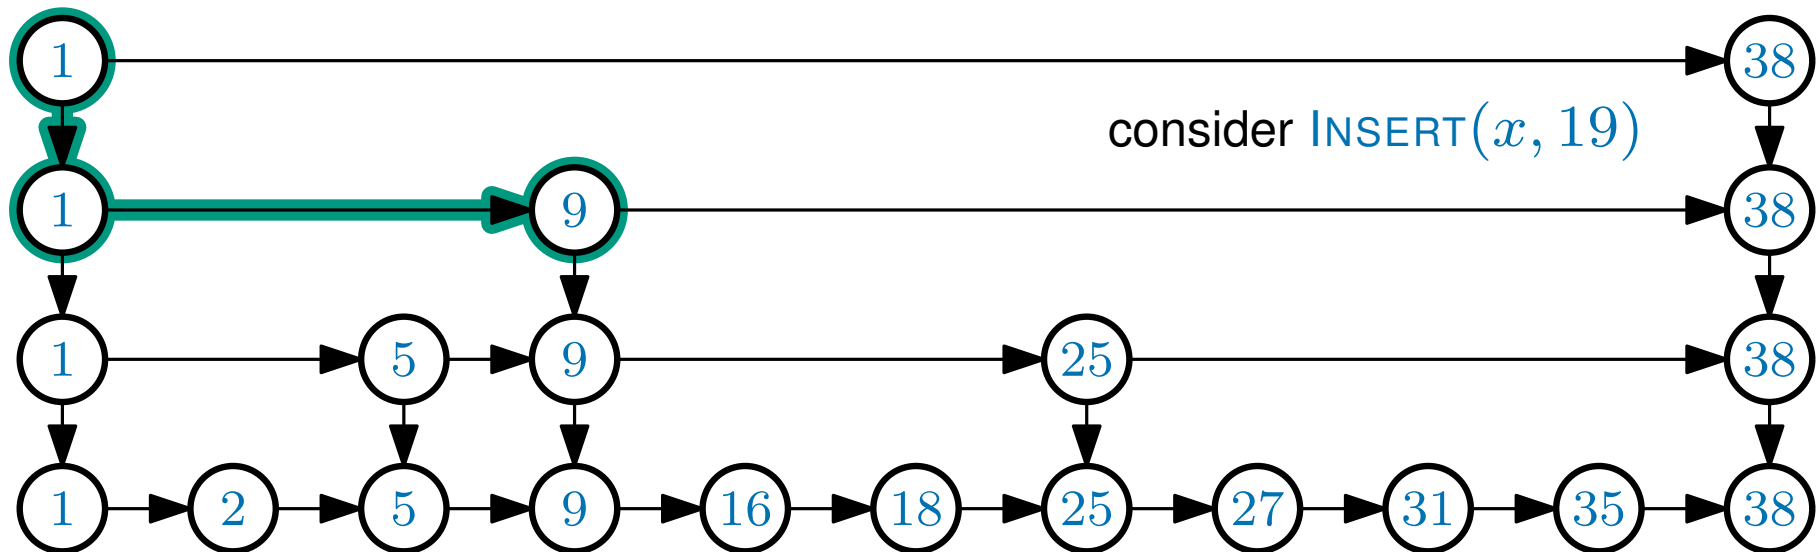
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform **INSERT**(x, k),

Step 1: Use **FIND**(k) to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

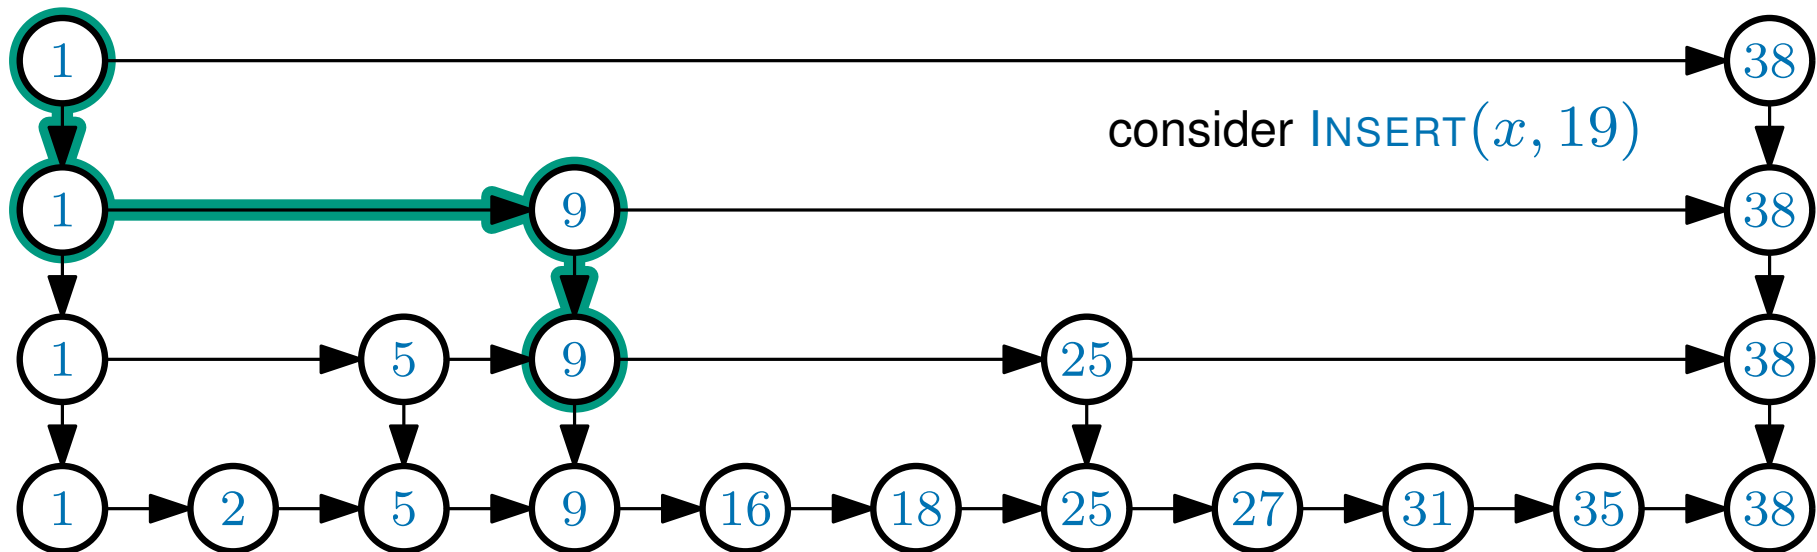
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

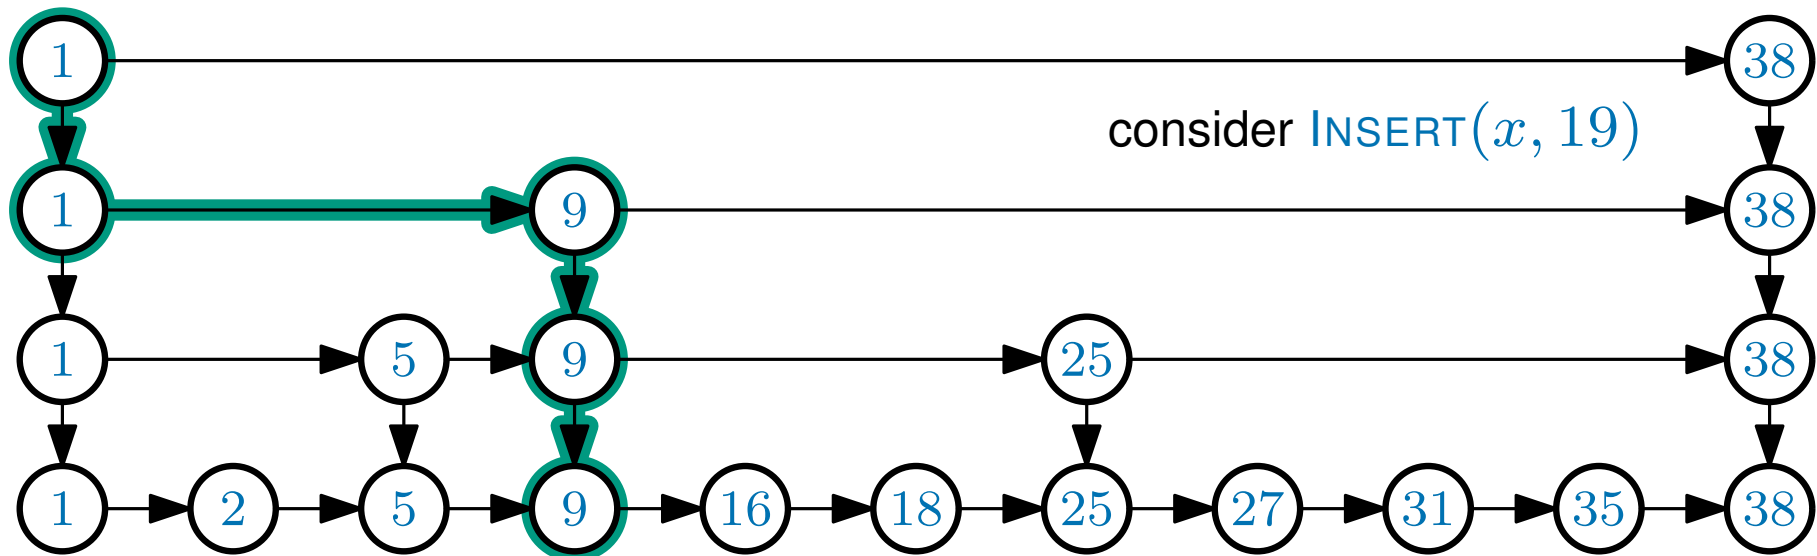
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

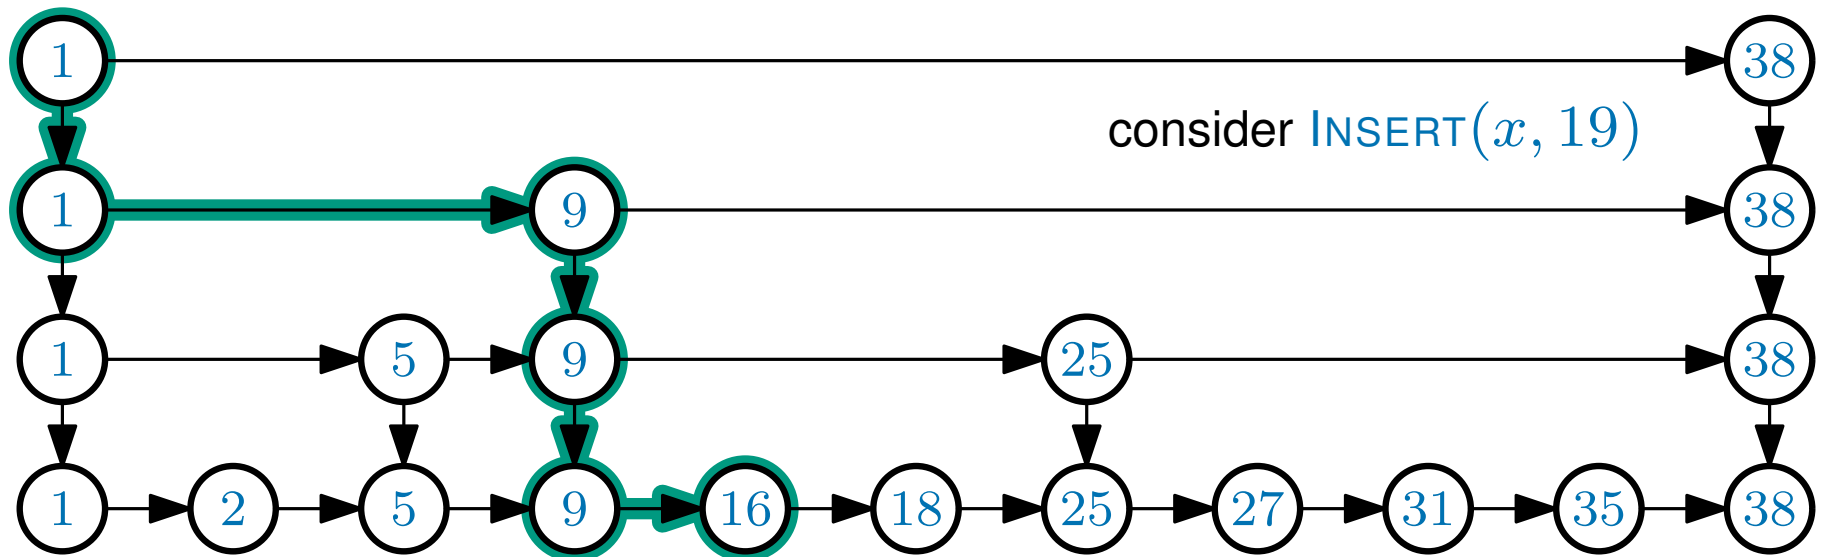
If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)


If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)


If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where

the **INSERTS** are done by flipping coins

i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:


If you get a **heads**, insert (x, k) into the next level up

(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level


Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level


Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

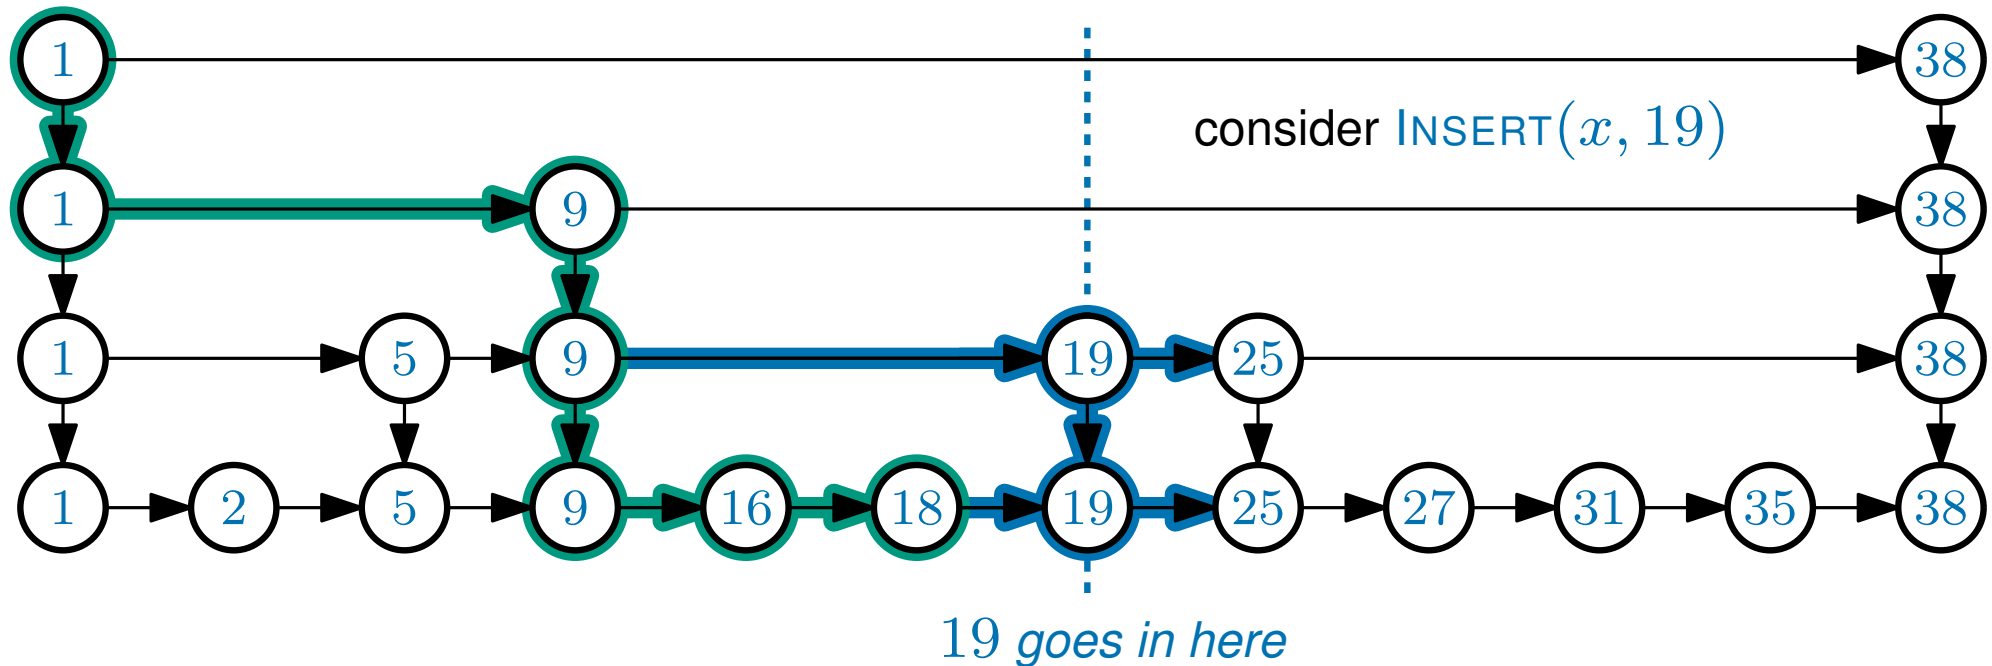
Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

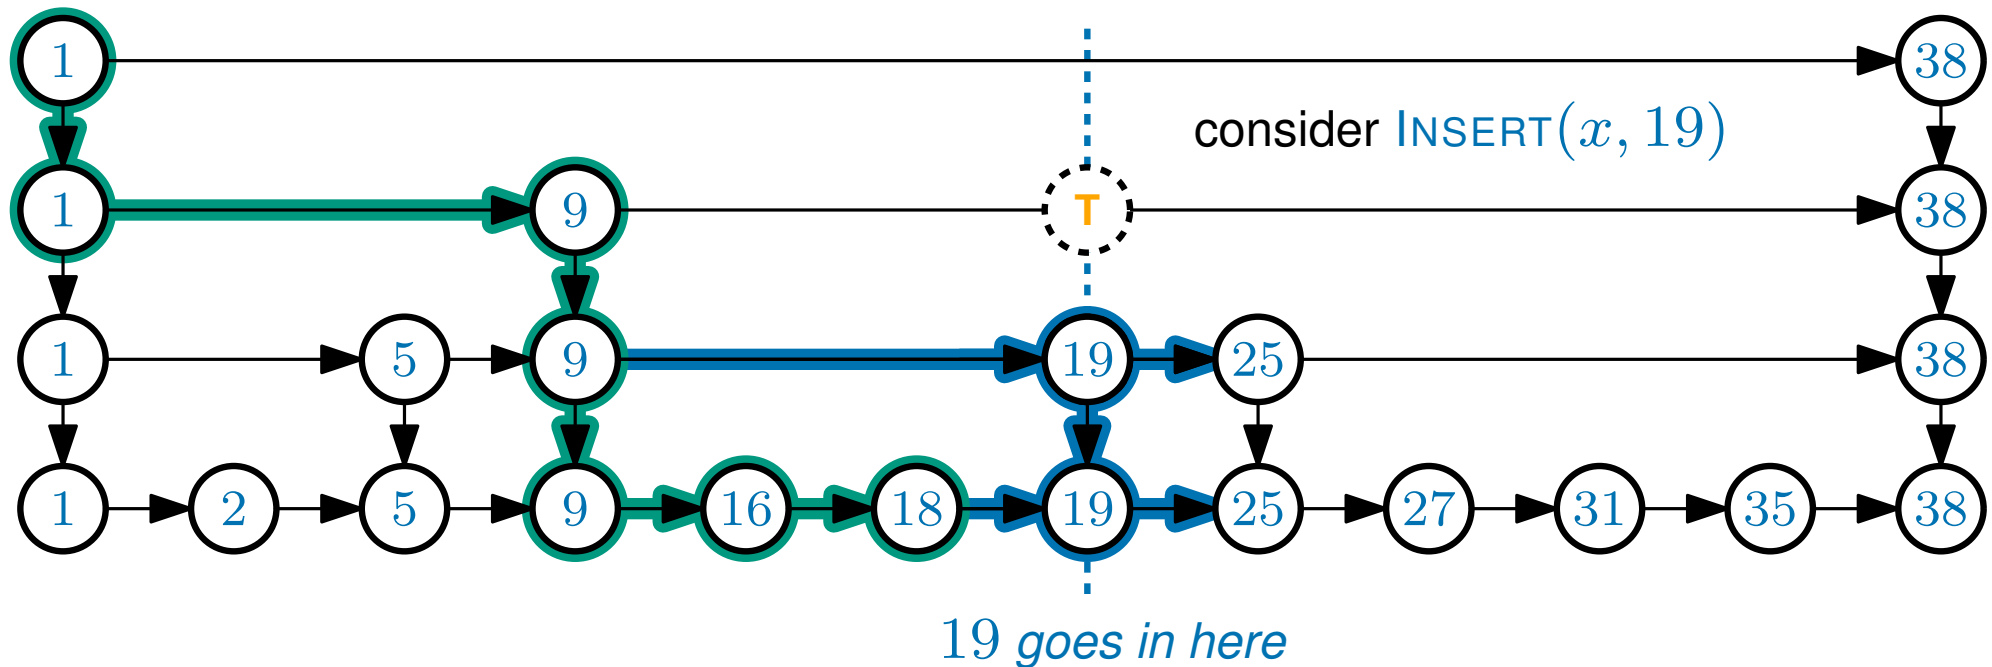
Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform $\text{INSERT}(x, k)$,

Step 1: Use $\text{FIND}(k)$ to insert (x, k) into the bottom level

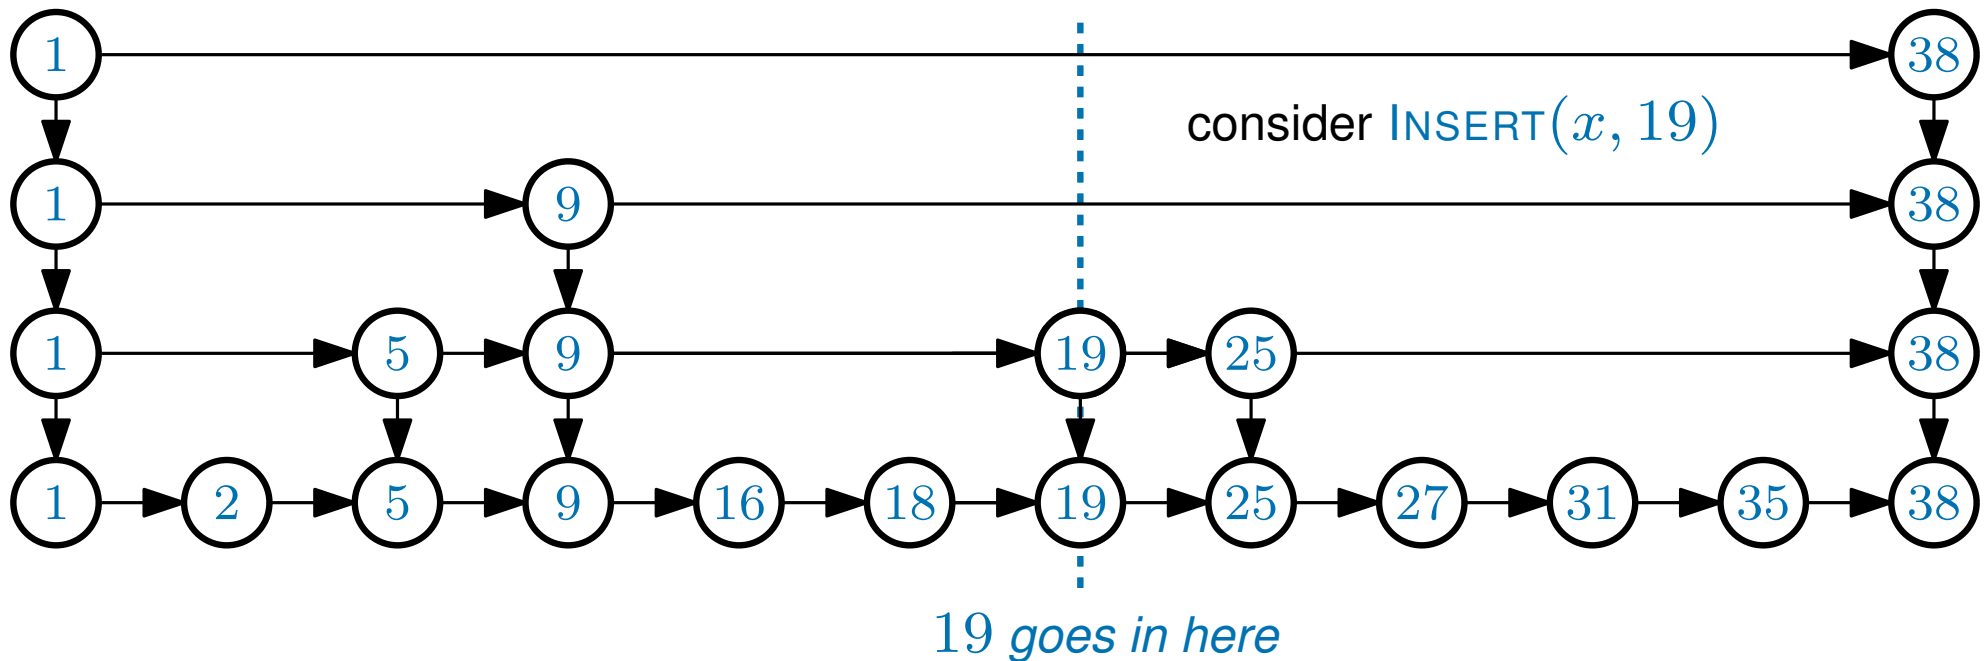
Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform **INSERT**(x, k),

Step 1: Use **FIND**(k) to insert (x, k) into the bottom level

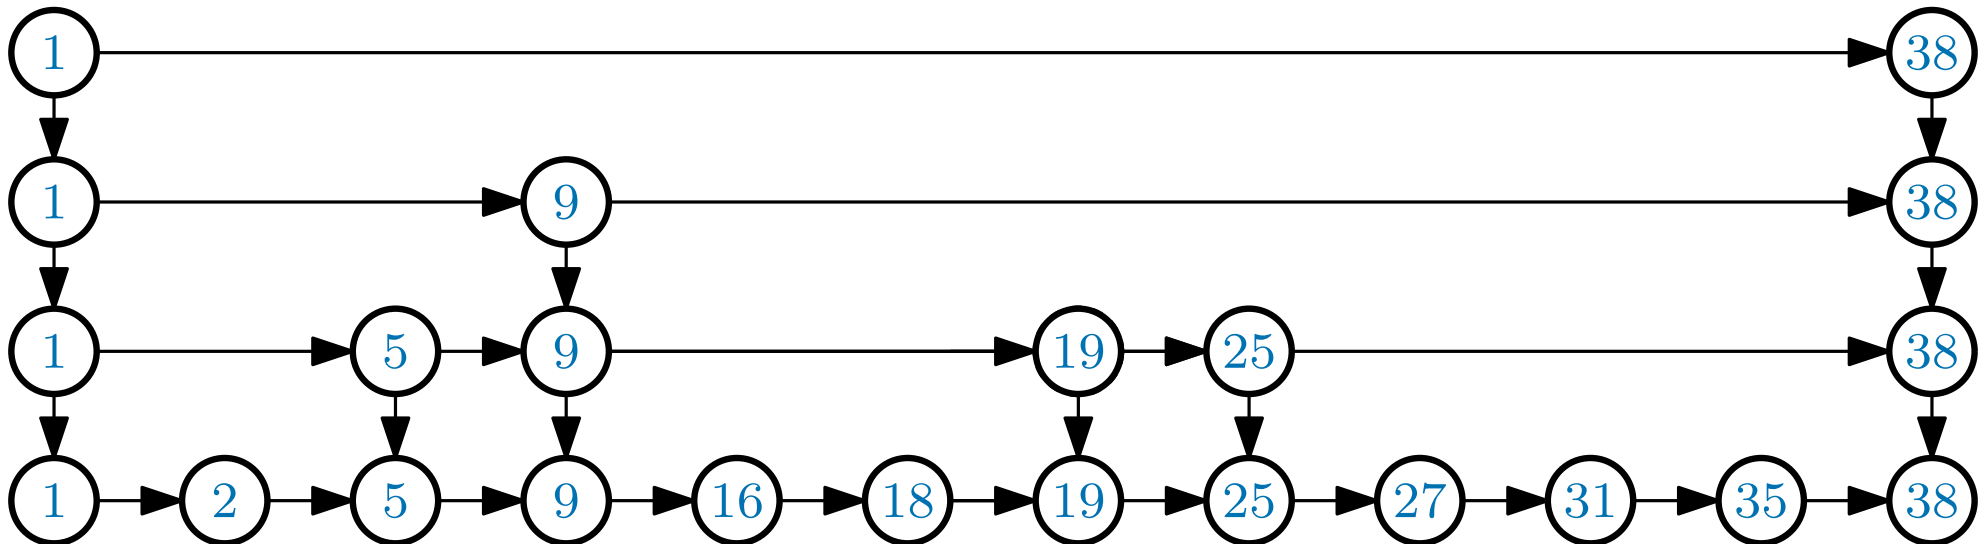
Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform **INSERT**(x, k),

Step 1: Use **FIND**(k) to insert (x, k) into the bottom level


Step 2: Flip a coin repeatedly:

If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

A skip list is a multi-level linked list where
the **INSERTS** are done by flipping coins
i.e. this is a skip list...



To perform **INSERT**(x, k),

Step 1: Use **FIND**(k) to insert (x, k) into the bottom level

Step 2: Flip a coin repeatedly:

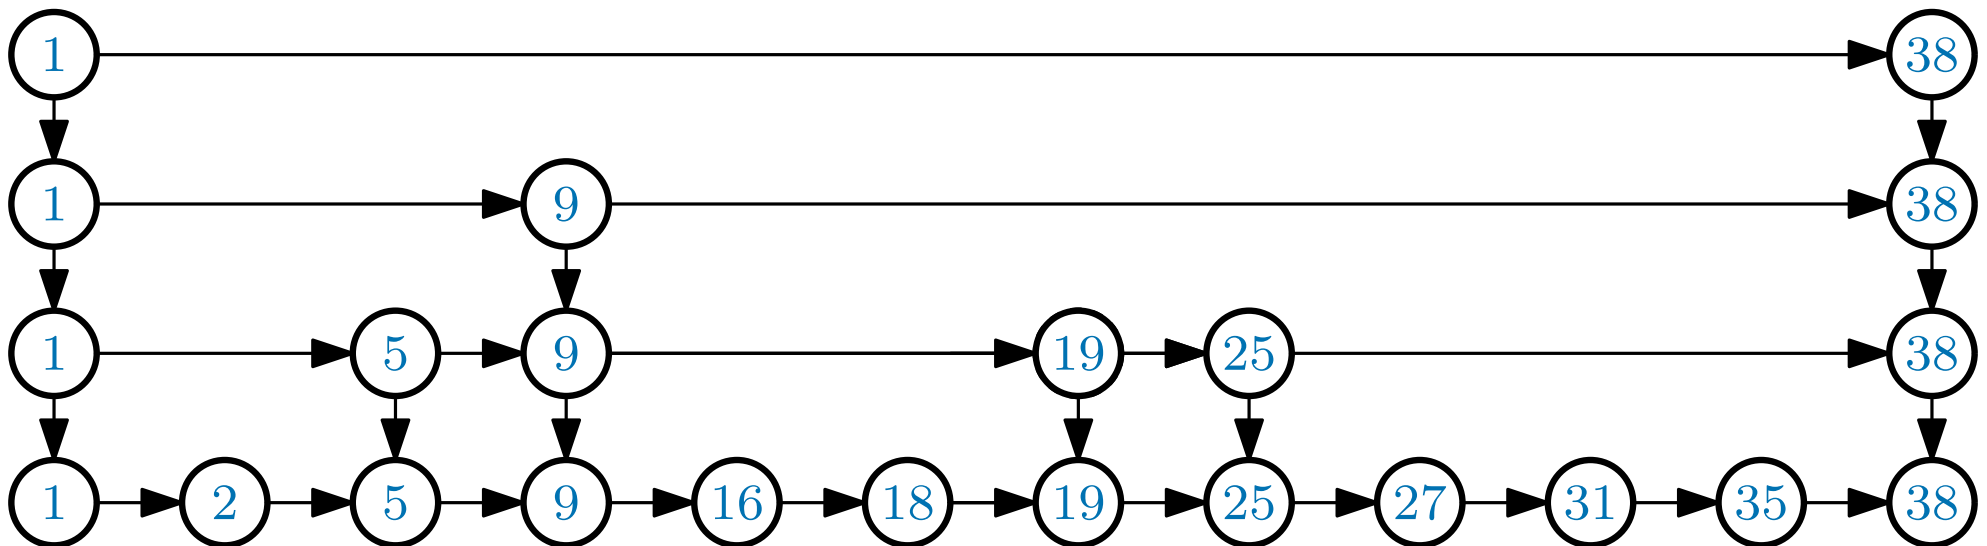
If you get a **heads**, insert (x, k) into the next level up
(if there is no 'next level up', create a new level at the top)

If you get a **tails**, **stop**

Skip Lists

That about **DELETES**?

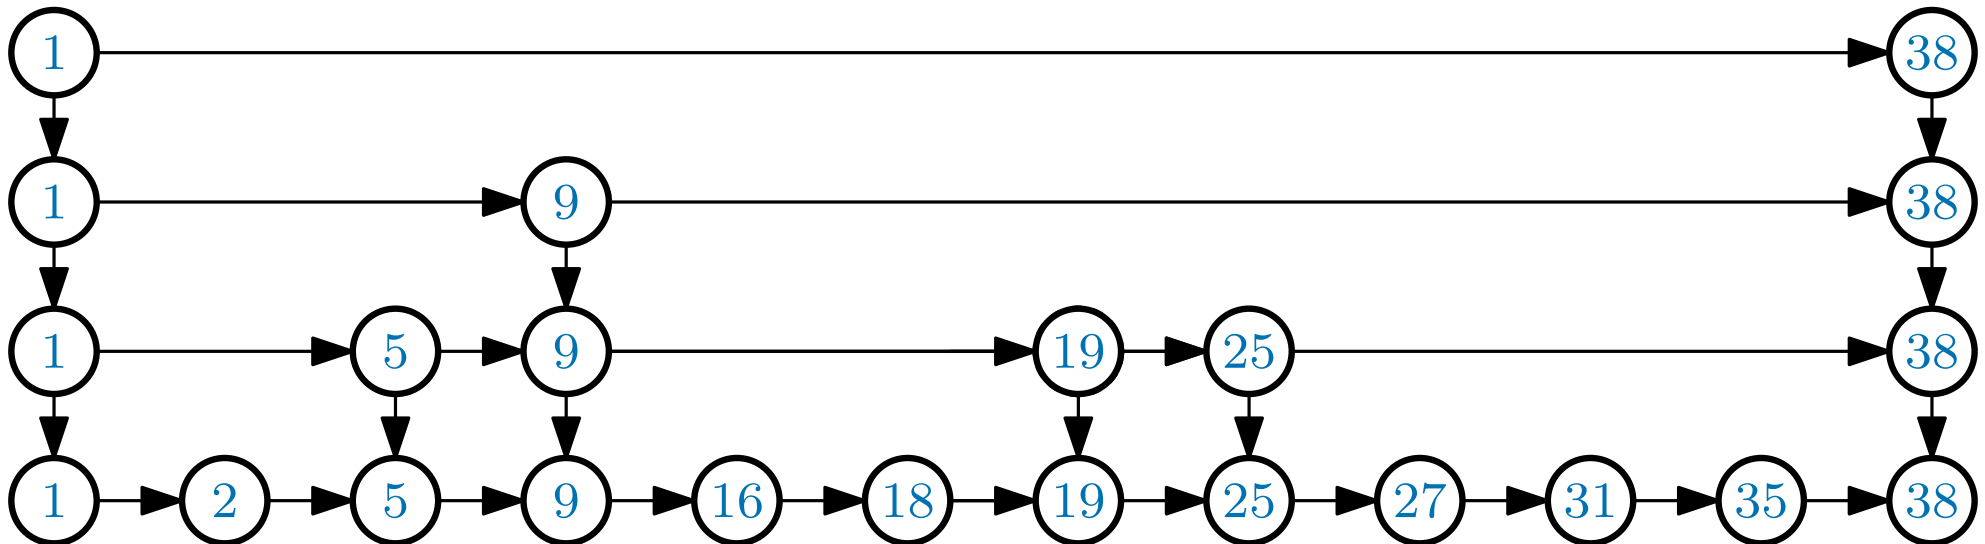
DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform **DELETE**(k),

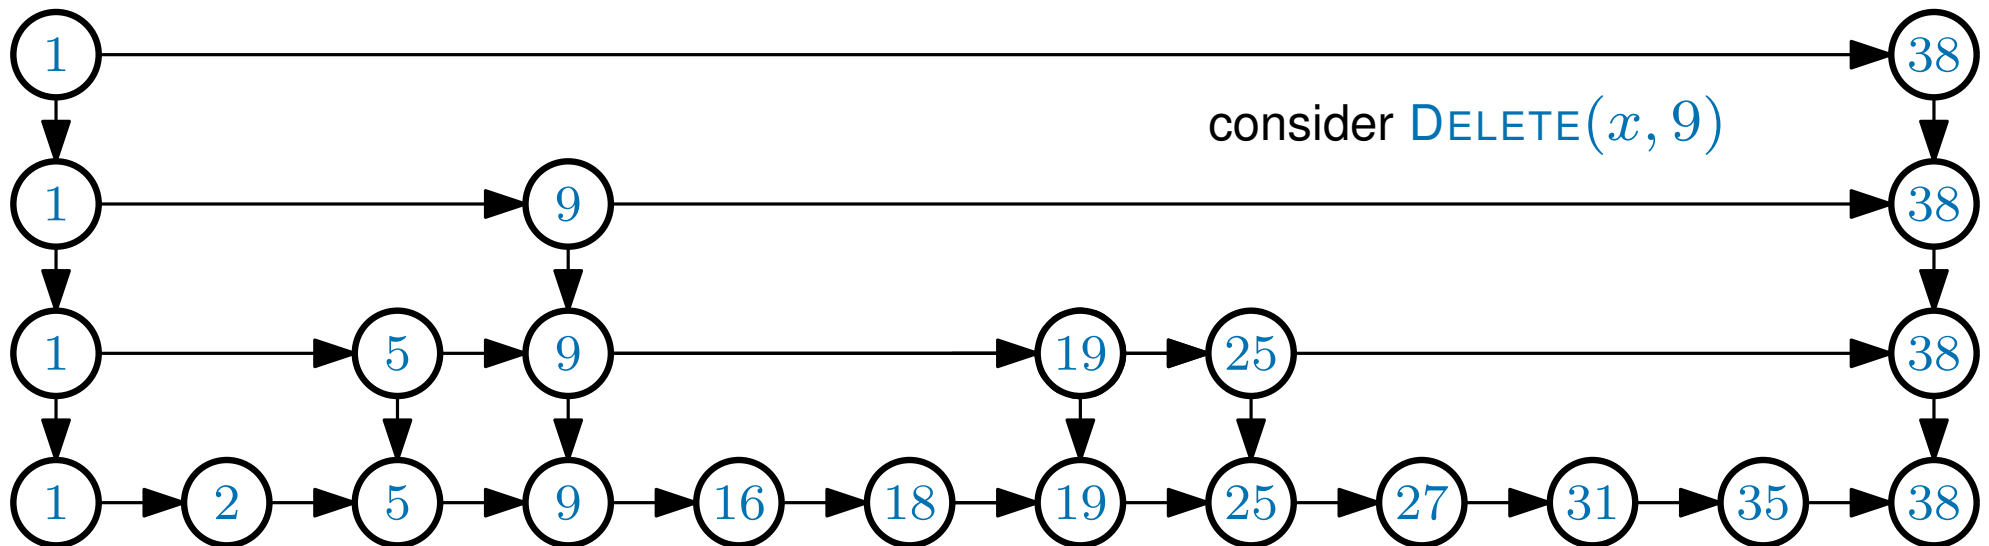
Step 1: Use **FIND**(k) to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform $\text{DELETE}(k)$,

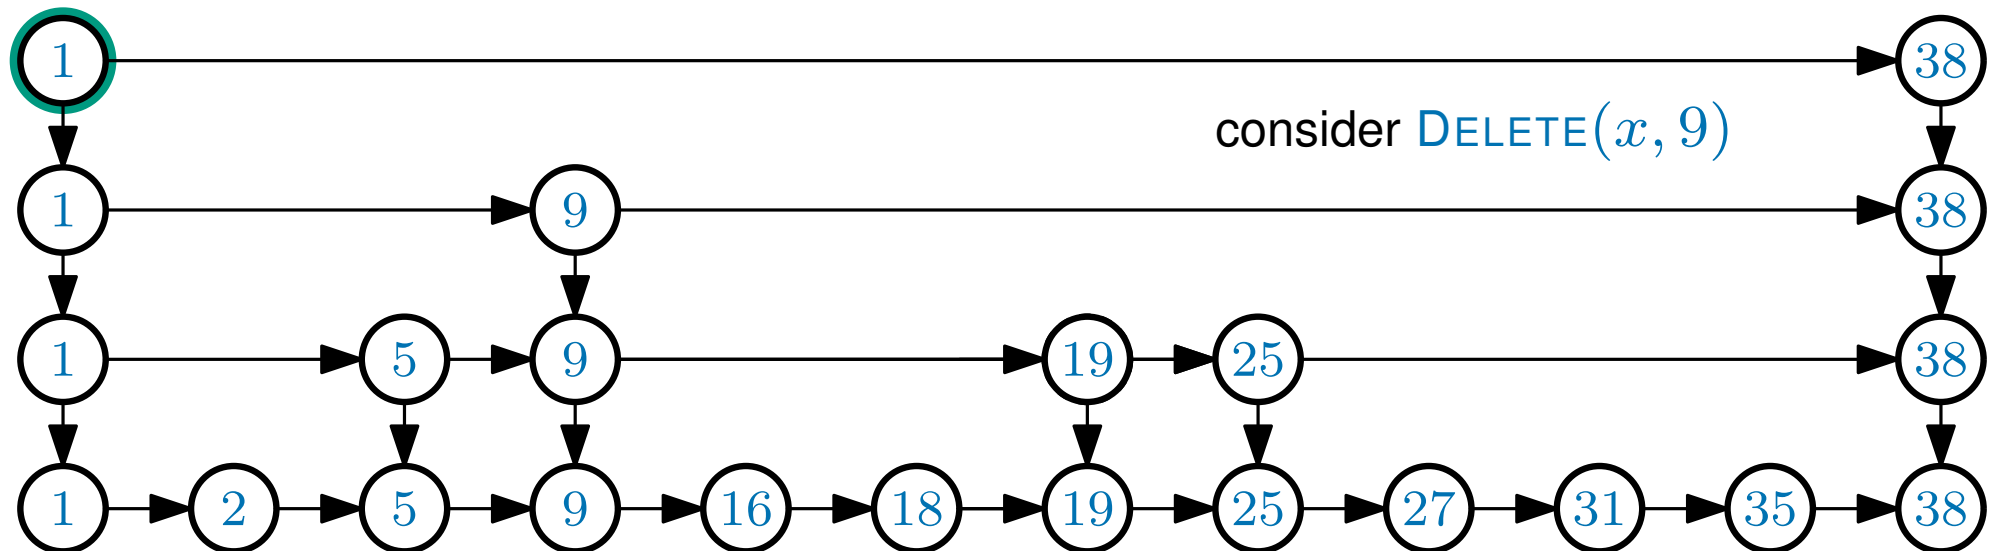
Step 1: Use $\text{FIND}(k)$ to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform $\text{DELETE}(k)$,

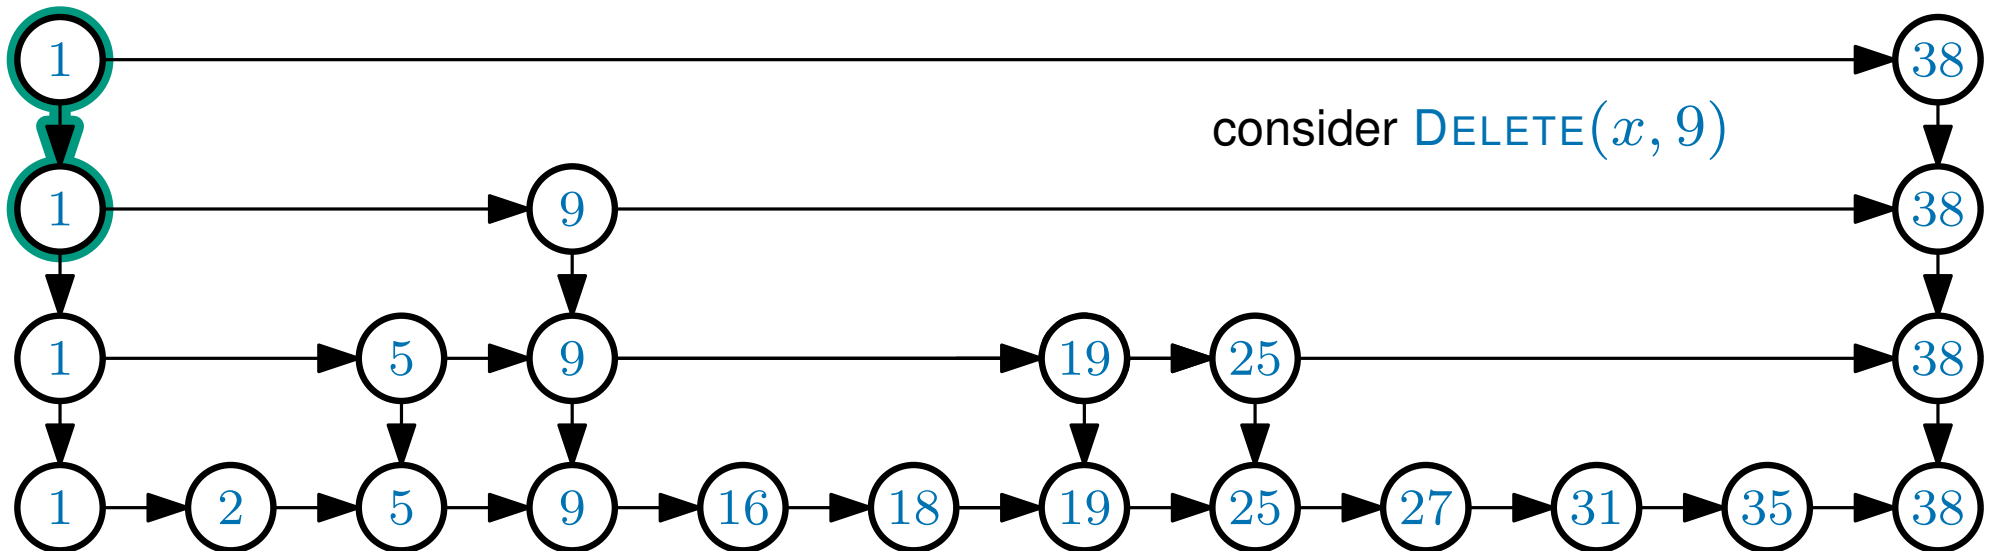
Step 1: Use $\text{FIND}(k)$ to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform $\text{DELETE}(k)$,

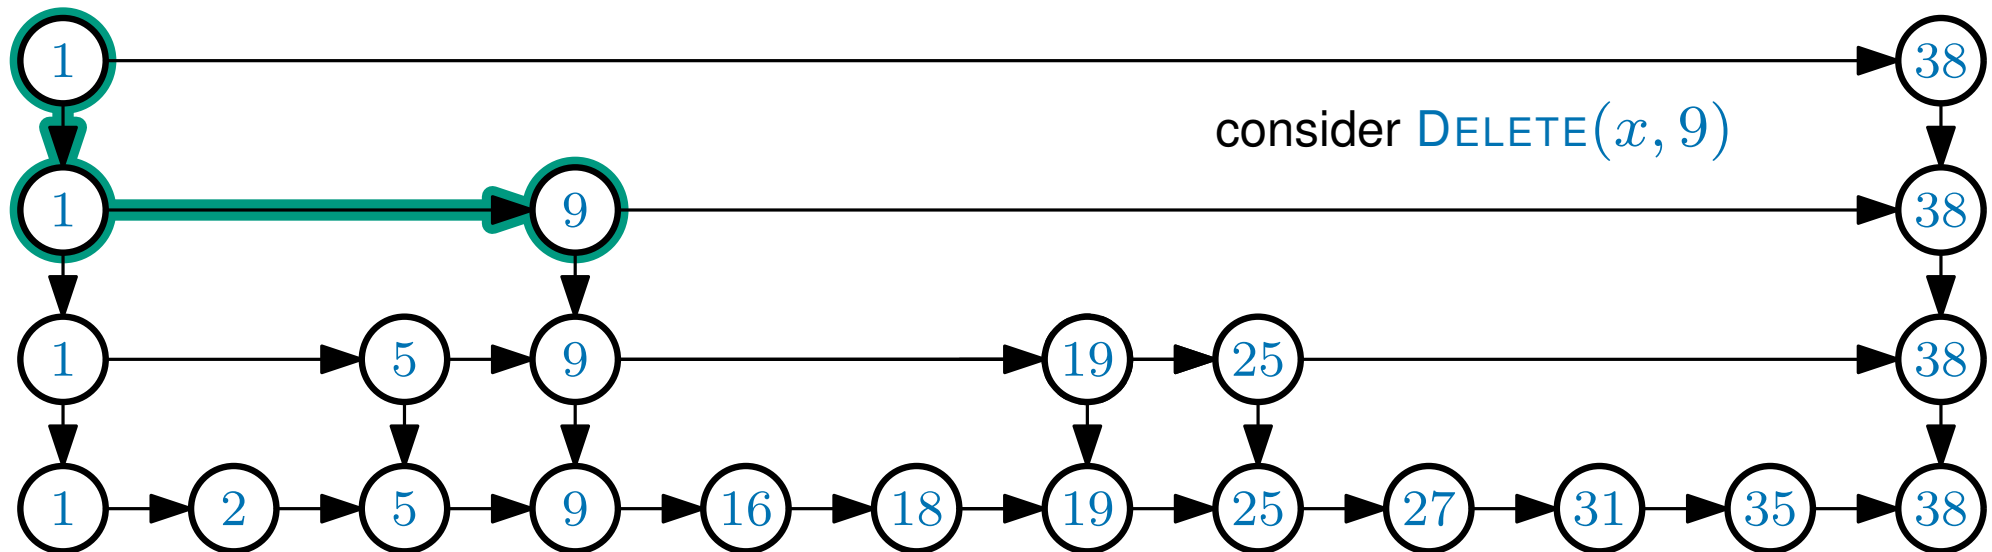
Step 1: Use $\text{FIND}(k)$ to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform $\text{DELETE}(k)$,

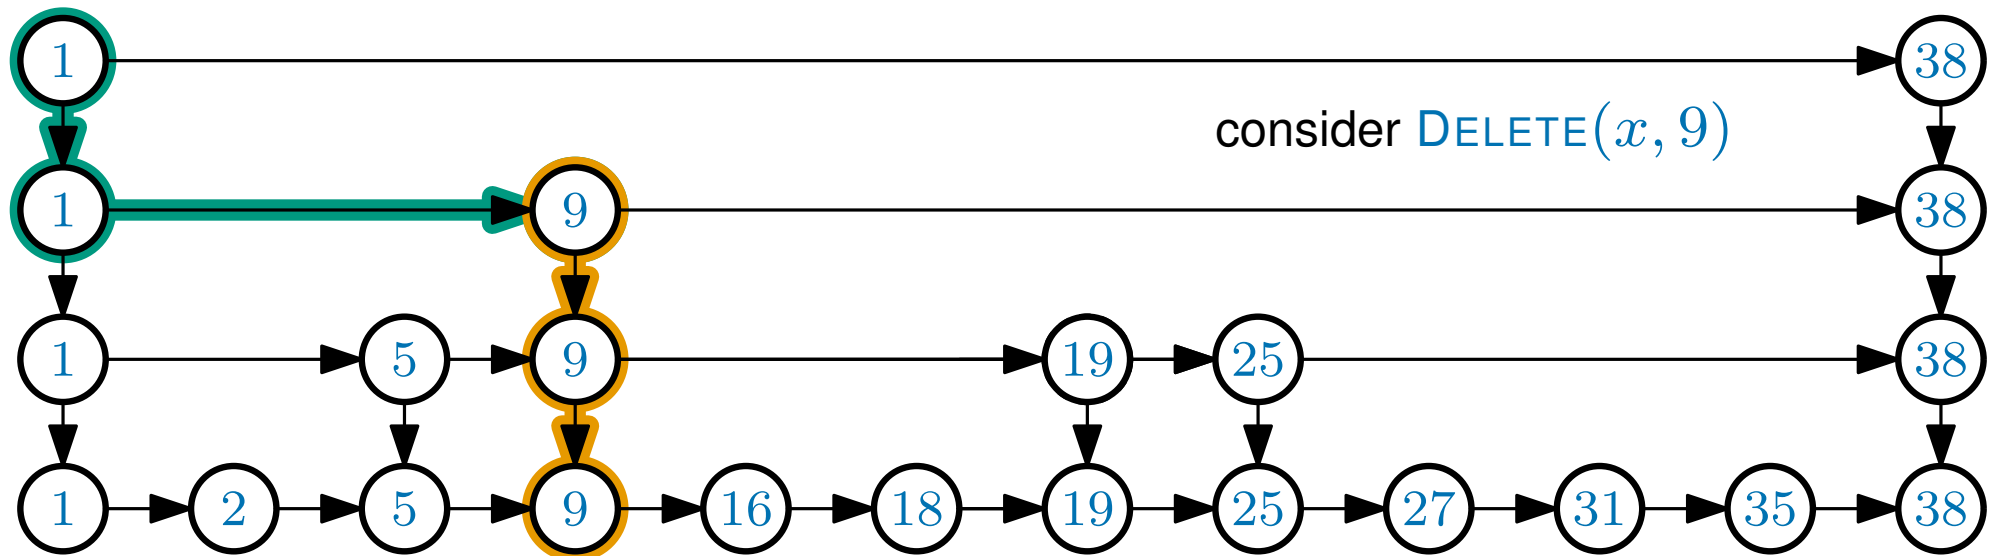
Step 1: Use $\text{FIND}(k)$ to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform $\text{DELETE}(k)$,

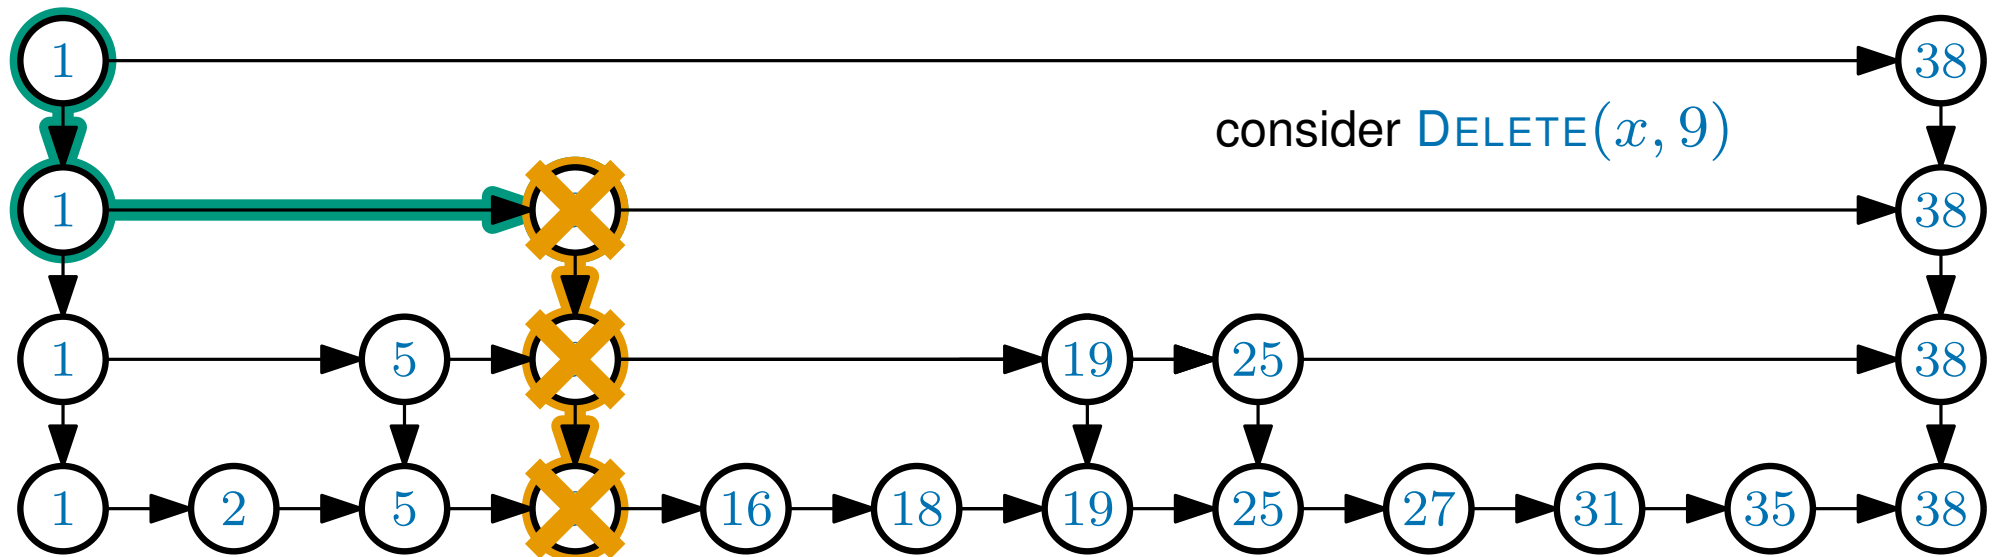
Step 1: Use $\text{FIND}(k)$ to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform **DELETE**(k),

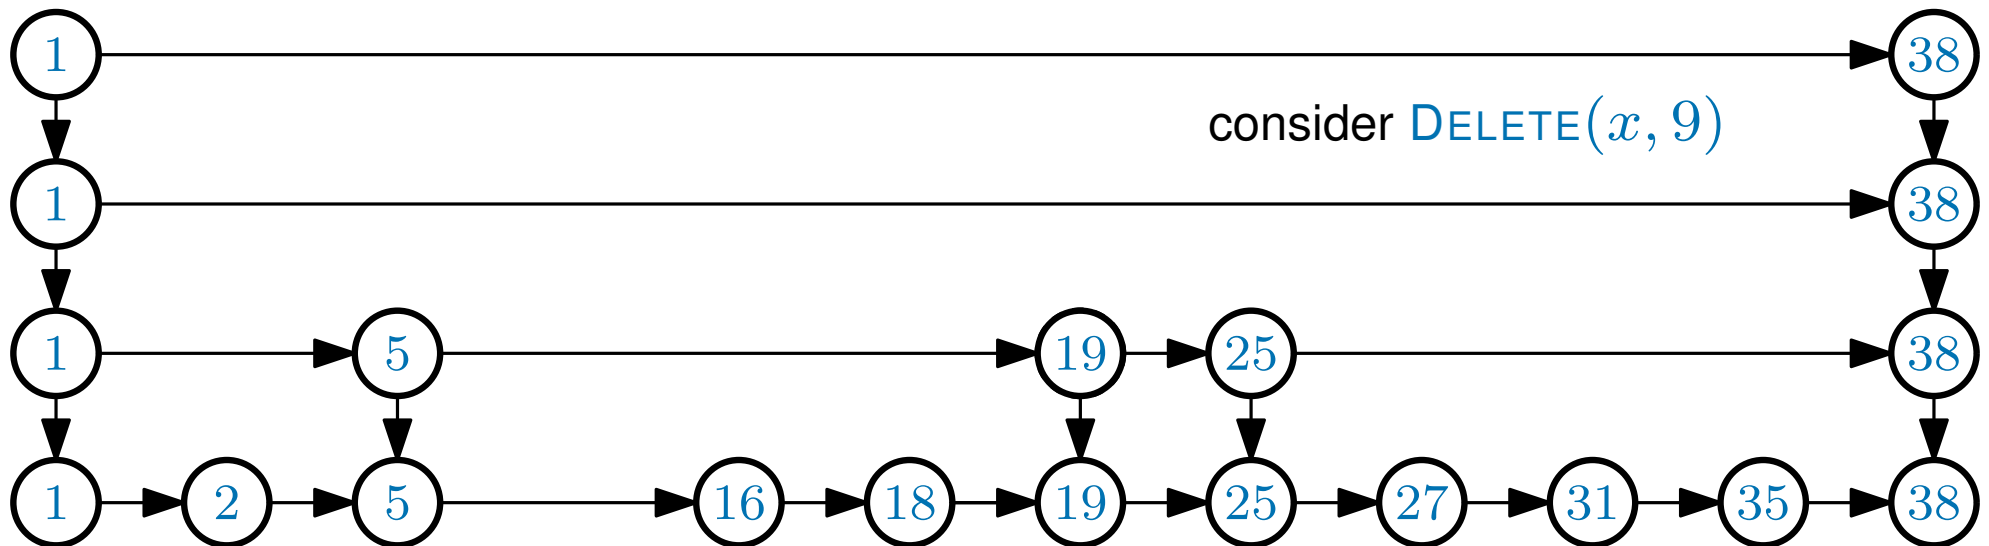
Step 1: Use **FIND**(k) to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform **DELETE**(k),

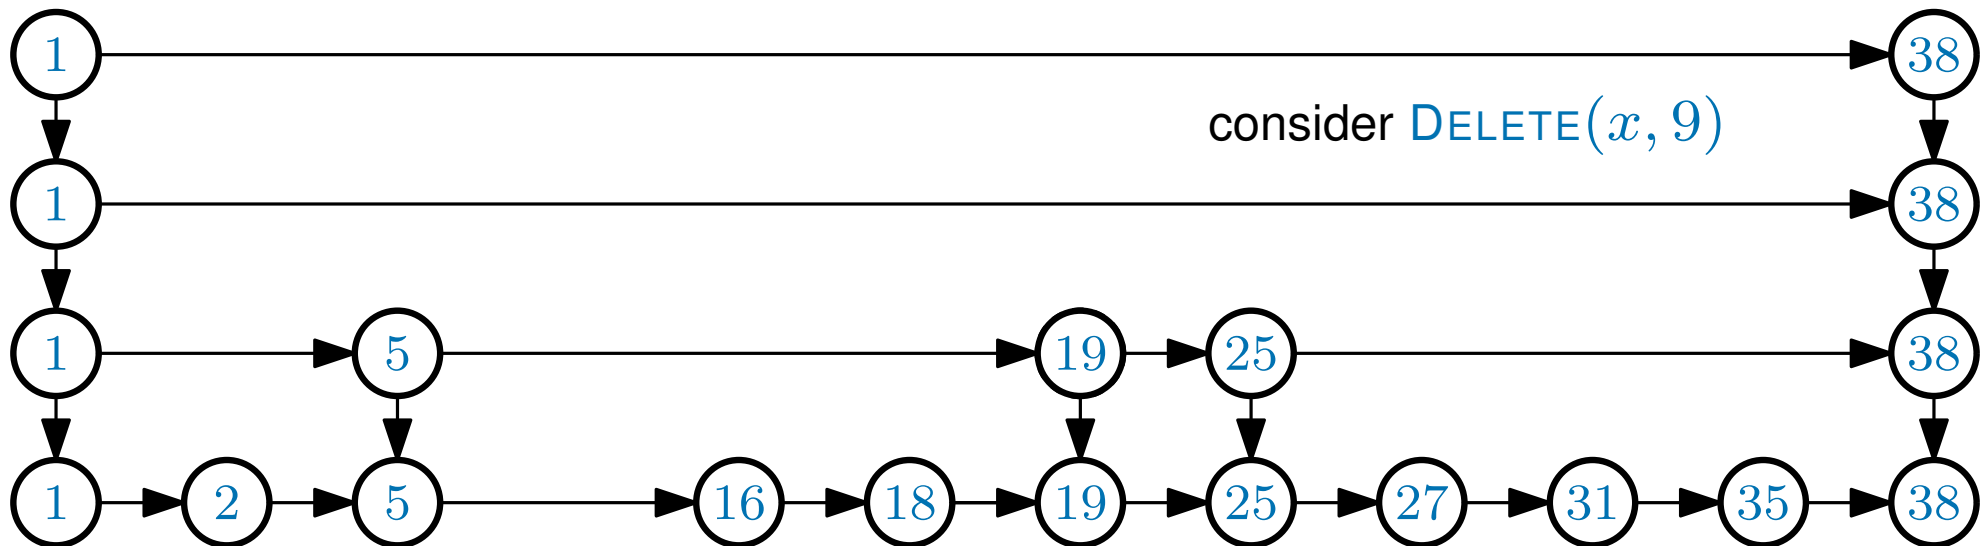
Step 1: Use **FIND**(k) to find (x, k)

Step 2: Delete (x, k) from all levels

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform **DELETE**(k),

Step 1: Use **FIND**(k) to find (x, k)

Step 2: Delete (x, k) from all levels

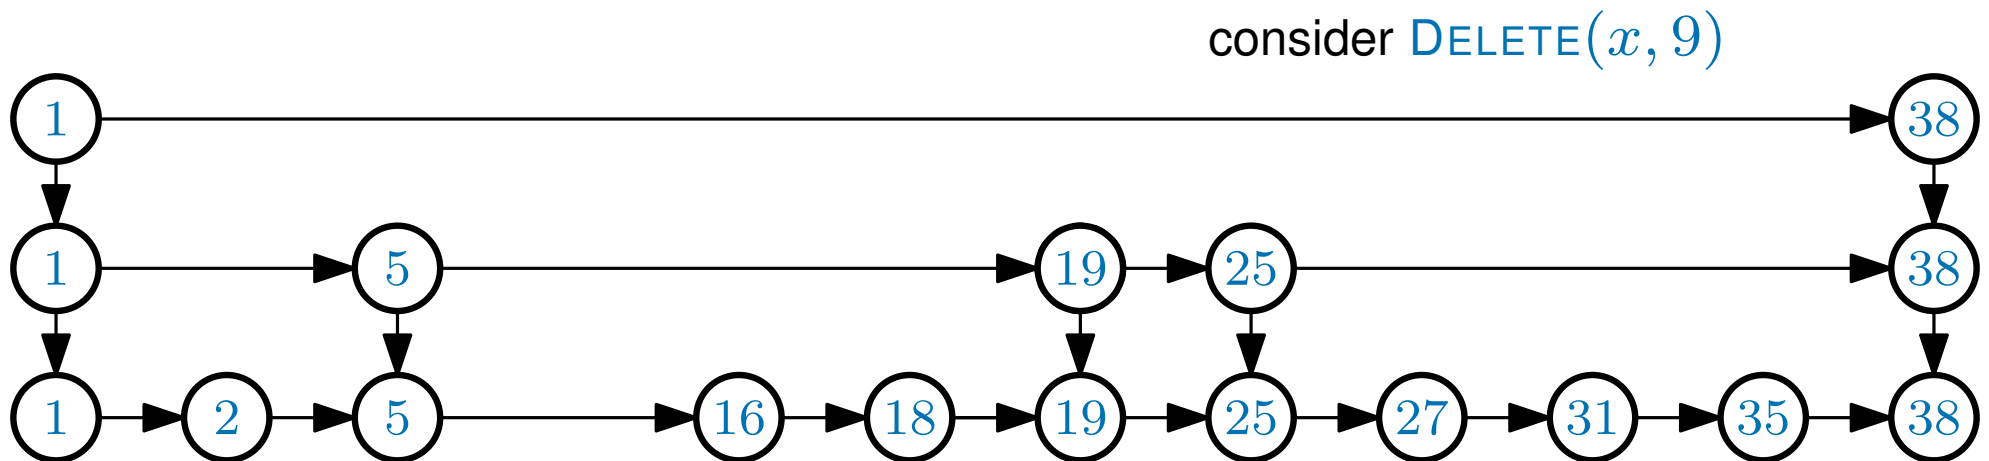
Step 3: Remove any empty levels

(ones containing only the smallest and largest keys)

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform **DELETE**(k),

Step 1: Use **FIND**(k) to find (x, k)

Step 2: Delete (x, k) from all levels

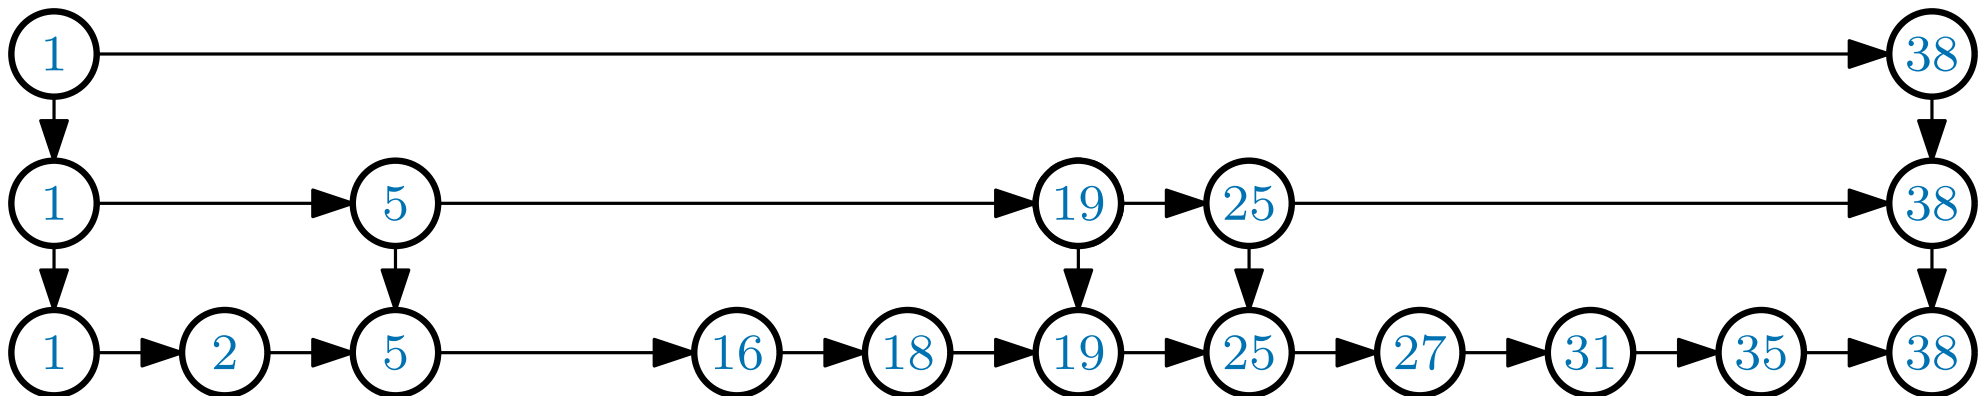
Step 3: Remove any empty levels

(ones containing only the smallest and largest keys)

Skip Lists

That about **DELETES**?

DELETING is straightforward, just **FIND** the key and **DELETE** it from all levels



To perform **DELETE**(k),

Step 1: Use **FIND**(k) to find (x, k)


Step 2: Delete (x, k) from all levels

Step 3: Remove any empty levels

(ones containing only the smallest and largest keys)

Skip Lists (pre-proof) summary

A skip list is a **randomised** data structure, based on link lists with **shortcuts**
which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



We will show that each of these operations takes *expected* $O(\log n)$ time

That is, they take $O(\log n)$ time ‘on average’

Important There is *no randomness in the data*,
the only randomness is in the coin flips

On *the worst case* input sequence, the expected time is $O(\log n)$

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an INSERT operation

How many levels are in a Skip list?

We begin by proving that after n **INSERT** operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an **INSERT** operation

Consider some **INSERT**(x, k) operation

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than 1 level is $\frac{1}{2}$
(the first coin flip is H)

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than 1 level is $\frac{1}{2}$
(the first coin flip is H)

The probability (x, k) is inserted into more than 2 levels is $\frac{1}{4}$
(we throw HH...)

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than 1 level is $\frac{1}{2}$
(the first coin flip is H)

The probability (x, k) is inserted into more than 2 levels is $\frac{1}{4}$
(we throw HH...)

The probability (x, k) is inserted into more than 3 levels is $\frac{1}{8}$
(we throw HHH...)

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than 1 level is $\frac{1}{2}$
(the first coin flip is H)

The probability (x, k) is inserted into more than 2 levels is $\frac{1}{4}$
(we throw HH...)

The probability (x, k) is inserted into more than 3 levels is $\frac{1}{8}$
(we throw HHH...)

The probability (x, k) is inserted into more than j levels is $\frac{1}{2^j}$

How many levels are in a Skip list?

We begin by proving that after n **INSERT** operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an **INSERT** operation

Consider some **INSERT**(x, k) operation

The probability (x, k) is inserted into more than j levels is $\frac{1}{2^j}$

How many levels are in a Skip list?

We begin by proving that after n **INSERT** operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an **INSERT** operation

Consider some **INSERT**(x, k) operation

The probability (x, k) is inserted into more than $2 \log n$ levels is $\frac{1}{2^{2 \log n}}$

How many levels are in a Skip list?

We begin by proving that after n **INSERT** operations, a skip list
is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level
and the only way this can increase is during an **INSERT** operation

Consider some **INSERT**(x, k) operation

The probability (x, k) is inserted into more than $2 \log n$ levels is $\frac{1}{2^{2 \log n}} = \frac{1}{n^2}$

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than $2 \log n$ levels is $\frac{1}{2^{2 \log n}} = \frac{1}{n^2}$

The **union** bound

Let $E_1, E_2 \dots E_n$ be events where E_j occurs with probability p_j

The probability of at least one E_j occurring is at most $\sum_j p_j$

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than $2 \log n$ levels is $\frac{1}{2^{2 \log n}} = \frac{1}{n^2}$

The **union** bound

Let $E_1, E_2 \dots E_n$ be events where E_j occurs with probability p_j

The probability of at least one E_j occurring is at most $\sum_j p_j$

Let E_j be the event that the j -th INSERT puts its element in more than $2 \log n$ levels

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than $2 \log n$ levels is $\frac{1}{2^{2 \log n}} = \frac{1}{n^2}$

The **union** bound

Let $E_1, E_2 \dots E_n$ be events where E_j occurs with probability p_j

The probability of at least one E_j occurring is at most $\sum_j p_j$

Let E_j be the event that the j -th INSERT puts its element in more than $2 \log n$ levels

The probability of at least one E_j occurring is at most $\sum_j \frac{1}{n^2}$

How many levels are in a Skip list?

We begin by proving that after n INSERT operations, a skip list is very unlikely to have more than $2 \log n$ levels...

An empty skip list contains only one level and the only way this can increase is during an INSERT operation

Consider some INSERT(x, k) operation

The probability (x, k) is inserted into more than $2 \log n$ levels is $\frac{1}{2^{2 \log n}} = \frac{1}{n^2}$

The **union** bound

Let $E_1, E_2 \dots E_n$ be events where E_j occurs with probability p_j

The probability of at least one E_j occurring is at most $\sum_j p_j$

Let E_j be the event that the j -th INSERT puts its element in more than $2 \log n$ levels

The probability of at least one E_j occurring is at most $\sum_j \frac{1}{n^2} = \frac{1}{n}$

How many levels are in a Skip list?

After n INSERT operations, the probability that a skip list
has more than $2 \log n$ levels...

is at most $\frac{1}{n}$

How many levels are in a Skip list?

After n INSERT operations, the probability that a skip list
has more than $2 \log n$ levels...

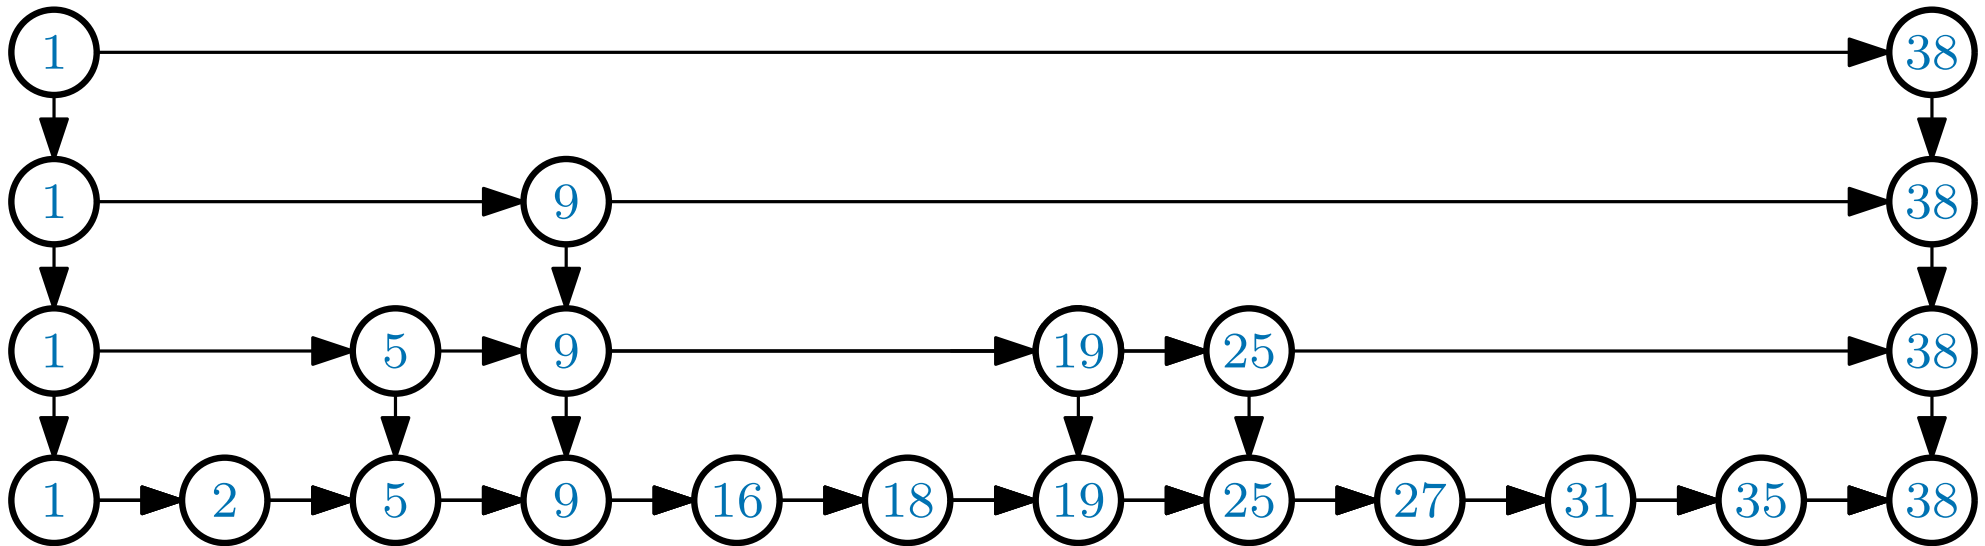
is at most $\frac{1}{n}$

It gets better as n increases!

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

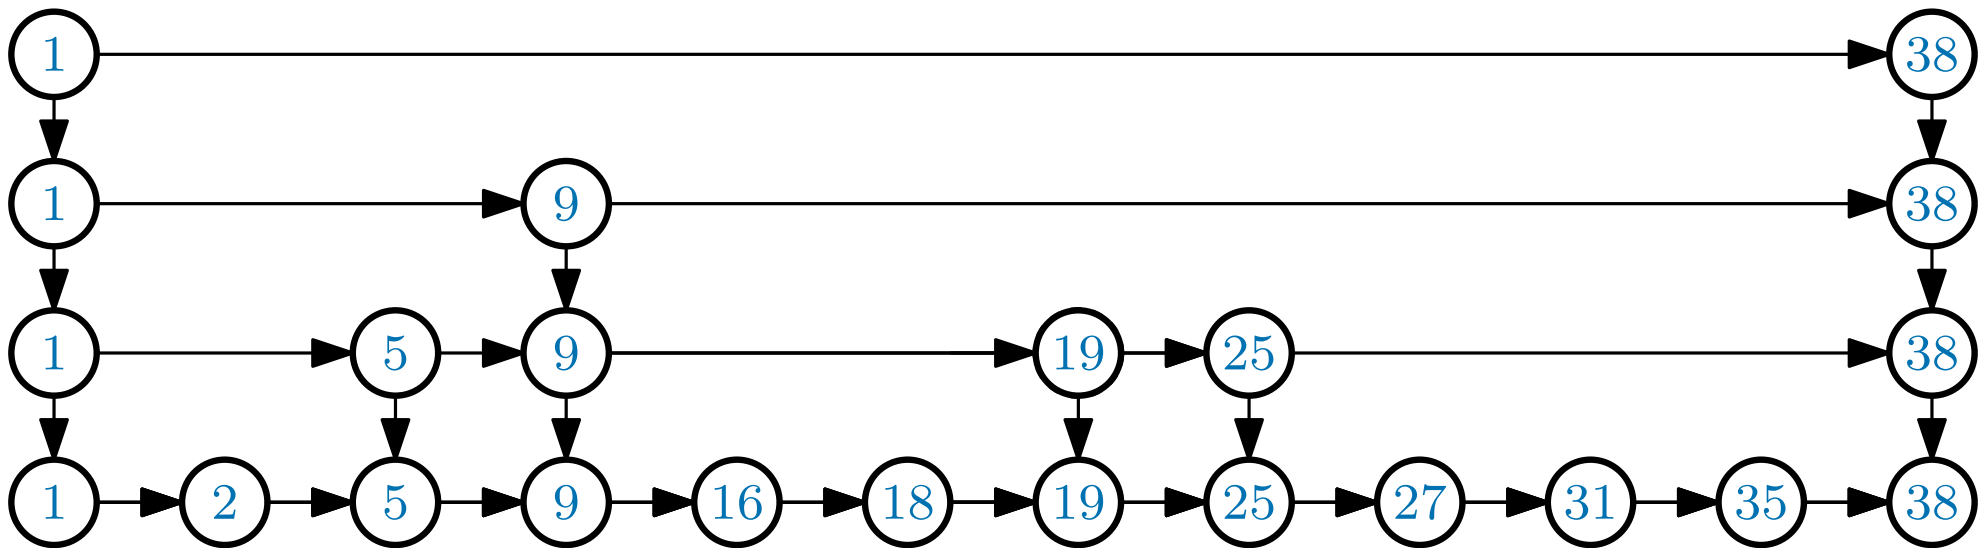
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

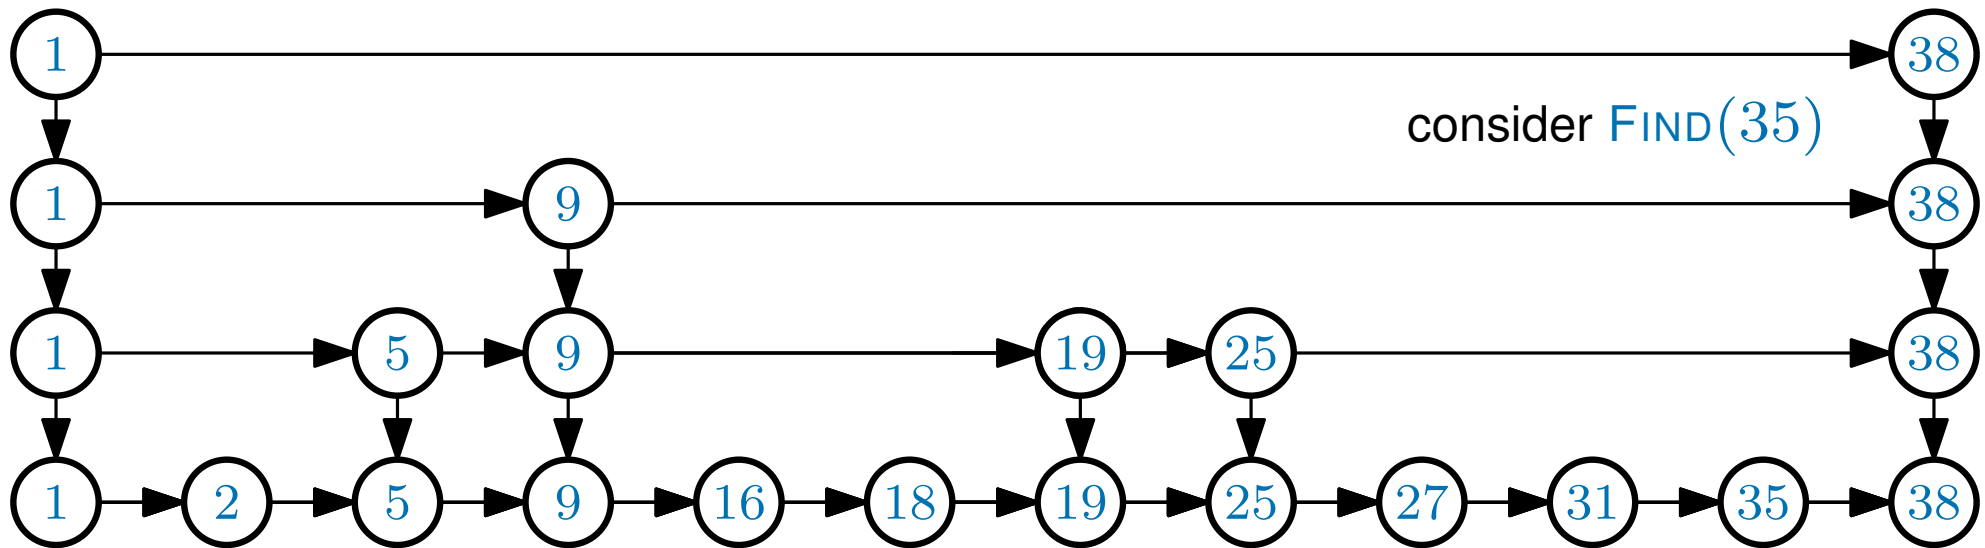
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

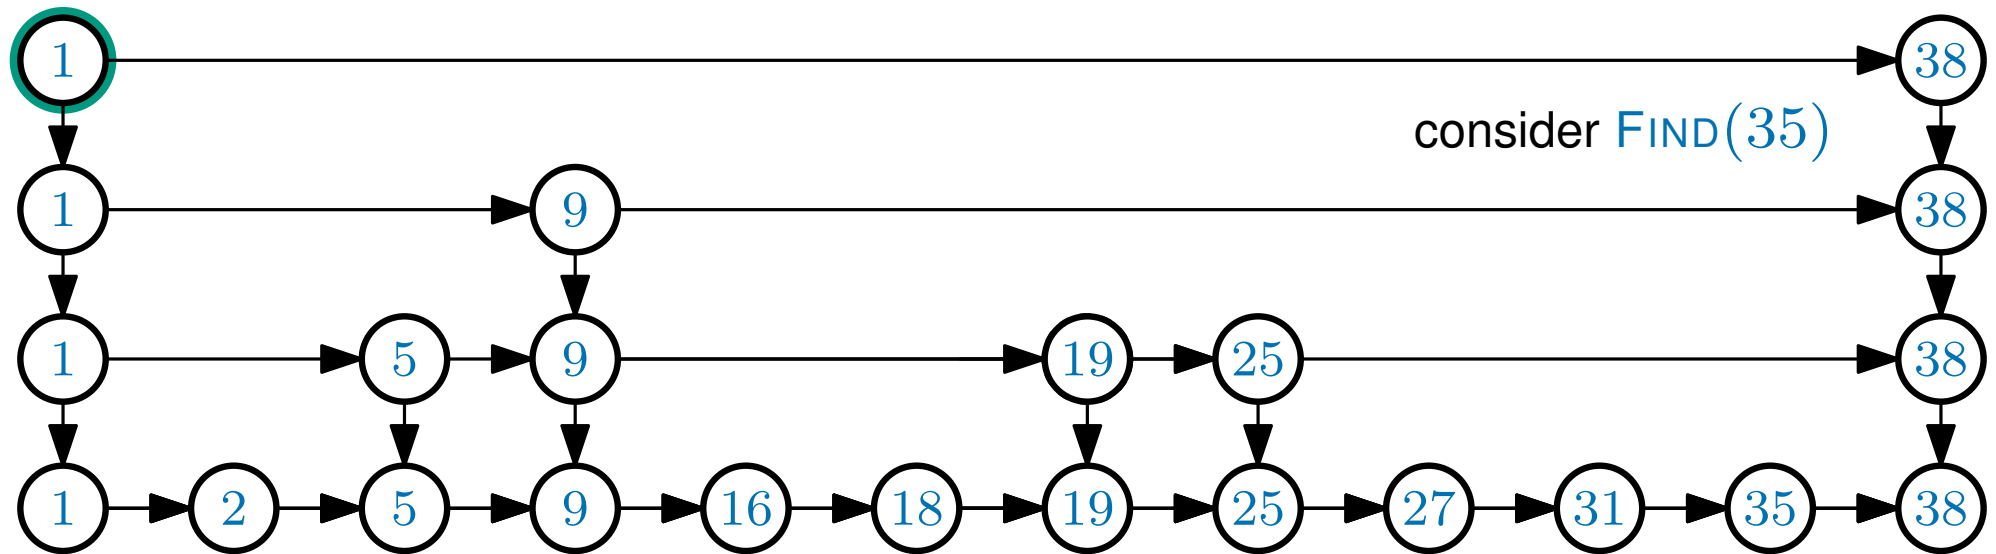
Else **Move down**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

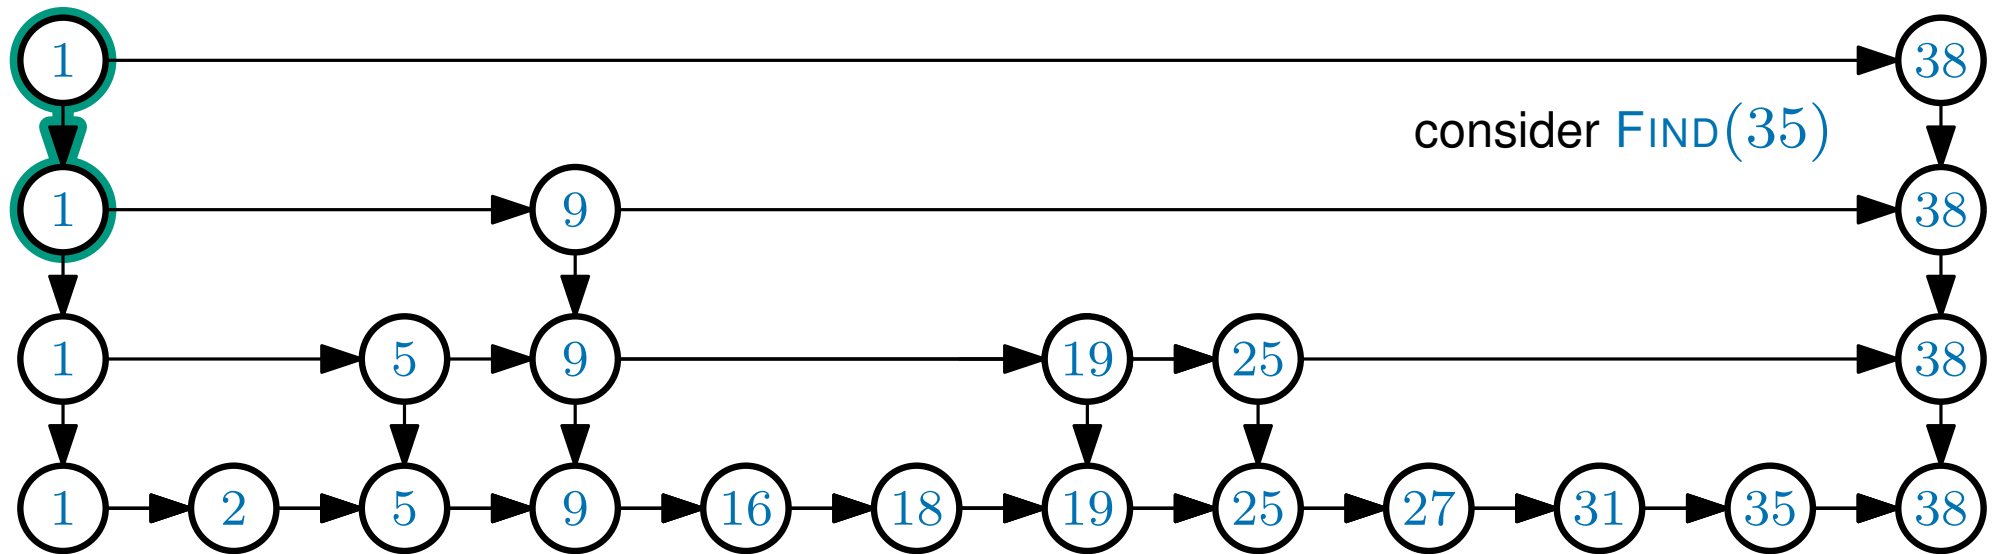
Else **Move down**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

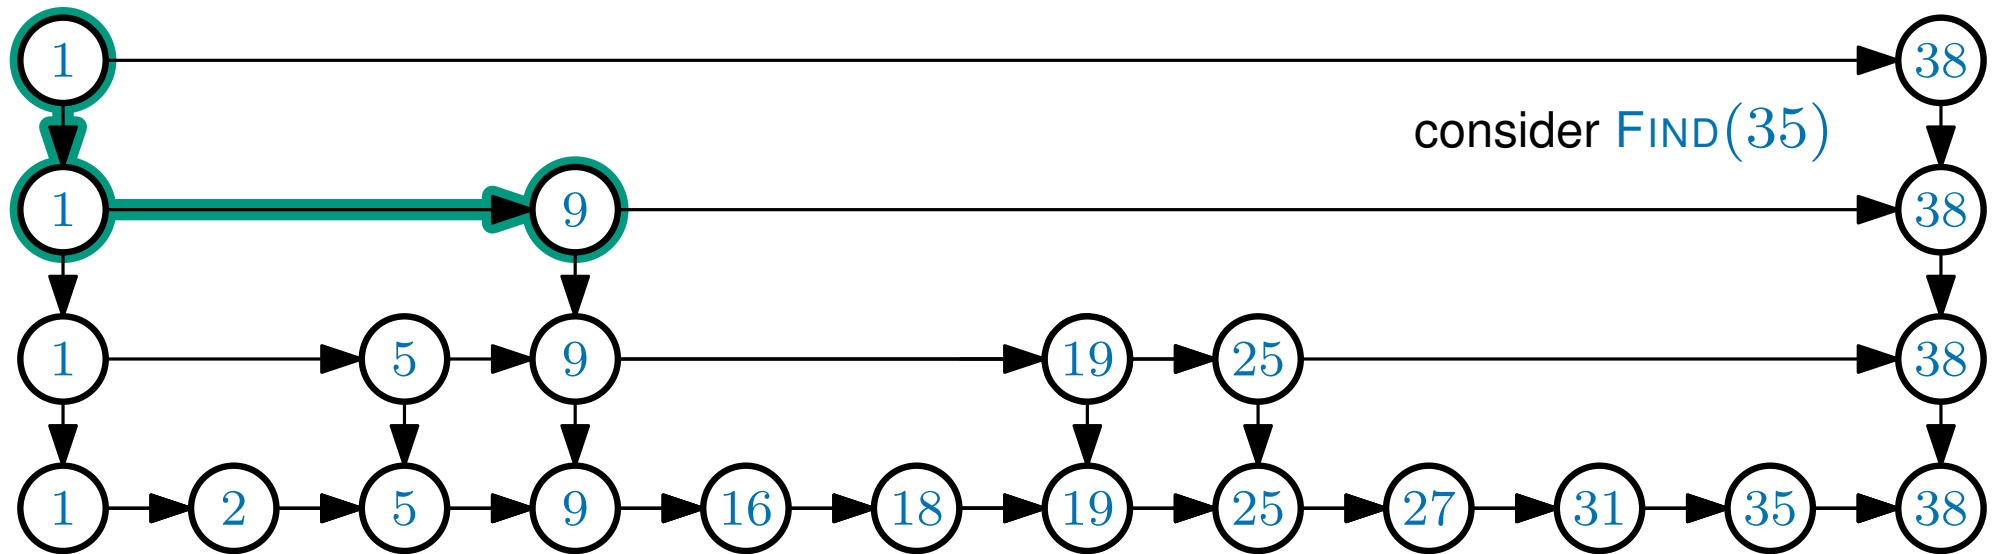
Else **Move down**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

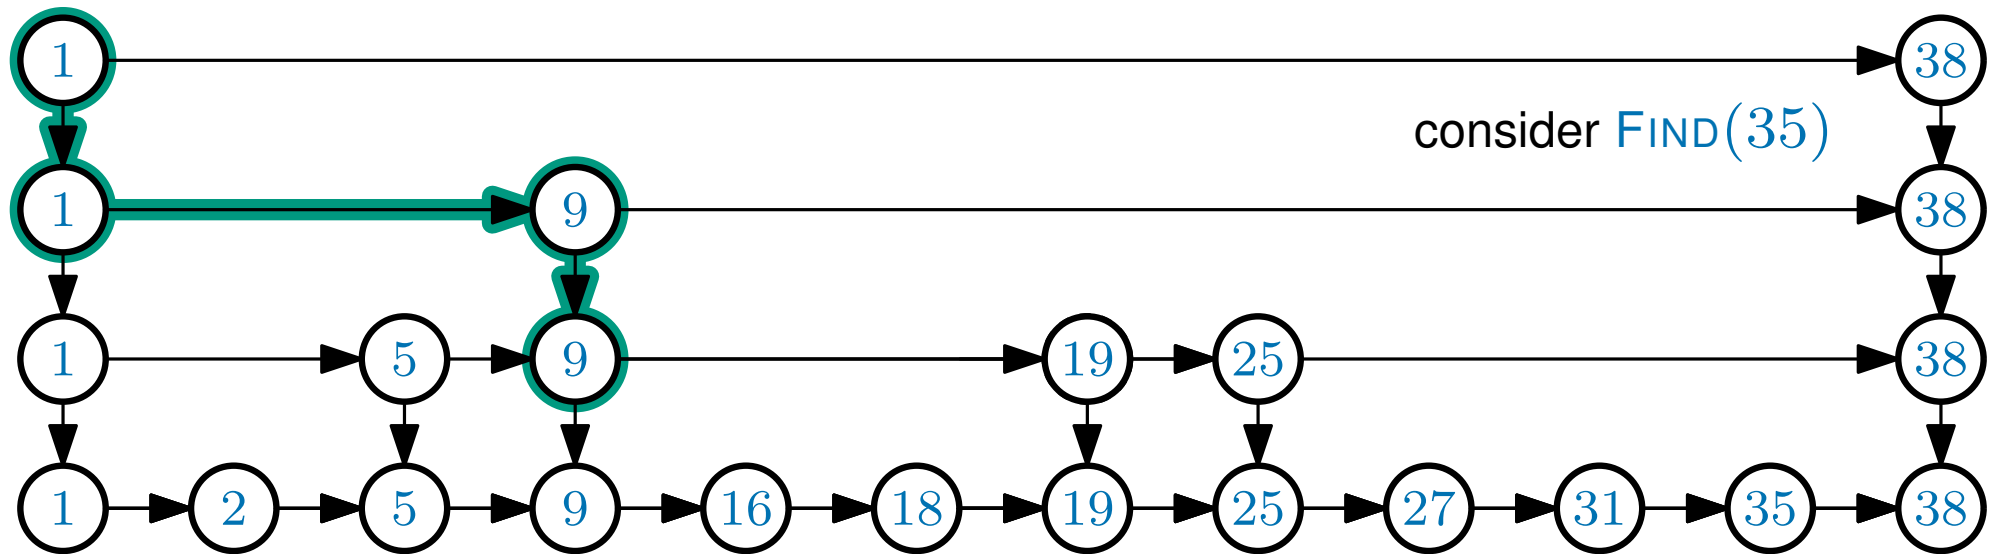
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

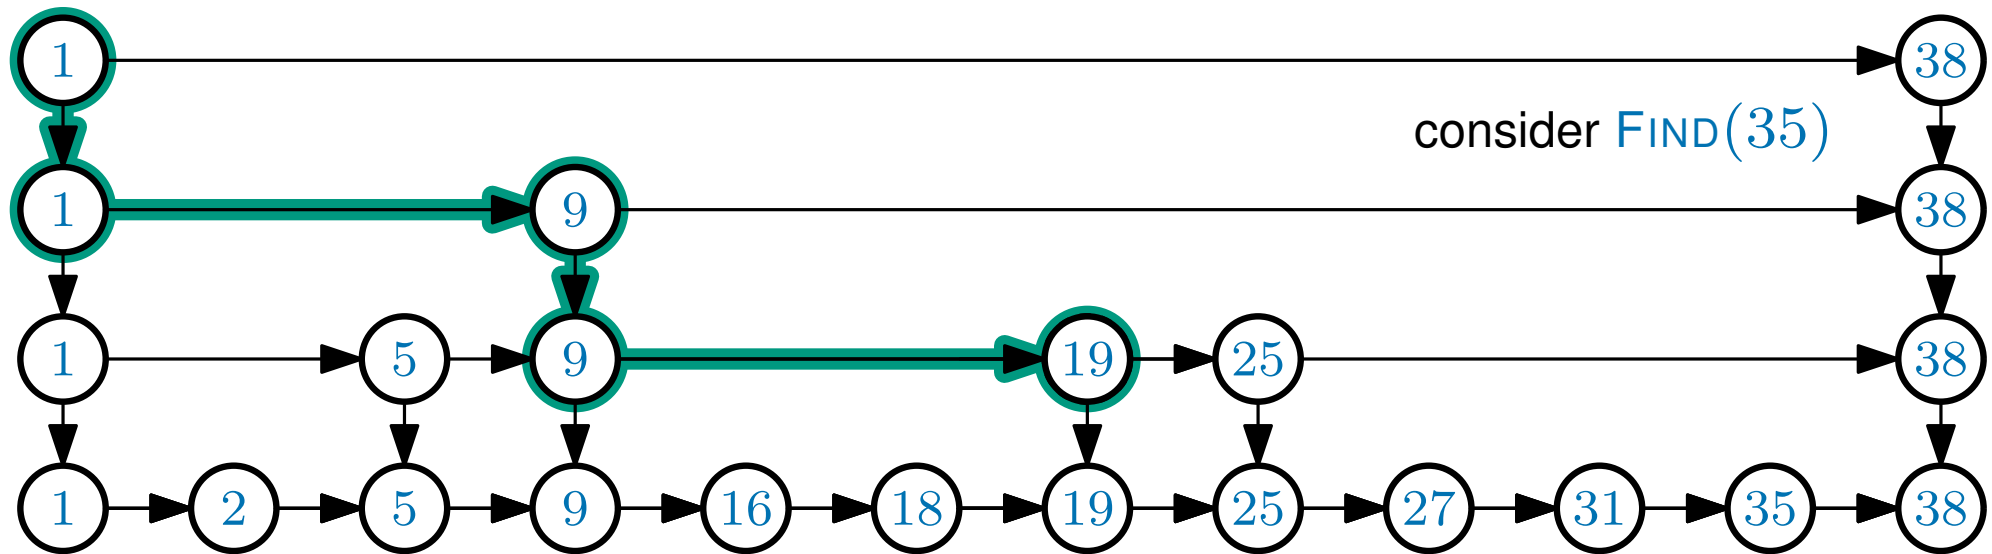
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

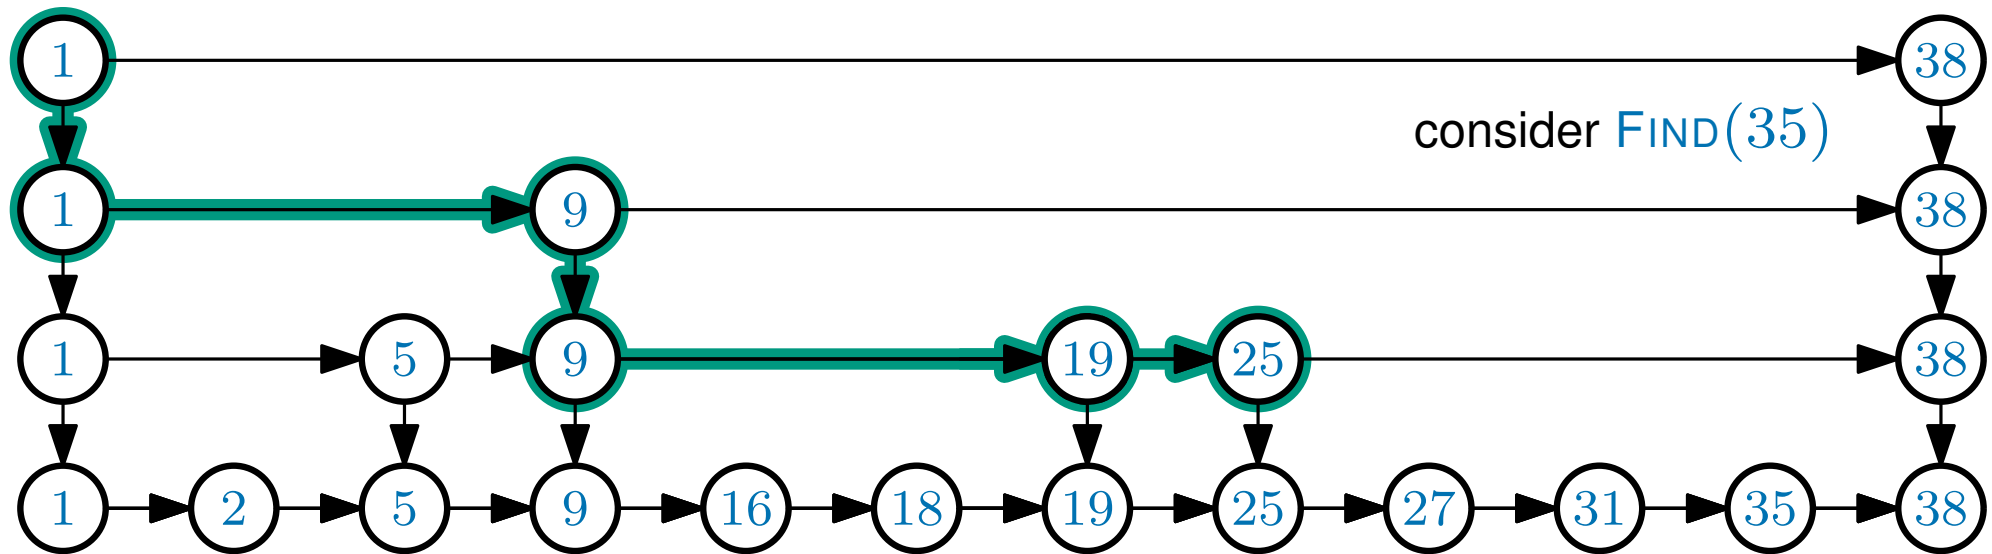
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (the head of the top level)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

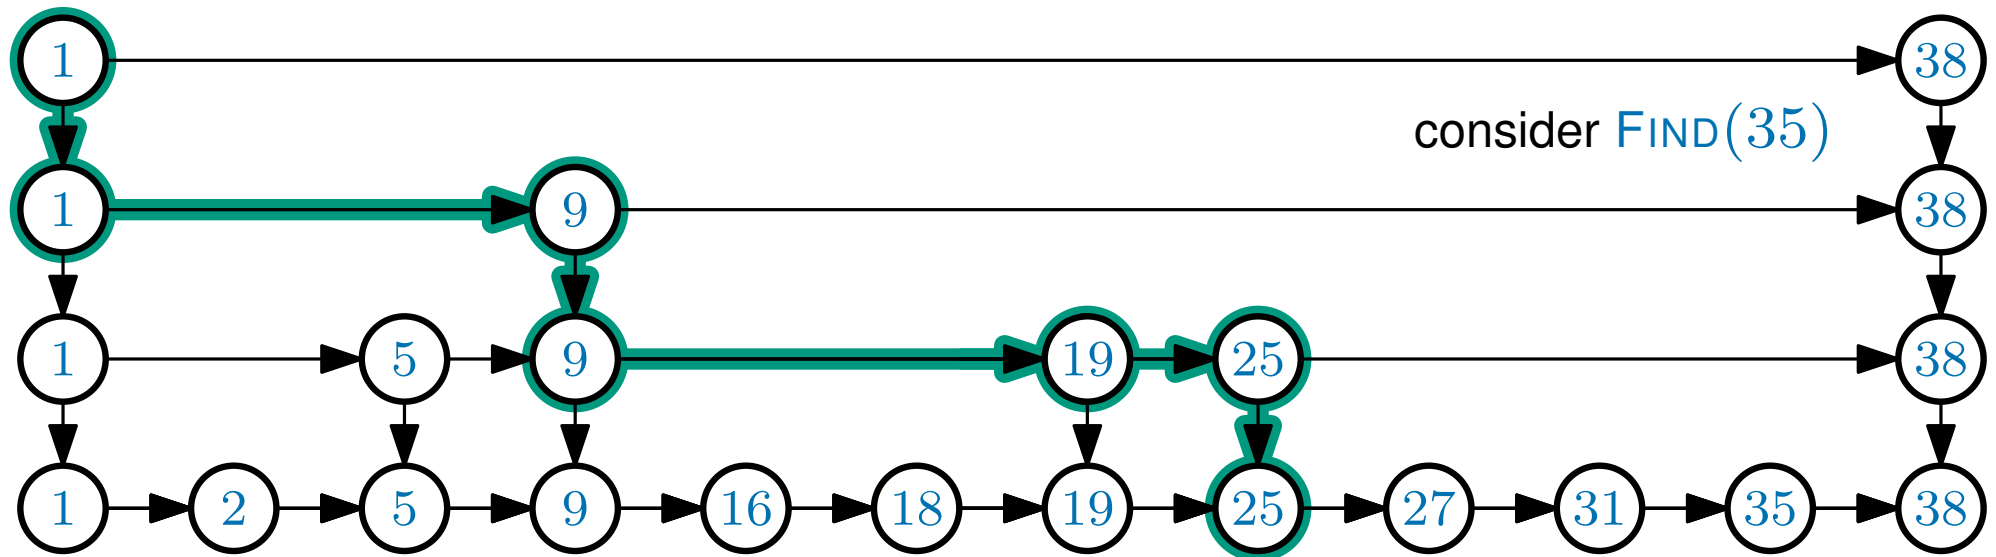
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right


Else **Move down**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right

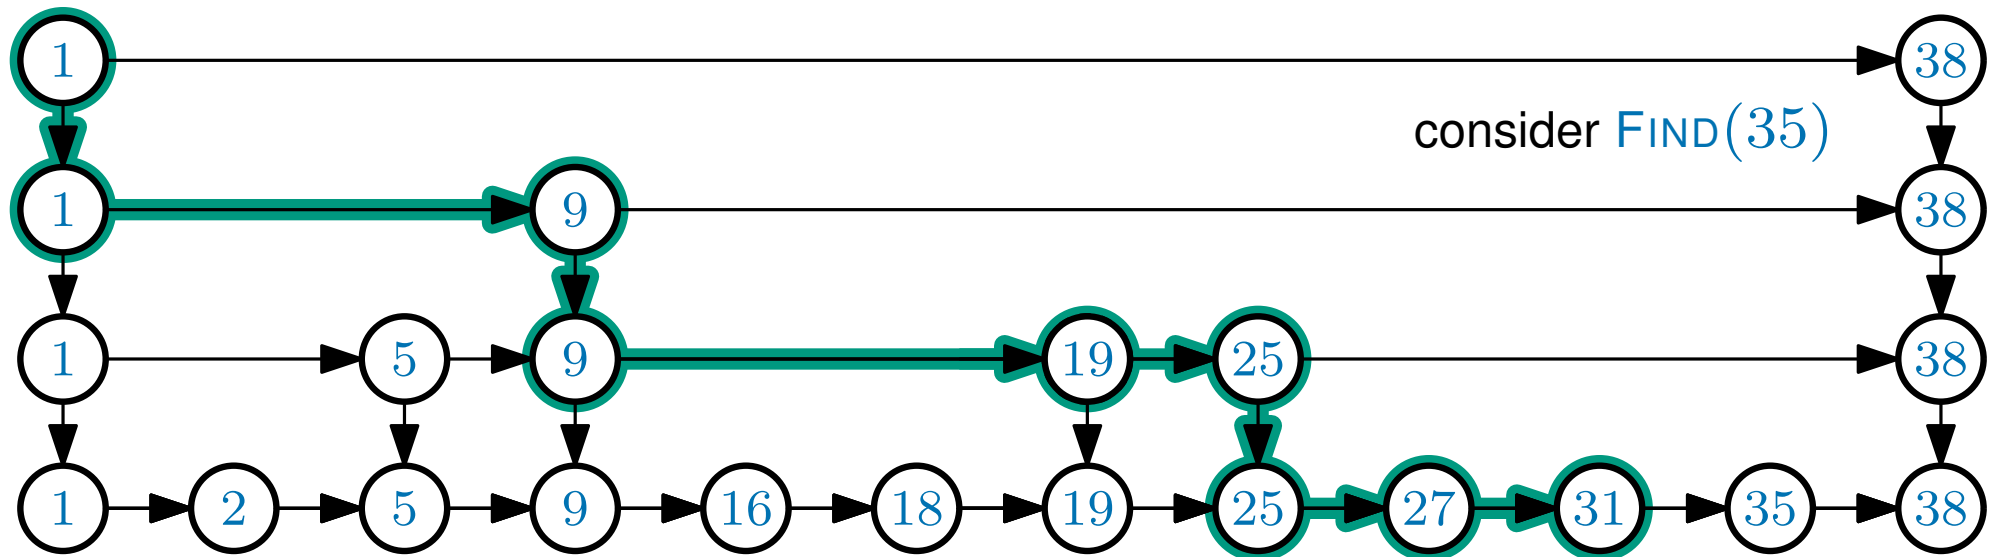
Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (the head of the top level)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right


Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



Start at the top-left (*the head of the top level*)

While you haven't found k :

To perform **FIND**(k),

If the node to the right's key, $k' \leq k$

Move right


Else Move down

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?

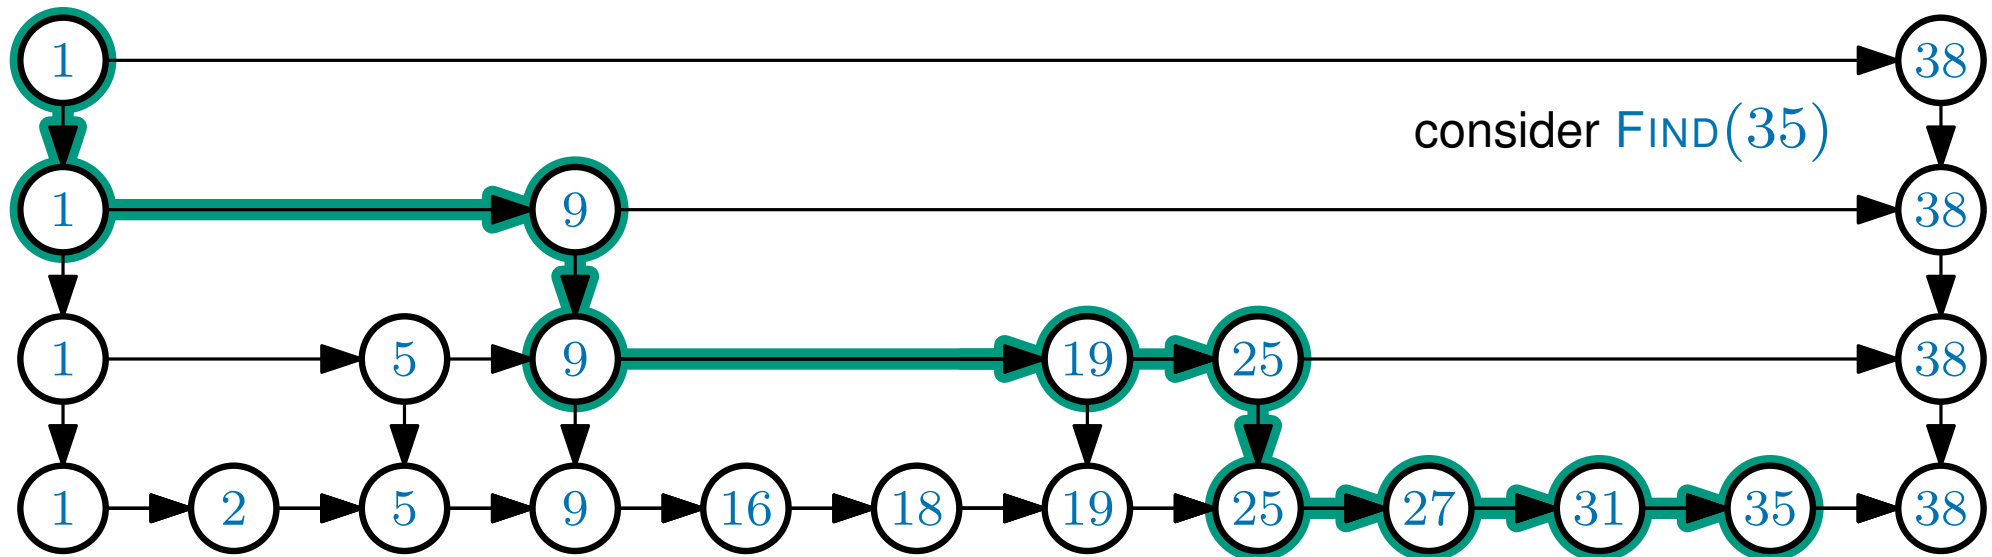


So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



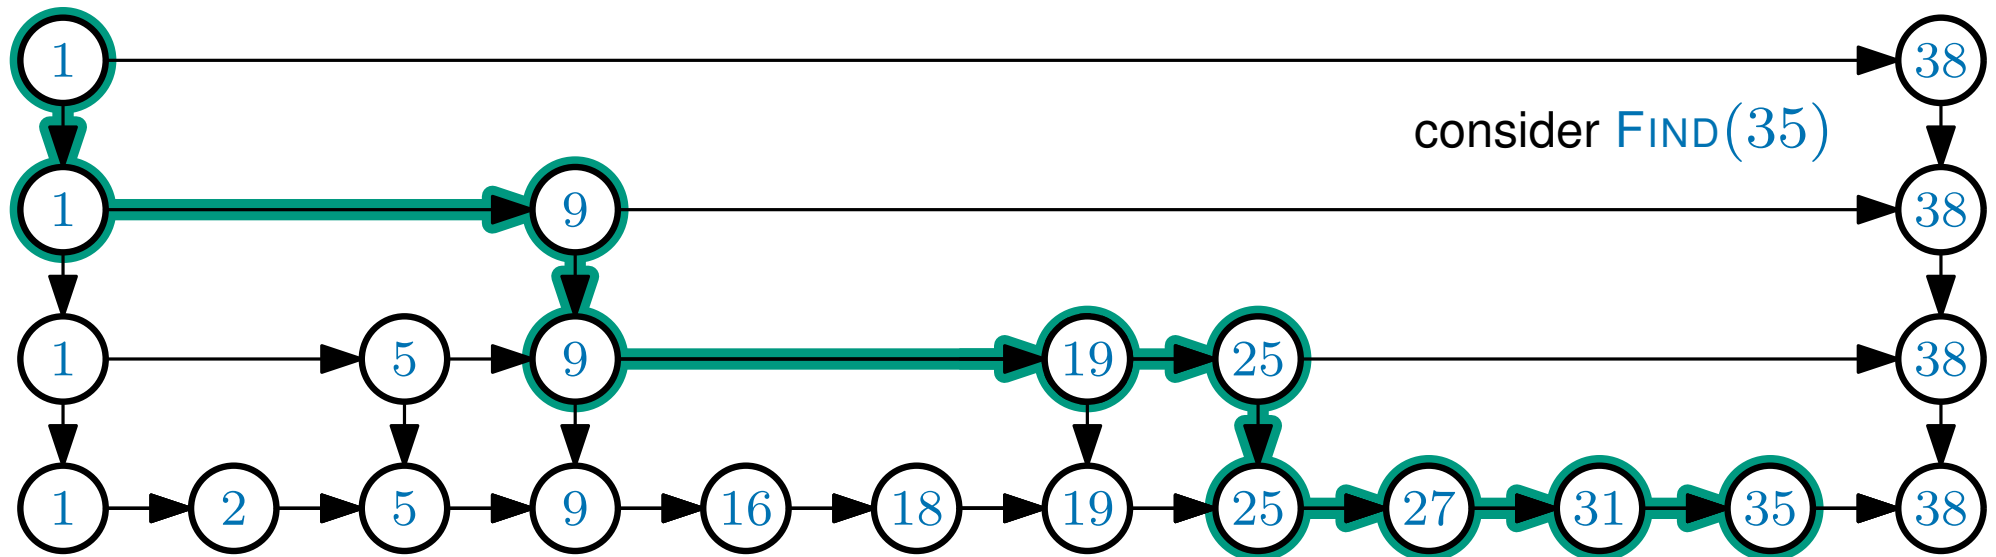
How long is this path?

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



How long is this path?

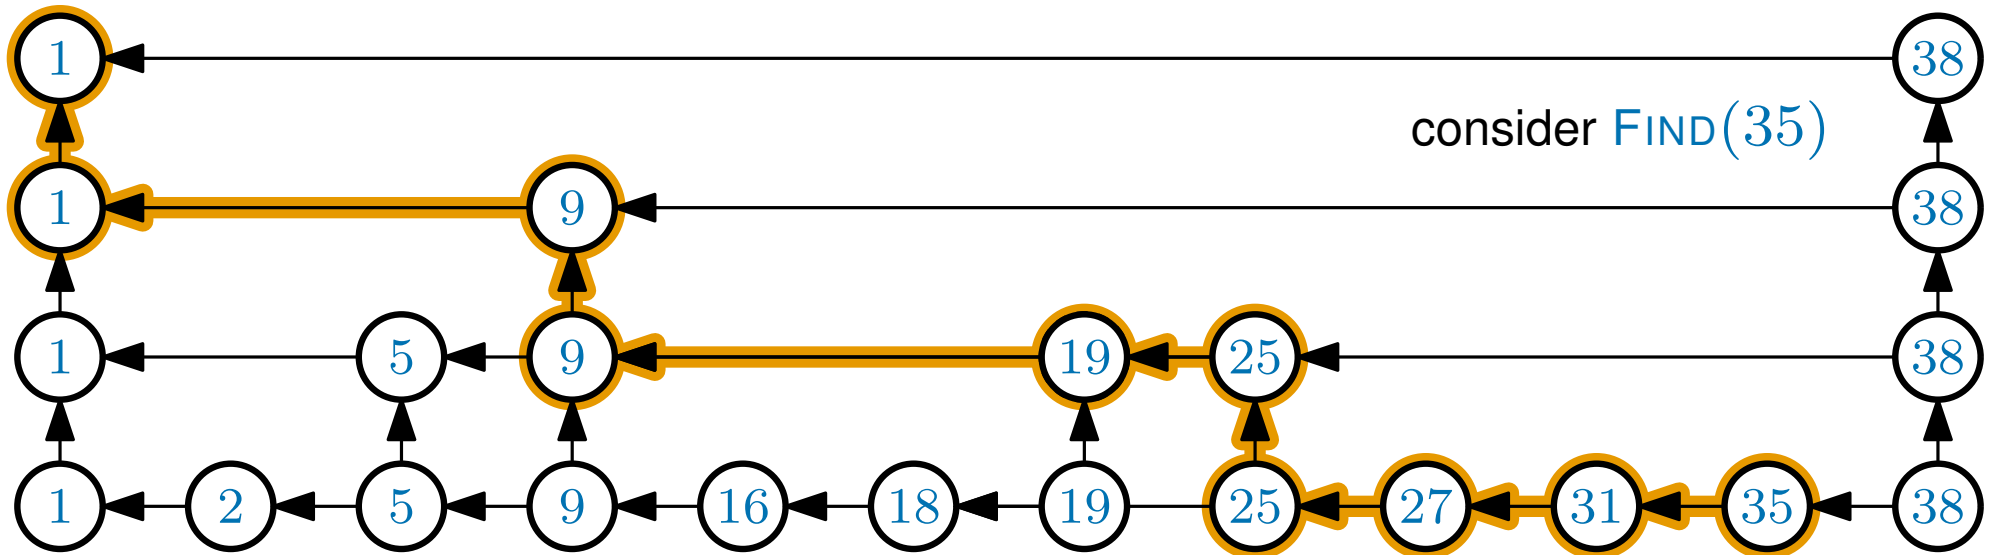
1. Reverse it

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?




1. Reverse it

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



How long is this path?

1. Reverse it
2. Convince yourself this is the same path:

Start at k

While not at the top-left:

If you can,
Move up

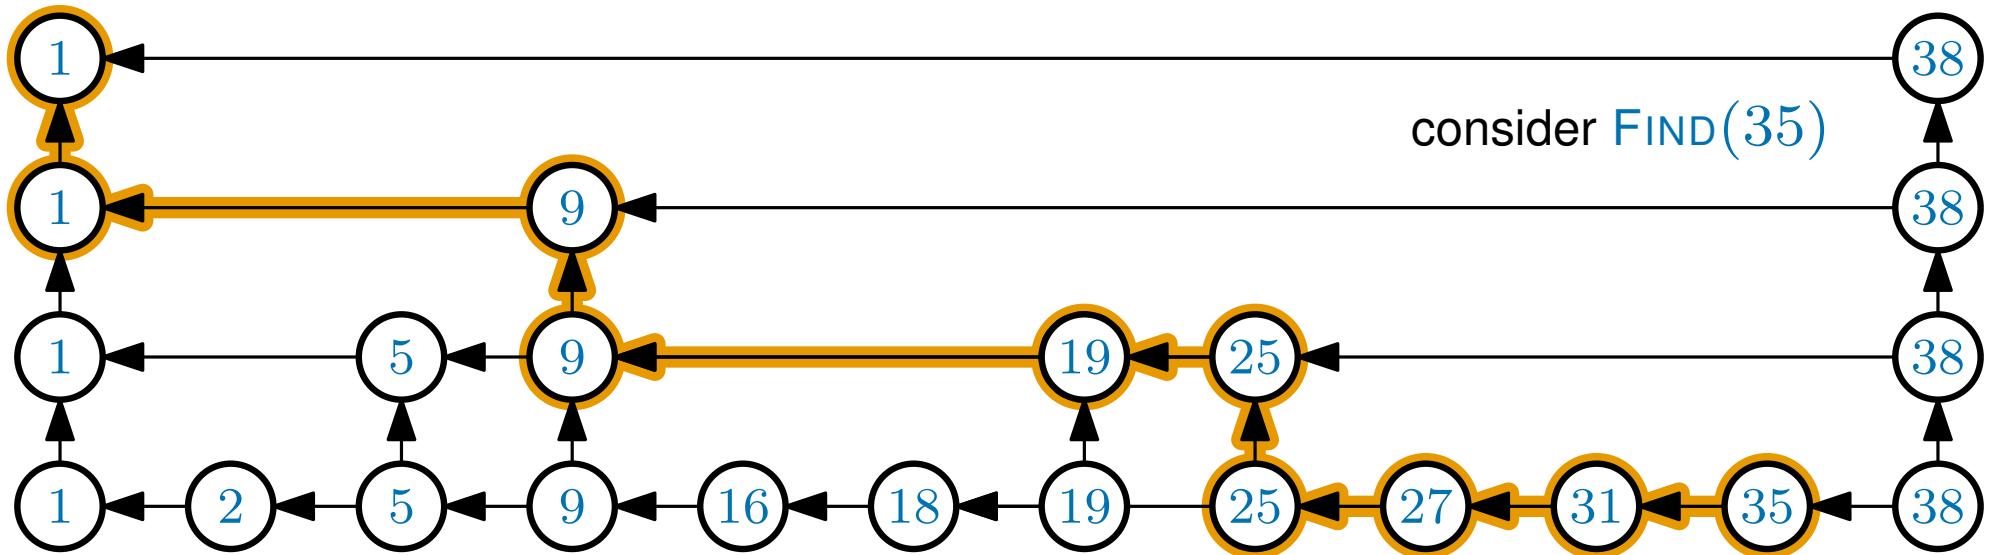
Else **Move left**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



How long is this path?

1. Reverse it
2. Convince yourself this is the same path:
3. Now convince yourself
it takes the same time as this:
(in expectation)

Start at k

While not at the top-left:

If (flip a coin)

Move up

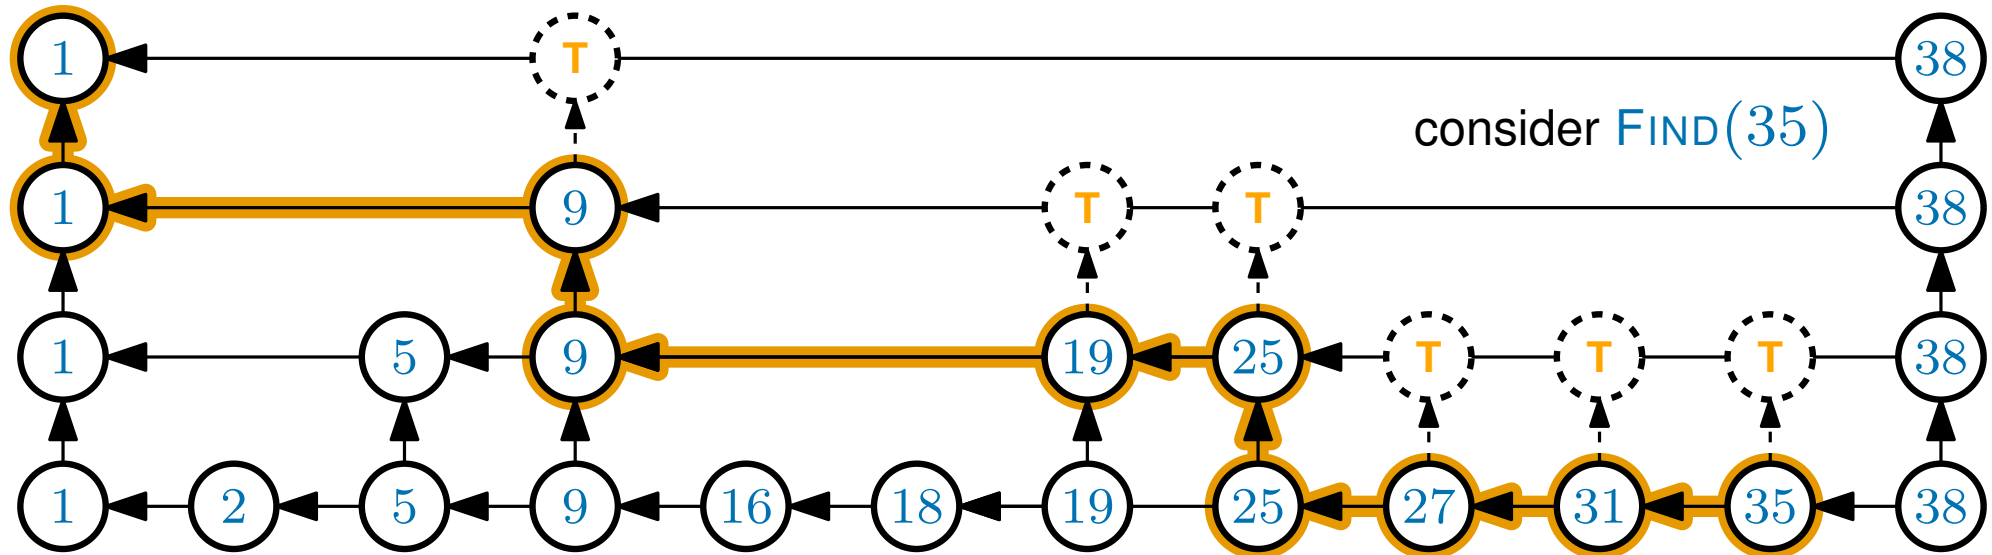
Else **Move left**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**?



How long is this path?

1. Reverse it
2. Convince yourself this is the same path:
3. Now convince yourself
it takes the same time as this:
(in expectation)

Start at k

While not at the top-left:

If (flip a coin)

Move up

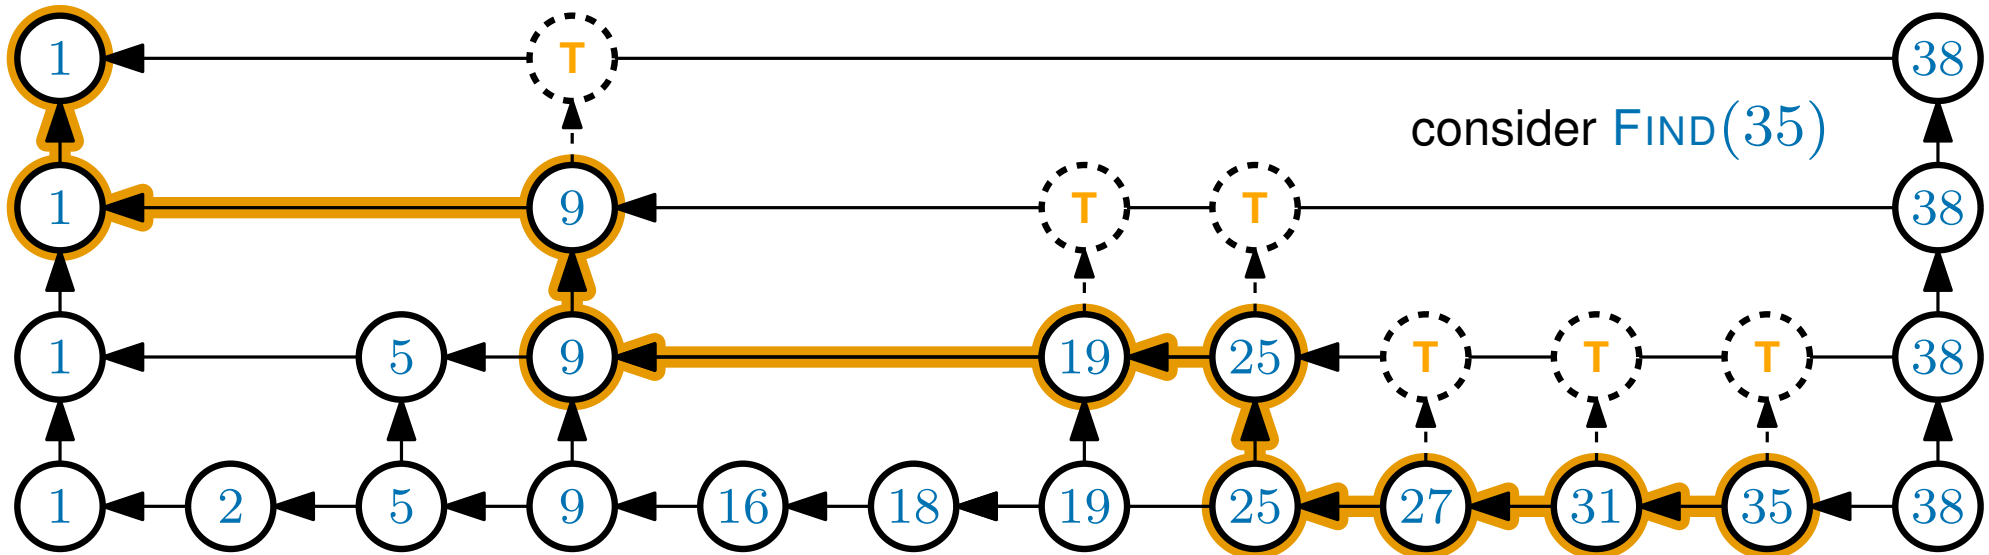
Else **Move left**

So how long does a **FIND** take? (sketch proof)

As the number of levels is $O(\log n)$ (with high probability),

we can conclude that the number of times we **move down** is very likely to be $O(\log n)$

but how many times do we **move right**? $O(\log n)$ in expectation



How long is this path?

1. Reverse it
2. Convince yourself this is the same path:
3. Now convince yourself
it takes the same time as this:
(in expectation)

Start at k

While not at the top-left:

If (flip a coin)

Move up

Else **Move left**

Time complexities

When performing a **FIND** operation, the number of moves is $O(\log n)$ in expectation

Time complexities

When performing a **FIND** operation, the number of moves is $O(\log n)$ in expectation
as each move takes $O(1)$ time, the expected time complexity is $O(\log n)$

Time complexities

When performing a **FIND** operation, the number of moves is $O(\log n)$ in expectation
as each move takes $O(1)$ time, the expected time complexity is $O(\log n)$

The number of levels is also $O(\log n)$ in expectation

Time complexities

When performing a **FIND** operation, the number of moves is $O(\log n)$ in expectation
as each move takes $O(1)$ time, the expected time complexity is $O(\log n)$

The number of levels is also $O(\log n)$ in expectation

Both **INSERT** and **DELETE** also take expected $O(\log n)$ time
*this is because they both call **FIND** and then spend $O(1)$ time per level*

Time complexities

When performing a **FIND** operation, the number of moves is $O(\log n)$ in expectation
 as each move takes $O(1)$ time, the expected time complexity is $O(\log n)$

The number of levels is also $O(\log n)$ in expectation

Both **INSERT** and **DELETE** also take expected $O(\log n)$ time
*this is because they both call **FIND** and then spend $O(1)$ time per level*

In fact, all three operations actually take $O(\log n)$ time
with high probability

i.e. the probability of an operation taking longer is at most $\frac{1}{n}$

Time complexities

When performing a **FIND** operation, the number of moves is $O(\log n)$ in expectation
as each move takes $O(1)$ time, the expected time complexity is $O(\log n)$

The number of levels is also $O(\log n)$ in expectation

Both **INSERT** and **DELETE** also take expected $O(\log n)$ time
*this is because they both call **FIND** and then spend $O(1)$ time per level*


In fact, all three operations actually take $O(\log n)$ time
with high probability

i.e. the probability of an operation taking longer is at most $\frac{1}{n}$

(this is a stronger claim but proving it is harder)

Skip Lists (post-proof) summary

A skip list is a **randomised** data structure, based on link lists with **shortcuts**
which supports $\text{INSERT}(x, k)$, $\text{FIND}(k)$ and $\text{DELETE}(k)$



each of these operations takes *expected* $O(\log n)$ time

That is, they take $O(\log n)$ time 'on average'

Important There is *no randomness in the data*,
the only randomness is in the coin flips

On *the worst case* input sequence, the expected time is $O(\log n)$

Dynamic Search Structure Summary

A **dynamic search structure** supports (at least) the following three operations

$\text{DELETE}(k)$ - deletes the (unique) element x with $x.\text{key} = k$

$\text{INSERT}(x, k)$ - inserts x with key $k = x.\text{key}$

$\text{FIND}(k)$ - returns the (unique) element x with $x.\text{key} = k$

Here are the time complexities of the structures we have seen...

	INSERT	DELETE	FIND
Unsorted Linked List	$O(1)$	$O(n)$	$O(n)$
Binary Search Tree	$O(n)$	$O(n)$	$O(n)$
2-3-4 Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$
Red-Black Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$
Skip list	$O(\log n)$	$O(\log n)$	$O(\log n)$

The time complexities for the Skip list are *expected*, for the others, they are *worst case*