

Distance-based methods


- L_k norm or *Minkowski metric* for d -dimensional vectors

$$L_k(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^d |x_i - y_i|^k \right)^{1/k}$$

- $k=1$: Manhattan distance $L_1(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^d |x_i - y_i|$

- $k=2$: Euclidean distance $L_2(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^d (x_i - y_i)^2} = \sqrt{(\mathbf{x} - \mathbf{y})^\top (\mathbf{x} - \mathbf{y})} = \|\mathbf{x} - \mathbf{y}\|$

- $k=\infty$: $L_\infty(\mathbf{x}, \mathbf{y}) = \max_i |x_i - y_i|$



Plot of all points at distance 1 from the origin for different values of k ($d=2$)

Distances and (in)variance

- Euclidean distance is invariant under translation and rotation, but sensitive to scaling
 - multiplying a dimension with a factor $f > 1$ ($g < 1$) will increase (decrease) its influence in the distance calculation
- Therefore we usually re-weight all dimensions to have unit variance

$$\sqrt{\sum_{i=1}^d \frac{(x_i - y_i)^2}{\sigma_i^2}} = \sqrt{(\mathbf{x} - \mathbf{y})^T \mathbf{W} (\mathbf{x} - \mathbf{y})} \text{ with } \mathbf{W} = \begin{bmatrix} 1/\sigma_1^2 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & 1/\sigma_d^2 \end{bmatrix}$$


- unit-distance circle becomes axis-parallel ellipse
- If we also want to eliminate correlations between features, we set \mathbf{W} to the inverse of the covariance matrix (**Mahalanobis distance**)
 - unit-distance ellipse is rotated

Symbolic distance

- For equal-length bit strings, the *Hamming distance* is the number of positions in which the two strings differ
 - 101100 and 011101 have Hamming distance 3
 - corresponds to the L_1 distance
 - can also do L_k , just a monotonic transformation in this case
- Can be generalised in various ways
 - symbolic unordered features: 0 if same, 1 if different
 - symbolic ordered features: represent by (unit-variance) numbers
 - strings of different lengths (e.g. words): edit distance
- In what follows we assume some distance metric on feature vectors with possibly mixed numeric and symbolic features


A close neighbour...

- **Nearest-neighbour classification:** store a set of exemplars (e.g. the training data), and assign a test instance to the class of its nearest exemplar
 - instance space divided by **Voronoi tessellation**
 - piecewise linear decision boundary



Increasing the number of neighbours

- ***k-Nearest-neighbour classification:*** assign a test instance to the majority class among its k nearest neighbours
 - smoother decision boundary than with $k=1$
 - $k \rightarrow \infty$: all instances are assigned most common class in training set (i.e., using prior distribution only)
 - k can be tuned using separate test set



To classify a test instance, draw the smallest hypersphere around it that contains k neighbours; then assign majority class among those

Example: 2-nearest-neighbour




Figure 8.8. (left) Voronoi tessellation for five exemplars. (middle) Taking the two nearest exemplars into account leads to a further subdivision of each Voronoi cell. (right) The shading indicates which exemplars contribute to which cell.

k-nearest neighbour: variations

- In *k*-NN, each of the *k* nearest neighbours \mathbf{x}_i contributes equally to the vote for the most likely class of the test instance \mathbf{x}
- **Distance-weighted *k*-NN**: weight the vote with a **kernel** $K(\mathbf{x}-\mathbf{x}_i)$ that decreases with the distance
 - e.g. Gaussian kernel $K(\mathbf{d}) \propto e^{-\|\mathbf{d}\|^2}$
- Using distance-weighting, we can also let $k \rightarrow \infty$ and have **all** training instances contribute to the vote
 - **global** rather than local method
- All these methods can also be used for **estimation**, where the dependent variable is a scalar rather than a class

Discussion

- k -nearest neighbour methods take **all** features into account. With large numbers of features this causes problems
 - irrelevant features dominate distance calculations
 - training set covers only a fraction of instance space, which causes overfitting: ‘curse of dimensionality’
- This can be addressed in several ways
 - stretch/shrink dimensions if it improves performance on a separate test set
 - perform **feature selection** in a pre-processing step → later lecture

Clustering

- **Clustering** involves segmenting the instance space into regions of similar objects
 - no guidance through labelled training instances -> *unsupervised learning*
- We can again use the idea of a distance metric -> **K-means clustering** (this lecture)
- We can also use the class as a 'hidden' variable -> **Gaussian mixture models** (later)


K-means clustering

```
function KMeans(Instances, K)  
  randomly initialise K vectors  $\mu_1 \dots \mu_K$ ;  
  repeat  
    assign each  $x \in \text{Instances}$  to the nearest  $\mu_j$ ;  
    recompute each  $\mu_j$  as the mean of the  
      instances assigned to it;  
  until no change in  $\mu_1 \dots \mu_K$ ;  
  return  $\mu_1 \dots \mu_K$ ;
```


- Notes:

- different initialisation schemes are possible for the centroids
- alternatively, we can partition the instances into K groups and calculate the initial means from those


K-means example



K-means example



K-means example



Analysis

- What does it do?
 - **within-cluster scatter**: total squared distance between points \mathbf{x}_i and centroid μ_k in k -th cluster: $\sum_i \|\mathbf{x}_i - \mu_k\|^2$
 - K -means attempts to find a configuration $\mu_1 \dots \mu_K$ that minimises within-cluster scatter summed over all clusters
 - this is equivalent to maximising the between-cluster scatter (move all data points to their centroid and calculate total squared distance to the global centroid)
- Does it work?
 1. The algorithm terminates.
 2. It finds a **local optimum** from which no further improvement is possible by making local changes.
 3. It does not necessarily find a **global optimum**.

Local optimum with K -means




Figure 8.12. (left) First iteration of 3-means on the same data as Figure 8.11 with differently initialised centroids. (right) 3-means has converged to a sub-optimal clustering.

Discussion

- To be reasonably sure that we're not too far from the global optimum, we should run *K*-means a number of times with different initial configurations
 - or use more clever and/or domain-dependent initialisation schemes
- *K*-medoids is a variant of *K*-means in which cluster center is a **data point** (rather than the mean of data points) that minimises within-cluster scatter
 - allows to use distance metrics other than Euclidean distance, which tends to be sensitive to outliers
- We can also use “soft” assignment of data points to cluster centers -> Gaussian mixture models